

# Bigsets...CRDT sets but BIGGER

Russell Brown, Basho Technologies  
[russelldb@basho.com](mailto:russelldb@basho.com)



**basho**







**riak KV**



**riak TS**



**riak S2**

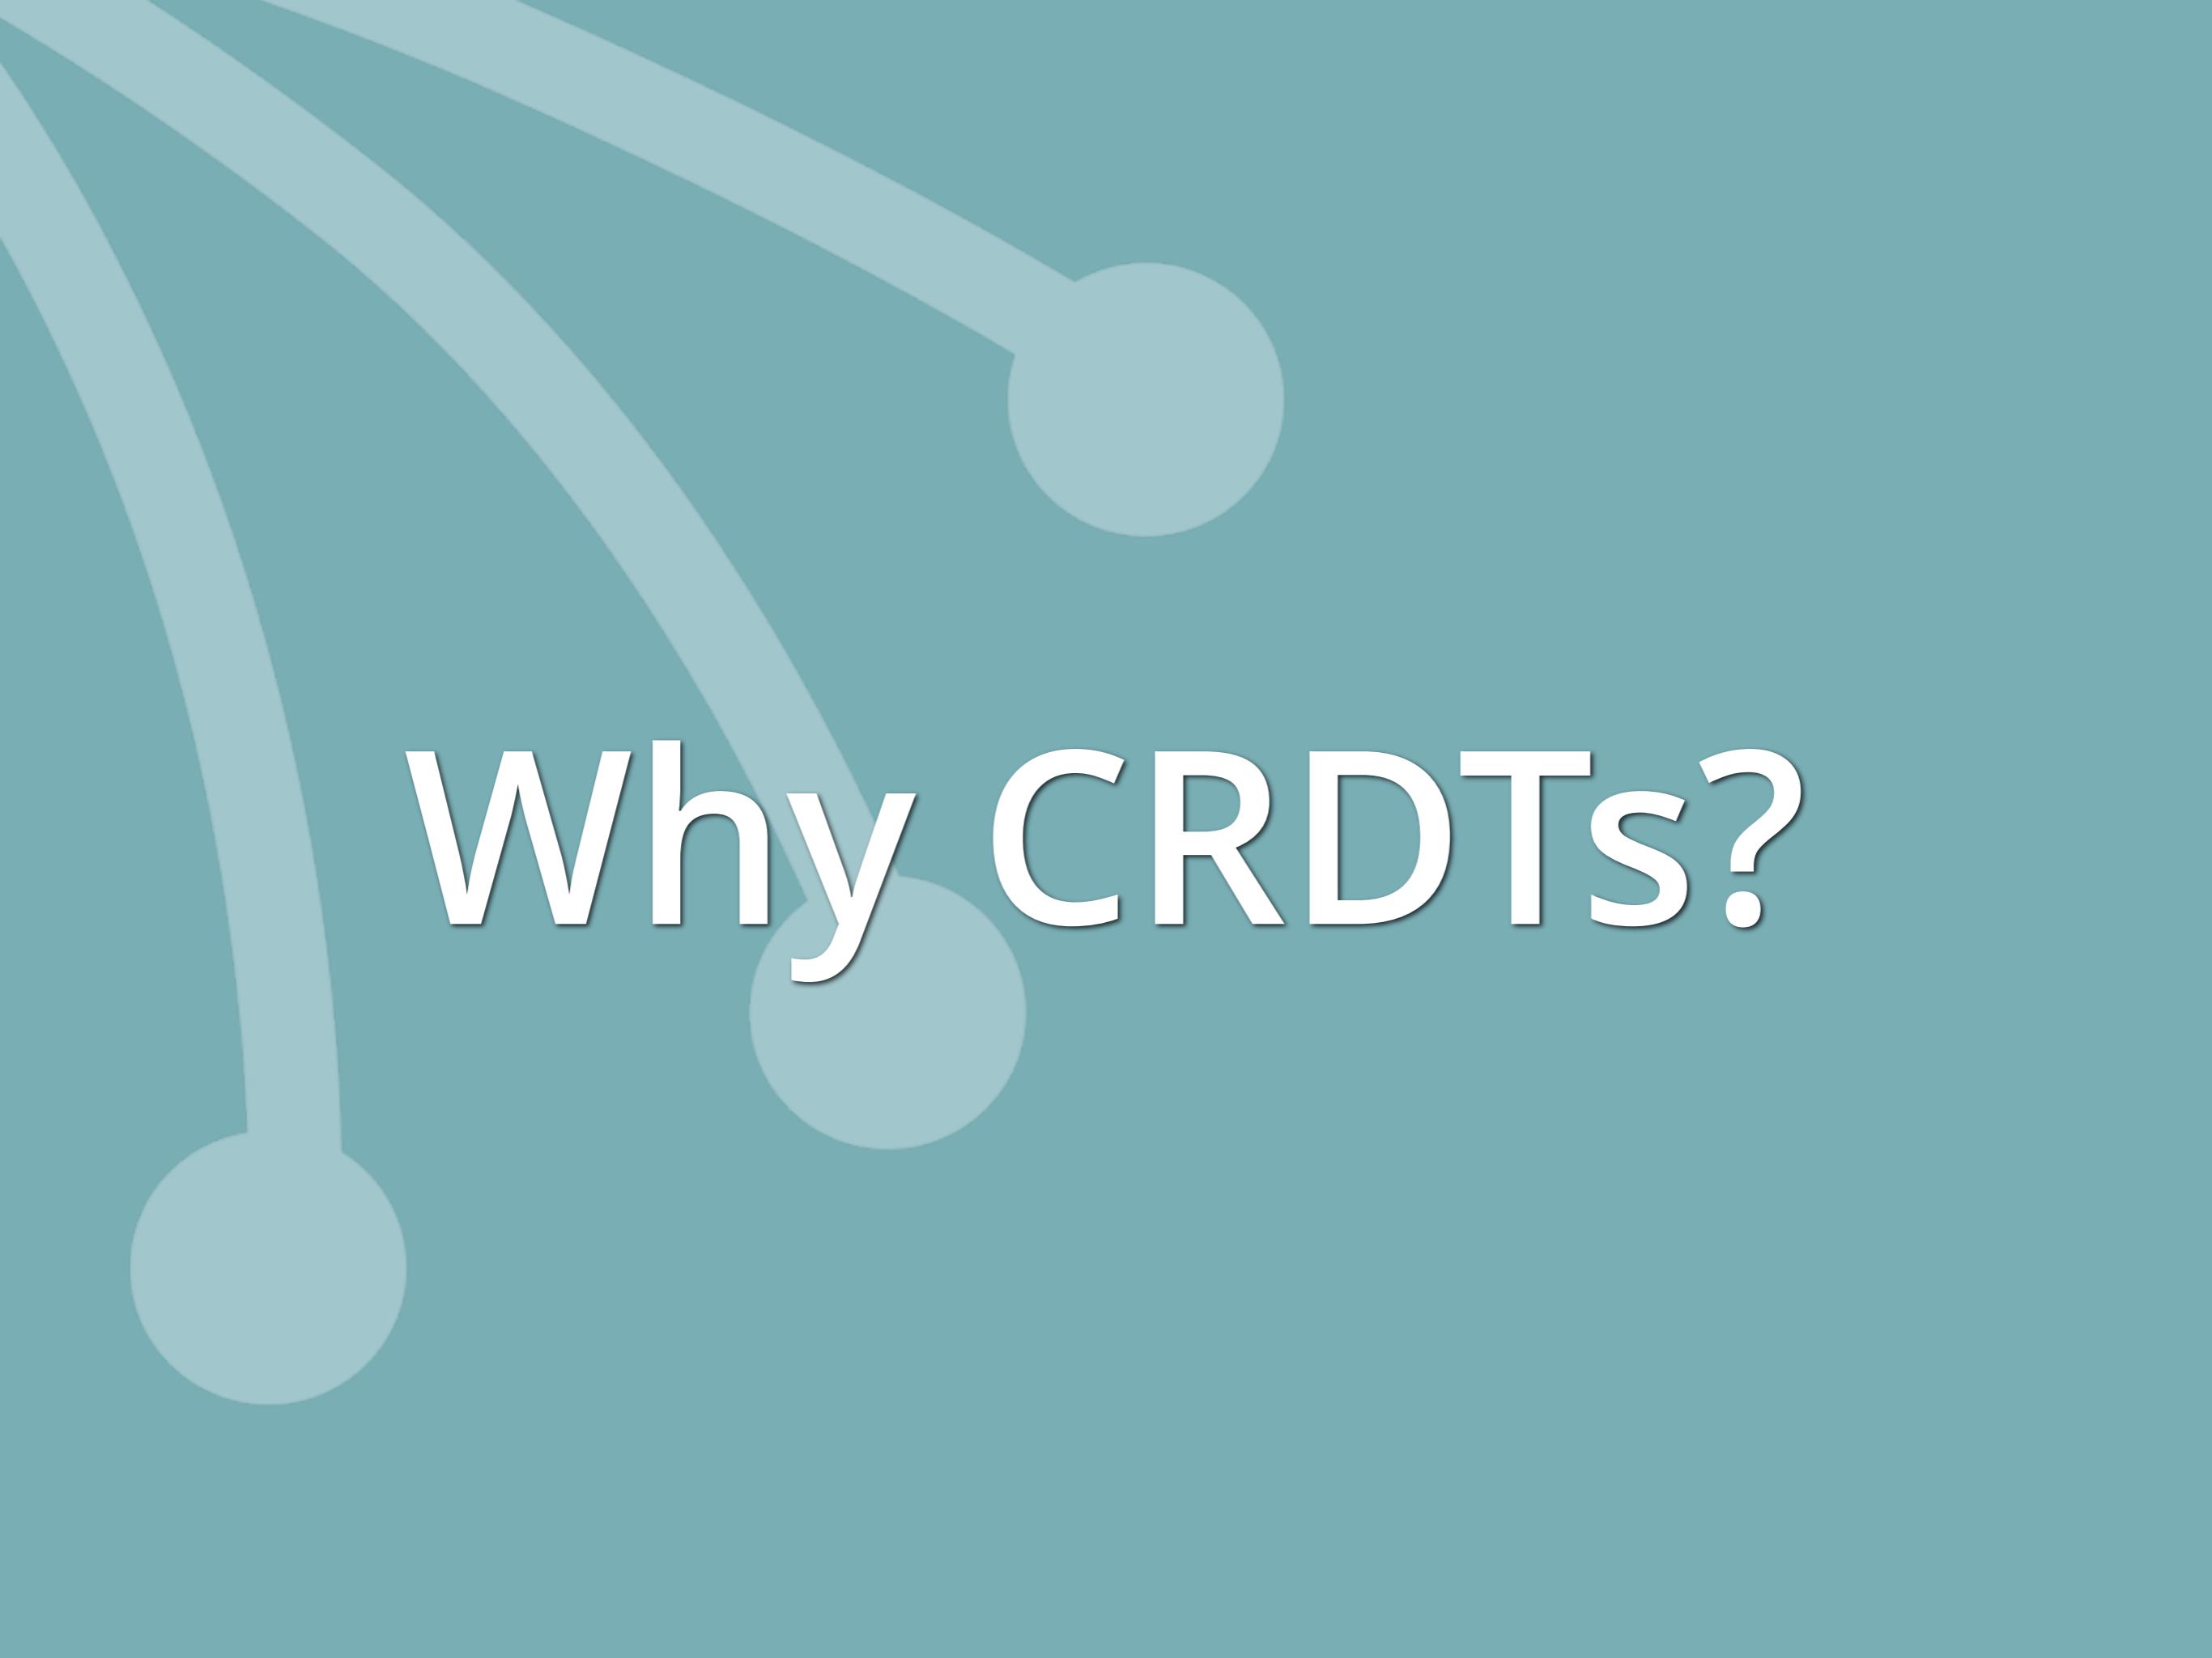


This project is funded by the European  
Union,  
7th Research Framework Programme, ICT  
call 10,  
grant agreement n°609551.



# 4 Sections

- 1.What are CRDTs (good for)?
- 2.History of CRDT Sets
- 3.Sets in Riak
- 4.Bigger Sets in Riak



# Why CRDTs?

# Fundamental Trade Off

- Lipton/Sandberg '88
- Attiya/Welch '94
- Gilbert/Lynch '02

## **Low Latency/Availability:**

- Increased Revenue
- User Engagement

## **Strong Consistency:**

- Easier for Programmers
- Less user “surprise”

# Dynamo: Amazon's Highly Available Key-value Store

Giuseppe DeCandia, Deniz Hastorun, Madan Jampani, Gunavardhan Kakulapati,  
Avinash Lakshman, Alex Pilchin, Swaminathan Sivasubramanian, Peter Vosshall  
and Werner Vogels

Amazon.com

## ABSTRACT

Reliability at massive scale is one of the biggest challenges we face at Amazon.com, one of the largest e-commerce operations in the world; even the slightest outage has significant financial consequences and impacts customer trust. The Amazon.com platform, which provides services for many web sites worldwide, is implemented on top of an infrastructure of tens of thousands of servers and network components located in many datacenters around the world. At this scale, small and large components fail continuously and the way persistent state is managed in the face of these failures drives the reliability and scalability of the software systems.

This paper presents the design and implementation of Dynamo, a highly available key-value storage system that some of Amazon's core services use to provide an "always-on" experience. To achieve this level of availability, Dynamo sacrifices consistency under certain failure scenarios. It makes extensive use of object versioning and application-assisted conflict resolution in a manner that provides a novel interface for developers to use.

## Categories and Subject Descriptors

D.4.2 [Operating Systems]: Storage Management; D.4.5 [Operating Systems]: Reliability; D.4.2 [Operating Systems]: Performance;

## General Terms

One of the lessons our organization has learned from operating Amazon's platform is that the reliability and scalability of a system is dependent on how its application state is managed. Amazon uses a highly decentralized, loosely coupled, service oriented architecture consisting of hundreds of services. In this environment there is a particular need for storage technologies that are always available. For example, customers should be able to view and add items to their shopping cart even if disks are failing, network routes are flapping, or data centers are being destroyed by tornados. Therefore, the service responsible for managing shopping carts requires that it can always write to and read from its data store, and that its data needs to be available across multiple data centers.

Dealing with failures in an infrastructure comprised of millions of components is our standard mode of operation; there are always a small but significant number of server and network components that are failing at any given time. As such Amazon's software systems need to be constructed in a manner that treats failure handling as the normal case without impacting availability or performance.

To meet the reliability and scaling needs, Amazon has developed a number of storage technologies, of which the Amazon Simple Storage Service (also available outside of Amazon and known as Amazon S3), is probably the best known. This paper presents the design and implementation of Dynamo, another highly available



Created by Creative Stall  
from Noun Project



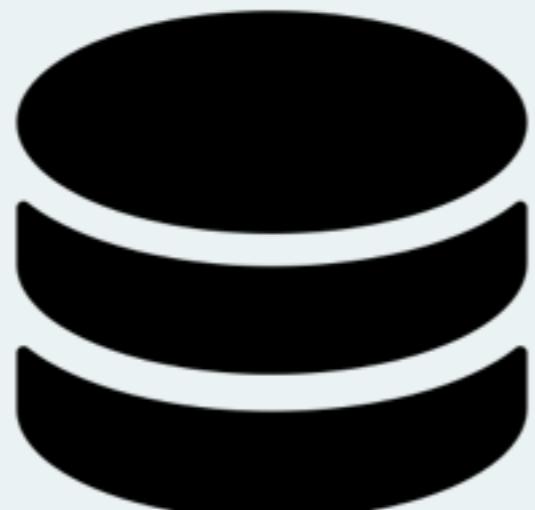
Created by Creative Stall  
from Noun Project



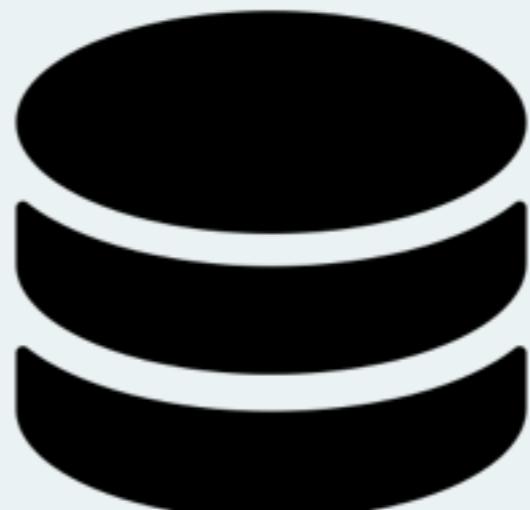
Created by Creative Stall  
from Noun Project



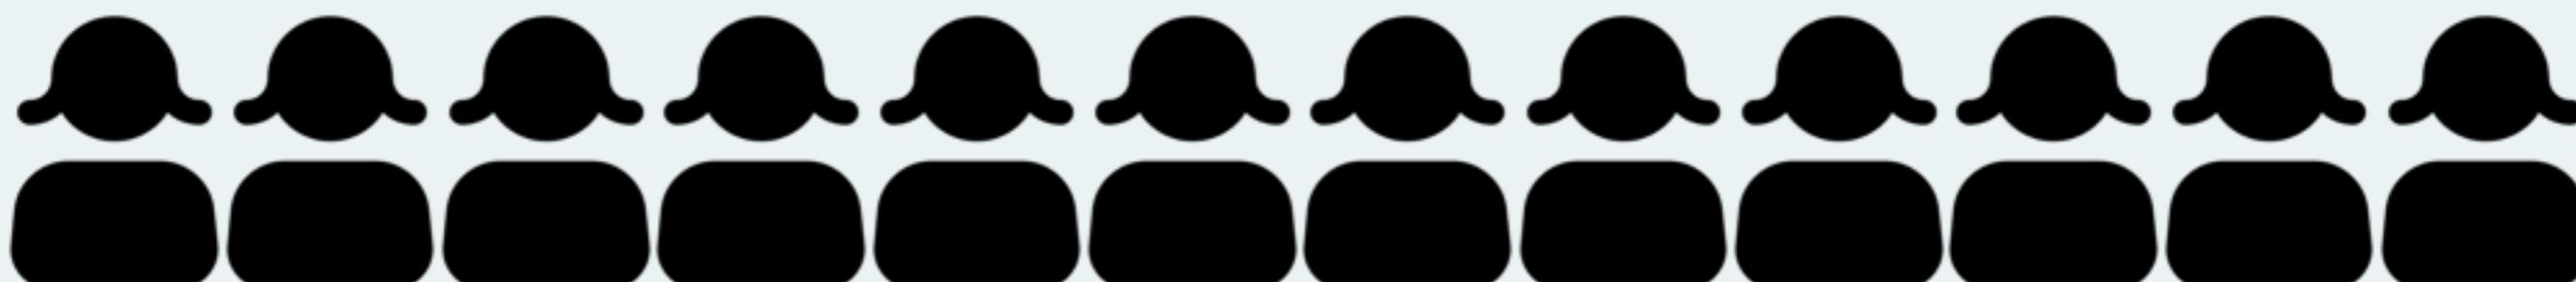
Created by Creative Stall  
from Noun Project



Created by Creative Stall  
from Noun Project



Created by Creative Stall  
from Noun Project



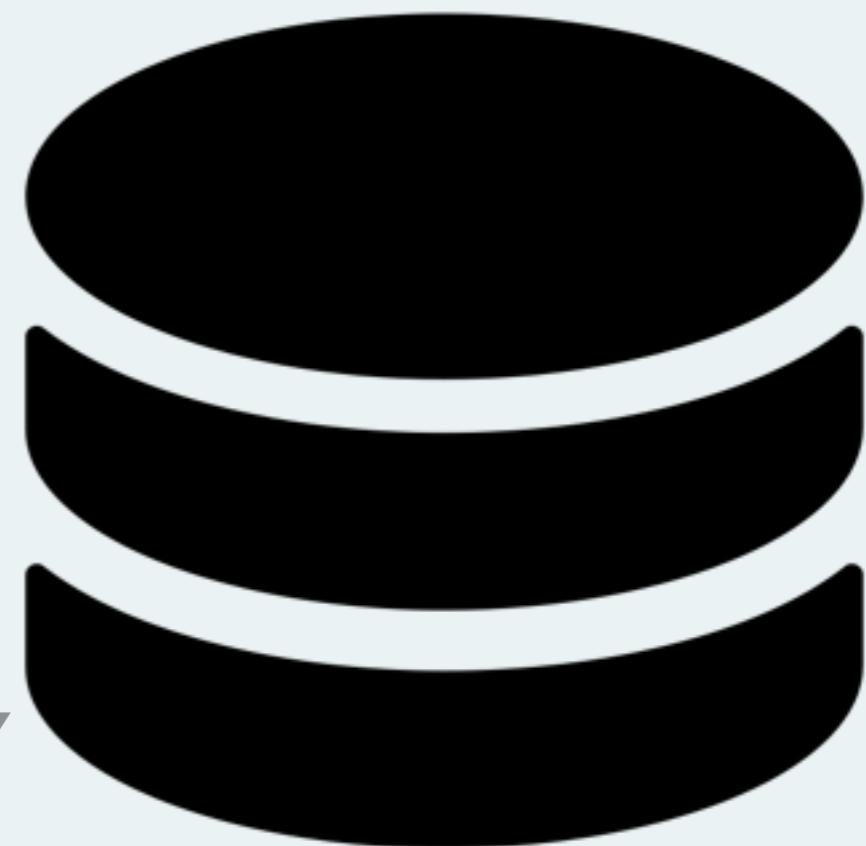
Created by Pumpkin Juice  
from Noun Project

# 2 REPLICAS



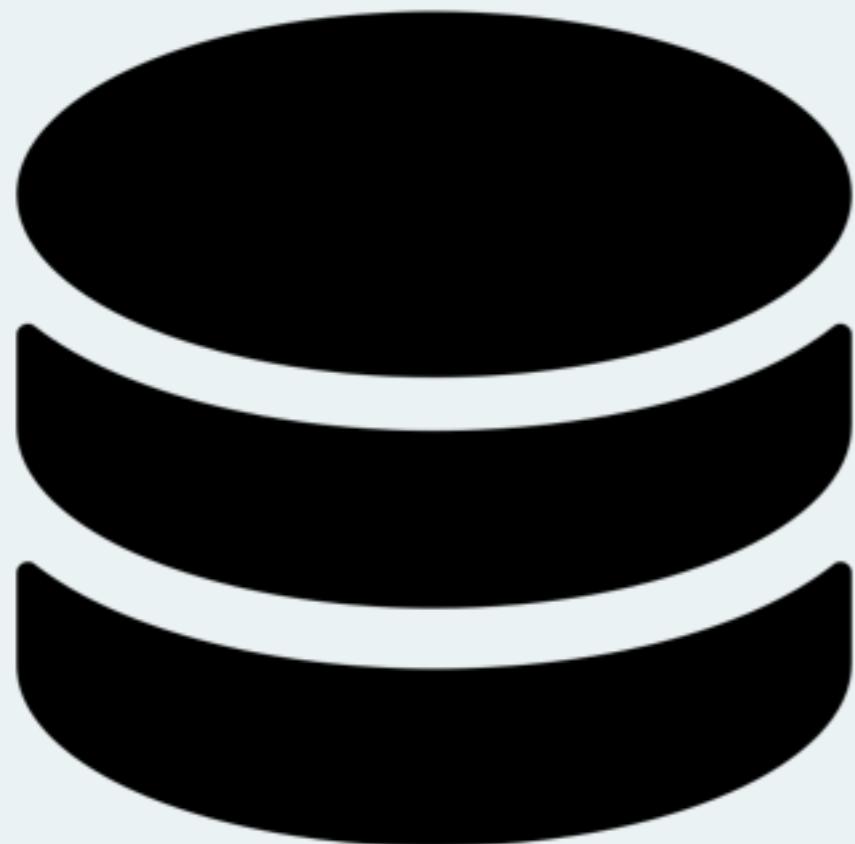
Created by Creative Stall  
from Noun Project

1 KEY



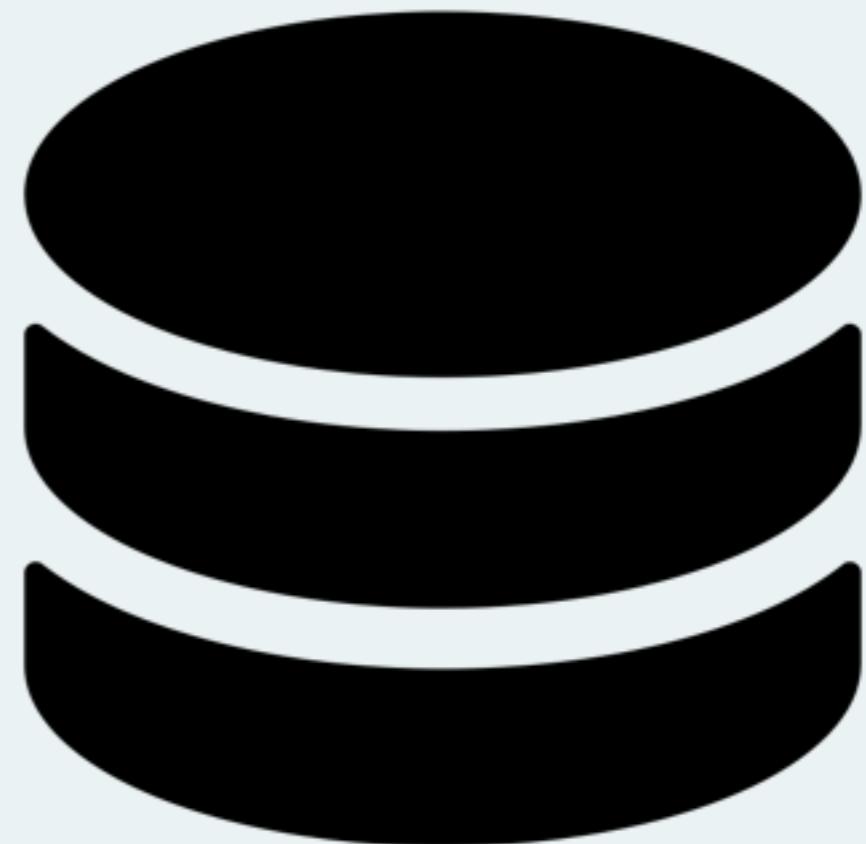
Created by Creative Stall  
from Noun Project

# 1 CLIENT



Created by Creative Stall  
from Noun Project

**REPLICATE**



Created by Creative Stall  
from Noun Project

**GET**



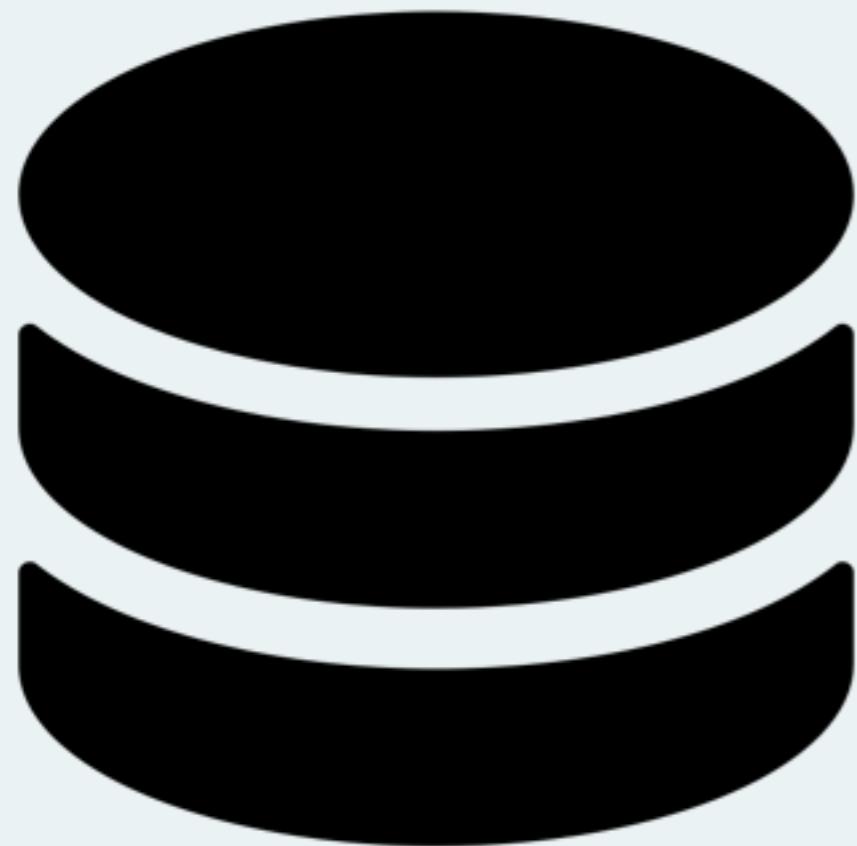
Created by Amy Schwartz  
from the Noun Project

**PUT**



Created by Amy Schwartz  
from the Noun Project

**UPDATE**

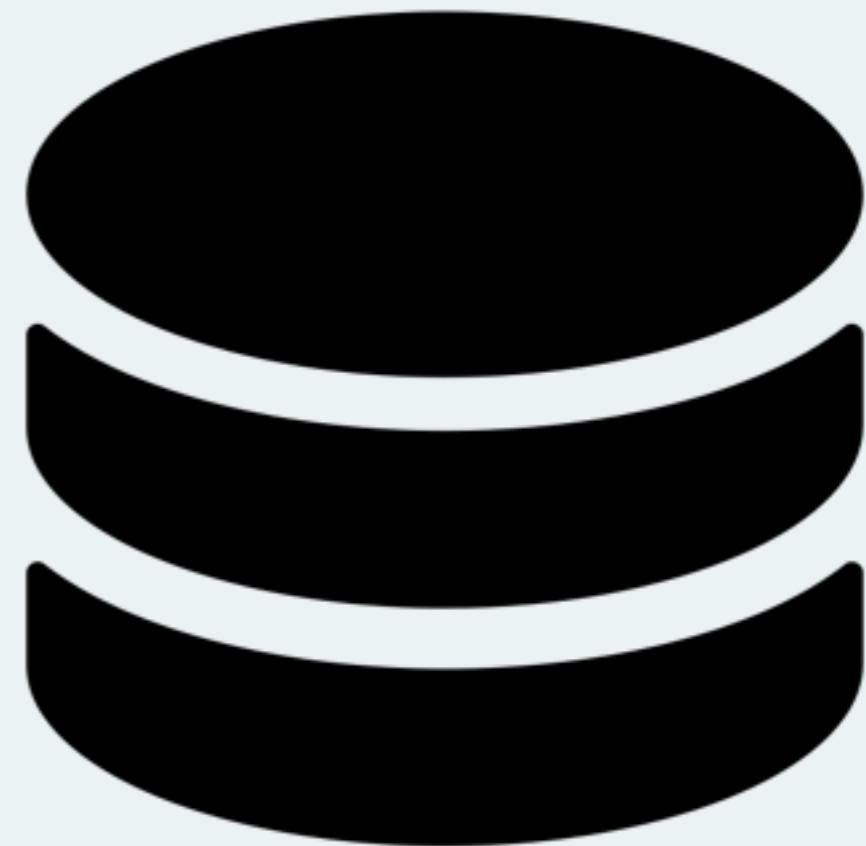
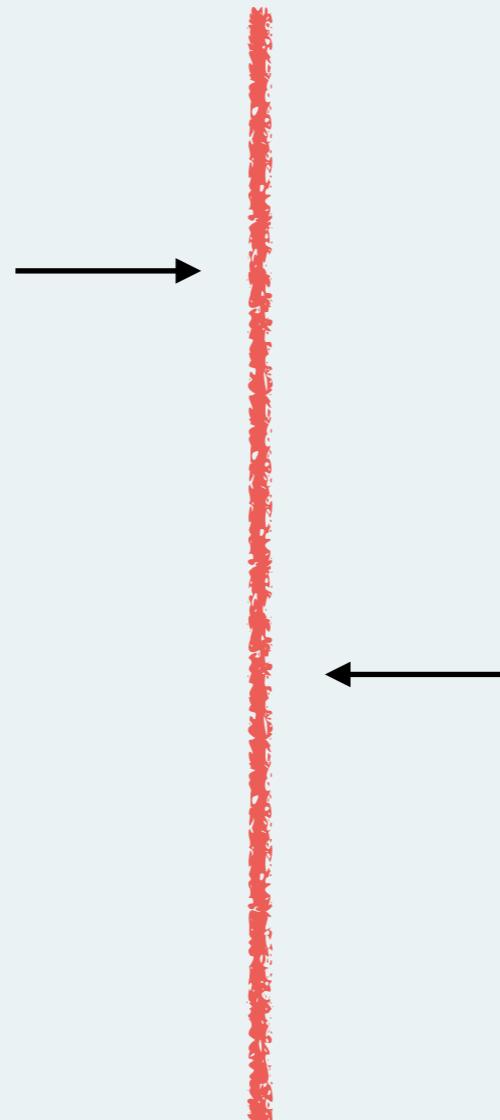


Created by Creative Stall  
from Noun Project

**PUT**



Created by Amy Schwartz  
from the Noun Project



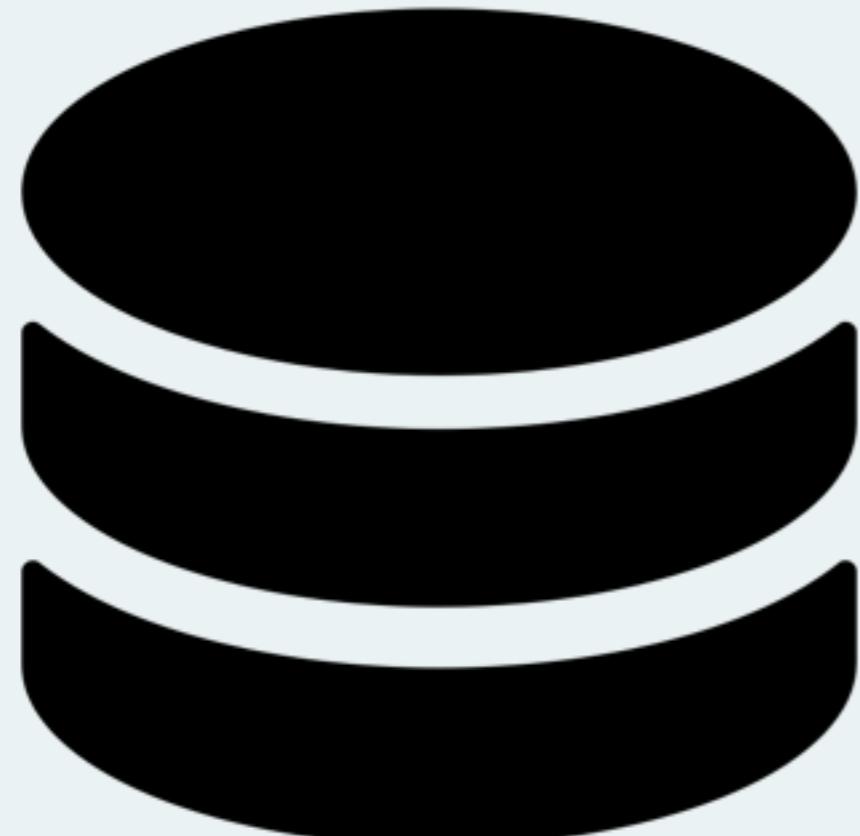
Created by Creative Stall  
from Noun Project

**PUT**

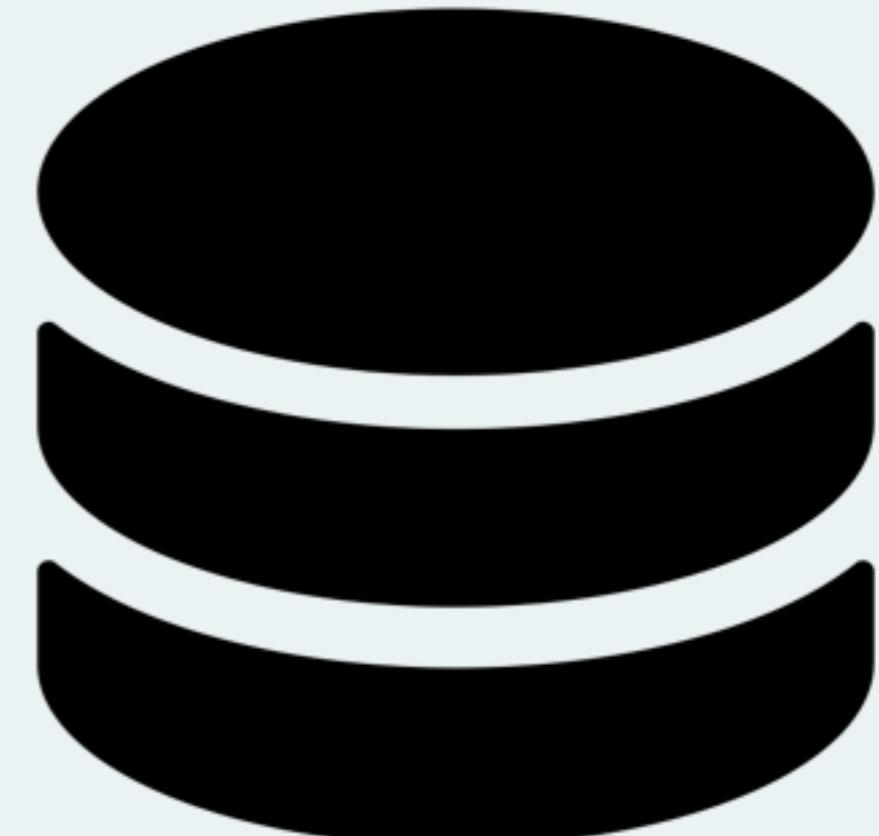


Created by Amy Schwartz  
from the Noun Project

# Quorum



Created by Creative Stall  
from Noun Project

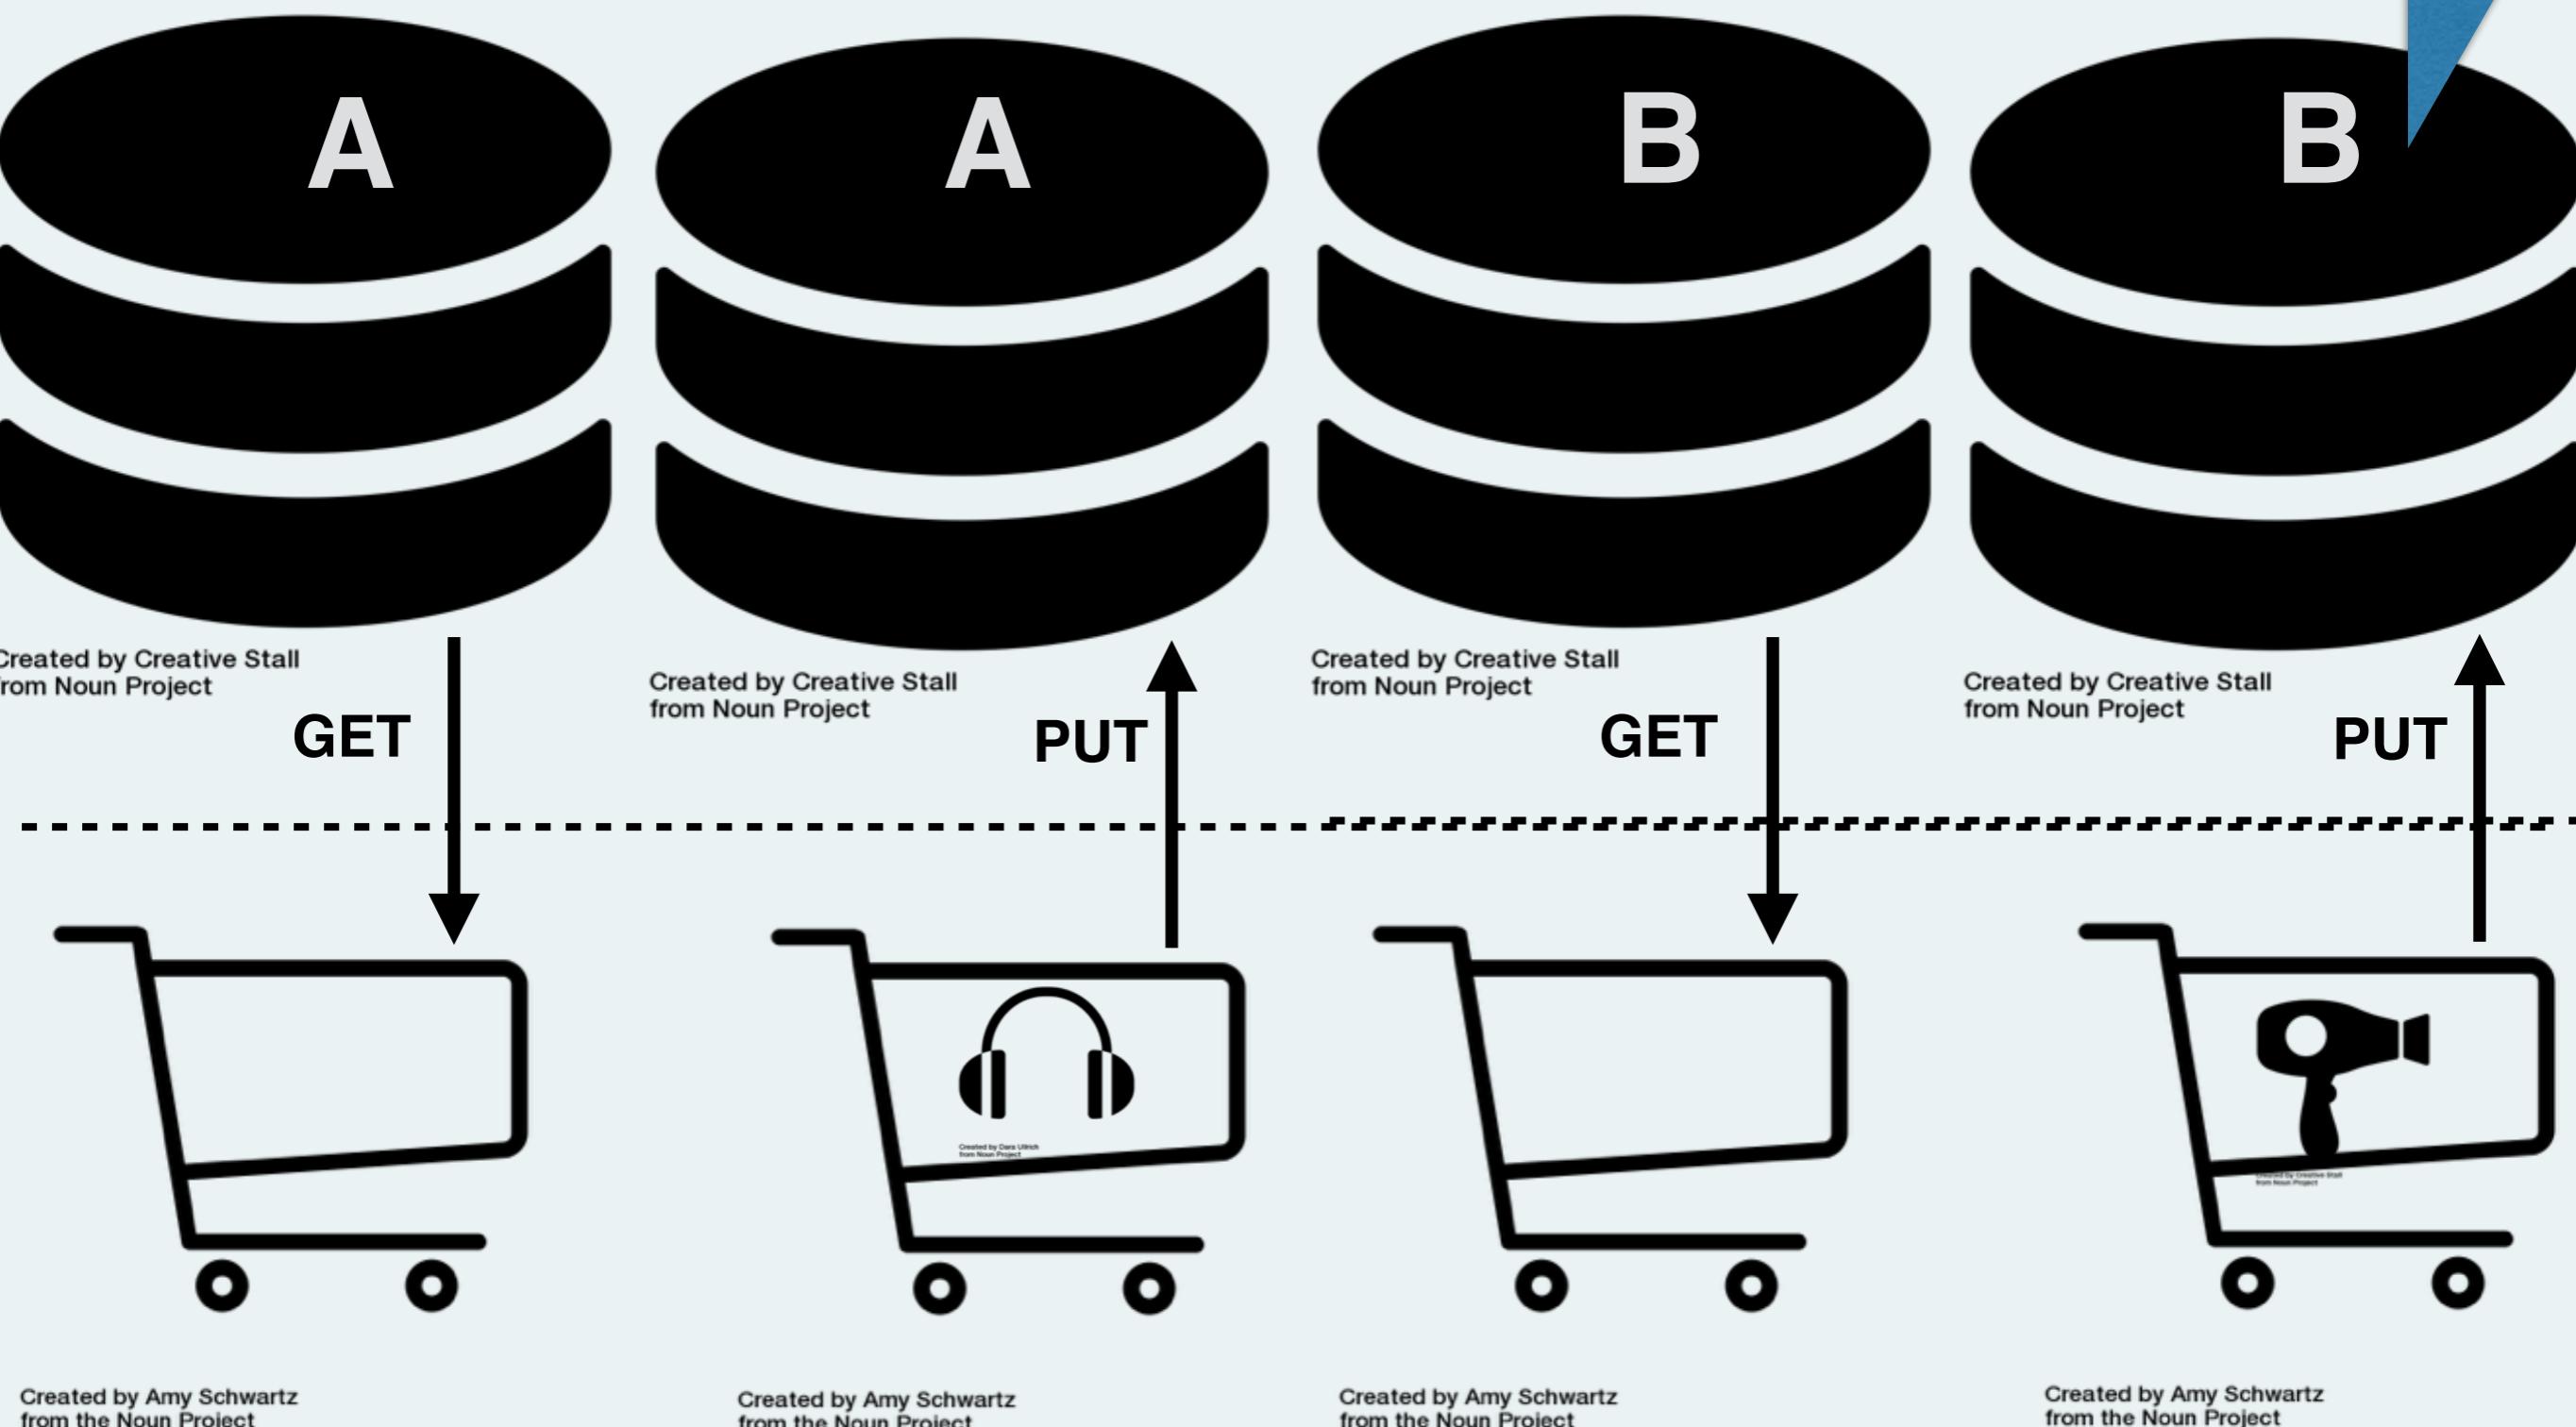


Created by Creative Stall  
from Noun Project



Created by Charlie Bob Gordon  
from Noun Project

# TEMPORAL TIME

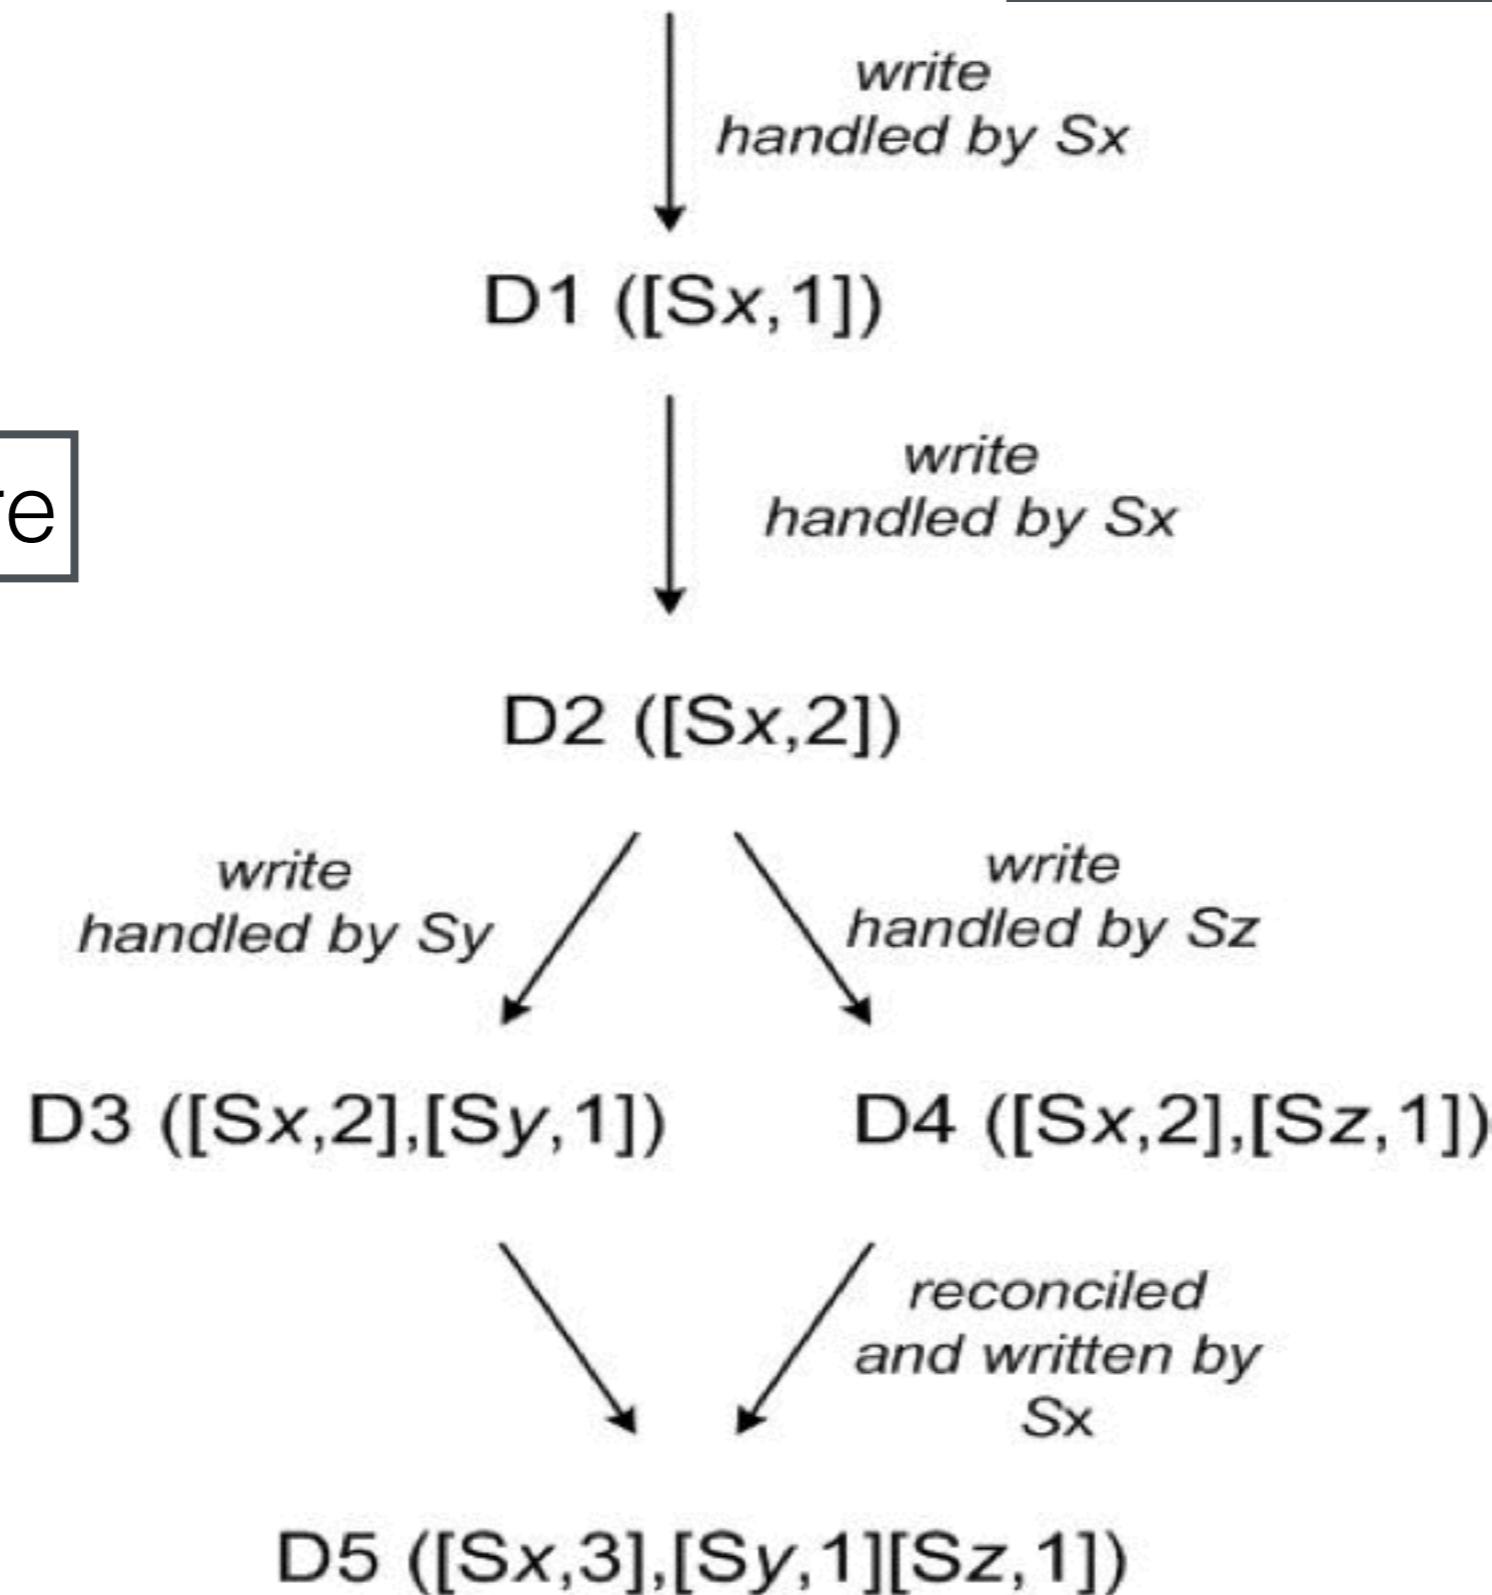


# Logical Clocks

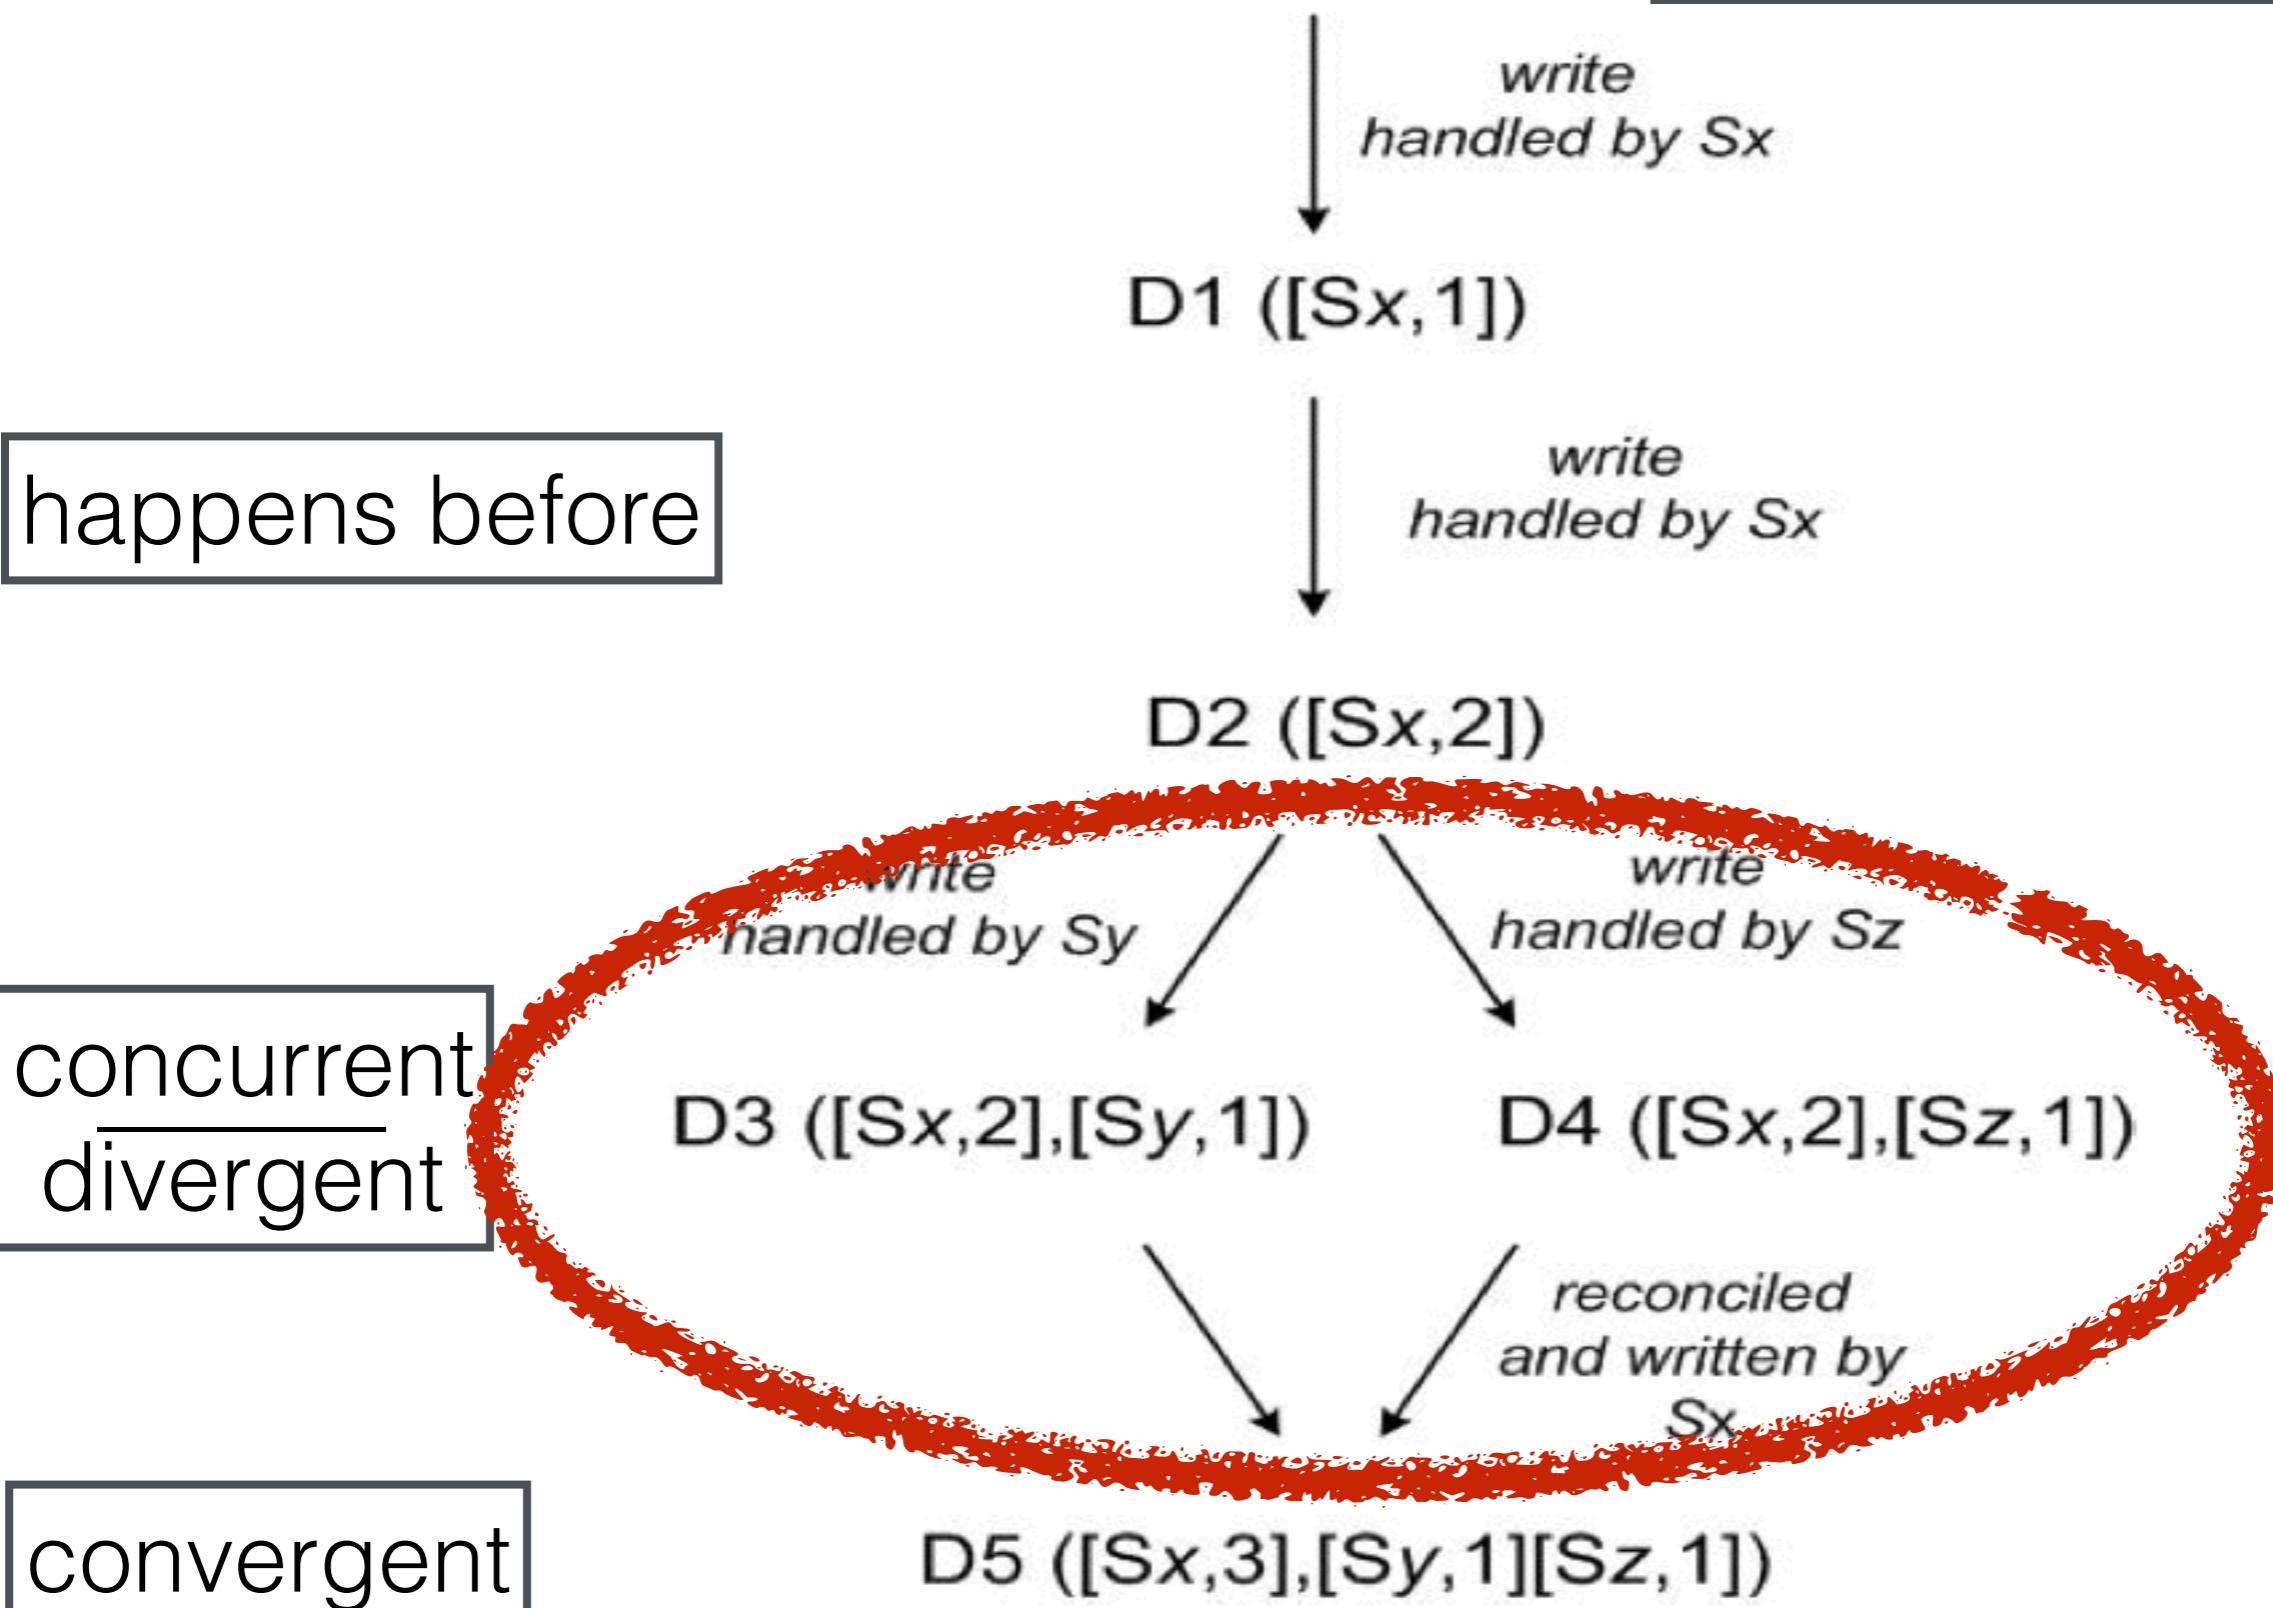
happens before

concurrent  
divergent

convergent



# Logical Clocks





```
if (result.hasConflicts()) {  
    // TODO: What should we do???
```

# Timestamp based reconciliation

155196119890

155196118001

>



Created by Amy Schwartz  
from the Noun Project

Created by Amy Schwartz  
from the Noun Project



**Created by Amy Schwartz  
from the Noun Project**

# Business Logic/Semantic Reconciliation



Created by Amy Schwartz  
from the Noun Project



Created by Amy Schwartz  
from the Noun Project



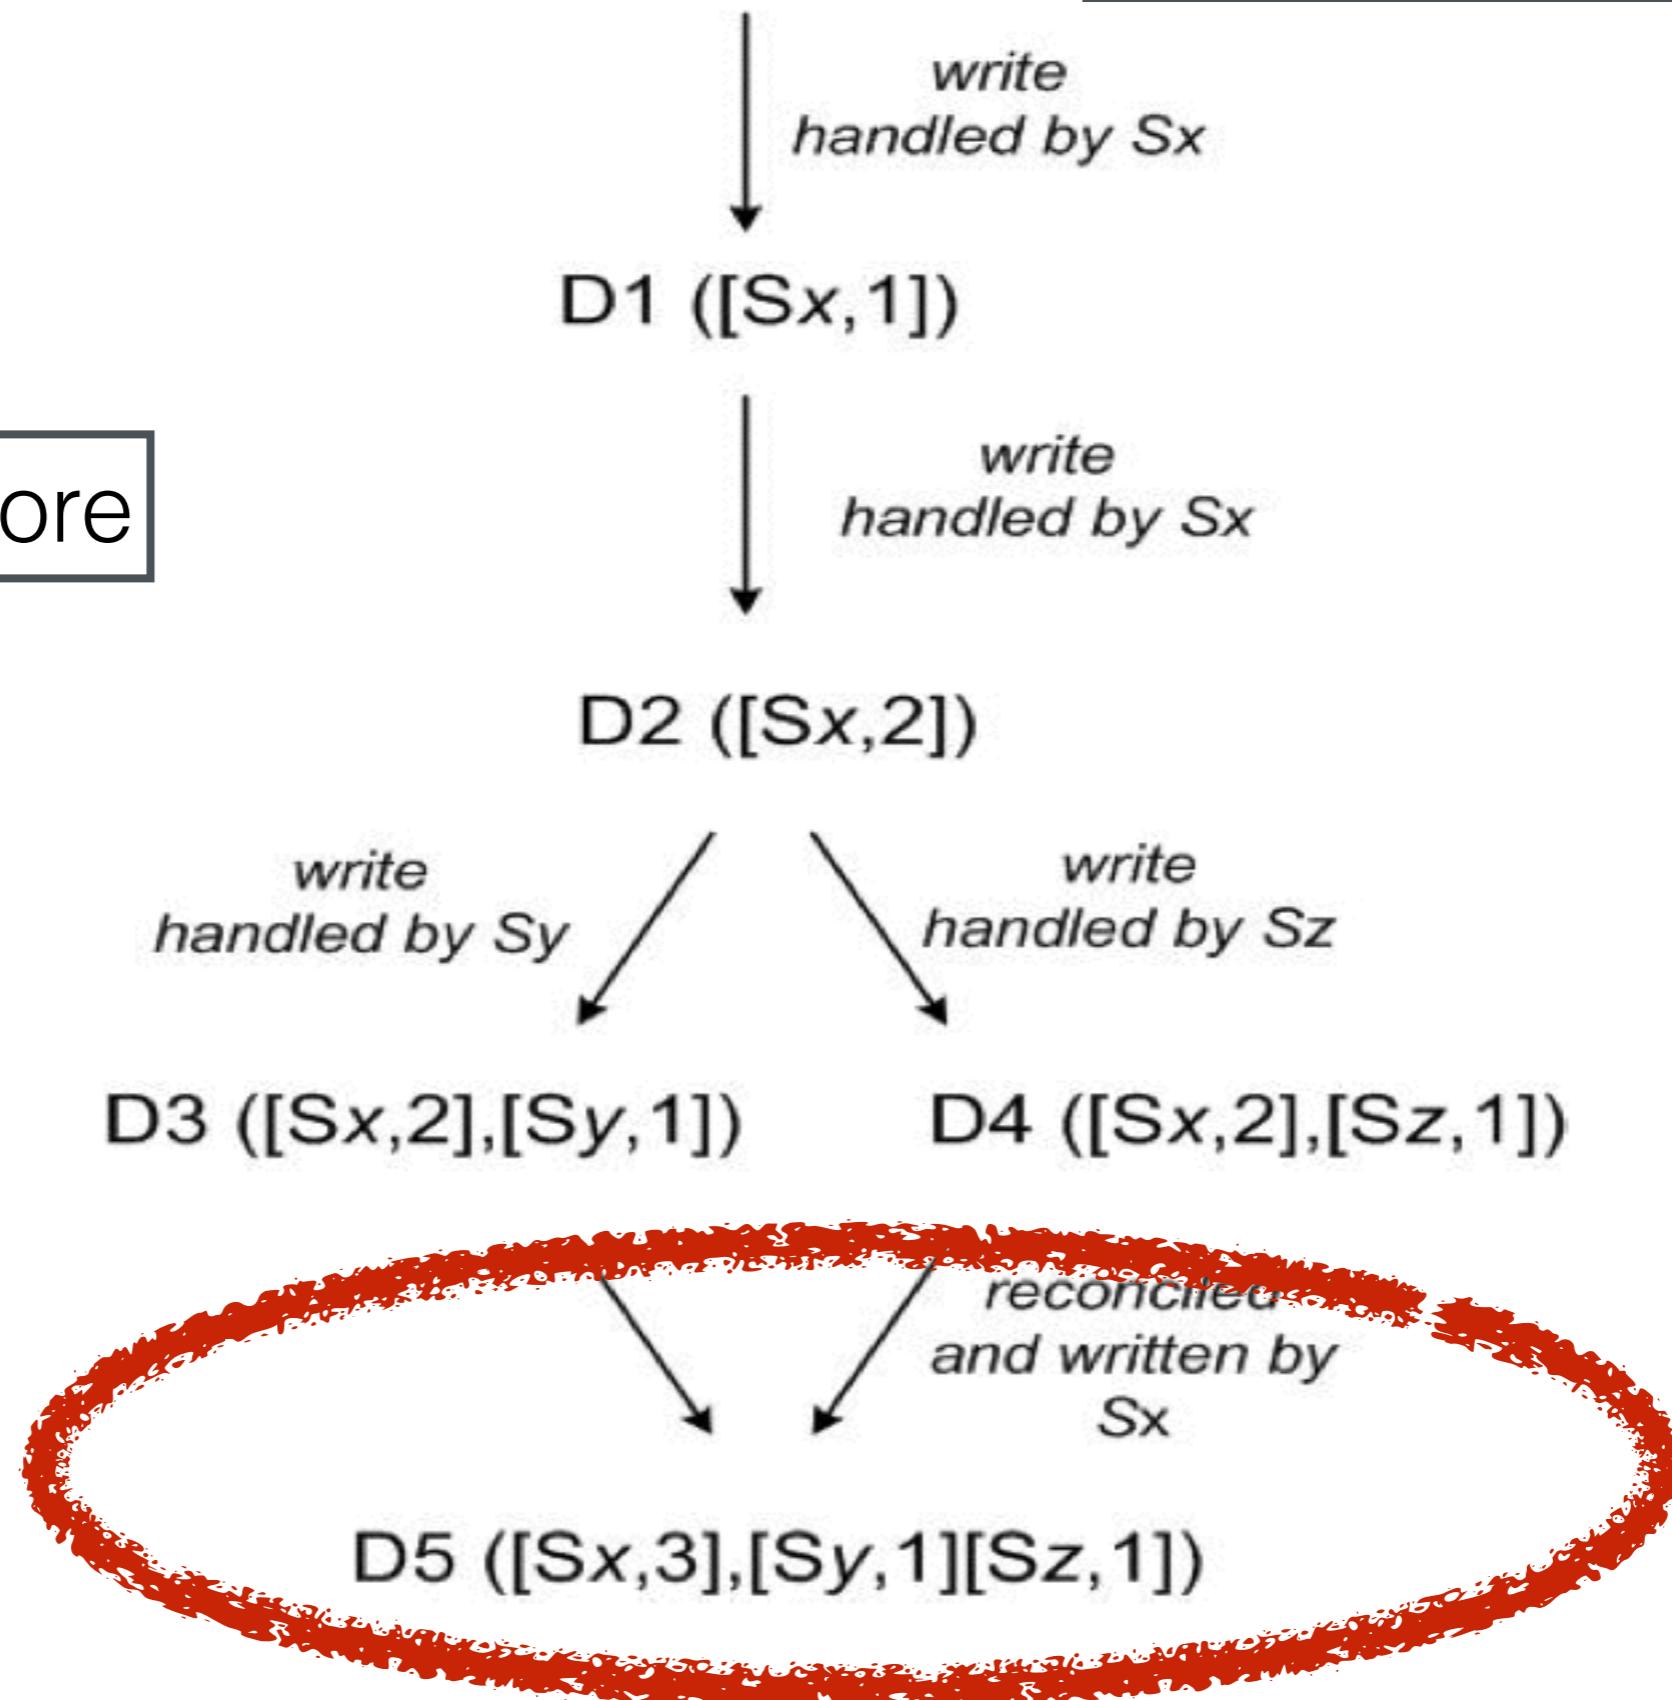
**Created by Amy Schwartz  
from the Noun Project**

# Logical Clocks

happens before

concurrent  
divergent

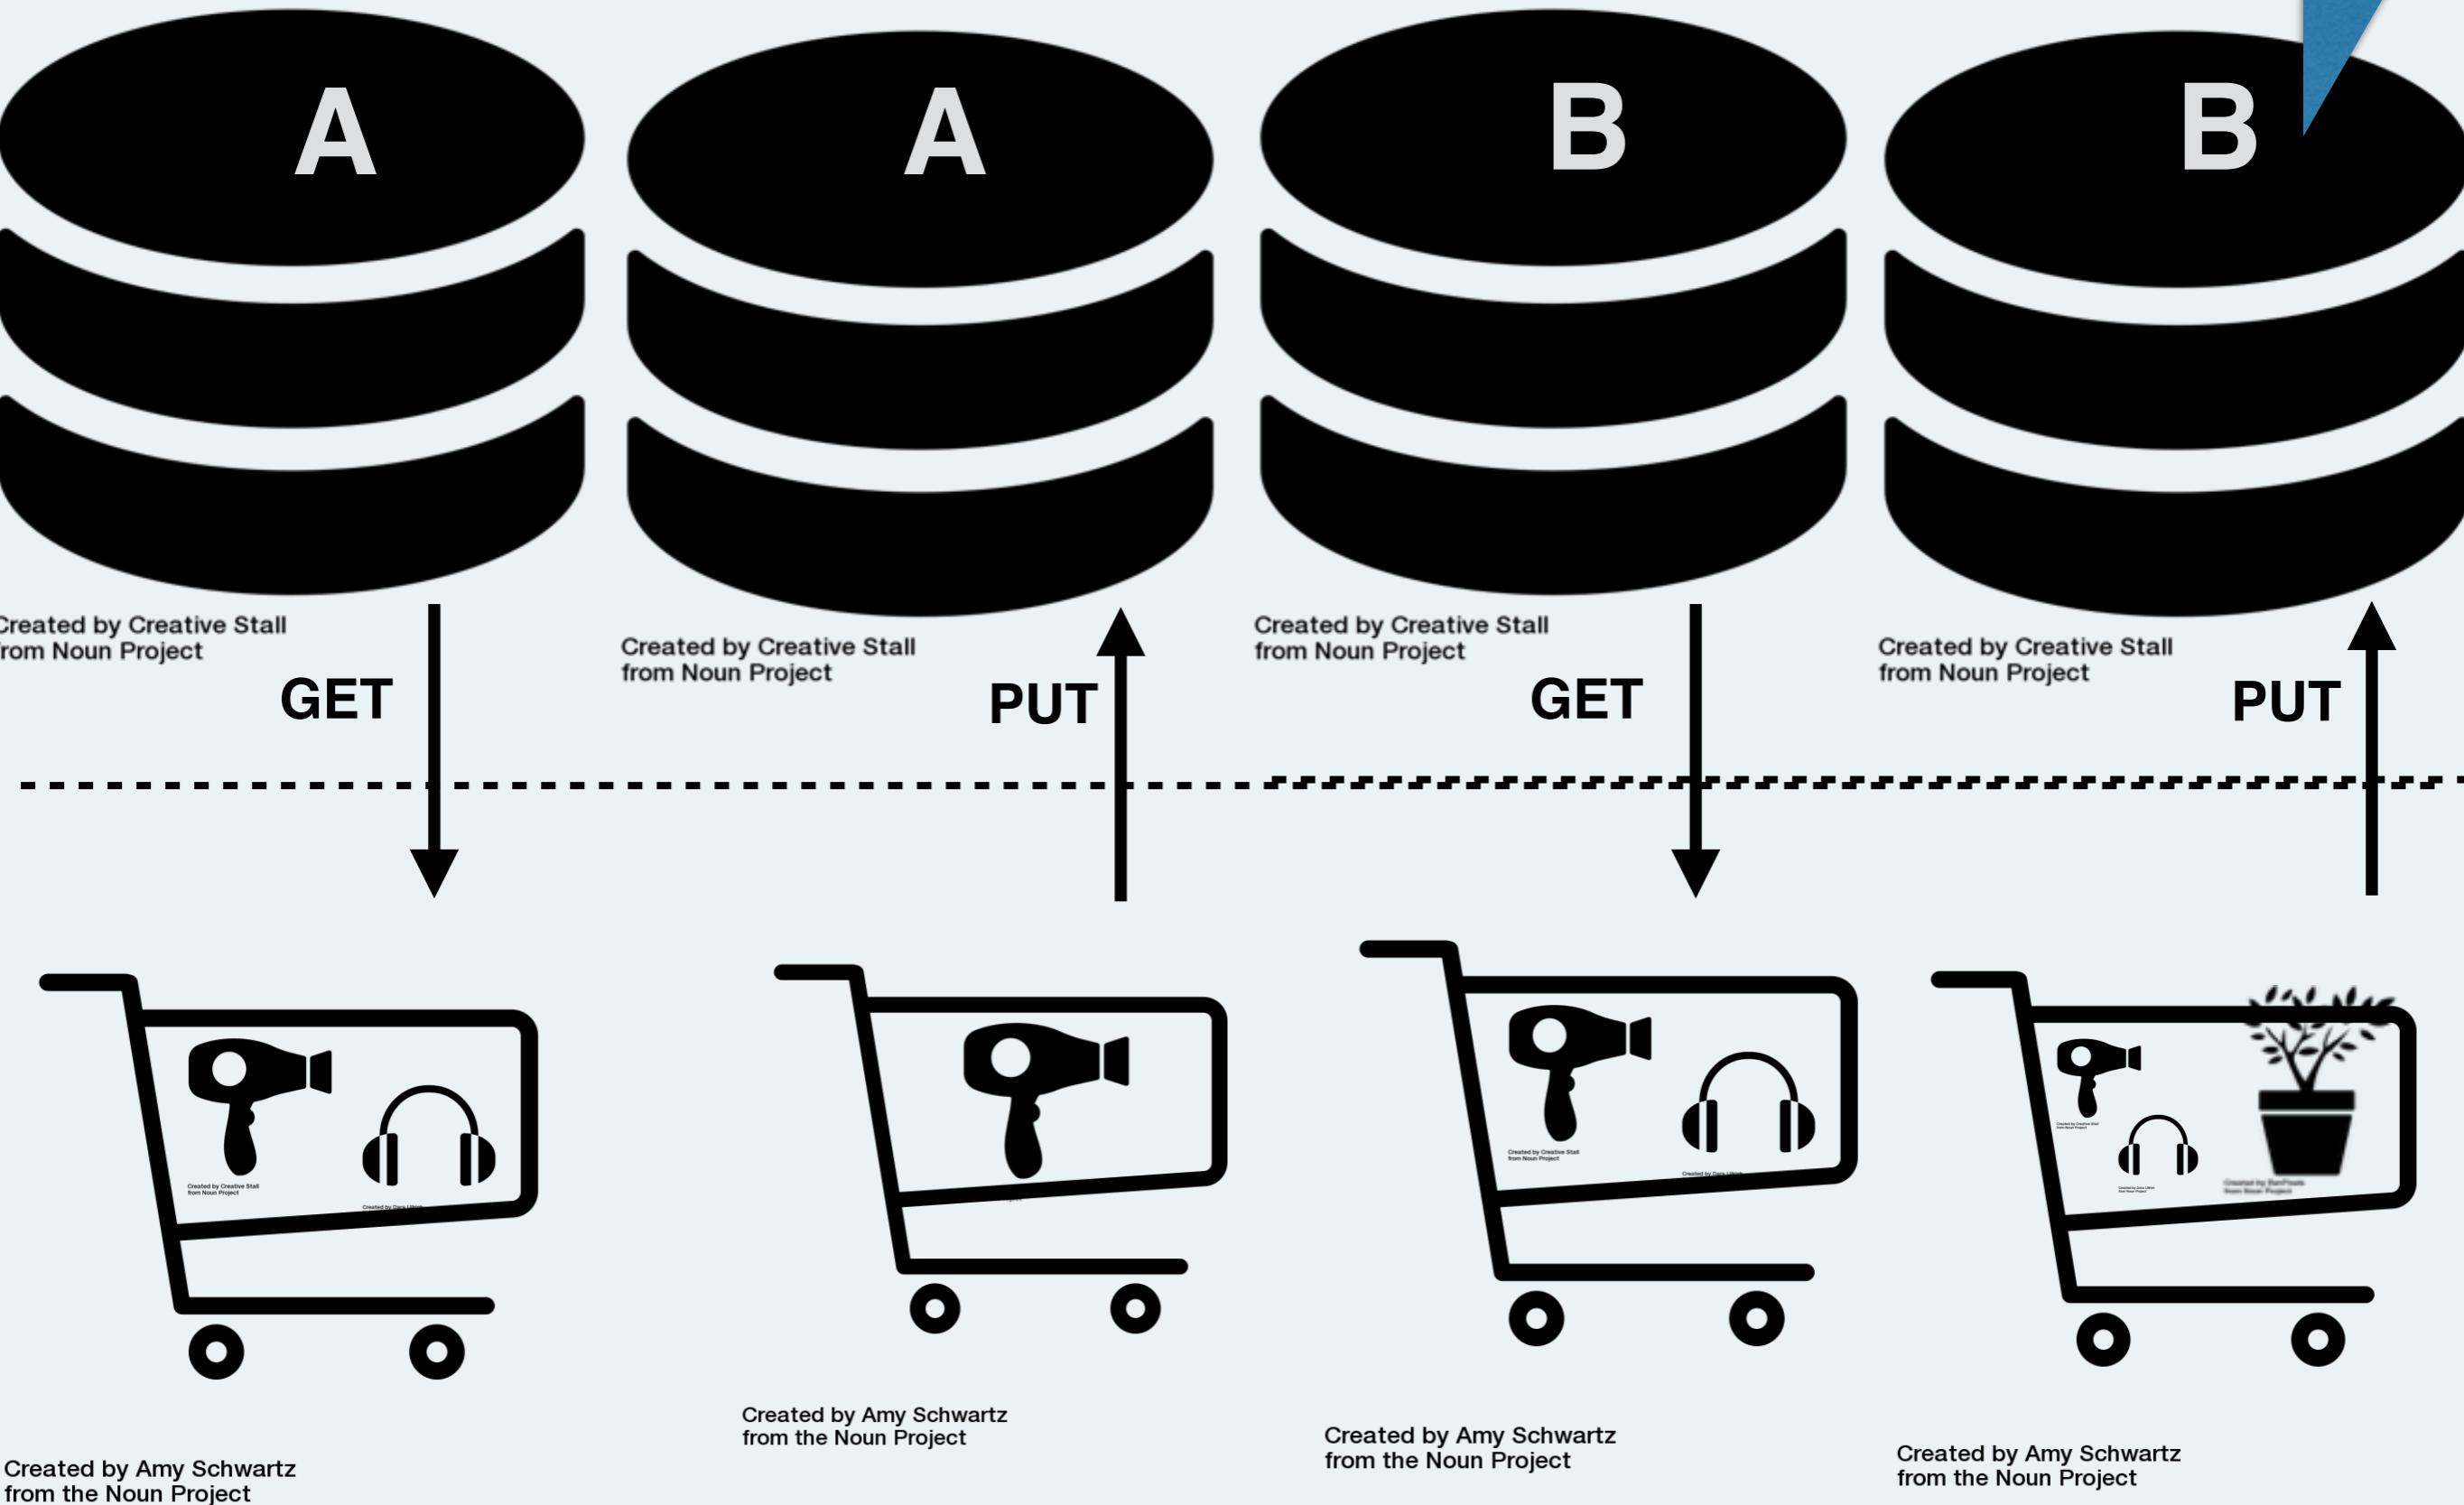
convergent

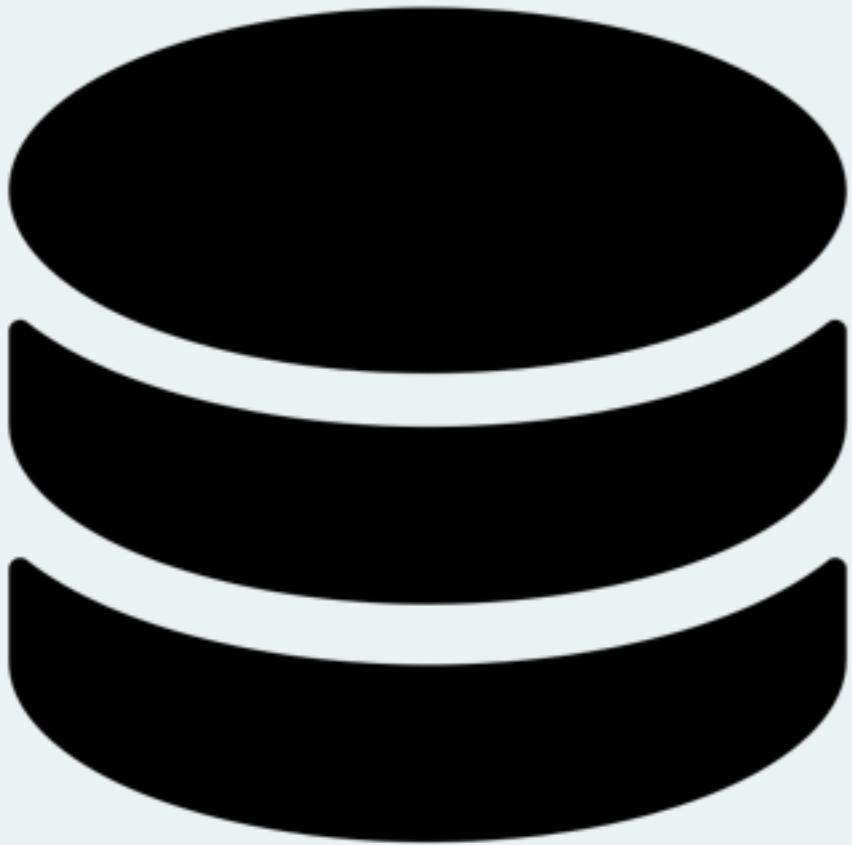




# Removes?

# TEMPORAL TIME





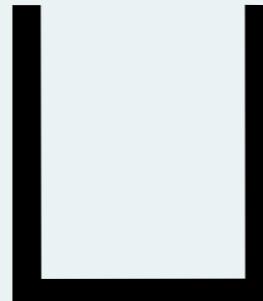
Created by Creative Stall  
from Noun Project



Created by Creative Stall  
from Noun Project



Created by Amy Schwartz  
from the Noun Project



Created by Amy Schwartz  
from the Noun Project



# Removes?

“merging” different versions of a customer’s shopping cart. Using this reconciliation mechanism, an “add to cart” operation is never lost. However, deleted items can resurface.

# Google F1

“Designing applications to cope with concurrency **anomalies** in their data is very **error-prone**, **time-consuming**, and ultimately *not worth the performance gains.*”

“...writing merge functions was likely to **confuse the hell** out of all our developers and ***slow down development.***”



<http://www.infoq.com/articles/key-lessons-learned-from-transition-to-nosql>

# CRDTs

DATA TYPES  
That CONVERGE

# CRDTs

Off the shelf  
MERGE functions

# CRDTs

CRDTs are Data Types  
(maps/sets/booleans/graphs/  
etc)

THAT CONVERGE



INSTITUT NATIONAL DE RECHERCHE EN INFORMATIQUE ET EN AUTOMATIQUE

*A comprehensive study of*  
***Convergent and Commutative Replicated Data Types***

Marc Shapiro, INRIA & LIP6, Paris, France

Nuno Preguiça, CITI, Universidade Nova de Lisboa, Portugal

Carlos Baquero, Universidade do Minho, Portugal

Marek Zawirski, INRIA & UPMC, Paris, France

# CRDT SETS

“...after some analysis we found that  
**much of our data could be modelled**  
**within sets** so by leveraging CRDT's our  
developers don't have to worry about  
writing bespoke merge functions for **95%**  
**of carefully selected use cases...**”



<http://www.infoq.com/articles/key-lessons-learned-from-transition-to-nosql>

# Evolution of a CRDT Set

# Evolution of a Set

## G-SET

# Replica A

Shelly

Bob

Pete

Anna

# Replica B

Alex

Shelly

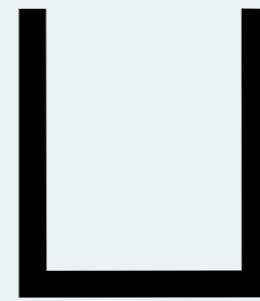
# Replica A

Shelly

Bob

Pete

Anna



# Replica B

Alex

Shelly

**Shelly**

**Bob**

**Pete**

**Anna**

**Alex**

**G-SET**



# Removes?

# Evolution of a Set

G-SET

2P-SET

**Adds**

**Shelly**

**Bob**

**Pete**

**Anna**

**Removes**

**Shelly**

**Bob**

**Pete**

**2P-SET**

# Adds

Shelly

Bob

Pete

Anna

# Removes

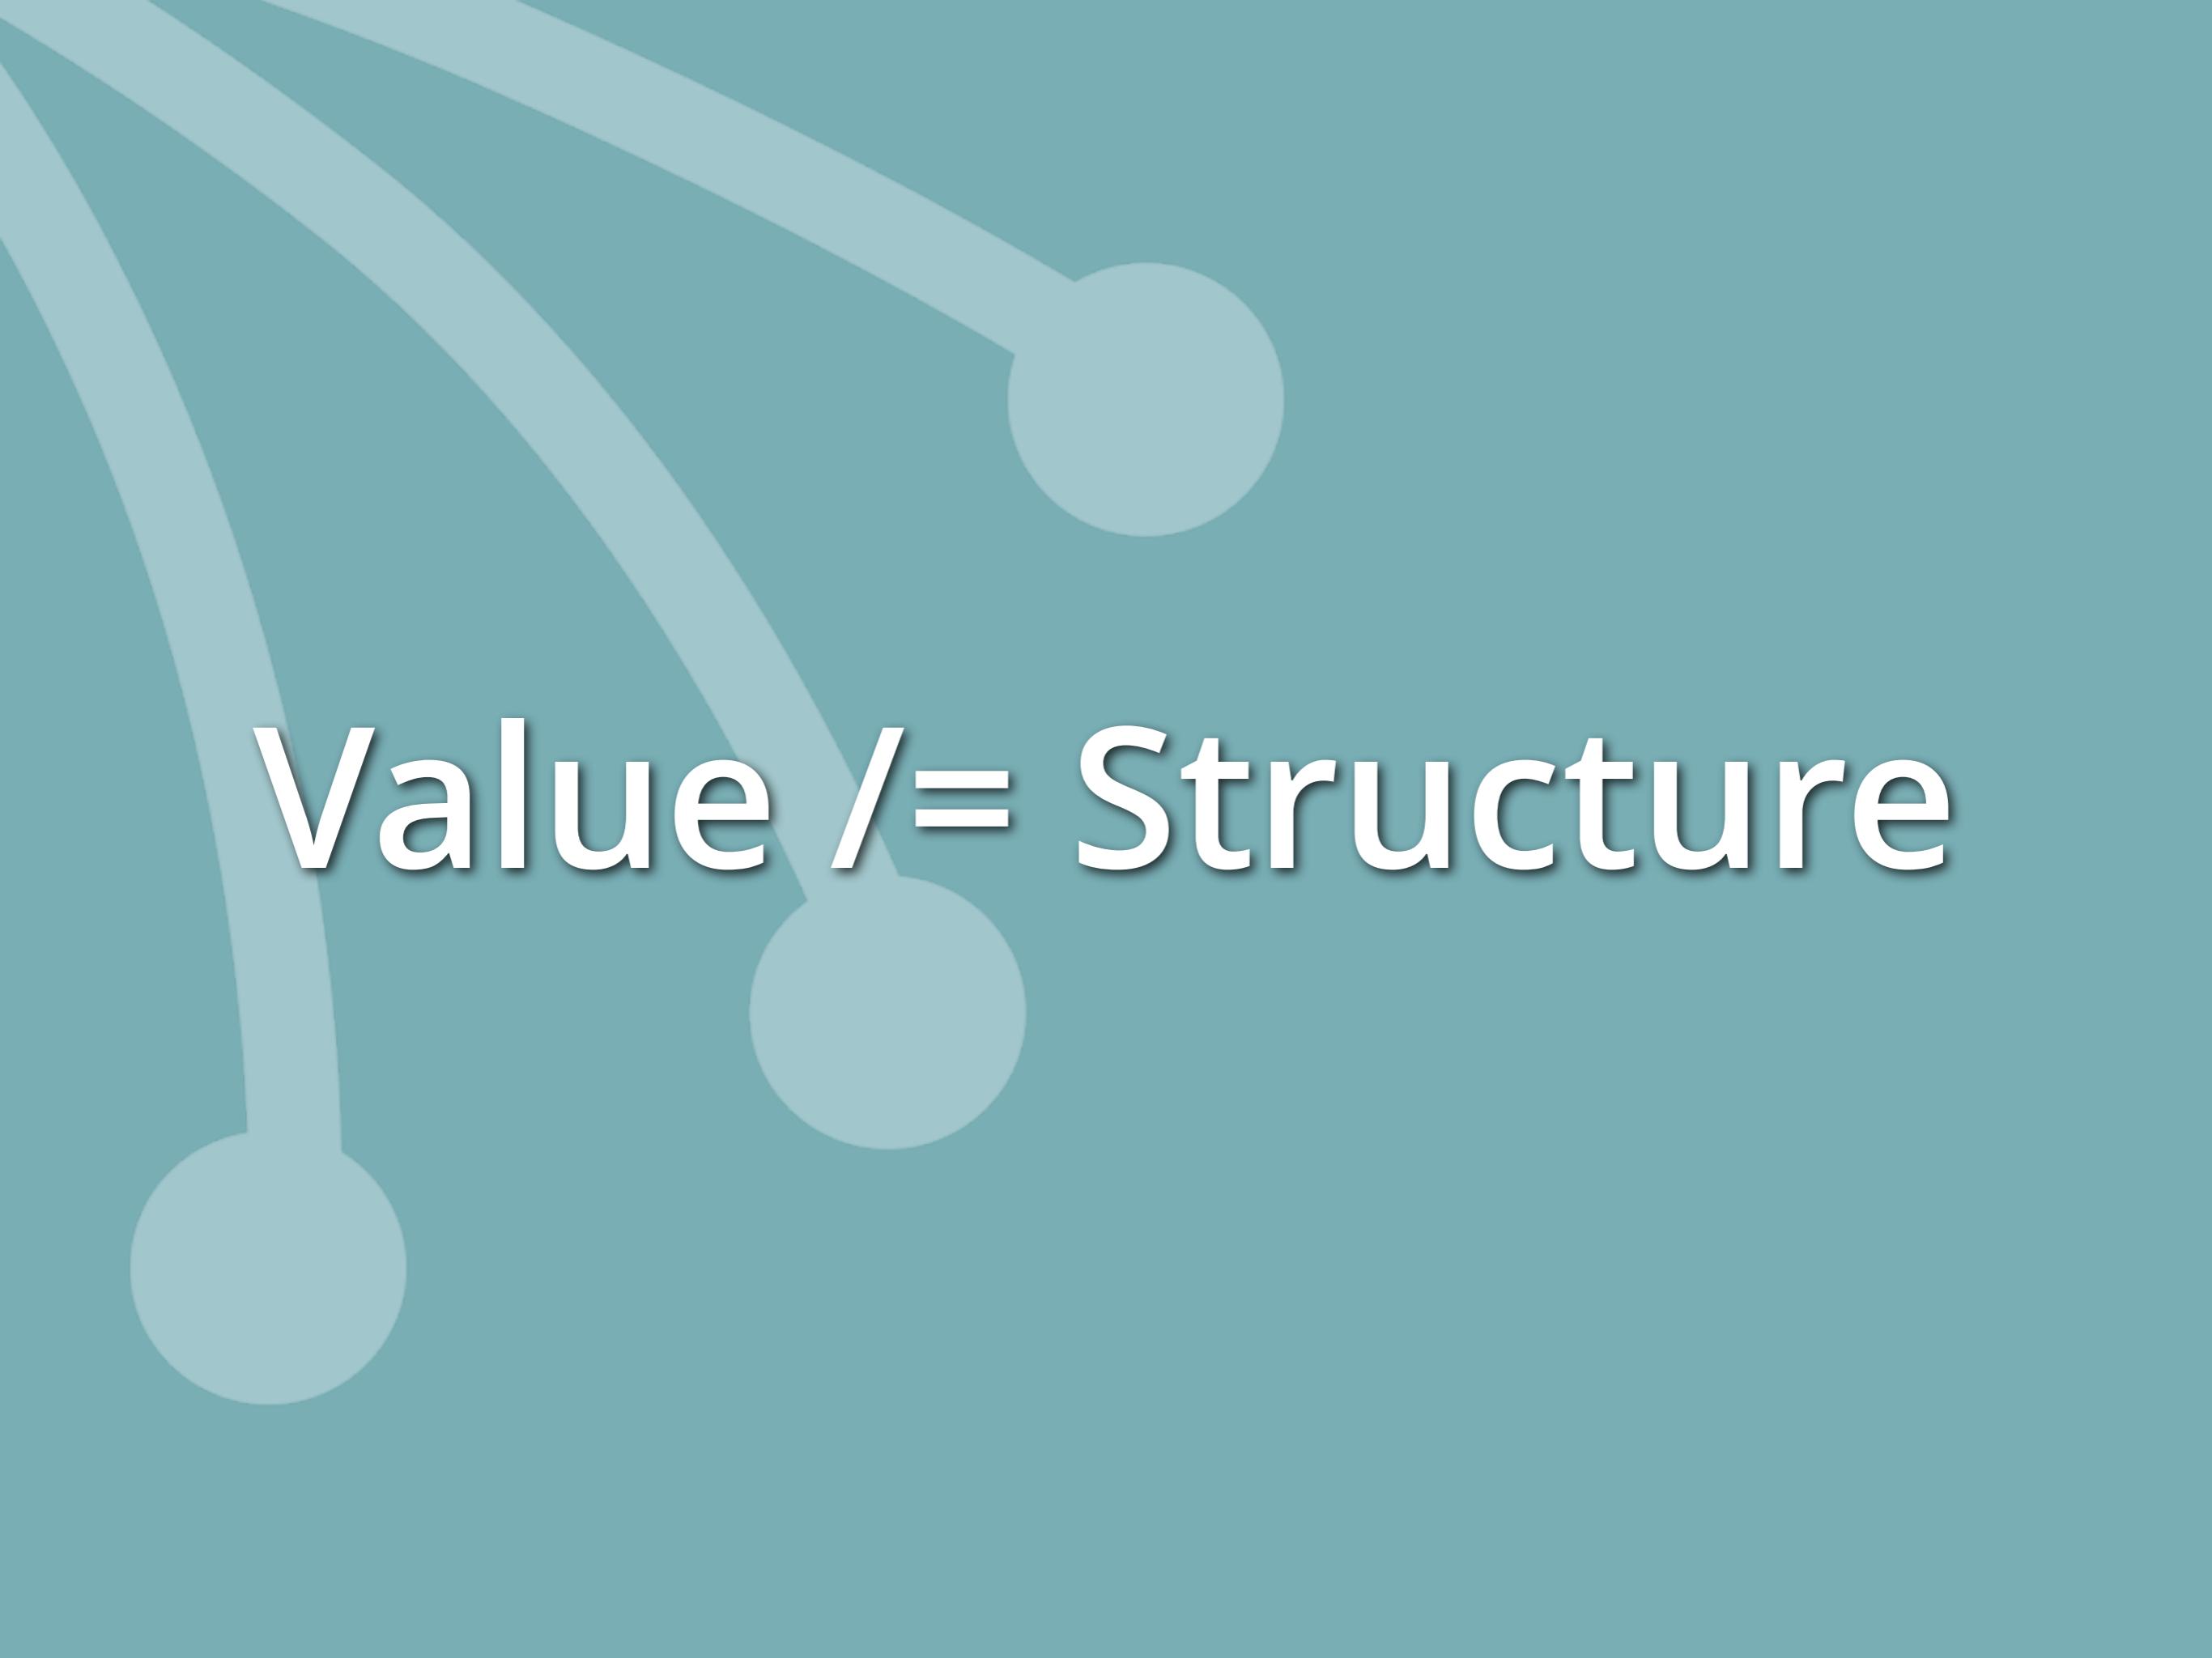
Shelly

Bob

Pete

==

Anna

The background features a teal gradient with three large, semi-transparent light-blue circles of varying sizes. Two circles are positioned vertically on the left side, and one is located in the upper right quadrant.

Value /= Structure

# Adds

Shelly

Bob

Pete

Anna

# Removes

Shelly

Bob

Pete

==

Anna

I changed  
my mind!

**Adds**

**Shelly**

**Bob**

**Pete**

**Anna**

**Shelly**

**Removes**

**Shelly**

**Bob**

**Pete**

**Anna**

**=**

# Evolution of a Set

# U-SET

# Replica A

1	Shelly
2	Bob
3	Pete
4	Anna

U

# Replica B

5	Alex
6	Shelly

1,6	Shelly
2	Bob
3	Pete
4	Anna
5	Alex

**U-SET**

# Evolution of a Set

**U-SET**

**OR-SET**

# Evolution of a Set

OR SET

AW-SET

## Adds

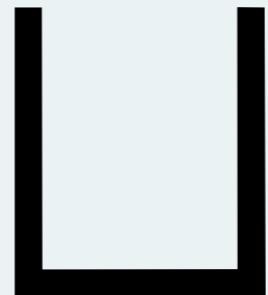
## Removes

1,5	Shelly	1	Shelly
2	Bob	2	Bob
3	Pete	3	Pete
4	Anna		

# AW-SET

# Replica A

Adds		Removes	
1	Shelly	1	Shelly
2	Bob	2	Bob
3	Pete	3	Pete



# Replica B

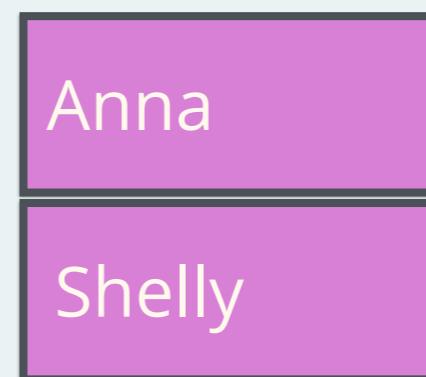
Adds	
4	Anna
5	Shelly

## Adds

1,5	Shelly
2	Bob
3	Pete
4	Anna

## Removes

1	Shelly
2	Bob
3	Pete





Observed  
Remove

Semantics  
Add  
Wins

# Evolution of a Set

# AW-SET

## Adds

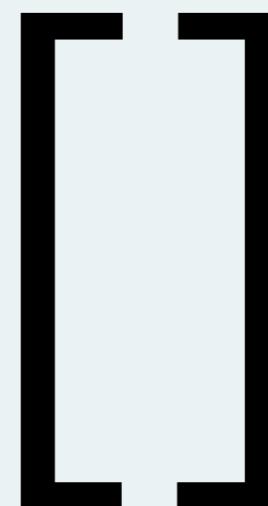
1,5	Shelly
2	Bob
3	Pete
4	Anna

## Removes

1,5	Shelly
2	Bob
3	Pete
4	Anna



Shelly





INSTITUT NATIONAL DE RECHERCHE EN INFORMATIQUE ET EN AUTOMATIQUE

## *An Optimized Conflict-free Replicated Set*

Annette Bieniusa, INRIA & UPMC, Paris, France

Marek Zawirski, INRIA & UPMC, Paris, France

Nuno Preguiça, CITI, Universidade Nova de Lisboa, Portugal

Marc Shapiro, INRIA & LIP6, Paris, France

Carlos Baquero, HASLab, INESC TEC & Universidade do Minho, Portugal

Valter Balegas, CITI, Universidade Nova de Lisboa, Portugal

Sérgio Duarte CITI, Universidade Nova de Lisboa, Portugal

# Evolution of a Set OR-SWOT

# Evolution of a Set

~~OR SWOT~~

# Optimised

# AW-SET

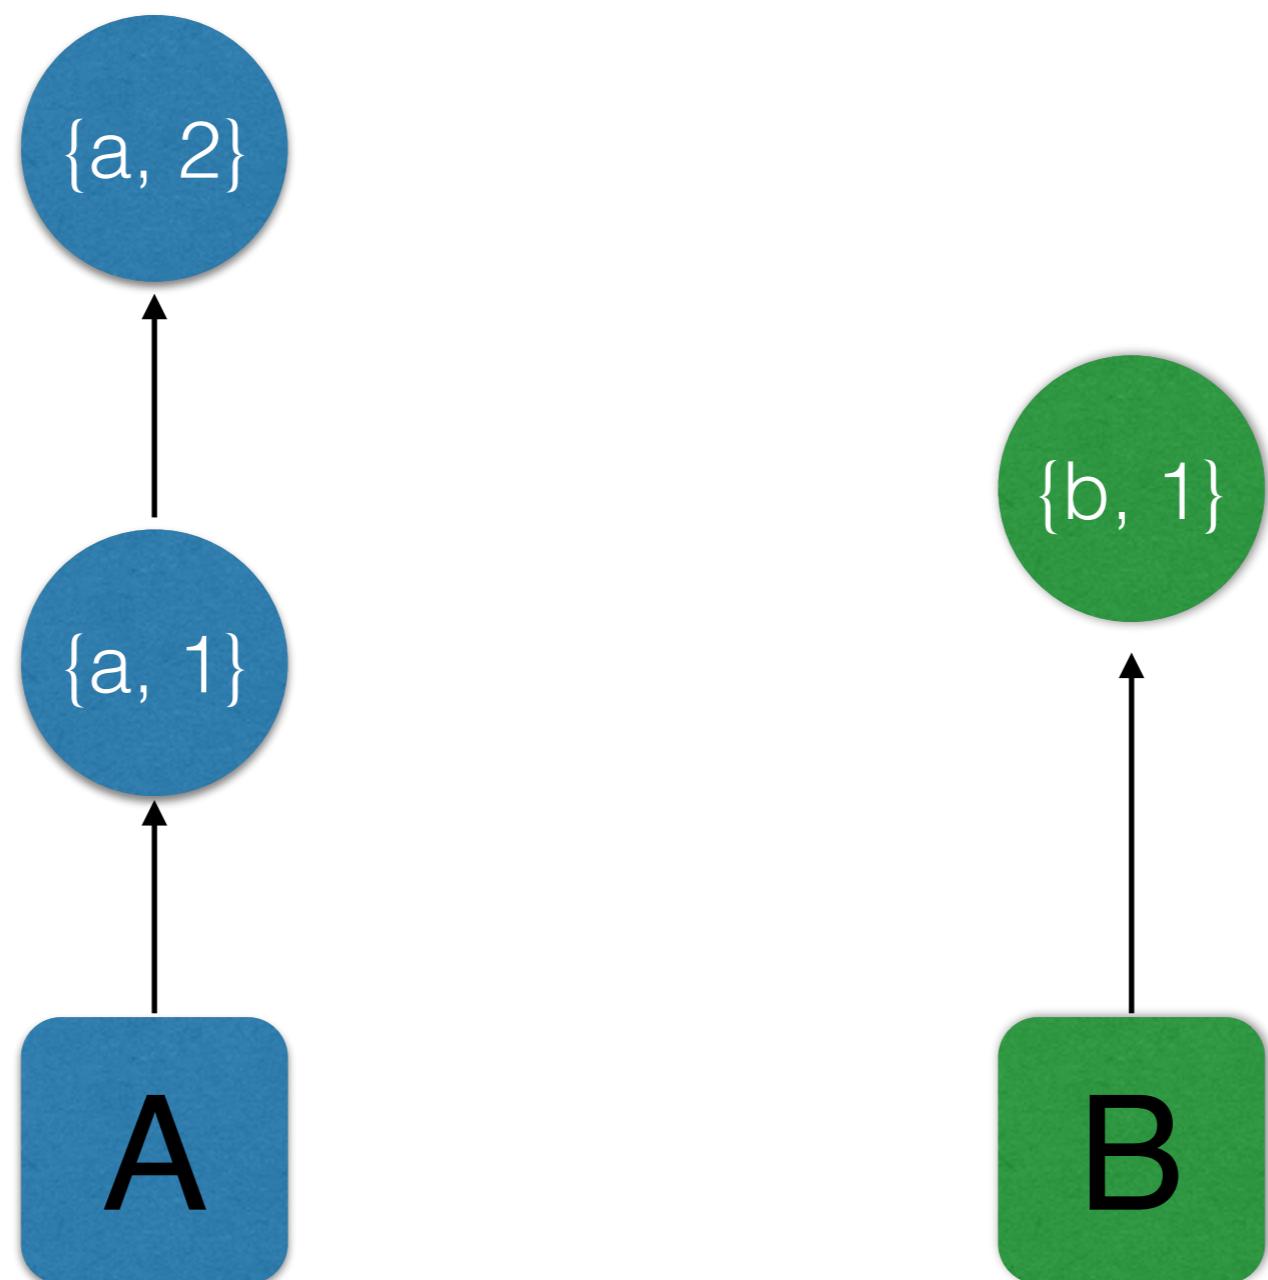
# Version Vectors

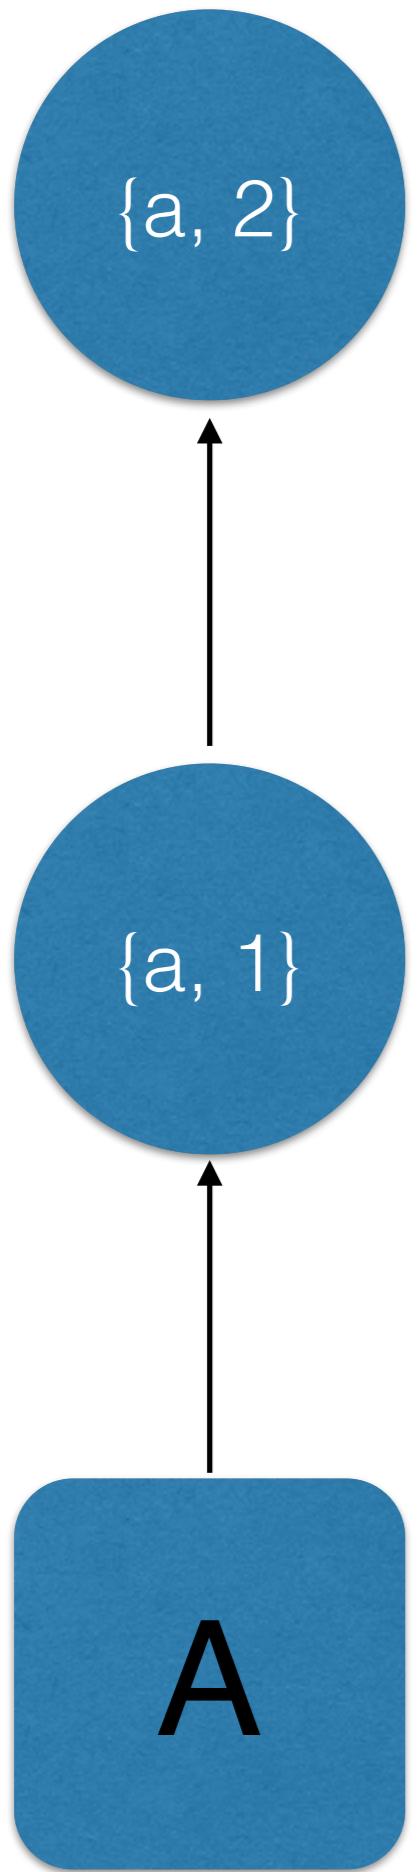
A

B

# Version Vectors

[  
    {a, 2}, {b, 1}  
]





EVENTS/TAGS

# Replica A

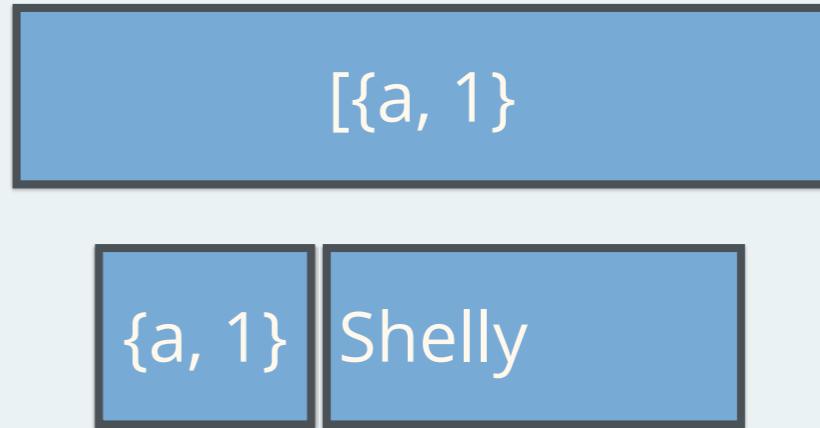
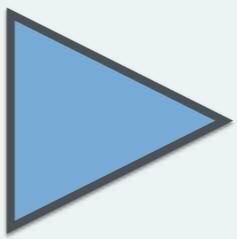
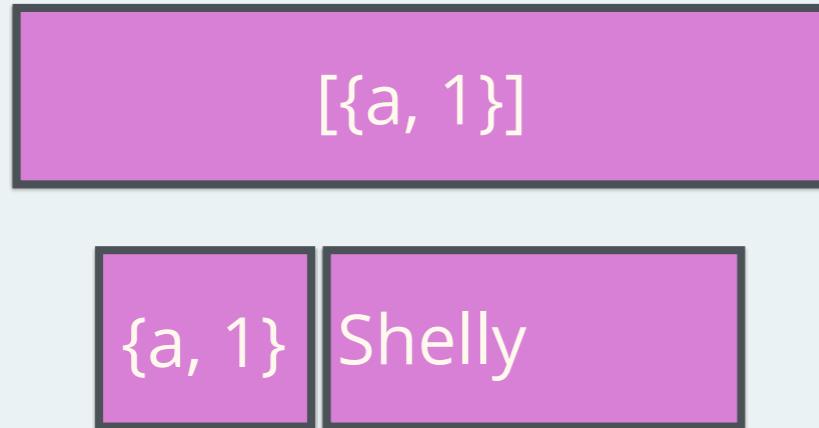
[{a, 1}]

{a, 1}

Shelly

# Replica A

# Replica B



# Replica A

[{a, 1}]

{a, 1} Shelly

# Replica B

[{a, 1}, {b, 3}]

{a, 1}	Shelly
{b, 1}	Bob
{b, 2}	Phil
{b, 3}	Pete

# Replica A

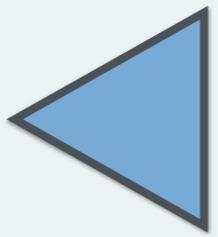
# Replica B

$\{\{a, 1\}, \{b, 3\}\}$

{a, 1}	Shelly
{b, 1}	Bob
{b, 2}	Phil
{b, 3}	Pete

$\{\{a, 1\}, \{b, 3\}\}$

{a, 1}	Shelly
{b, 1}	Bob
{b, 2}	Phil
{b, 3}	Pete



# Replica A

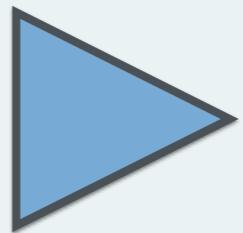
$\{\{a, 2\}, \{b, 3\}\}$	
{a, 1}	Shelly
{b, 1}	Bob
{b, 2}	Phil
{b, 3}	Pete
{a, 2}	Anna

# Replica B

$\{\{a, 1\}, \{b, 4\}\}$	
{a, 1}	Shelly
{b, 1}	Bob
{b, 2}	Phil
{b, 3}	Pete
{b, 4}	John

# Replica A

$\{\{a, 2\}, \{b, 3\}\}$	
{a, 1}	Shelly
{b, 1}	Bob
{b, 3}	Pete
{a, 2}	Anna



# Replica B

$\{\{a, 1\}, \{b, 4\}\}$	
{a, 1}	Shelly
{b, 1}	Bob
{b, 2}	Phil
{b, 3}	Pete
{b, 4}	John

# Replica A

$\{\{a, 2\}, \{b, 3\}\}$

{a, 1}	Shelly
{b, 1}	Bob

{b, 3}	Pete
{a, 2}	Anna

# Replica B

$\{\{a, 1\}, \{b, 4\}\}$

{a, 1}	Shelly
{b, 1}	Bob
{b, 2}	Phil
{b, 3}	Pete
{b, 4}	John

# MERGE

# Replica A

$\{\{a, 2\}, \{b, 3\}\}$

{a, 1}	Shelly
{b, 1}	Bob

{b, 3}	Pete
{a, 2}	Anna

# Replica B

$\{\{a, 1\}, \{b, 4\}\}$

{a, 1}	Shelly
{b, 1}	Bob

{b, 2}	Phil
{b, 3}	Pete
{b, 4}	John

# MERGE

{a, 1}	Shelly
{b, 1}	Bob

**{b, 2}**

**∈**

**[{a, 2}, {b, 3}]**

Phil

$\{b, 2\}$

Phil

$\in$

$[\{a, 2\}, \{b, 3\}]$

# Replica A

$\{\{a, 2\}, \{b, 3\}\}$

{a, 1}	Shelly
{b, 1}	Bob

{b, 1}	Bob
--------	-----

{b, 3}	Pete
--------	------

{a, 2}	Anna
--------	------

# Replica B

$\{\{a, 1\}, \{b, 4\}\}$

{a, 1}	Shelly
{b, 1}	Bob

{b, 1}	Bob
--------	-----

{b, 2}	Phil
--------	------

{b, 3}	Pete
--------	------

{b, 4}	John
--------	------

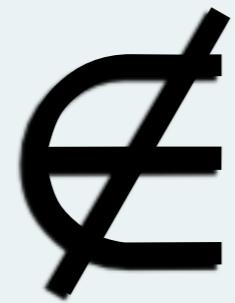
# MERGE

$\{\{a, 1\}, \{b, 3\}\}$

$\{\{b, 1\}, \{b, 4\}\}$

$\{\{b, 3\}, \{b, 4\}\}$

**{a, 2}**



**[{a, 1}, {b, 4}]**

Anna

# Replica A

$\{\{a, 2\}, \{b, 3\}\}$

{a, 1}	Shelly
{b, 1}	Bob

{b, 1}	Bob
--------	-----

{b, 3}	Pete
--------	------

{a, 2}	Anna
--------	------

# Replica B

$\{\{a, 1\}, \{b, 4\}\}$

{a, 1}	Shelly
{b, 1}	Bob

{b, 1}	Bob
--------	-----

{b, 2}	Phil
--------	------

{b, 3}	Pete
--------	------

{b, 4}	John
--------	------

# MERGE

$\{\{a, 1\}, \{b, 3\}\}$

$\{\{b, 1\}, \{b, 4\}\}$

$\{\{b, 3\}, \{a, 2\}\}$

$\{\{a, 2\}, \{b, 4\}\}$

$\{\{b, 4\}, \{b, 4\}\}$

$[\{a, 2\}, \{b, 4\}]$

{a, 1}	Shelly
{b, 1}	Bob
{b, 3}	Pete
{a, 2}	Anna
{b, 4}	John

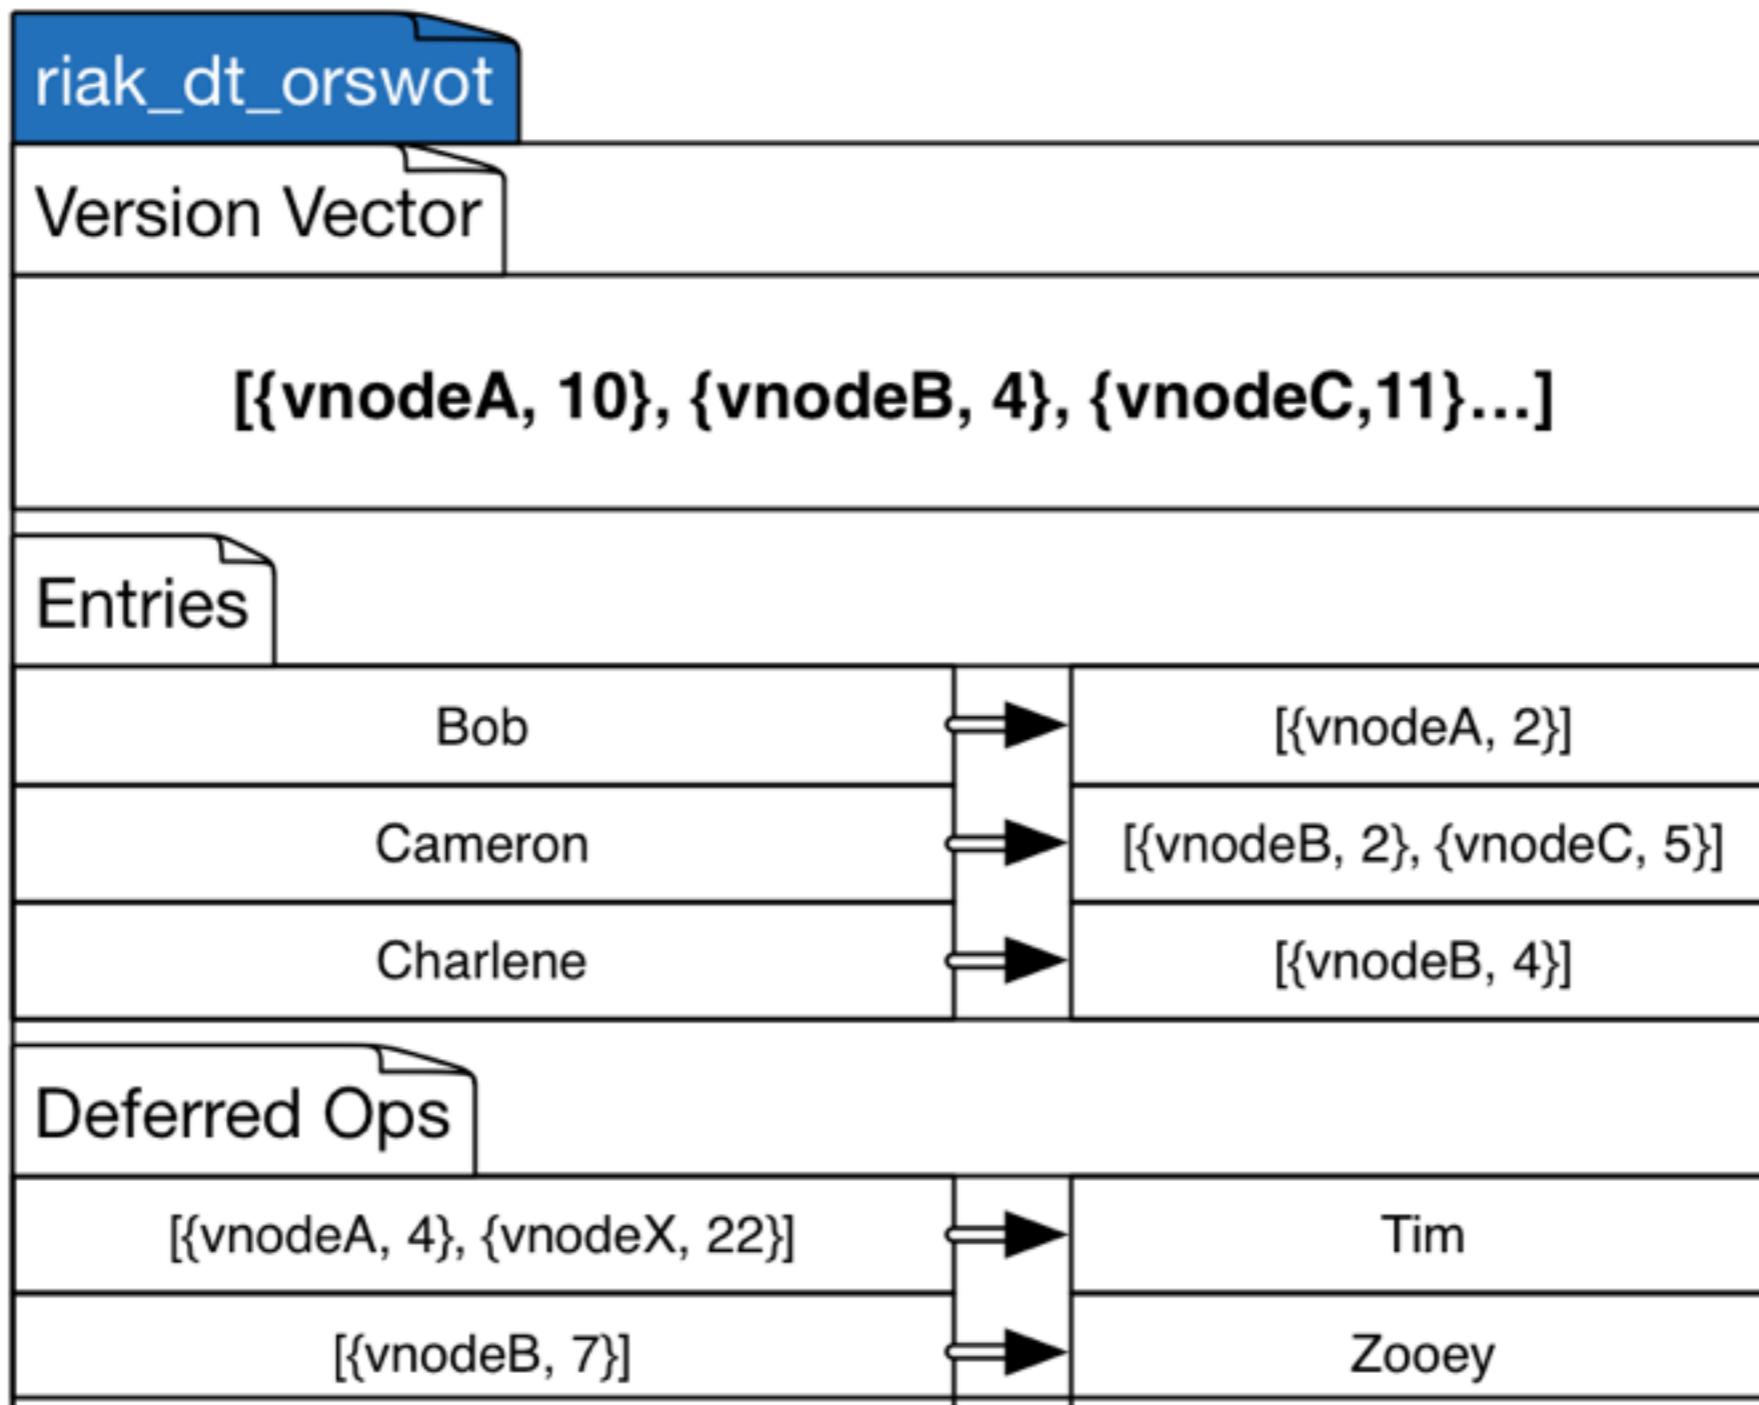


Shelly  
Bob  
  
Pete  
Anna  
John

# CRDT Sets

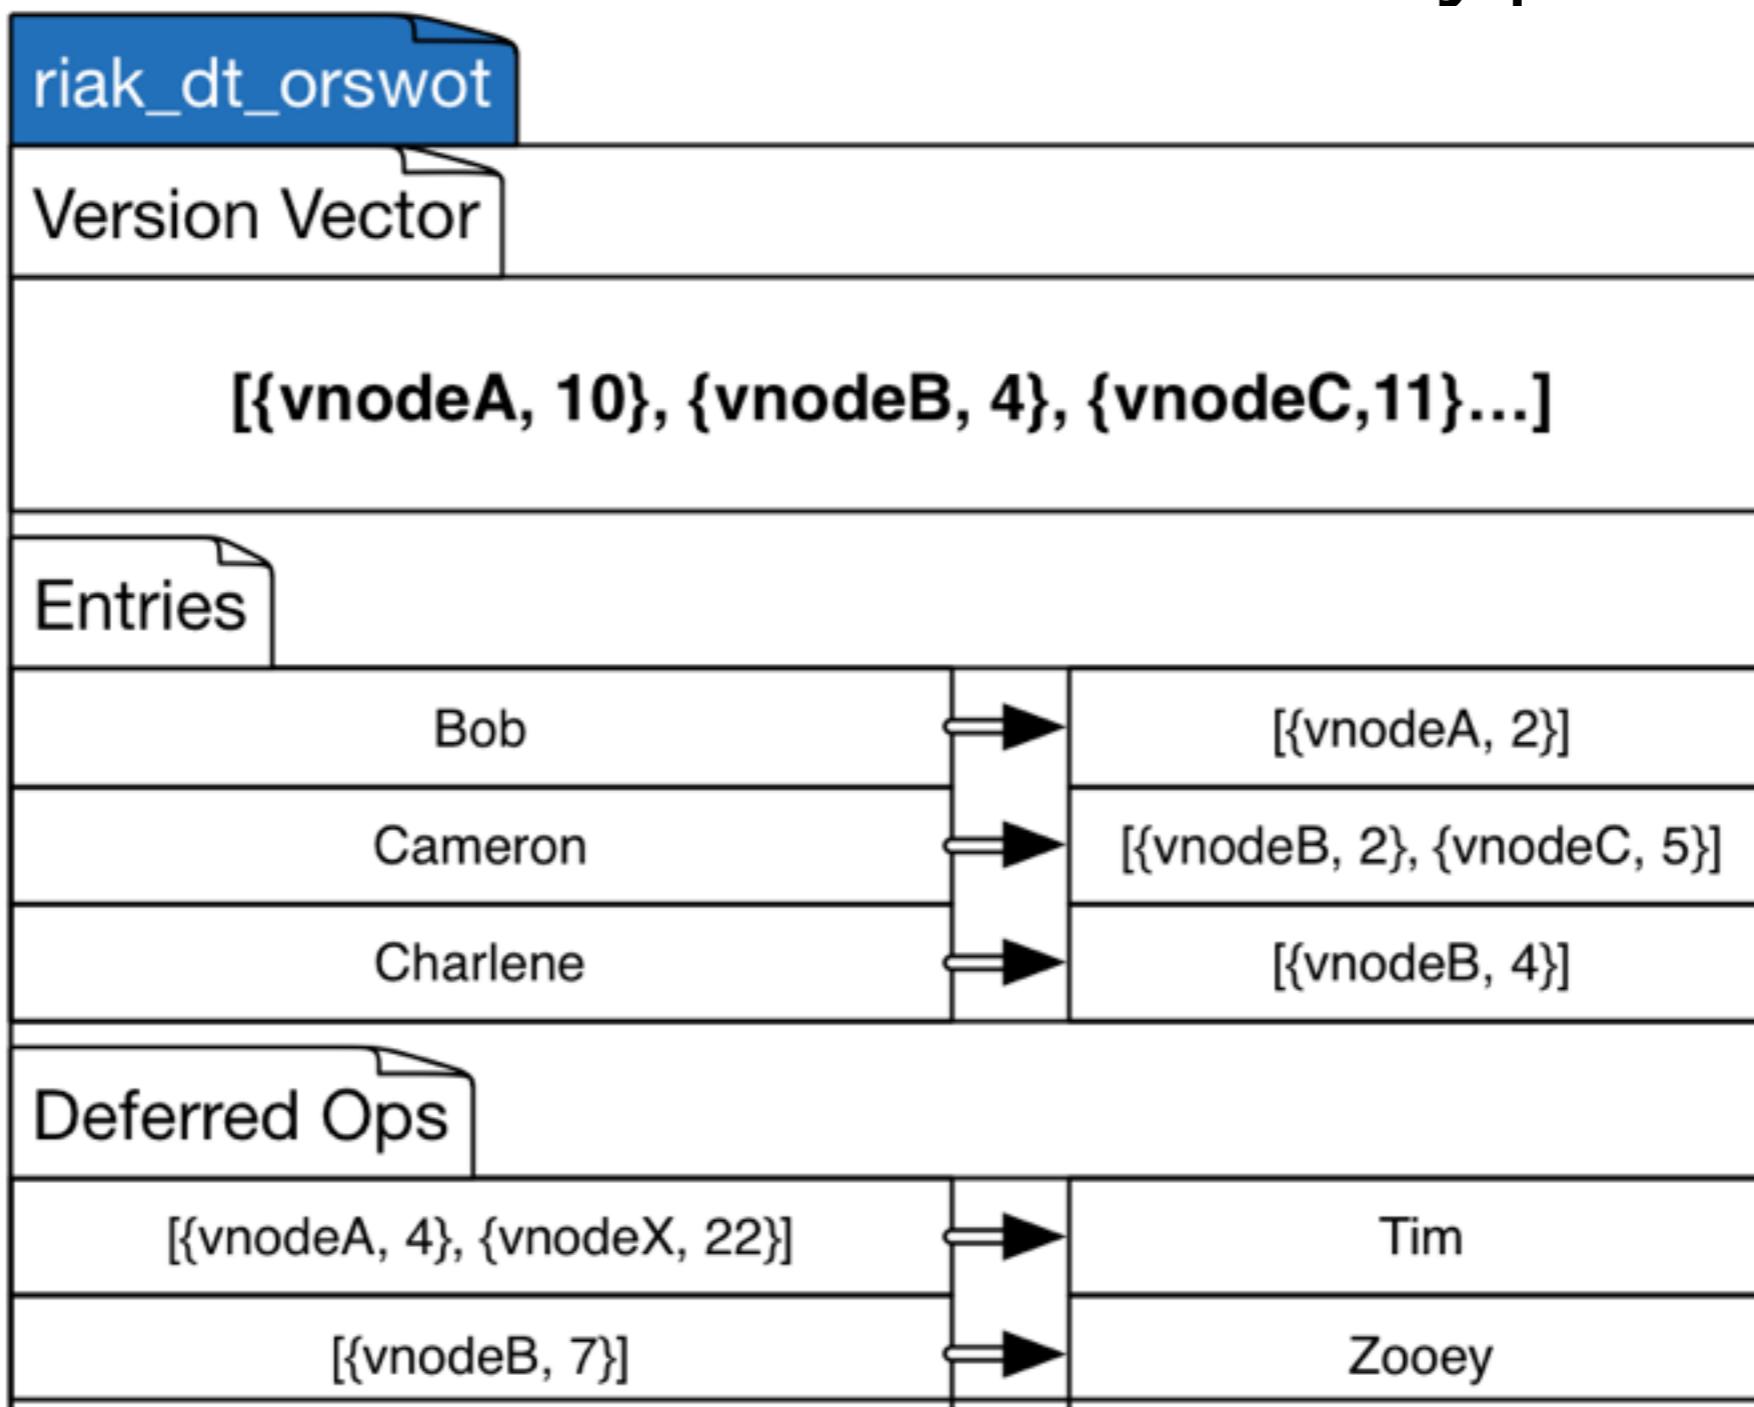
a semantic of “Add-Wins”  
via  
“Observed Remove”

# SETS in RIAK 2.0+

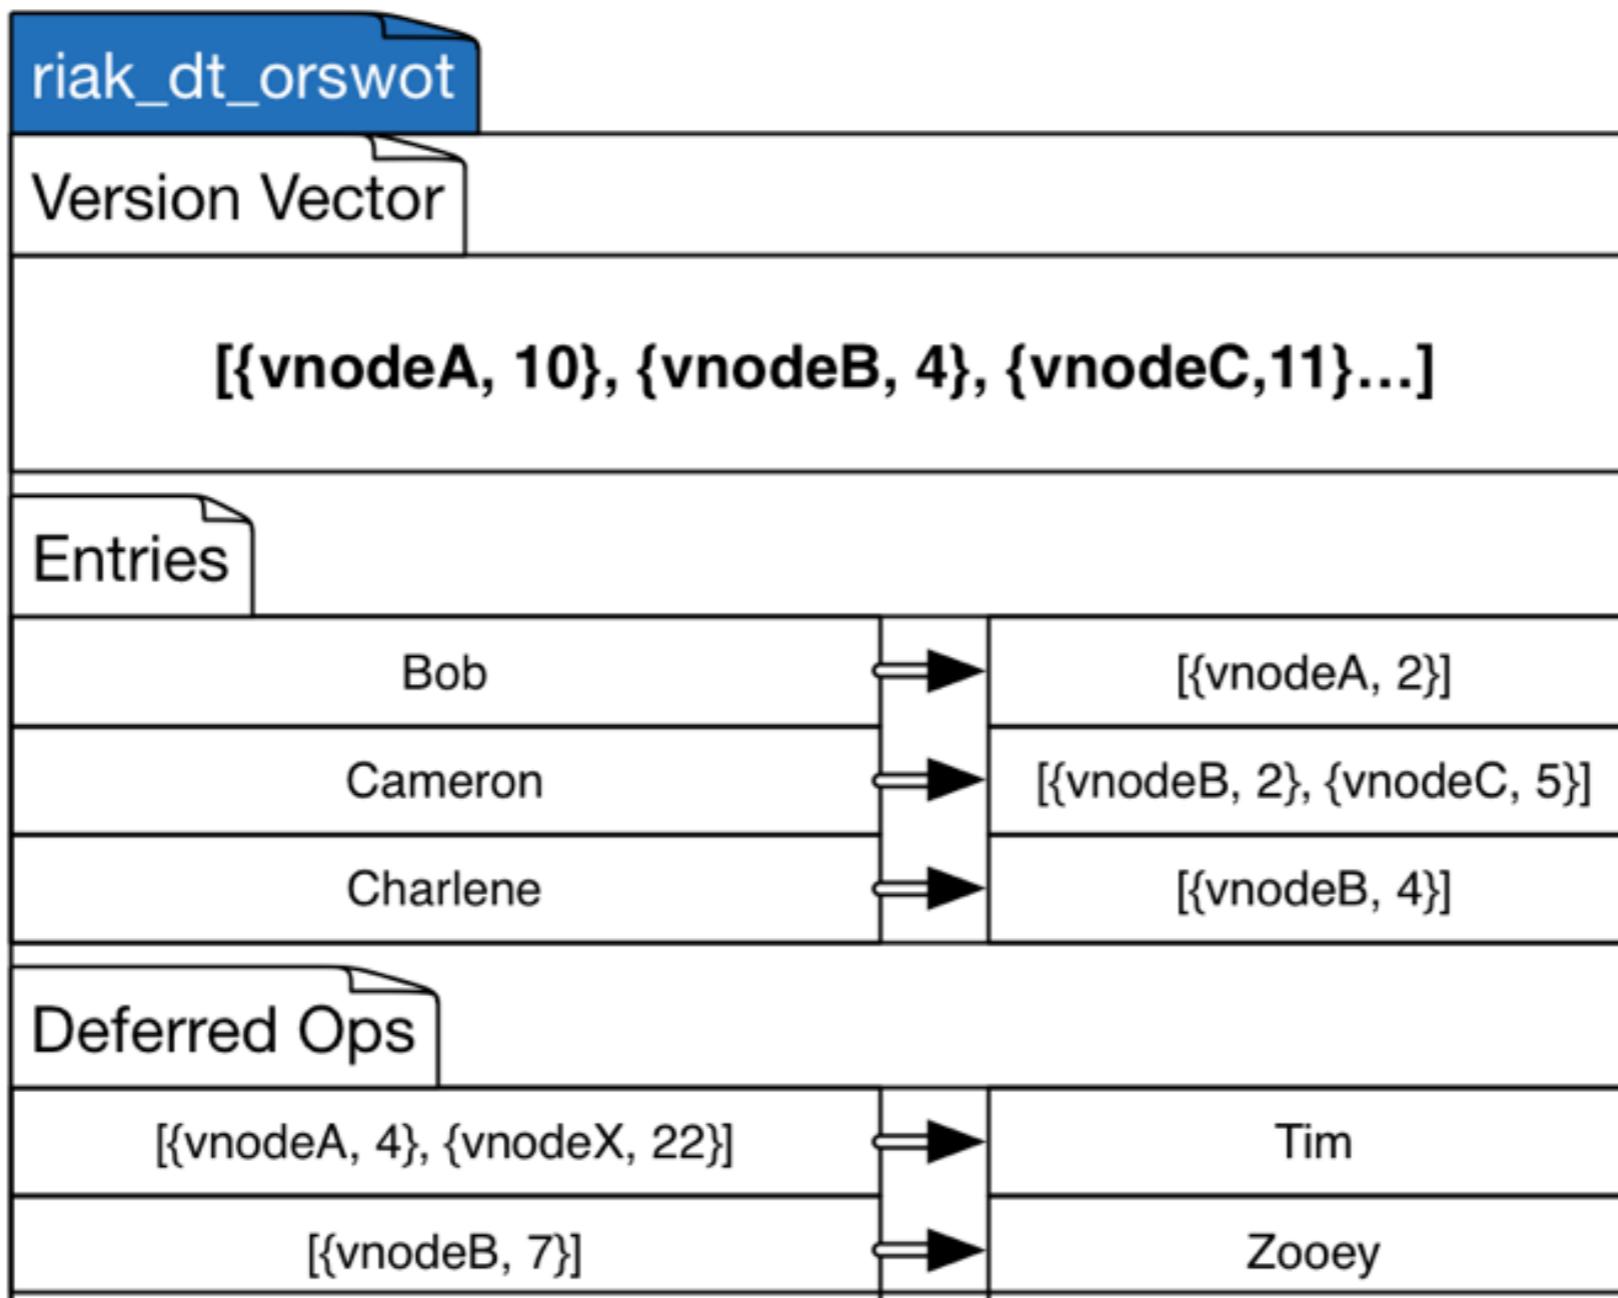


# Riak 2.0

## riak\_dt -> Riak Data Types



# Sets in Riak

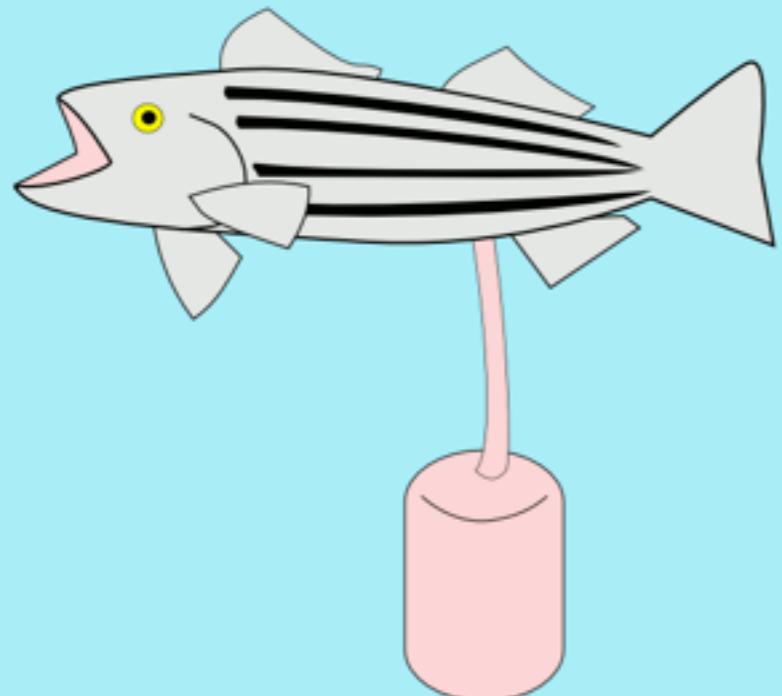


An optimized conflict-free replicated set

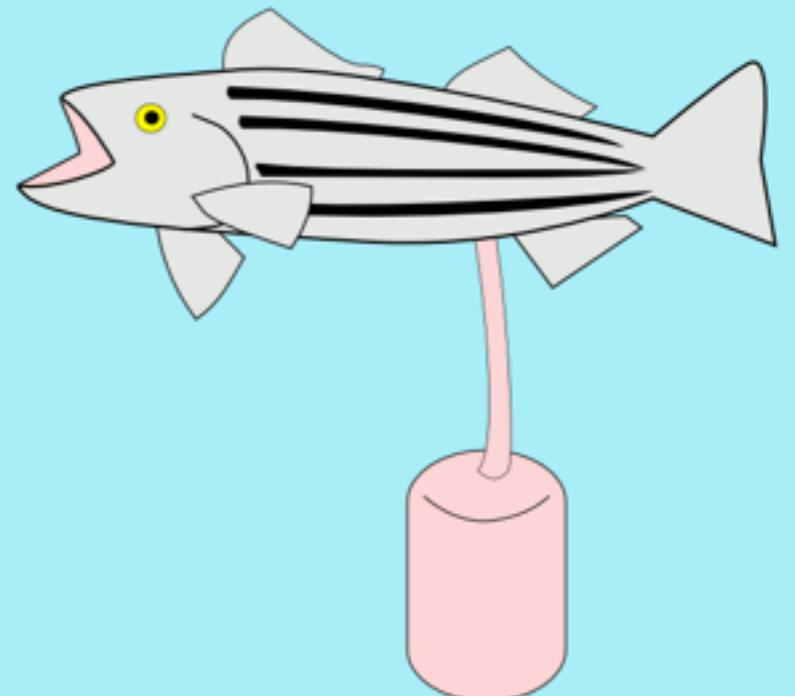
Annette Bieniusa et al

<http://arxiv.org/abs/1210.3368>

**WHO USES THE LIB?**



A



B

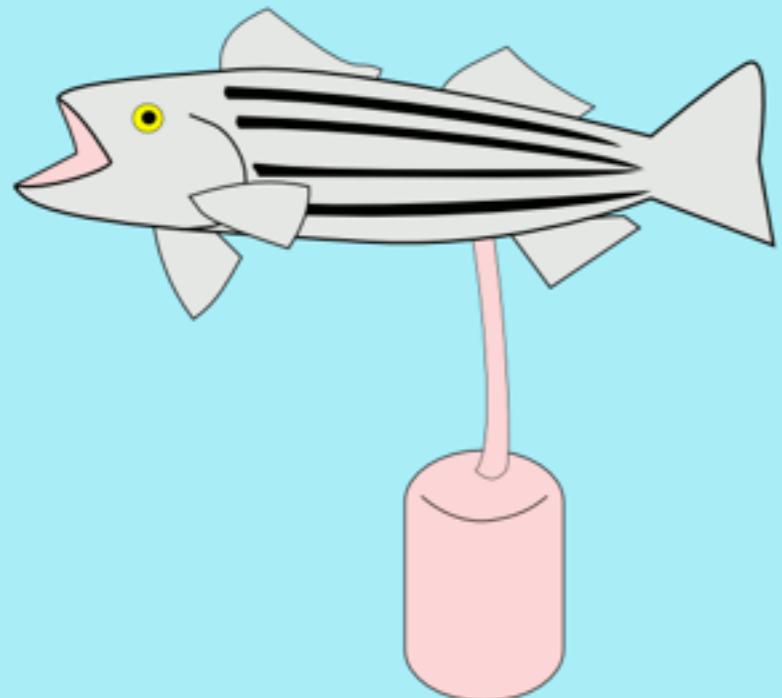
# who's the actor?

Client 1

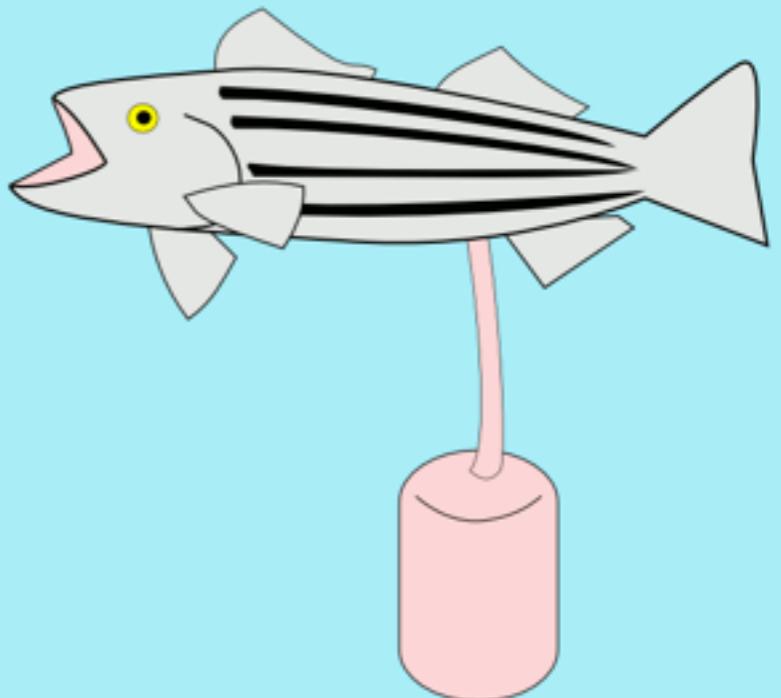
Client 2

.....

Client 10000



A



B

{c<sub>1</sub>, 1}

{c<sub>2</sub>, 1}

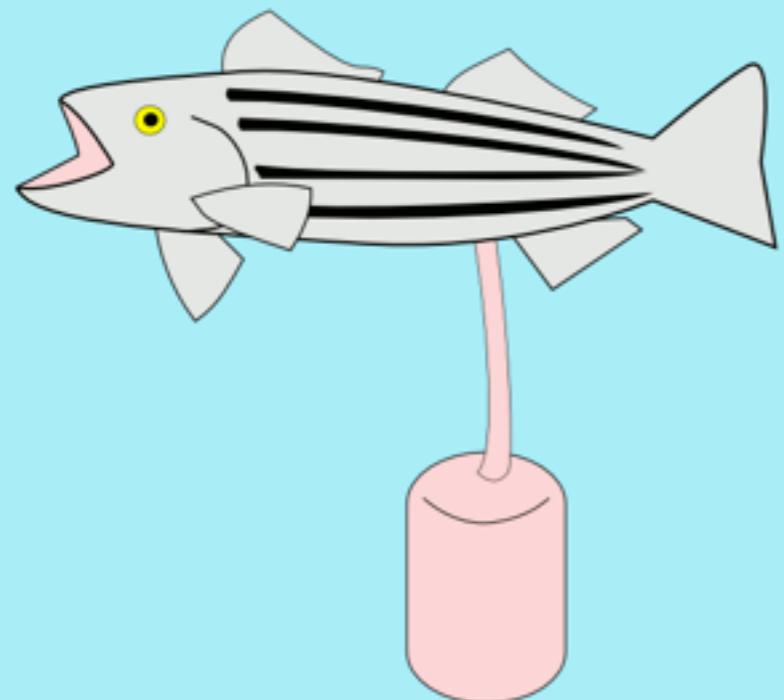
Client 1

Client 2

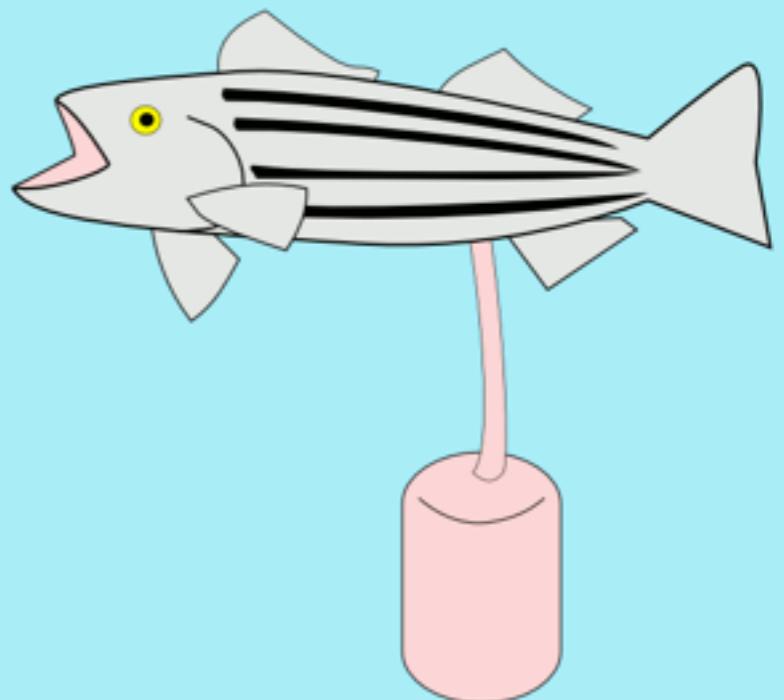
.....

{c<sub>X</sub>, 1}

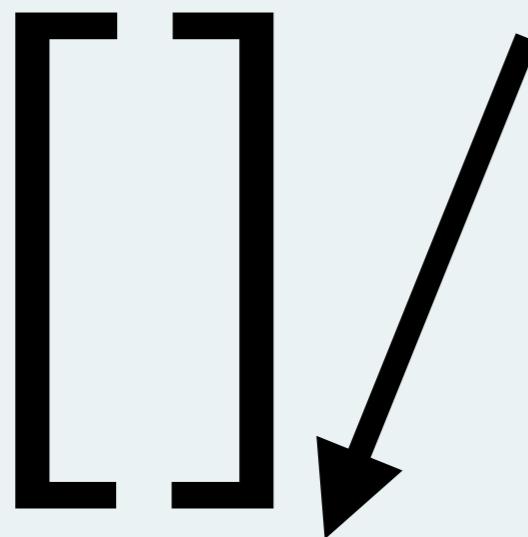
Client 10000



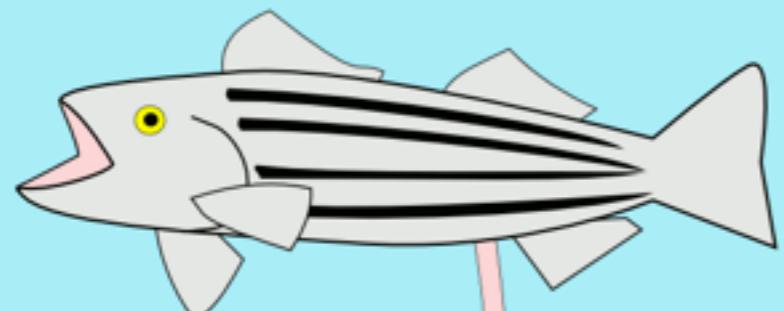
A



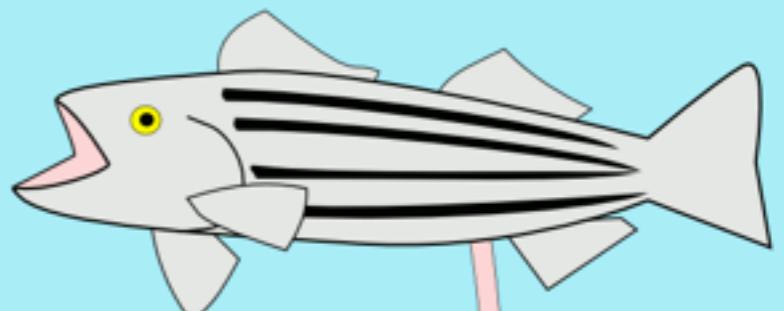
B



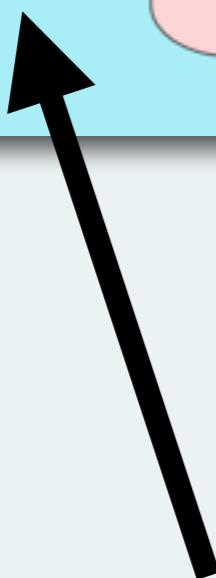
Client 1



A

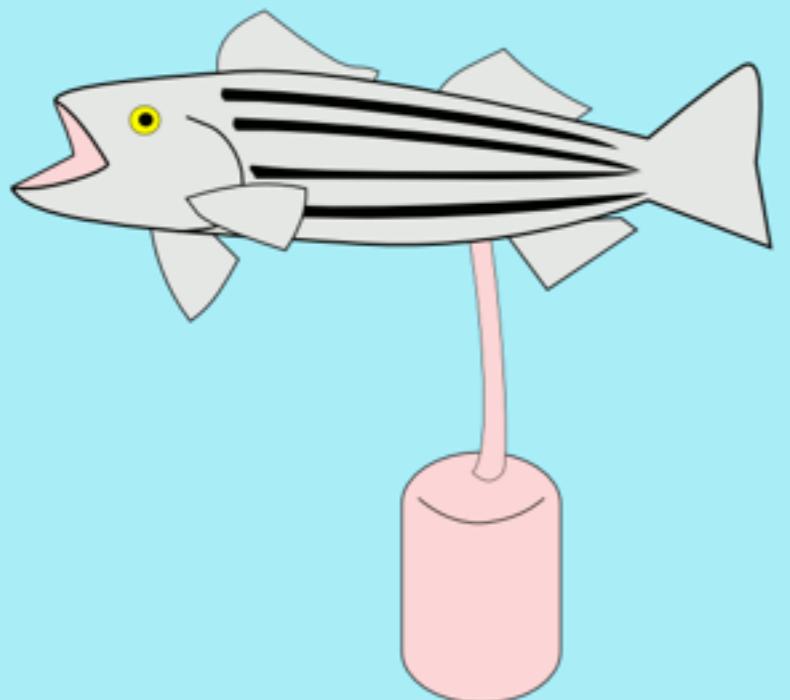


B

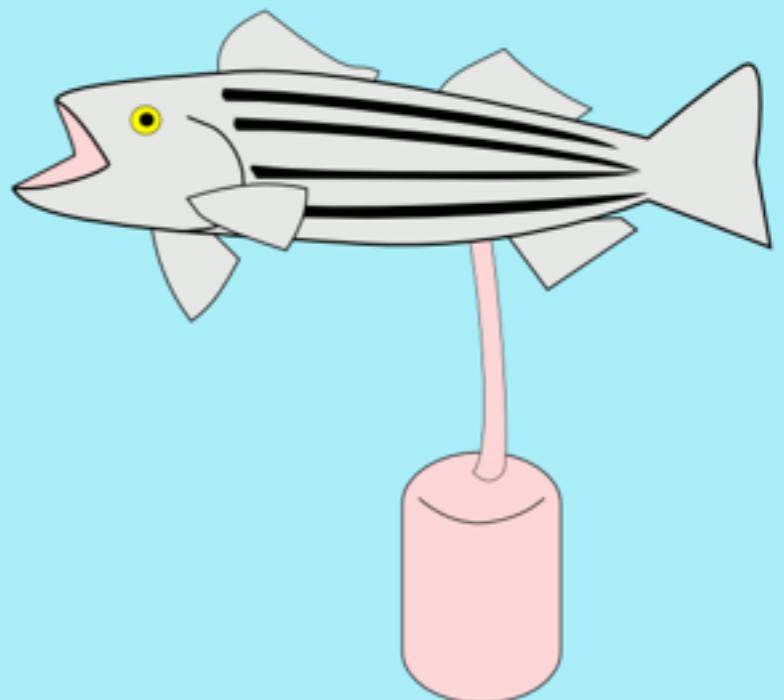


{c1, 1} Shelly

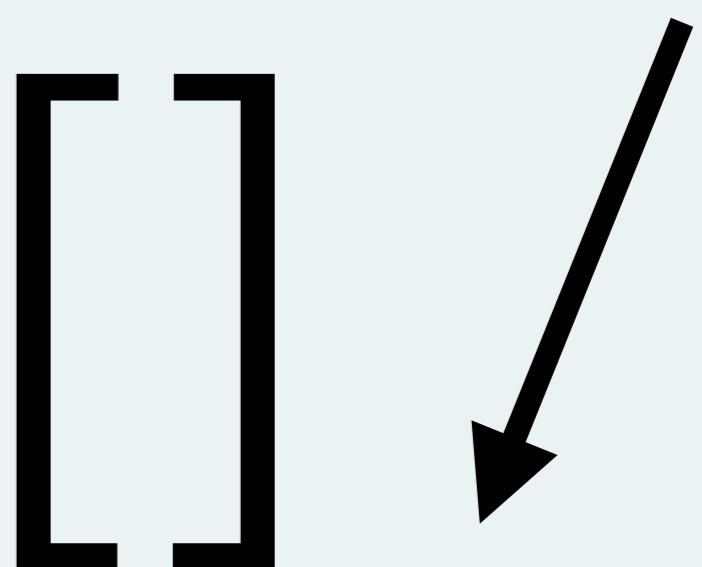
Client 1



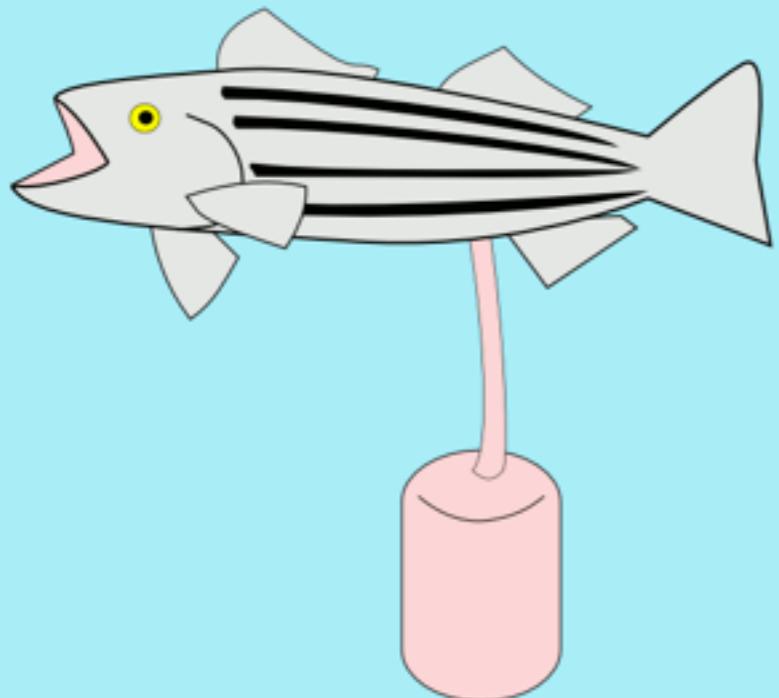
A



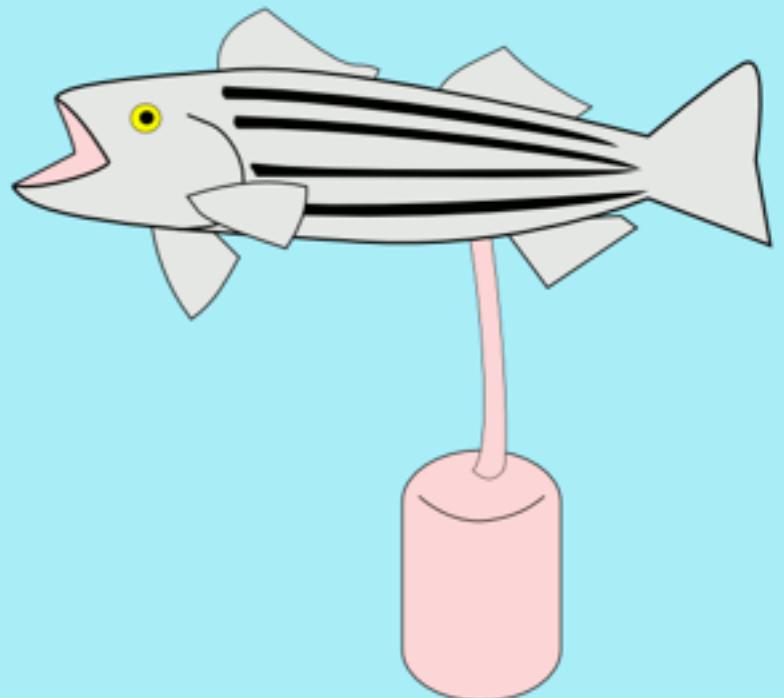
B



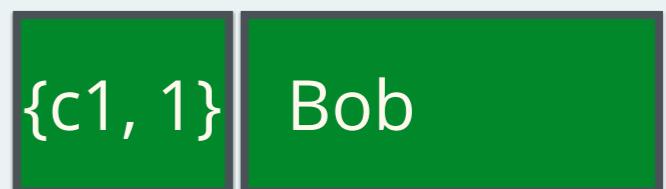
Client 1



A

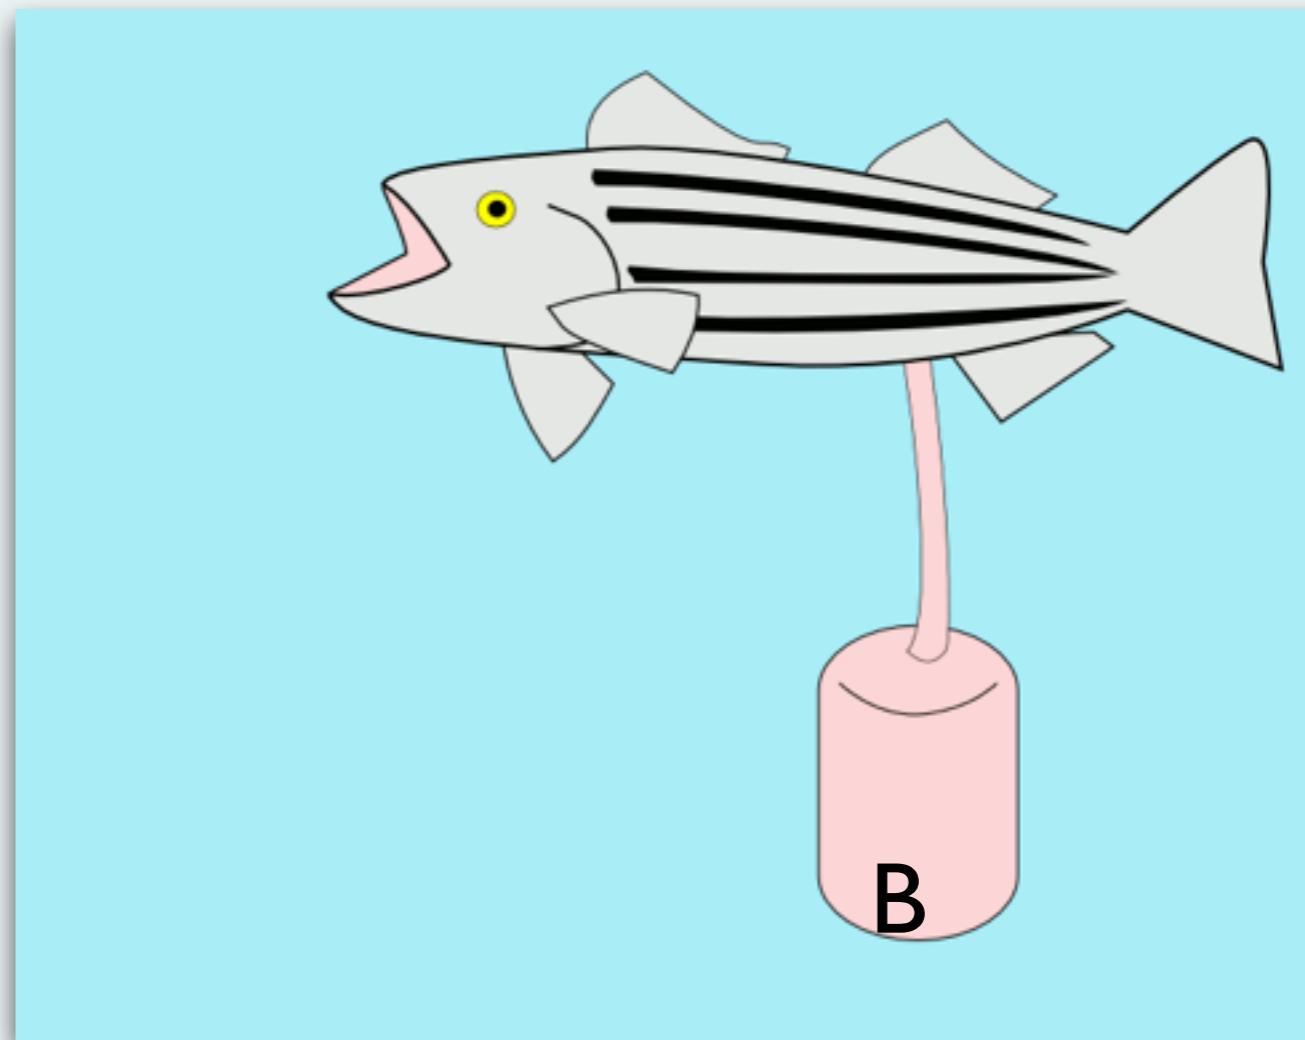
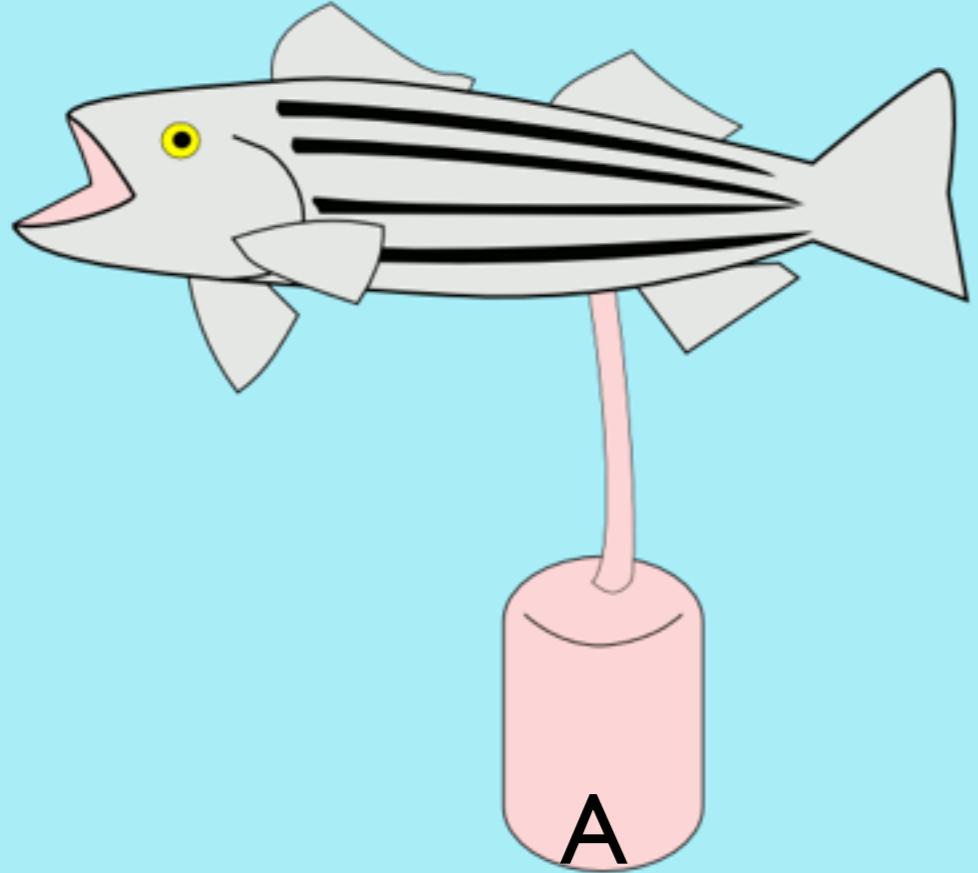


B



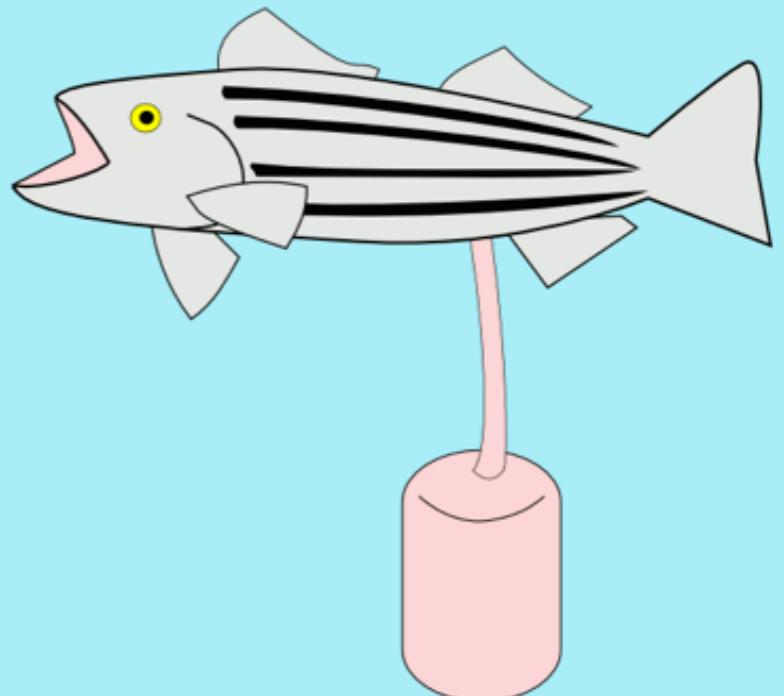
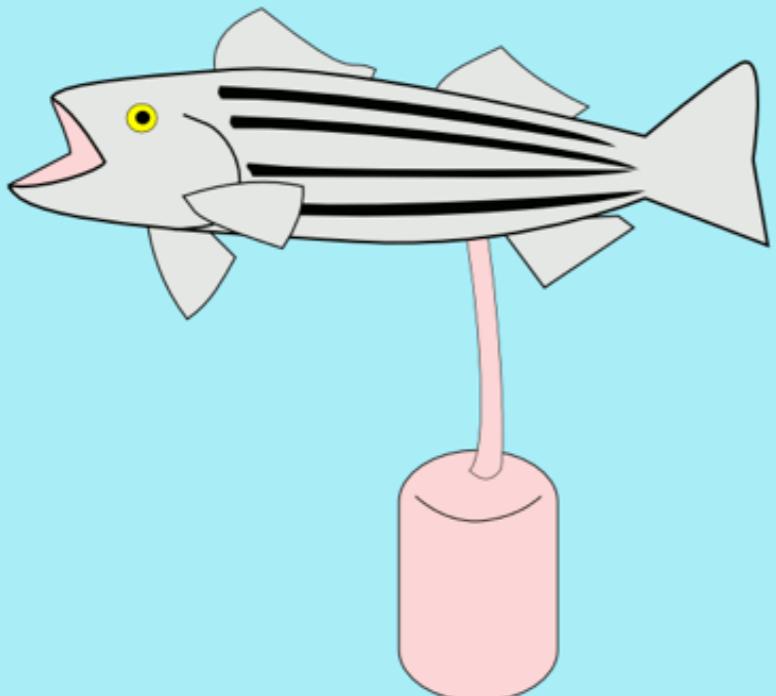
Client 1

Read  
Your  
Own  
Writes



**REPLICAS  
ACTORS**

**HOW TO USE  
THE LIB?**



A

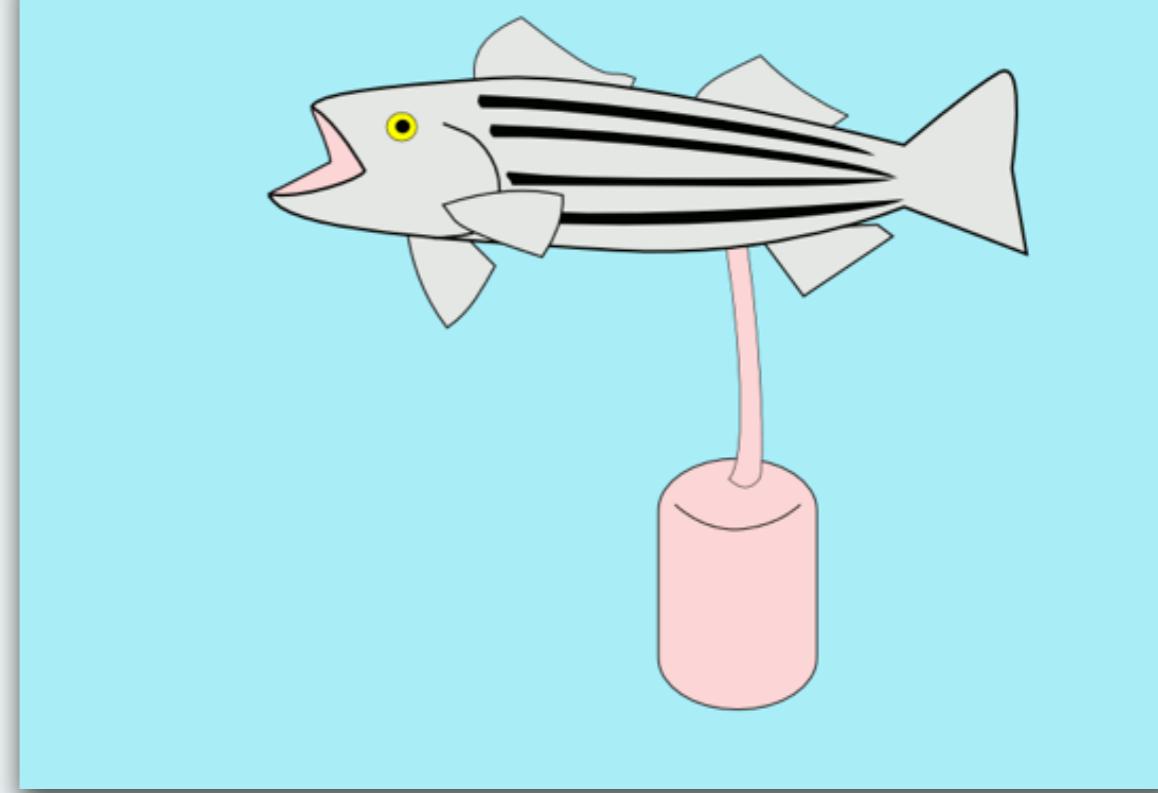
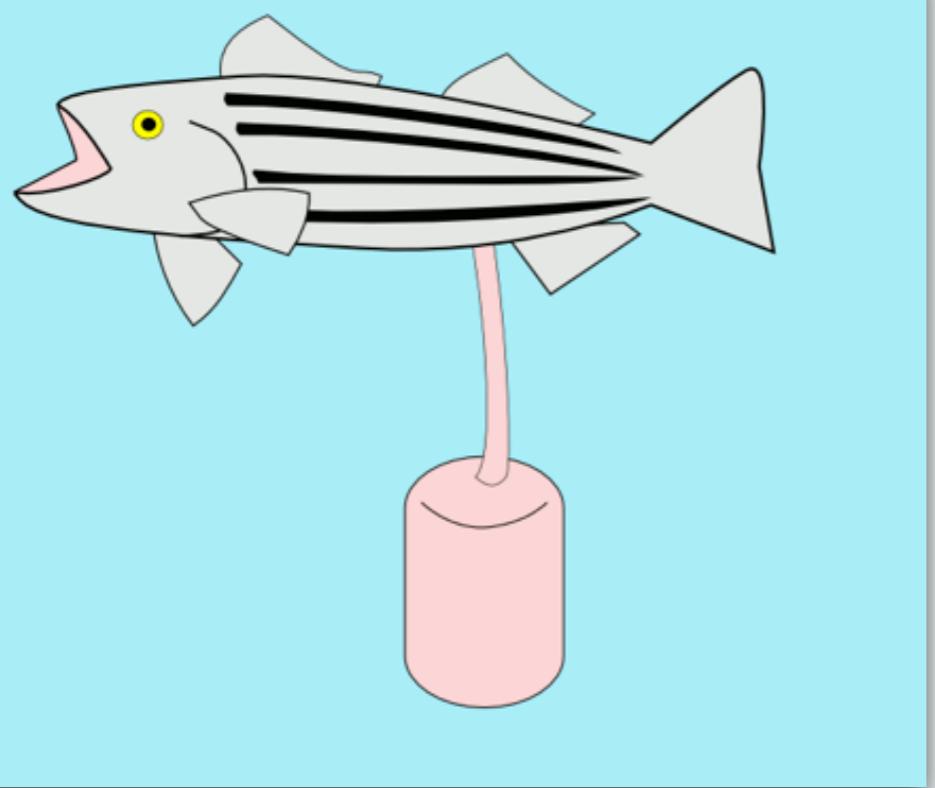


B



SHOPPING CART

[HAIRDRYER, PENCIL CASE]

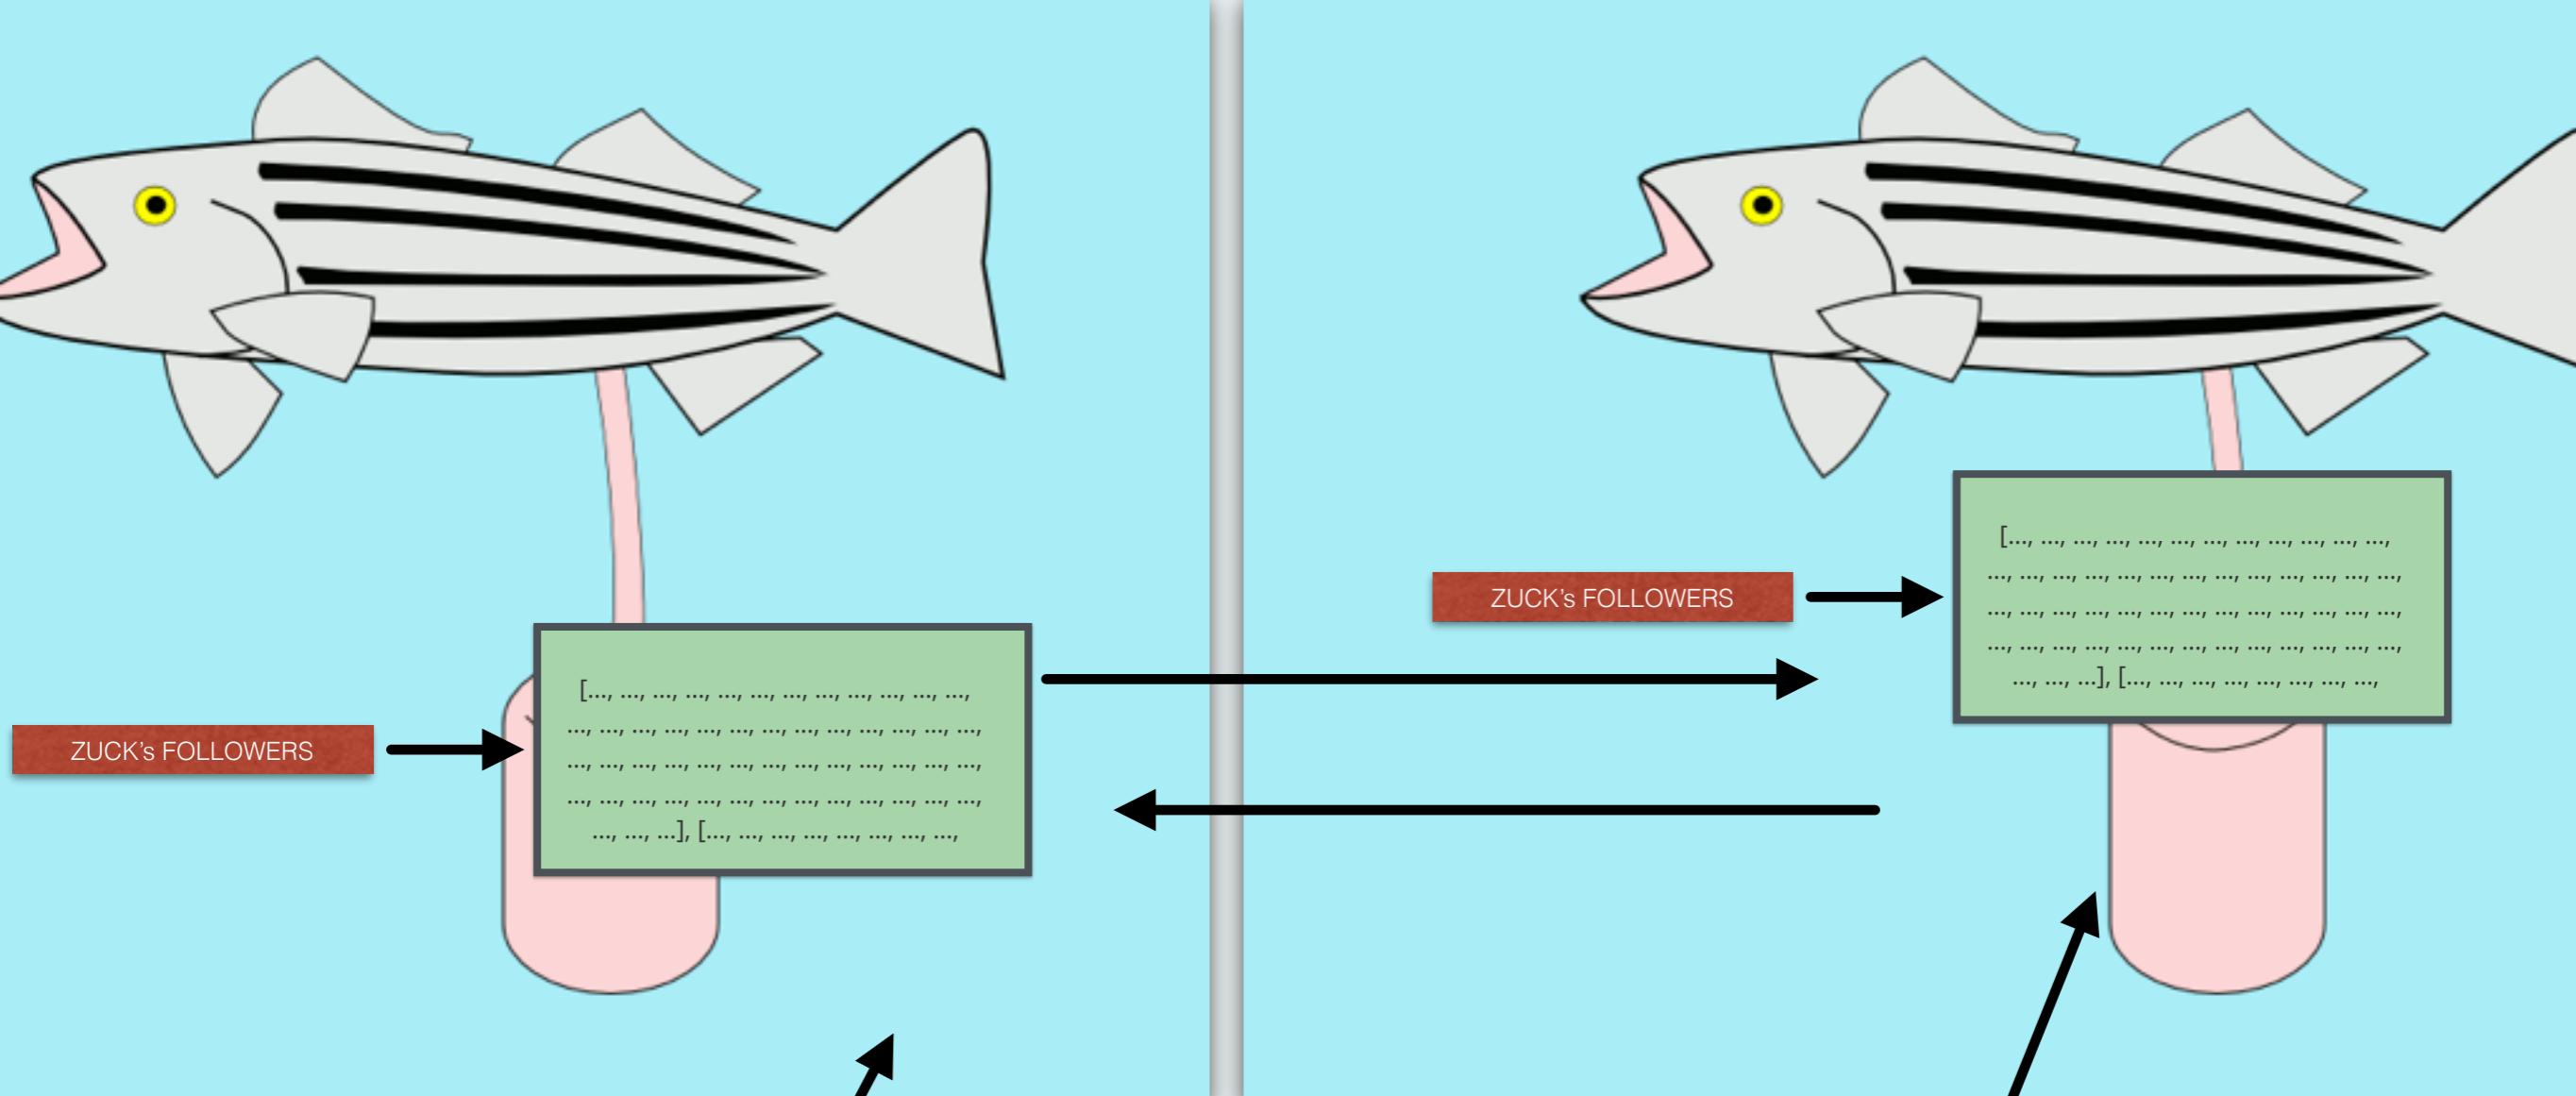


A

B

## ZUCK's FOLLOWERS?

[..., ..., ..., ..., ..., ..., ..., ..., ..., ..., ..., ...,  
..., ..., ..., ..., ..., ..., ..., ..., ..., ..., ..., ..., ...,  
..., ..., ..., ..., ..., ..., ..., ..., ..., ..., ..., ..., ...,  
..., ..., ..., ..., ..., ..., ..., ..., ..., ..., ..., ..., ...,  
..., ..., ..., ..., ..., ..., ..., ..., ..., ..., ..., ..., ...], [...,

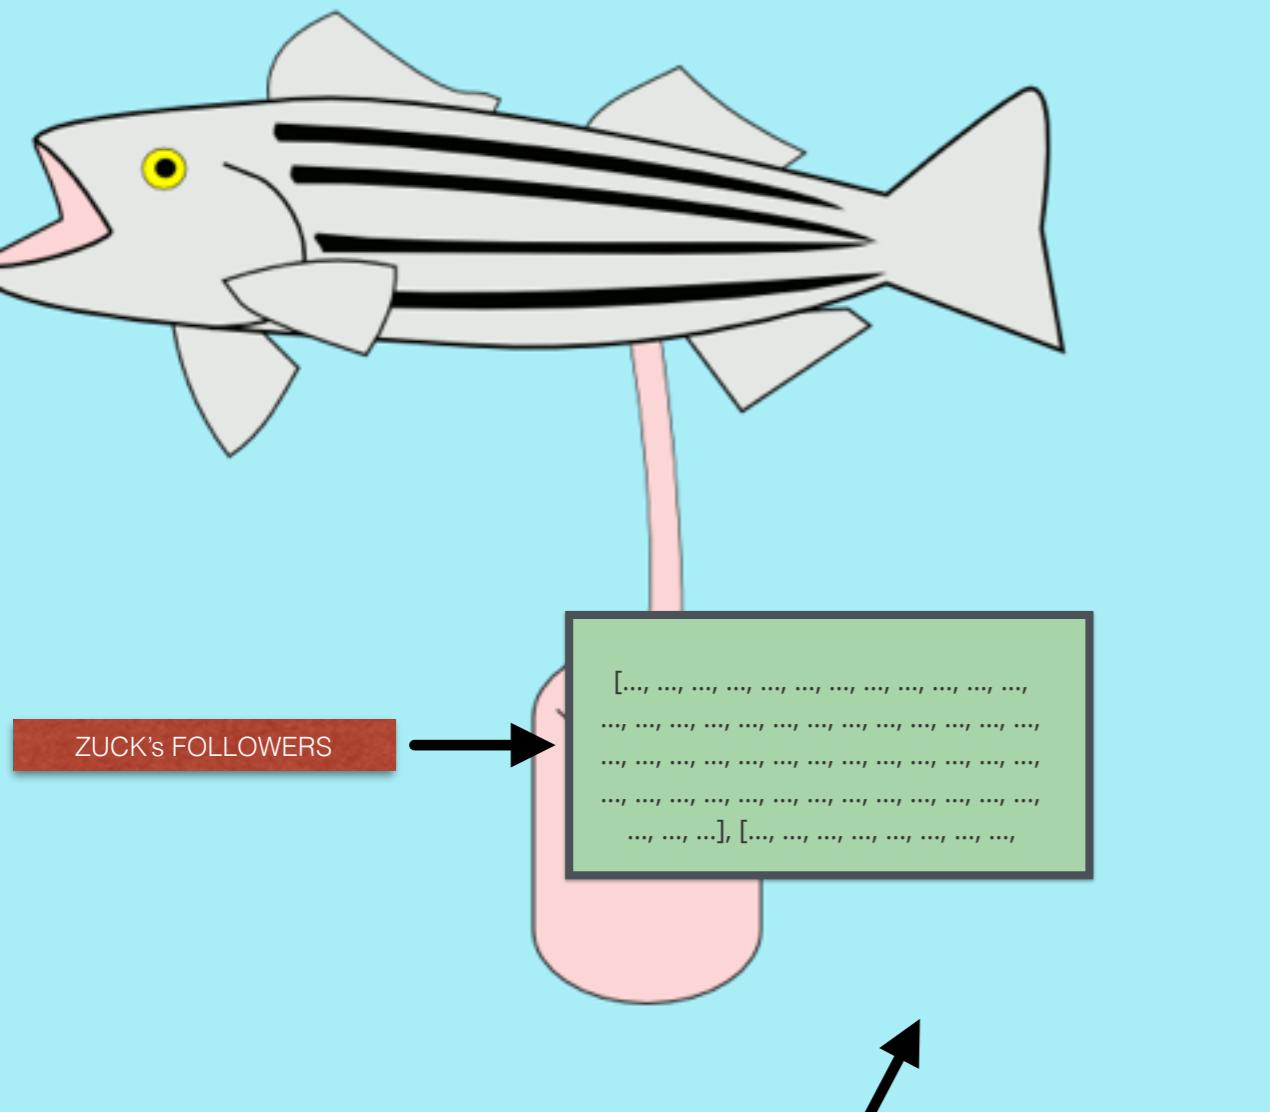


Add “Shelly”

Client X

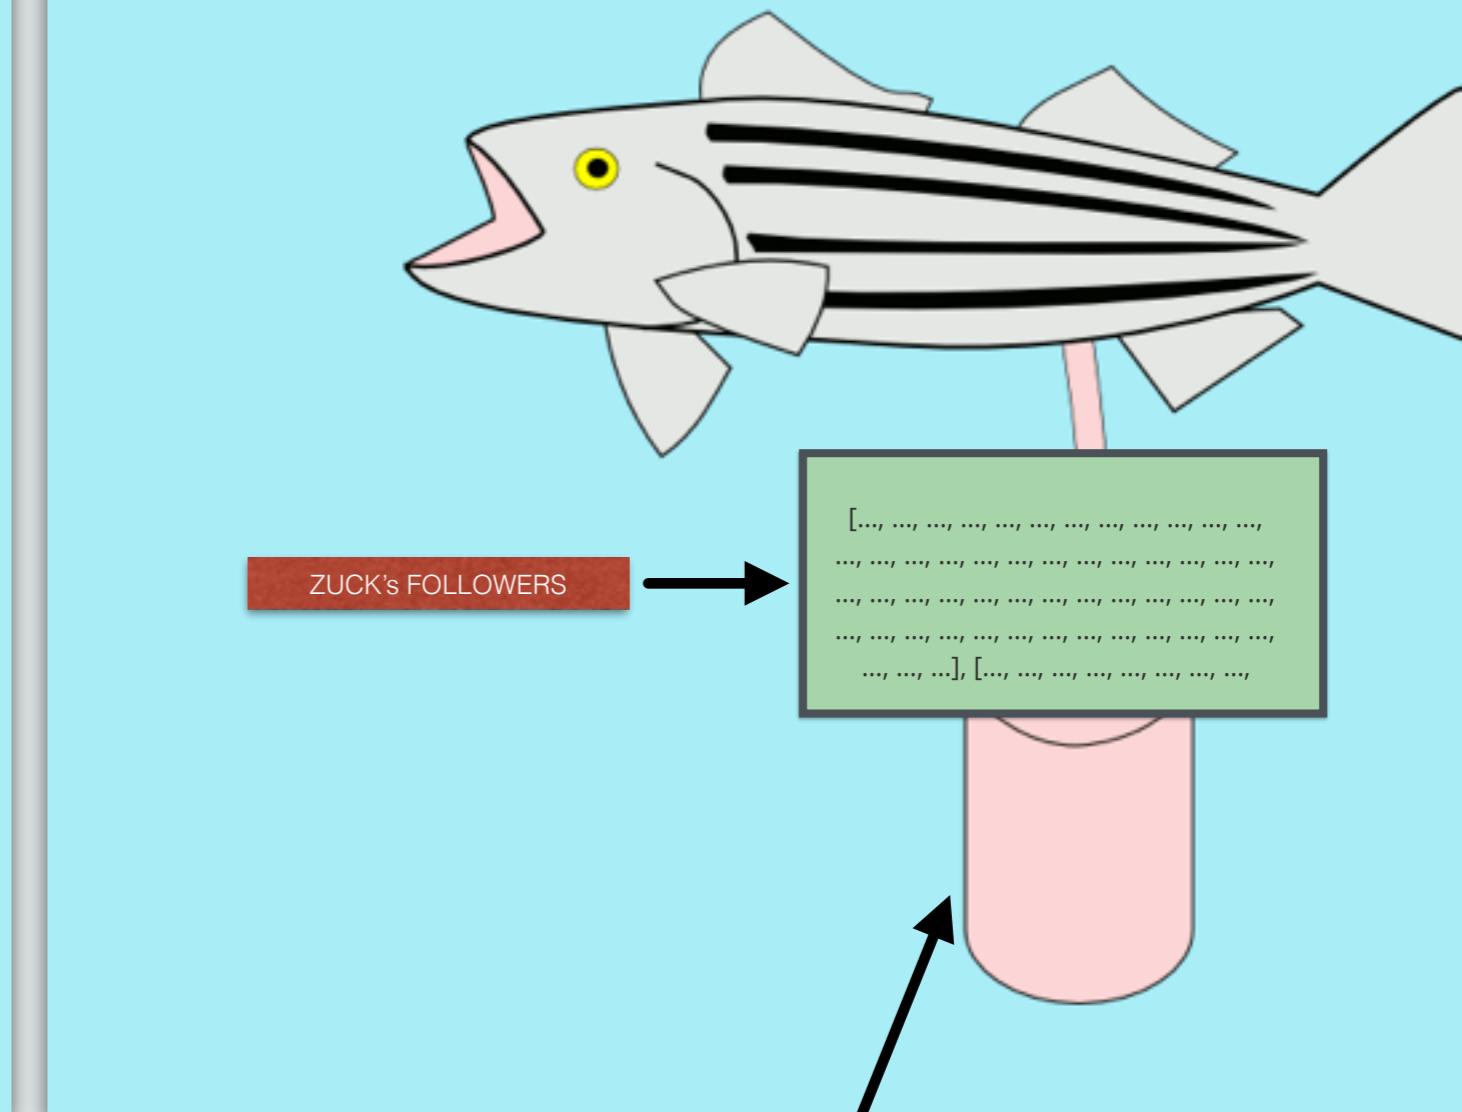
Add “Bob”

Client Y



# remove “Shelly”

# Client X

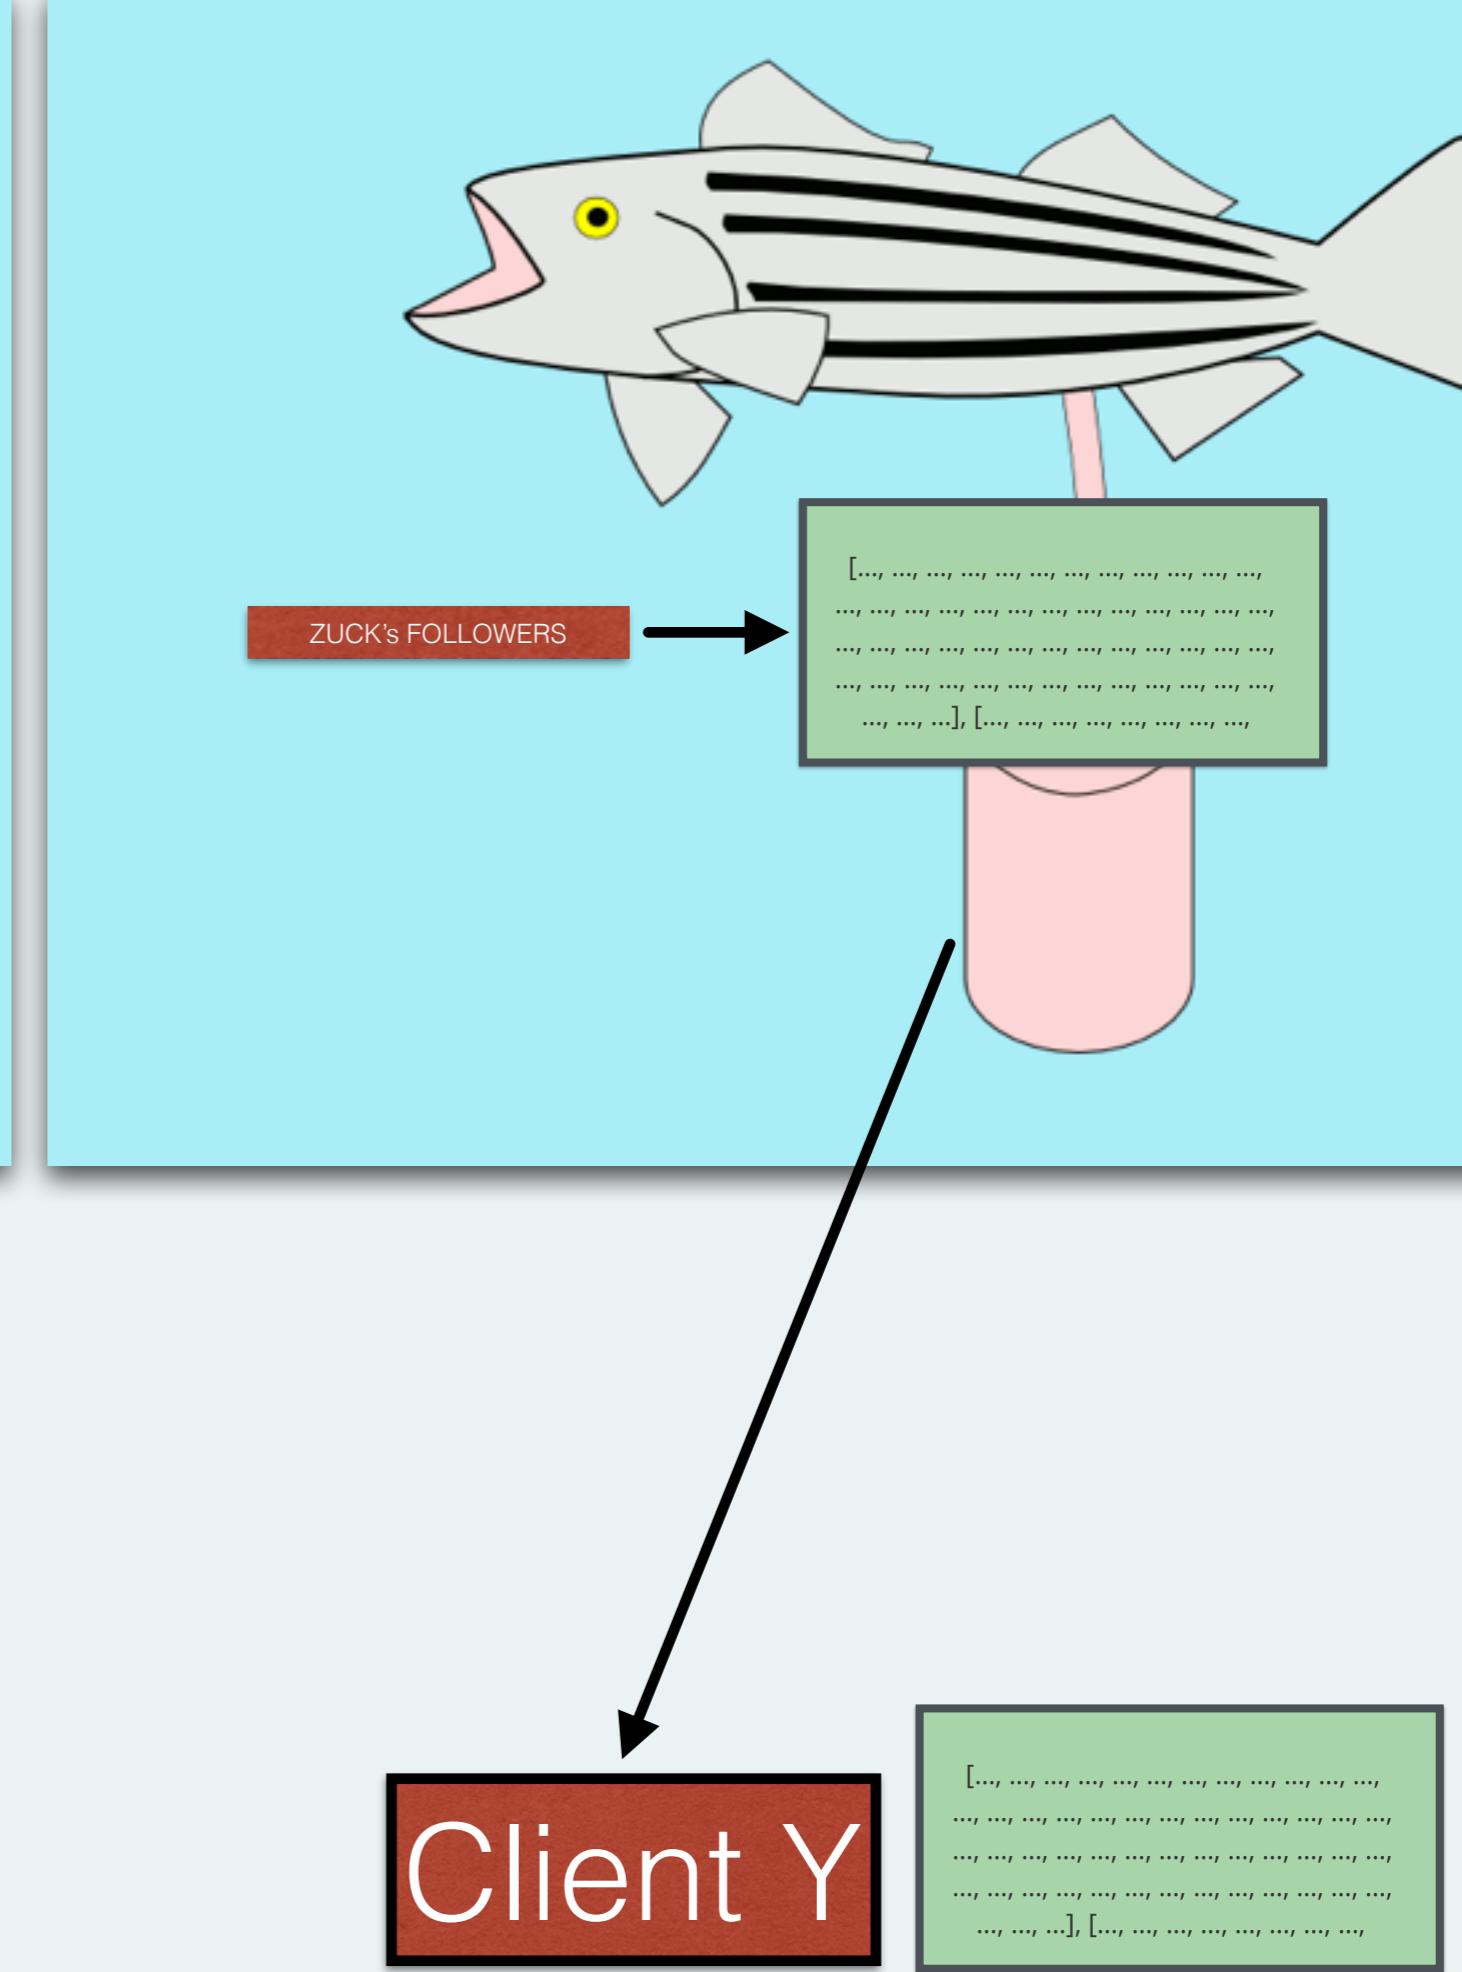
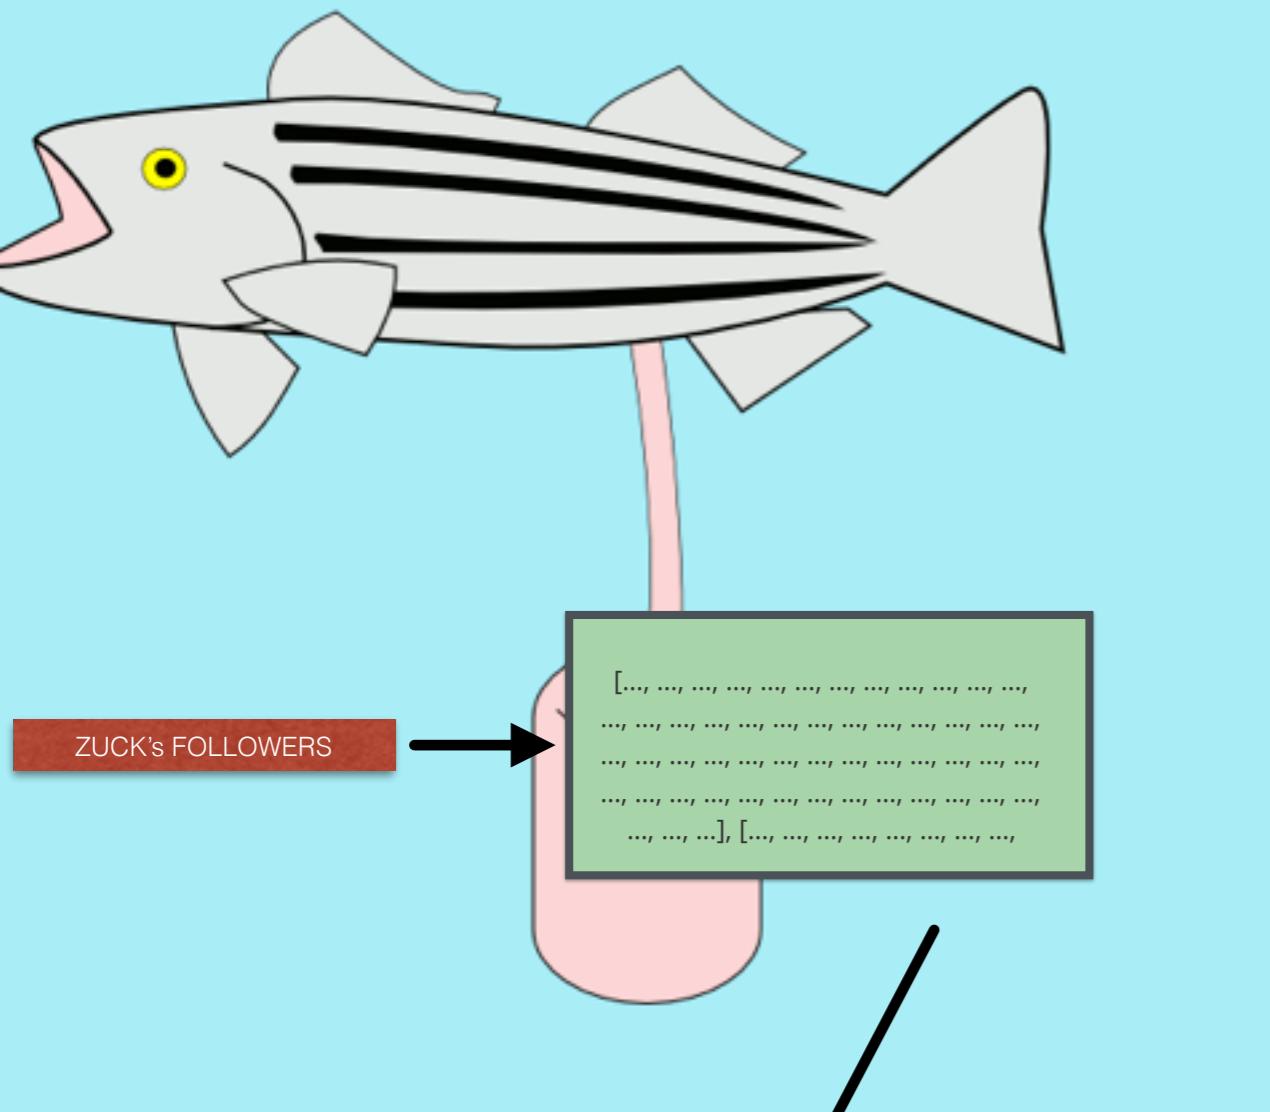


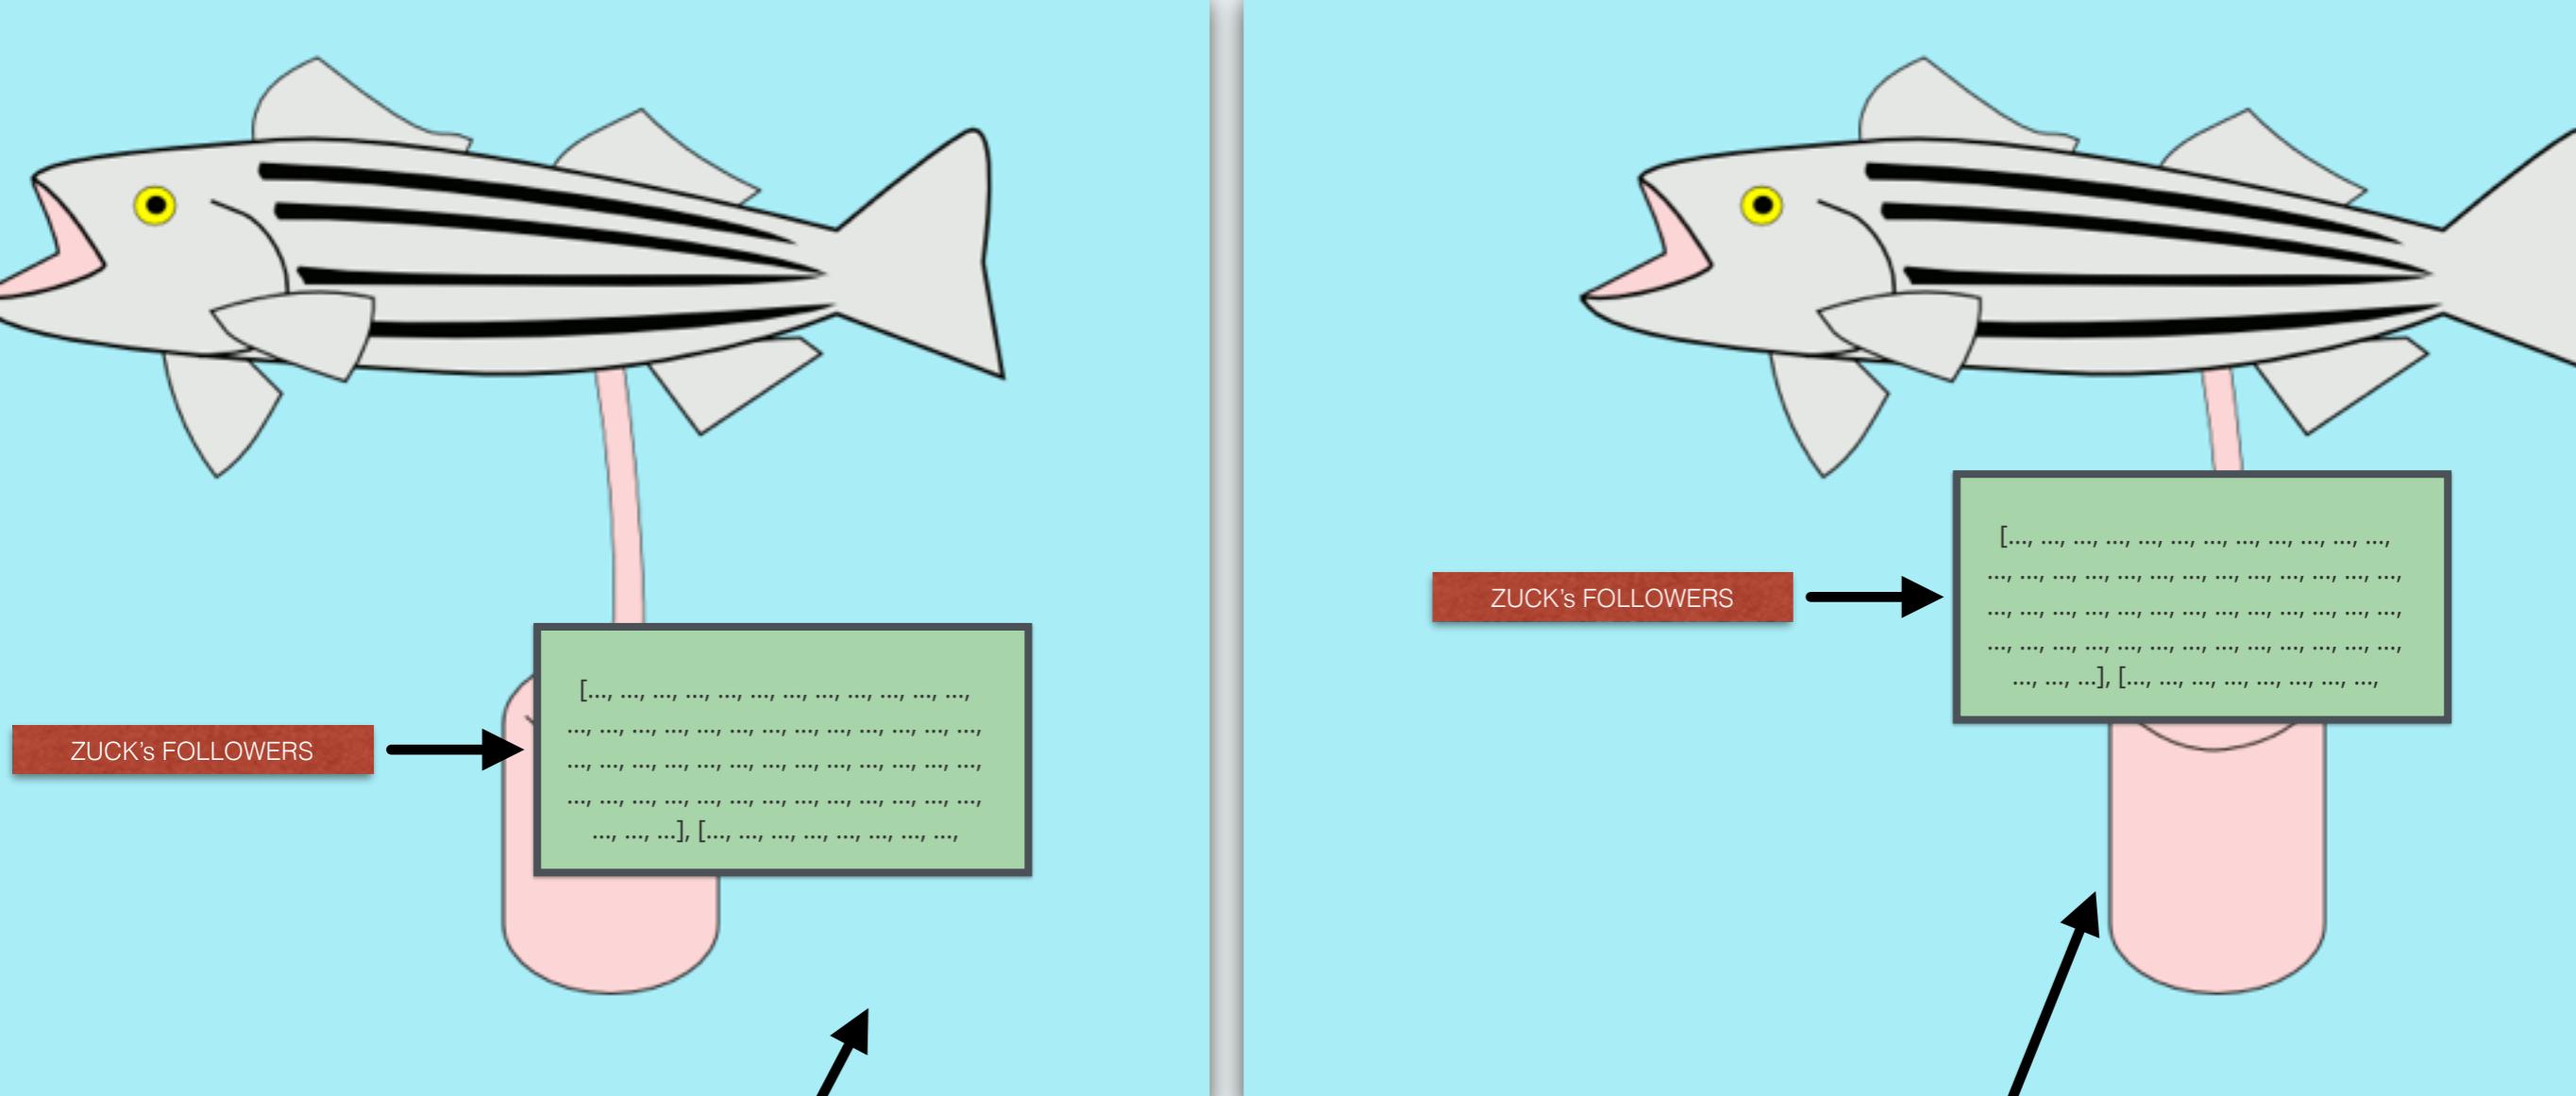
# remove “Bob”

# Client Y



Observed  
Remove





**remove “Shelly”**  
[{A3, B4}]

Client X

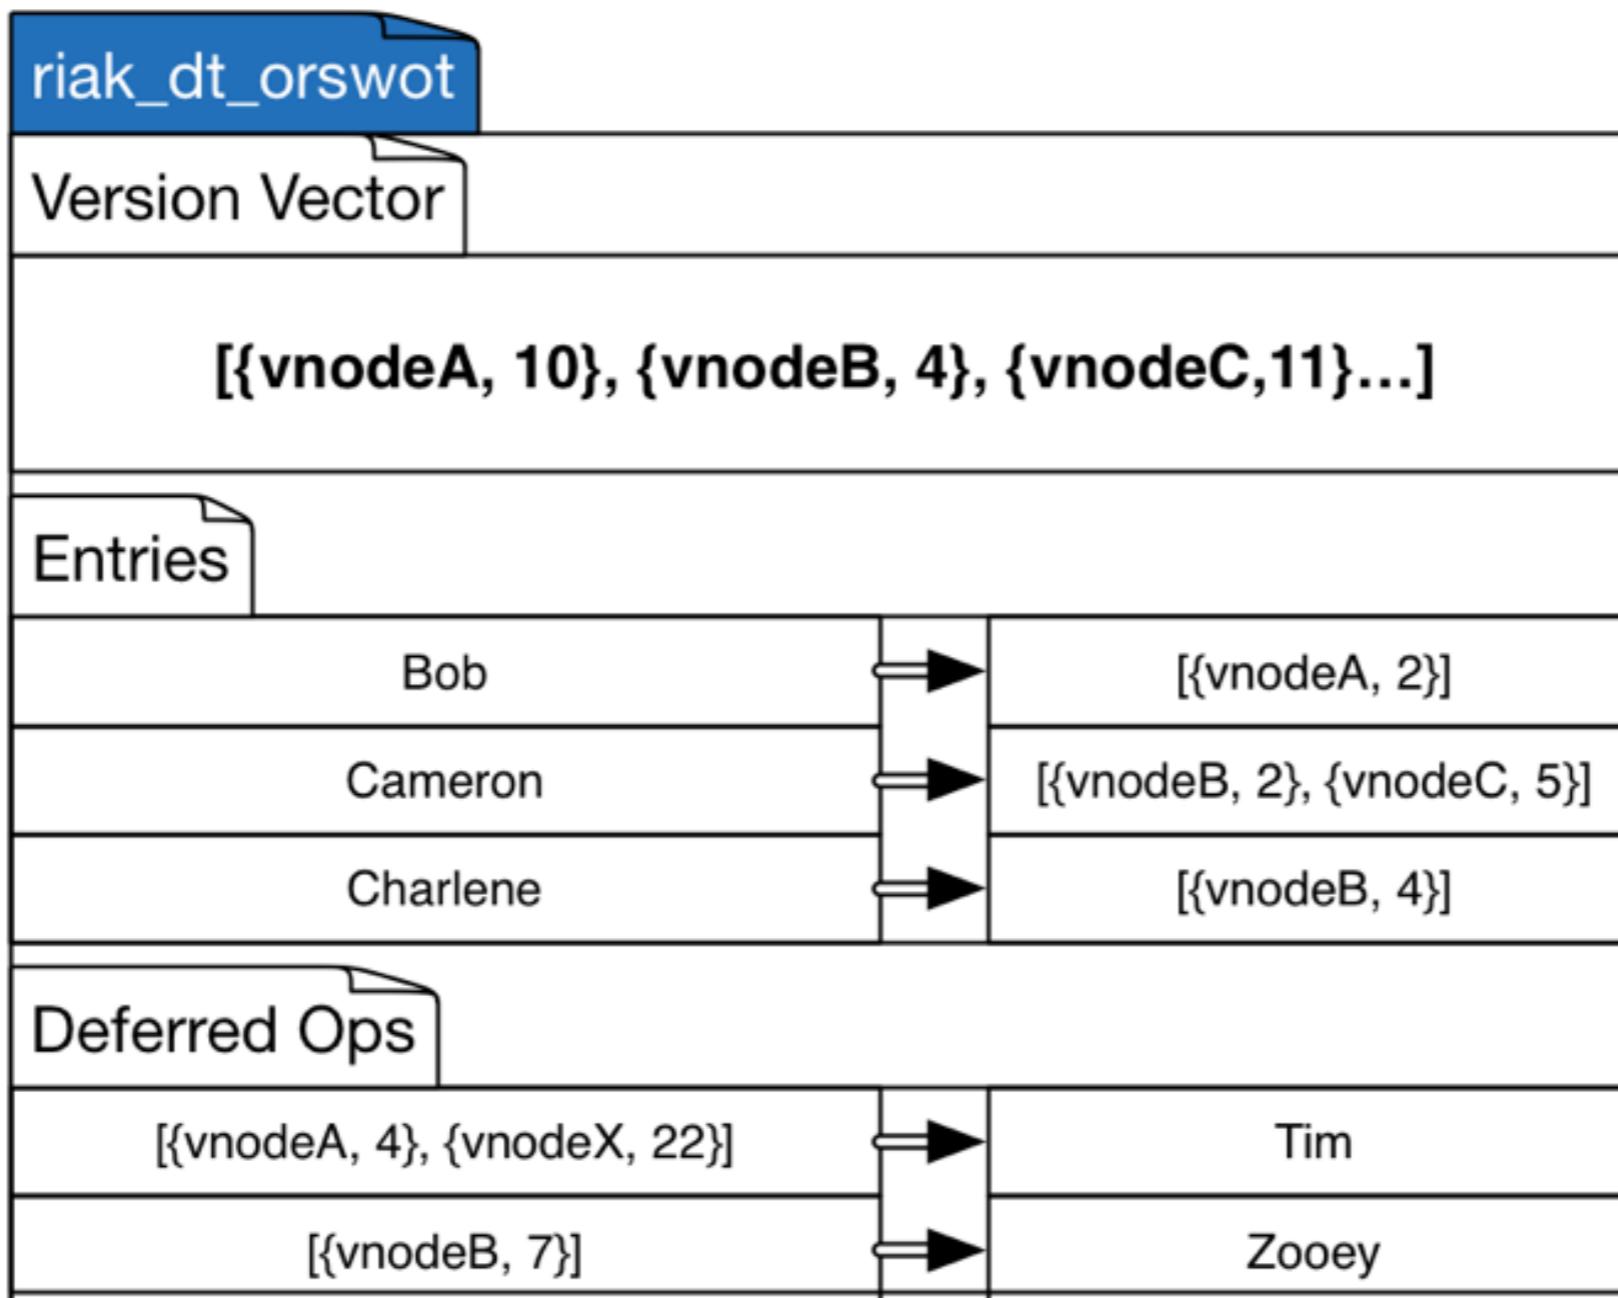
**remove “Bob”**  
[{A1, B5}]

Client Y

# Sets in Riak

- Operation Based API
  - With causal Context for removes!
- Vnode As Actor/Replica
  - Action-at-a-distance
- Full state replication

# Sets in Riak

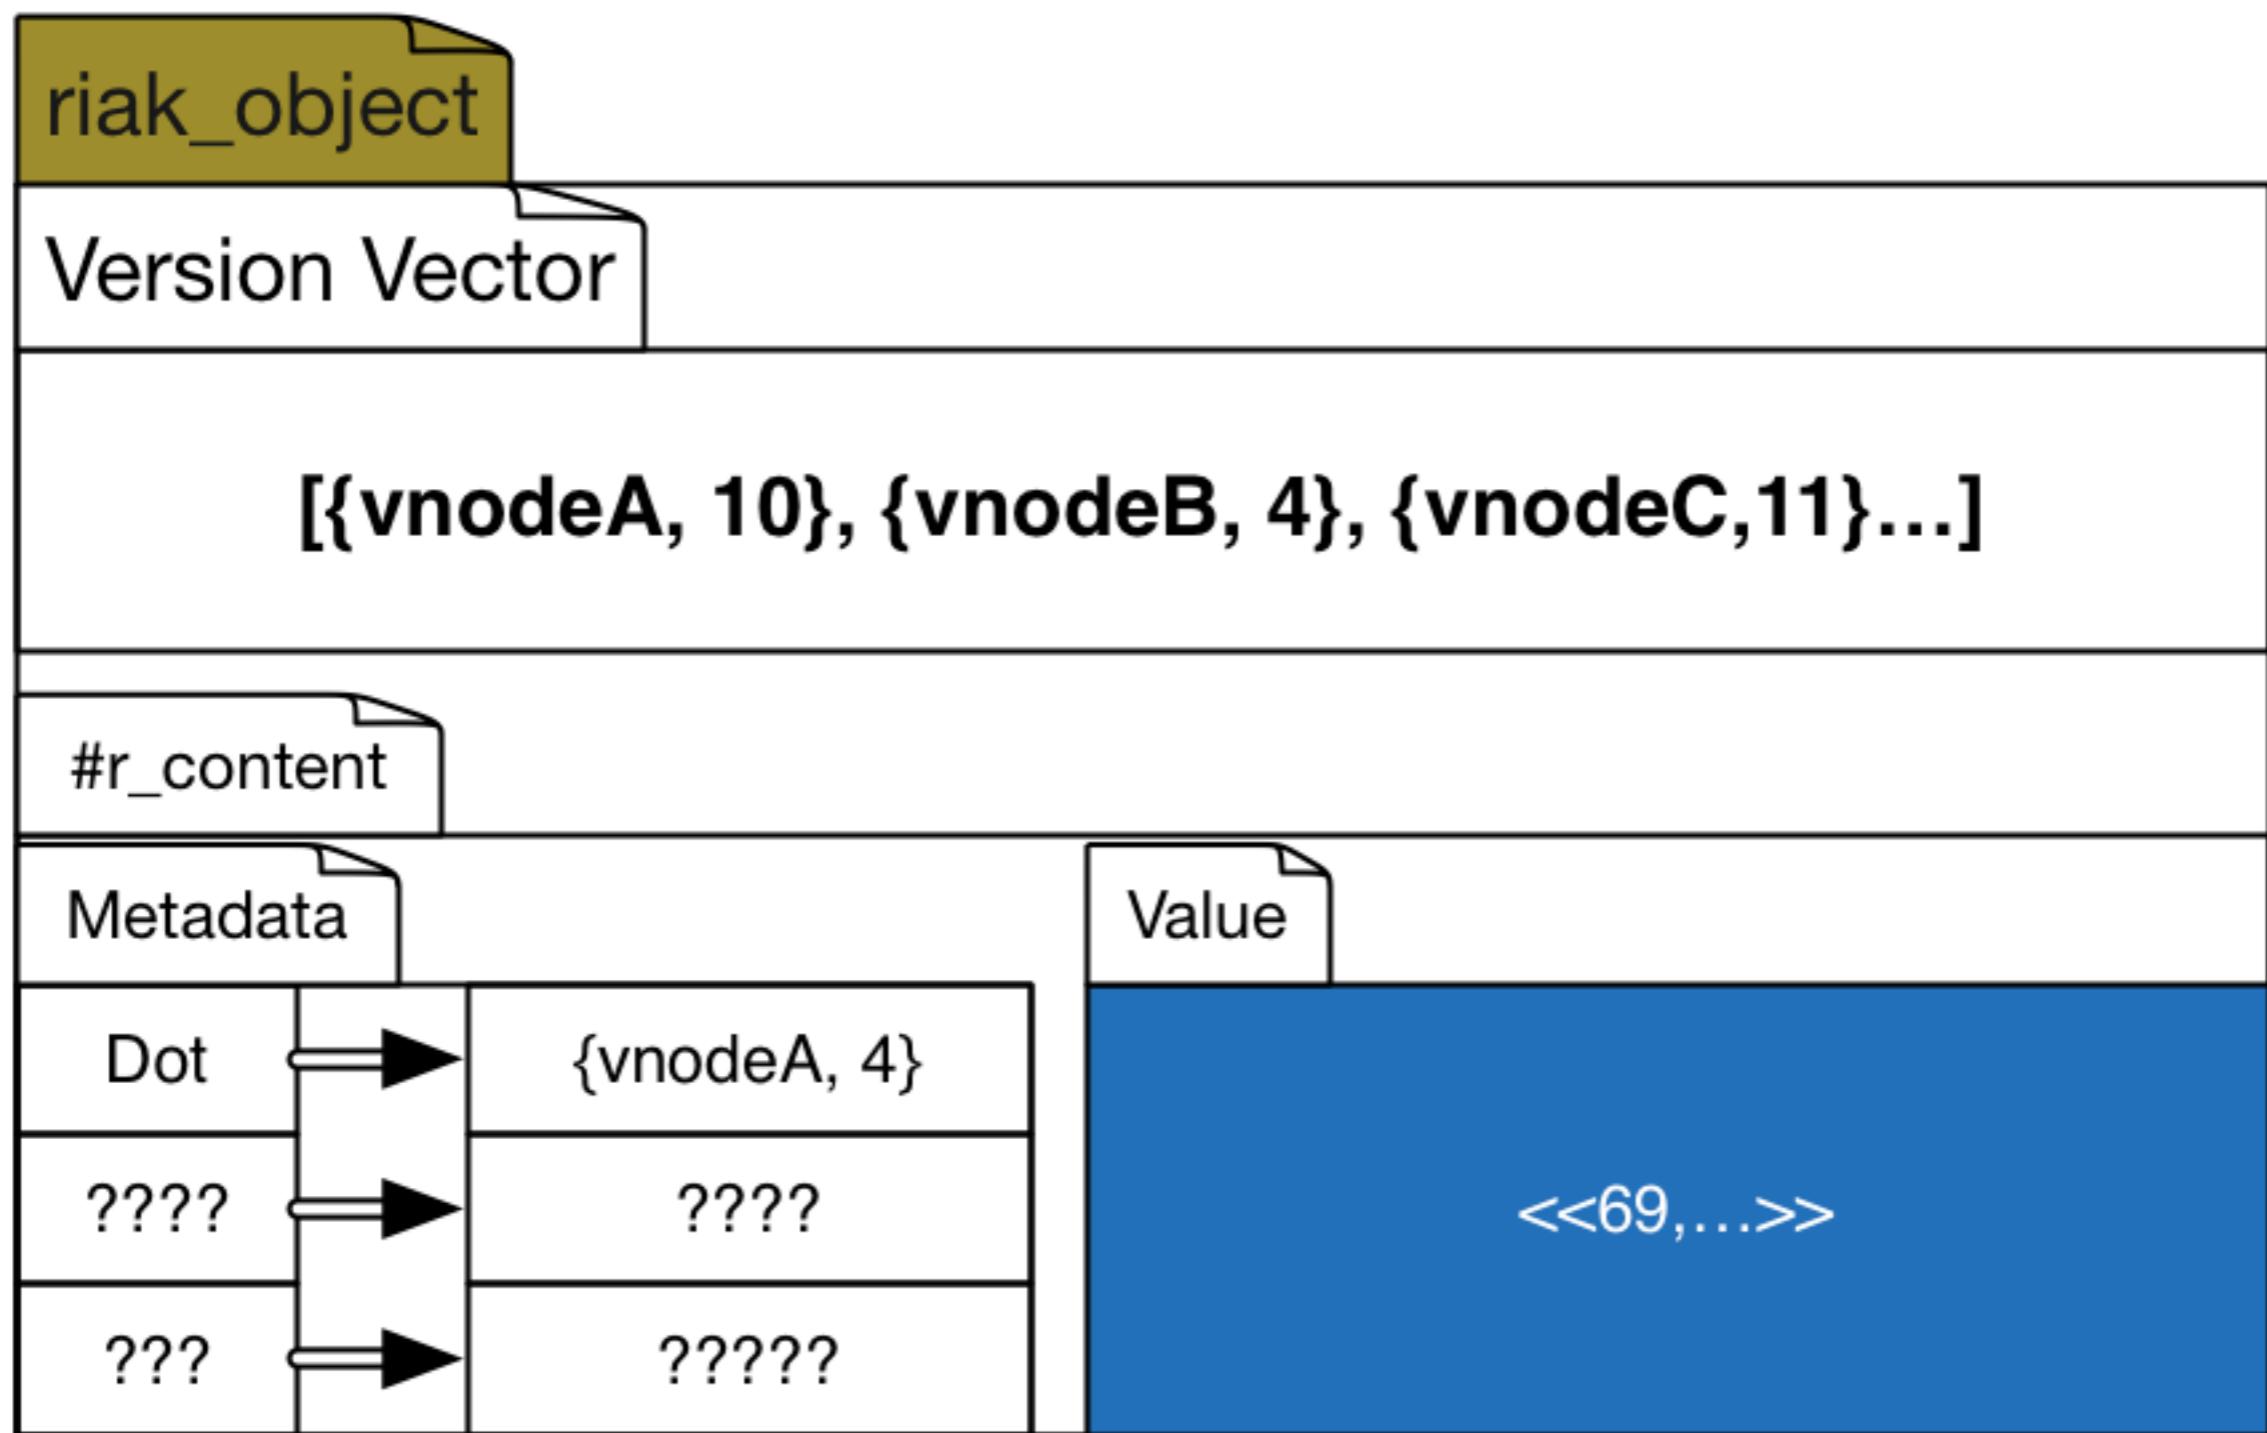


An optimized conflict-free replicated set

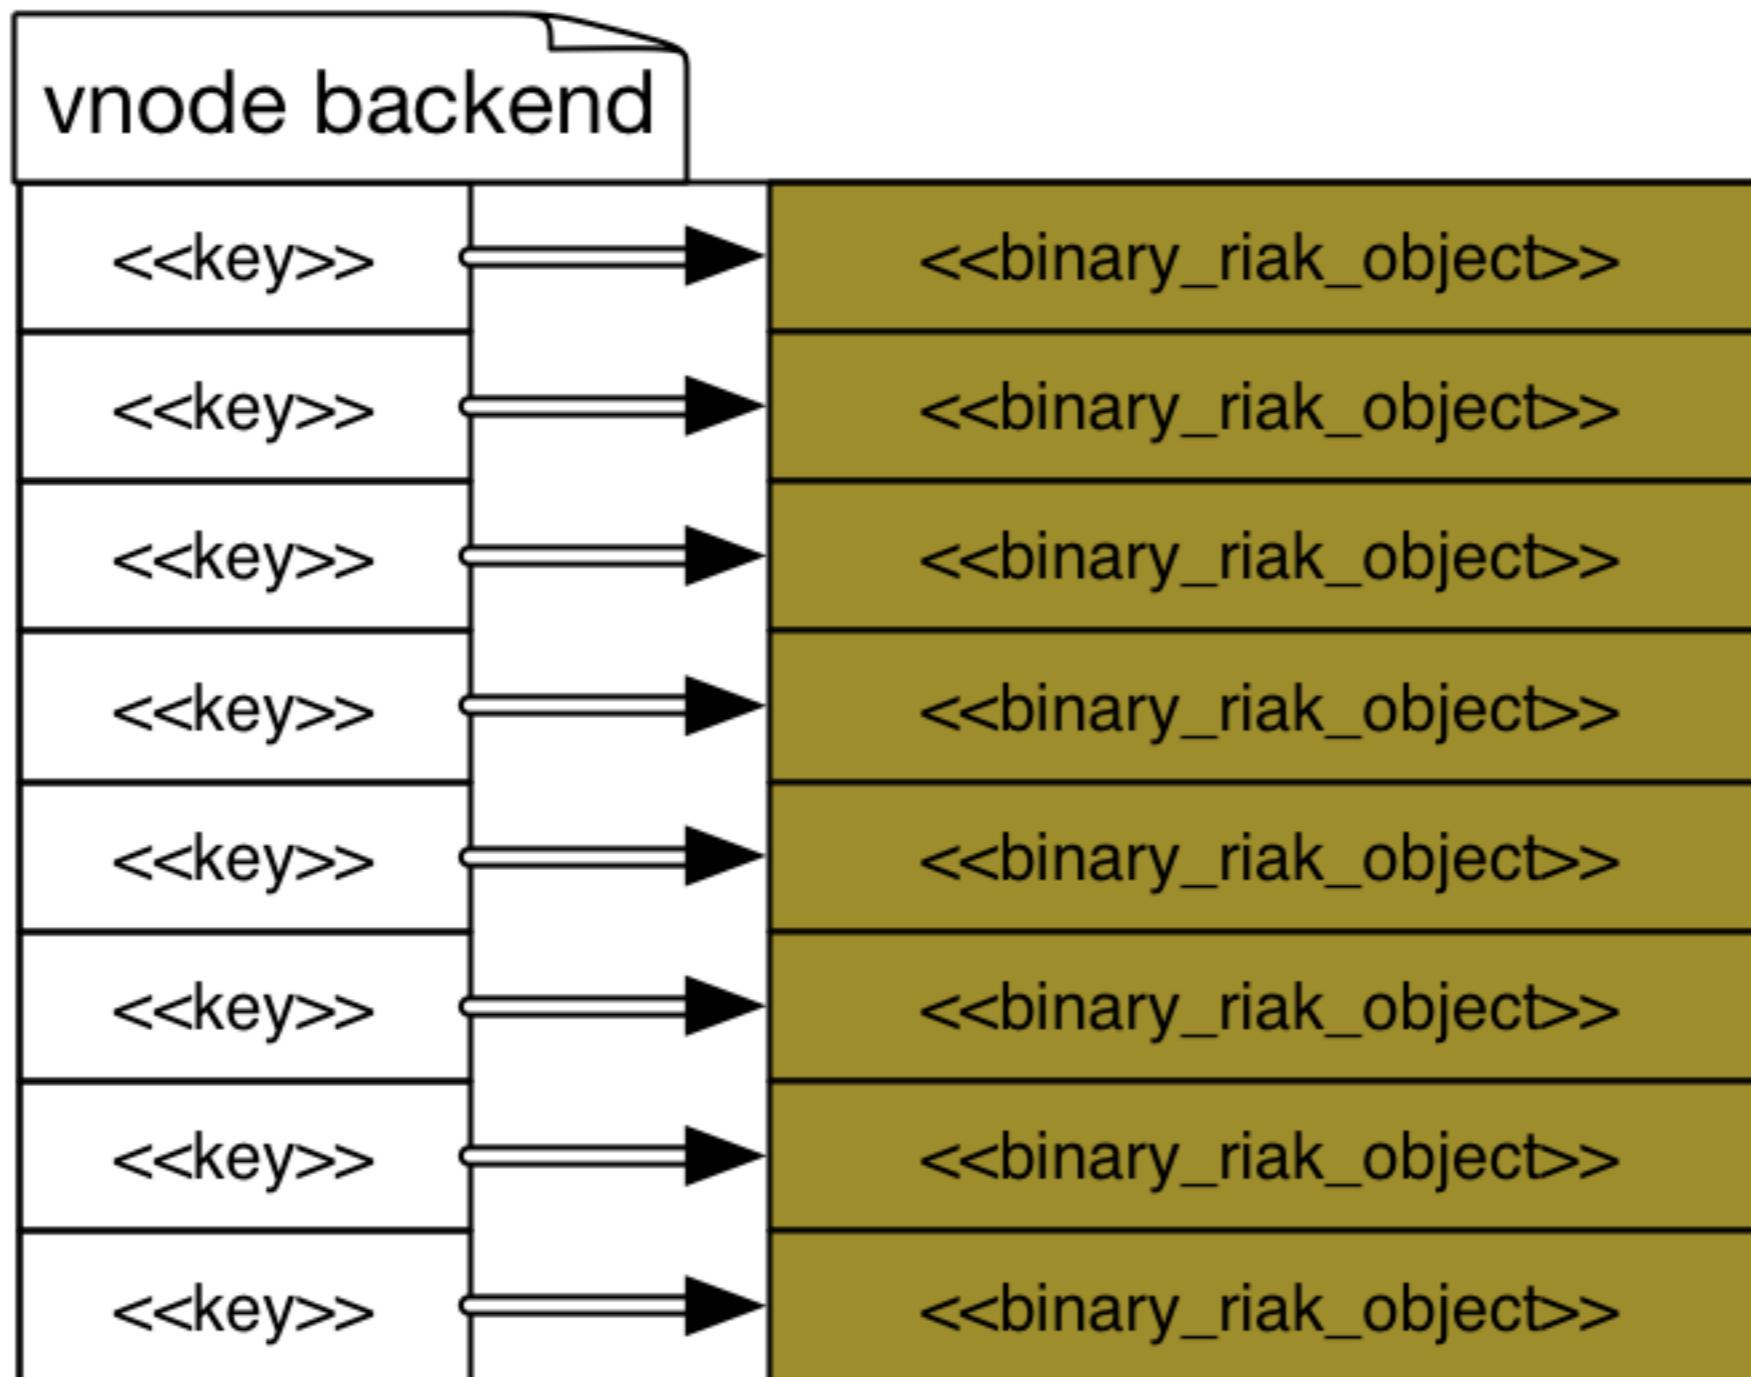
Annette Bieniusa et al

<http://arxiv.org/abs/1210.3368>

# Sets in Riak



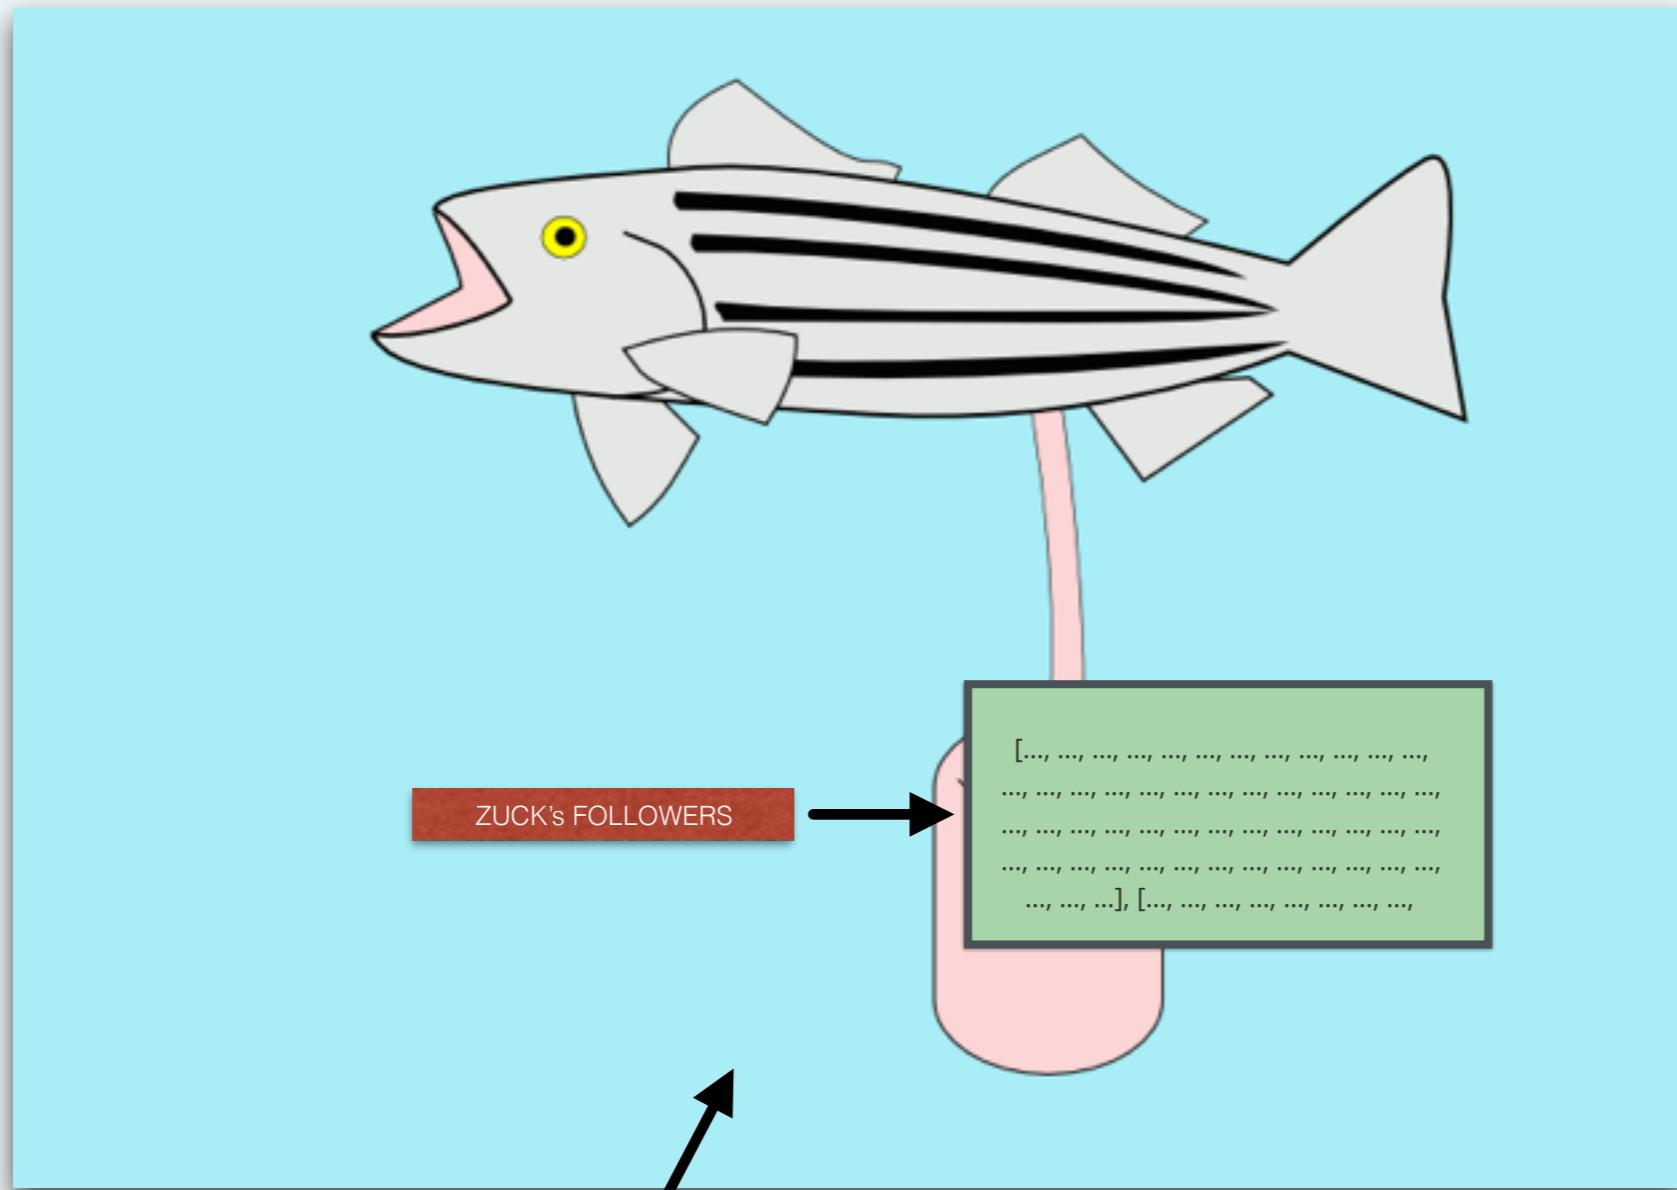
# Sets in Riak



# Sets in Riak

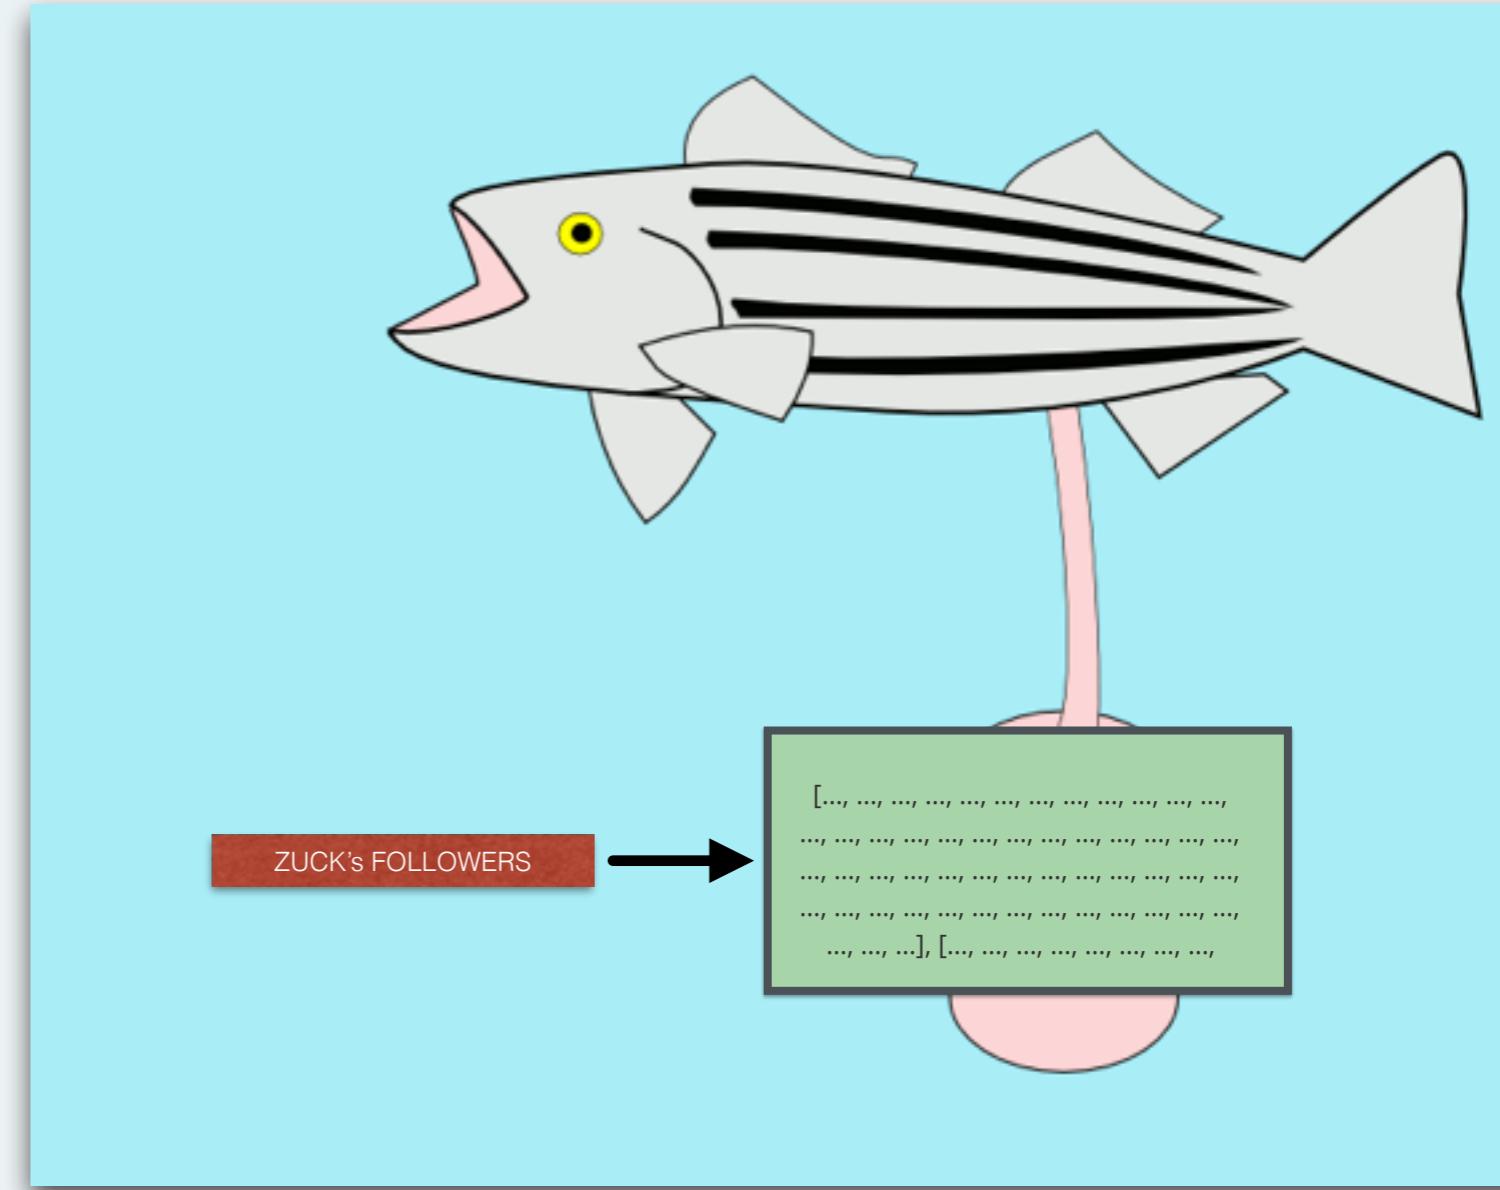


PHOTO © 2011 J. RONALD LEE, CC ATTRIBUTION 3.0.  
<https://www.flickr.com/photos/jronaldlee/5566380424>

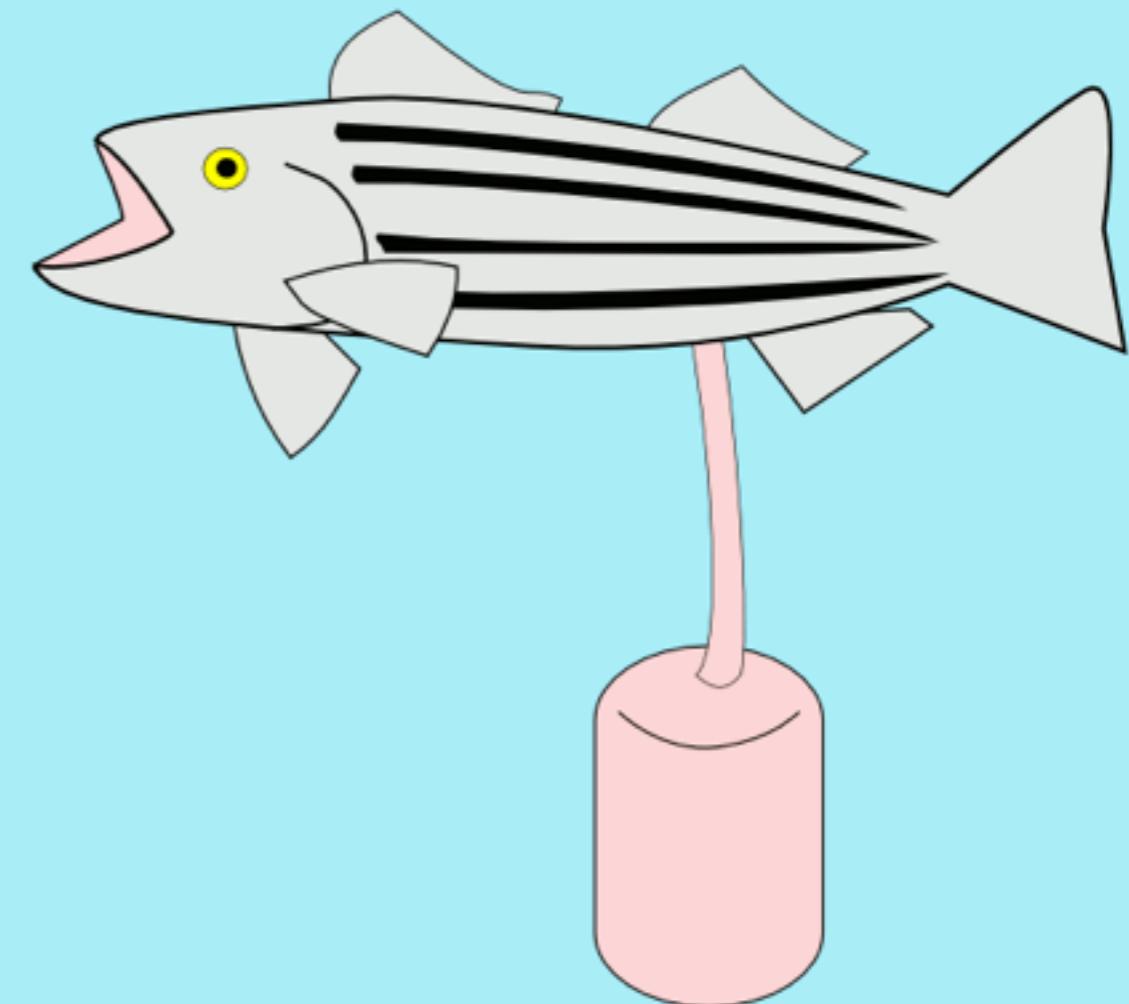


# Add “Shelly”

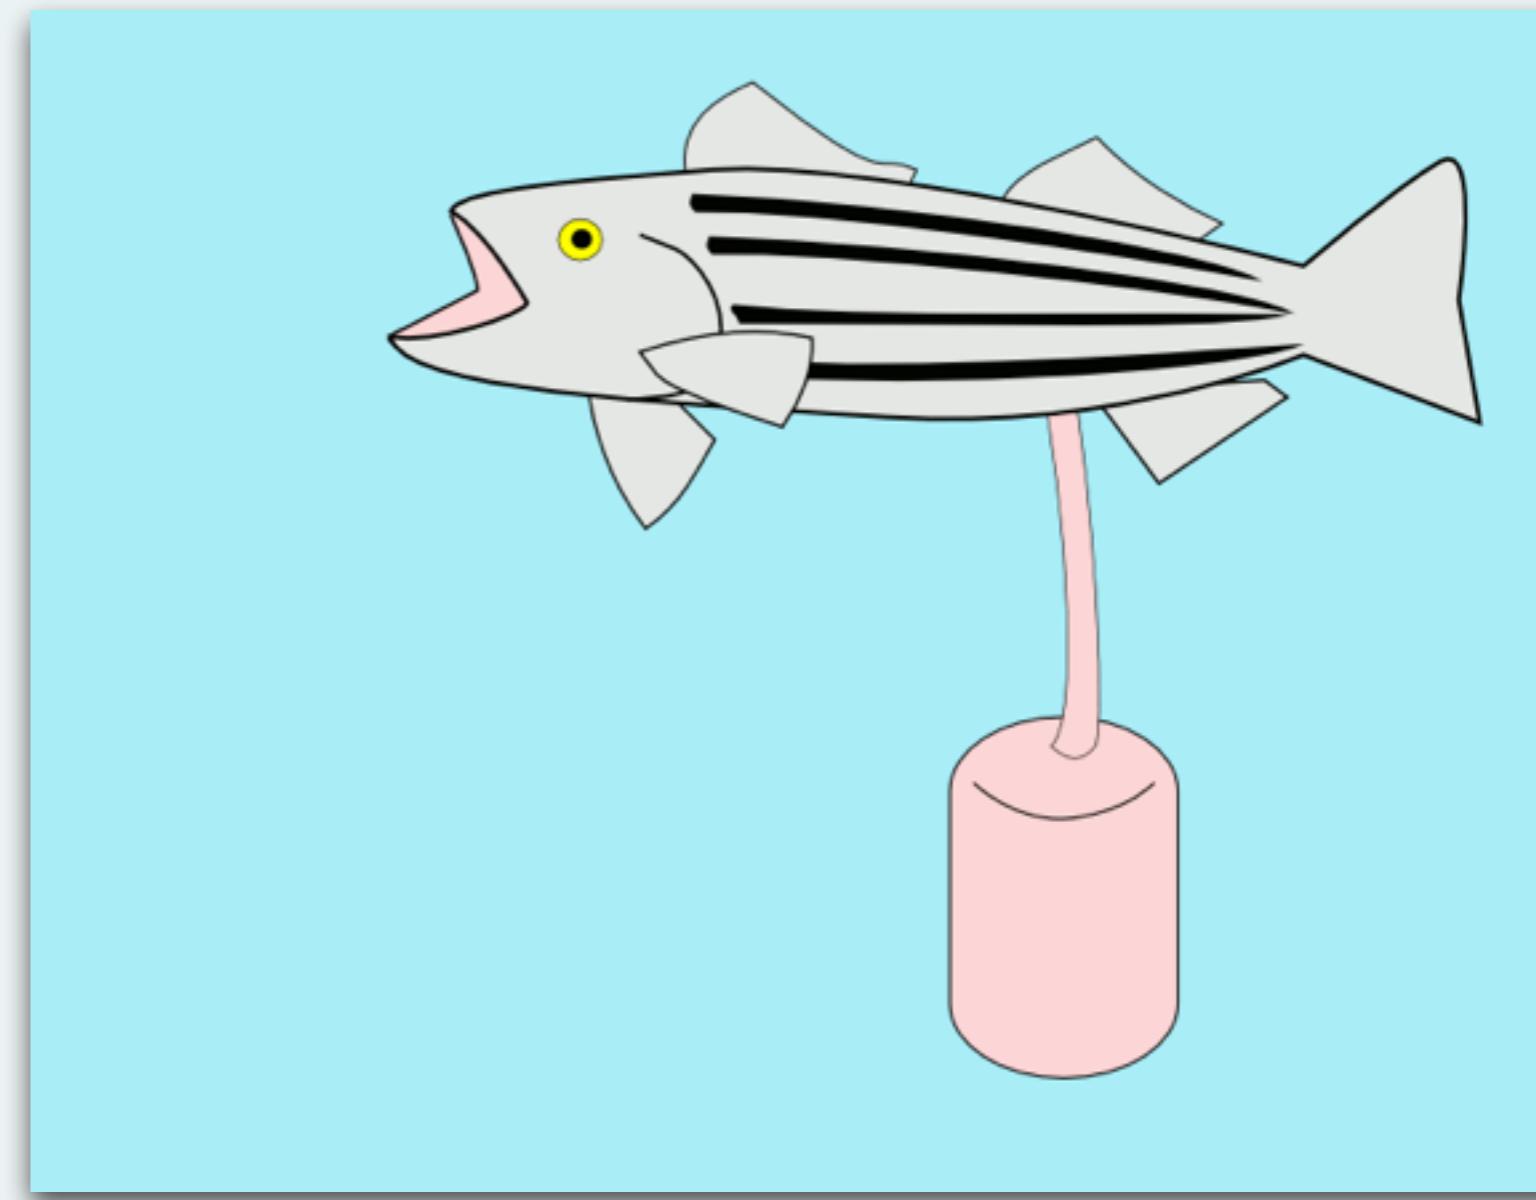
# Client X



ZUCK's FOLLOWERS



## ZUCK's FOLLOWERS



ZUCK's FOLLOWERS



```
[..., ..., ..., ..., ..., ..., ..., ..., ..., ...]  
..., ..., ..., ..., ..., ..., ..., ..., ..., ...  
..., ..., ..., ..., ..., ..., ..., ..., ..., ...  
..., ..., ..., ..., ..., ..., ..., ..., ..., ...  
..., ..., ...], [..., ..., ..., ..., ..., ..., ...]
```

[{a, 34}, {b,  
1000}...]

ZUCK's FOLLOWERS



```
[..., ..., ..., ..., ..., ..., ..., ..., ..., ..., ...,
..., ..., ..., ..., ..., ..., ..., ..., ..., ..., ...,
..., ..., ..., ..., ..., ..., ..., ..., ..., ..., ...,
..., ..., ..., ..., ..., ..., ..., ..., ..., ..., ...,
..., ..., ..., ..., ..., ..., ..., ..., ..., ..., ...]
```

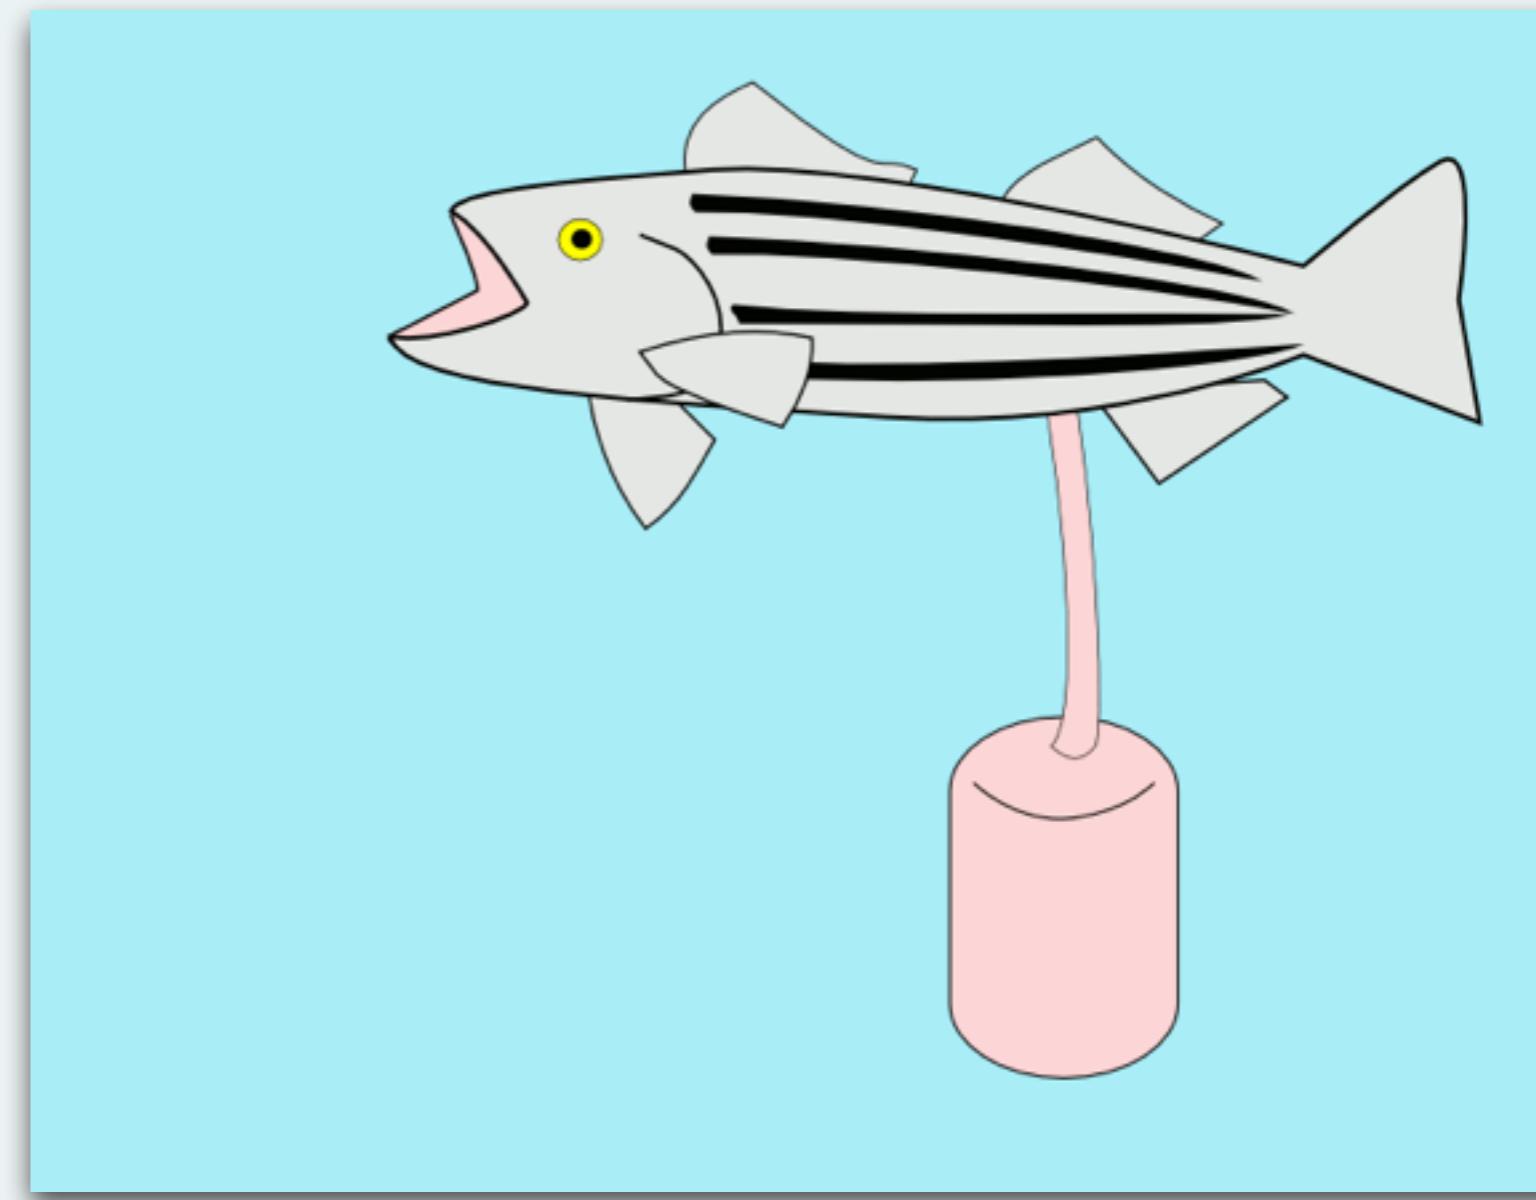
[{a, 35}, {b,  
1000}...]

ZUCK's FOLLOWERS

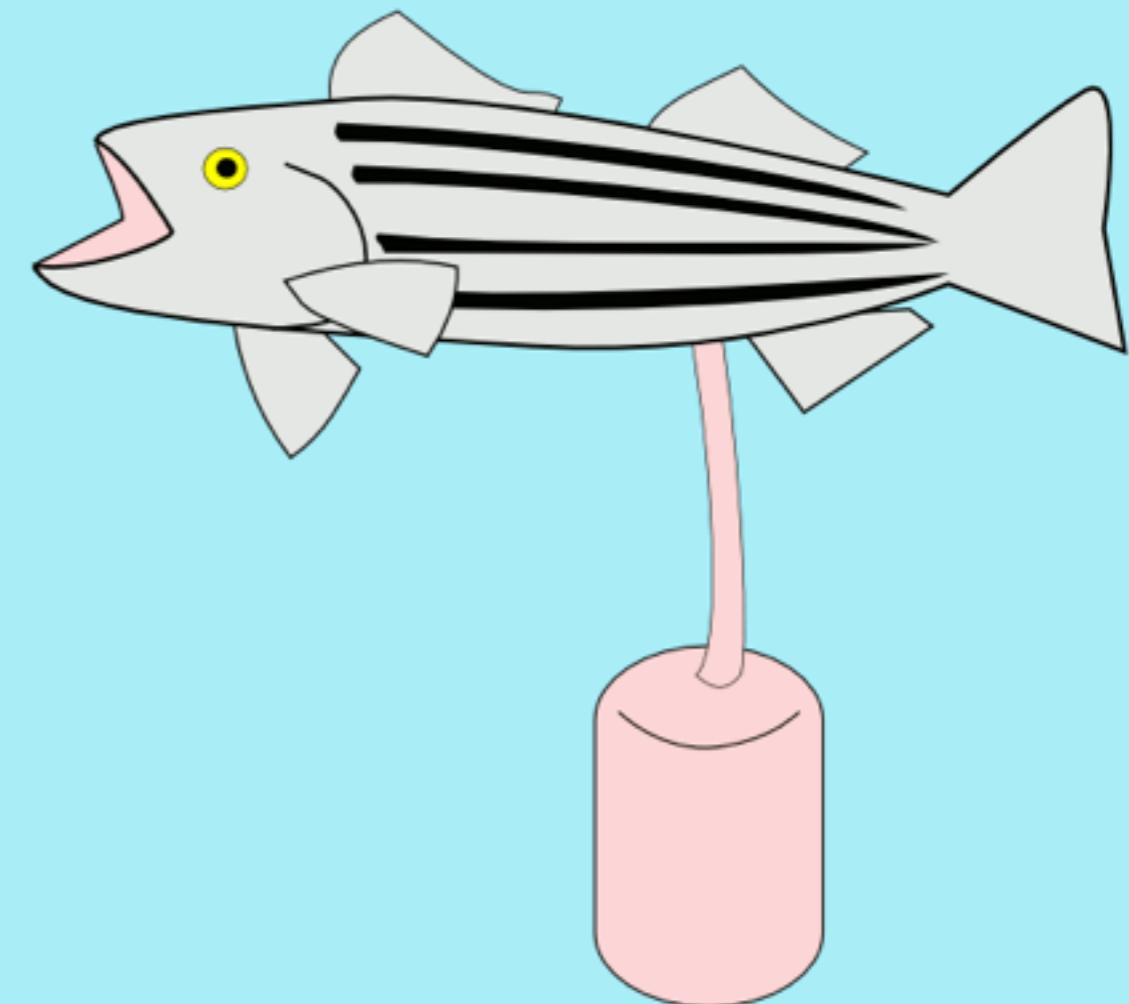
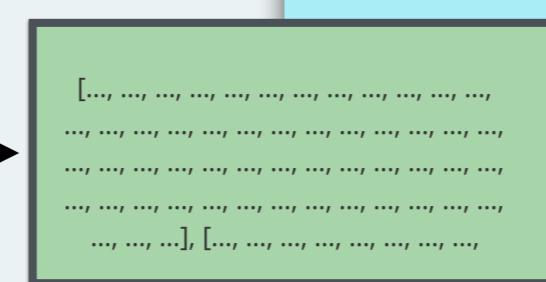


[{a35}shelly, ..., ..., ..., ..., ..., ..., ...,  
..., ..., ..., ..., ..., ..., ..., ..., ..., ..., ...,  
..., ..., ..., ..., ..., ..., ..., ..., ..., ..., ...,  
..., ..., ..., ..., ..., ..., ..., ..., ..., ..., ...,

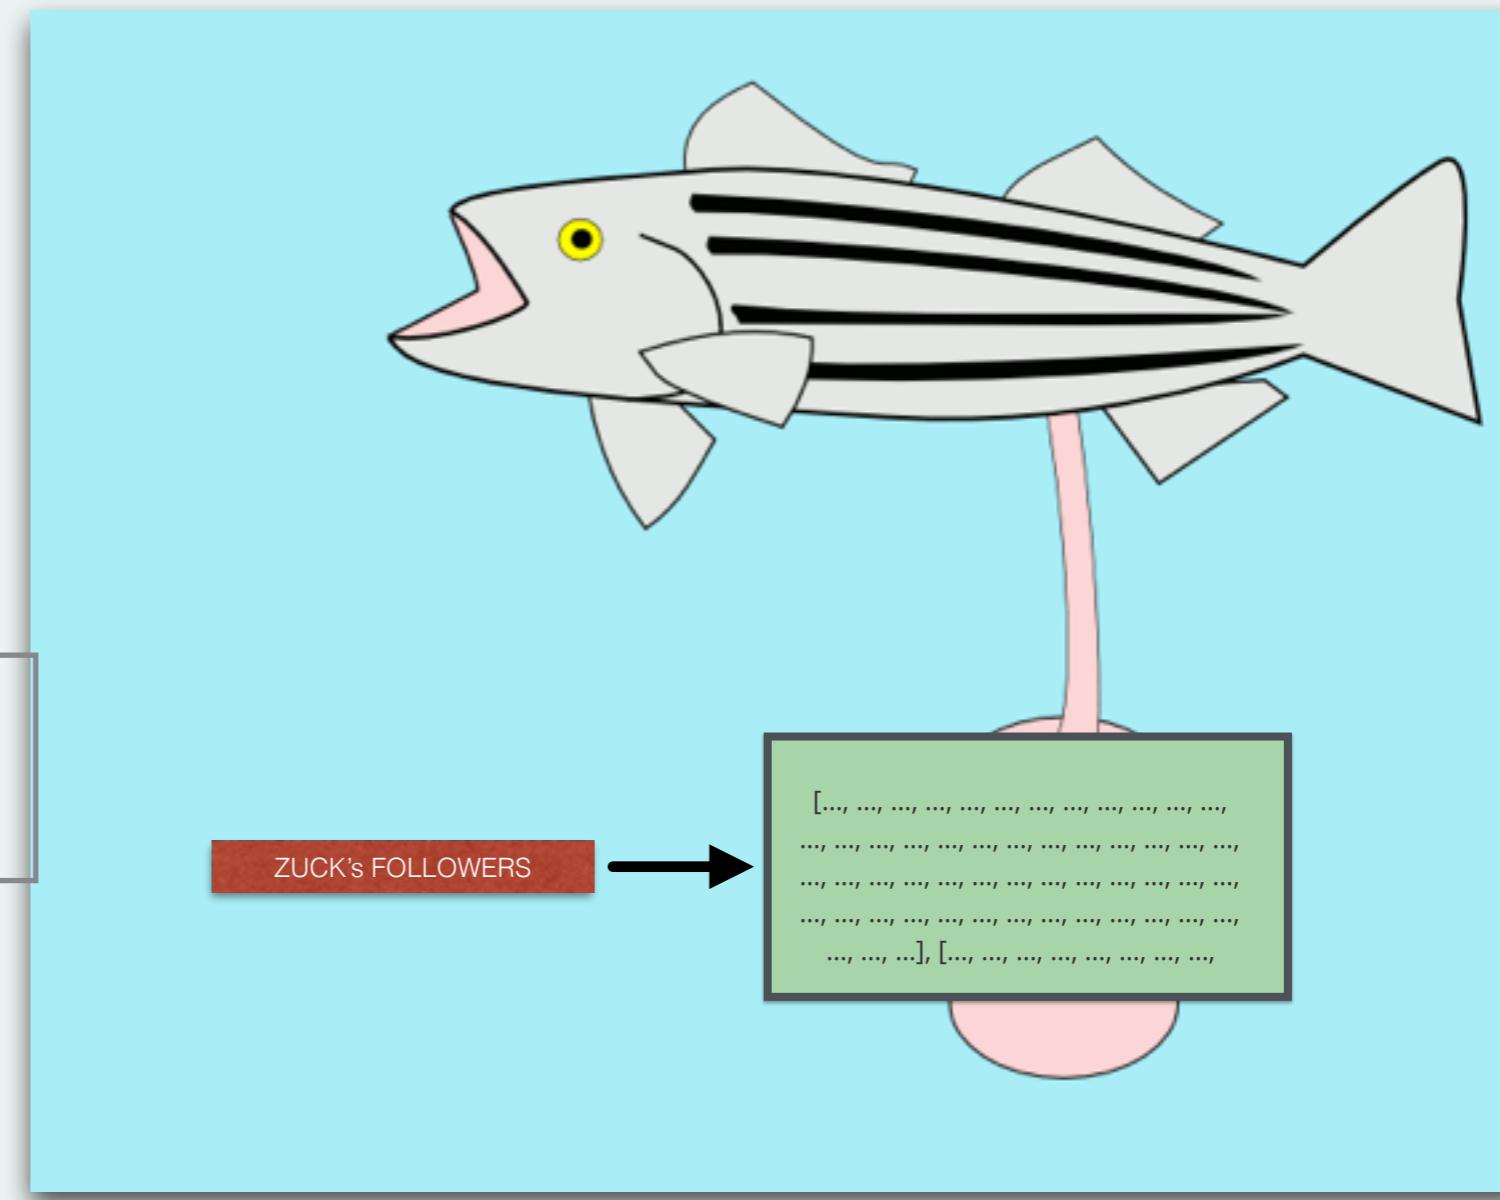
## ZUCK's FOLLOWERS

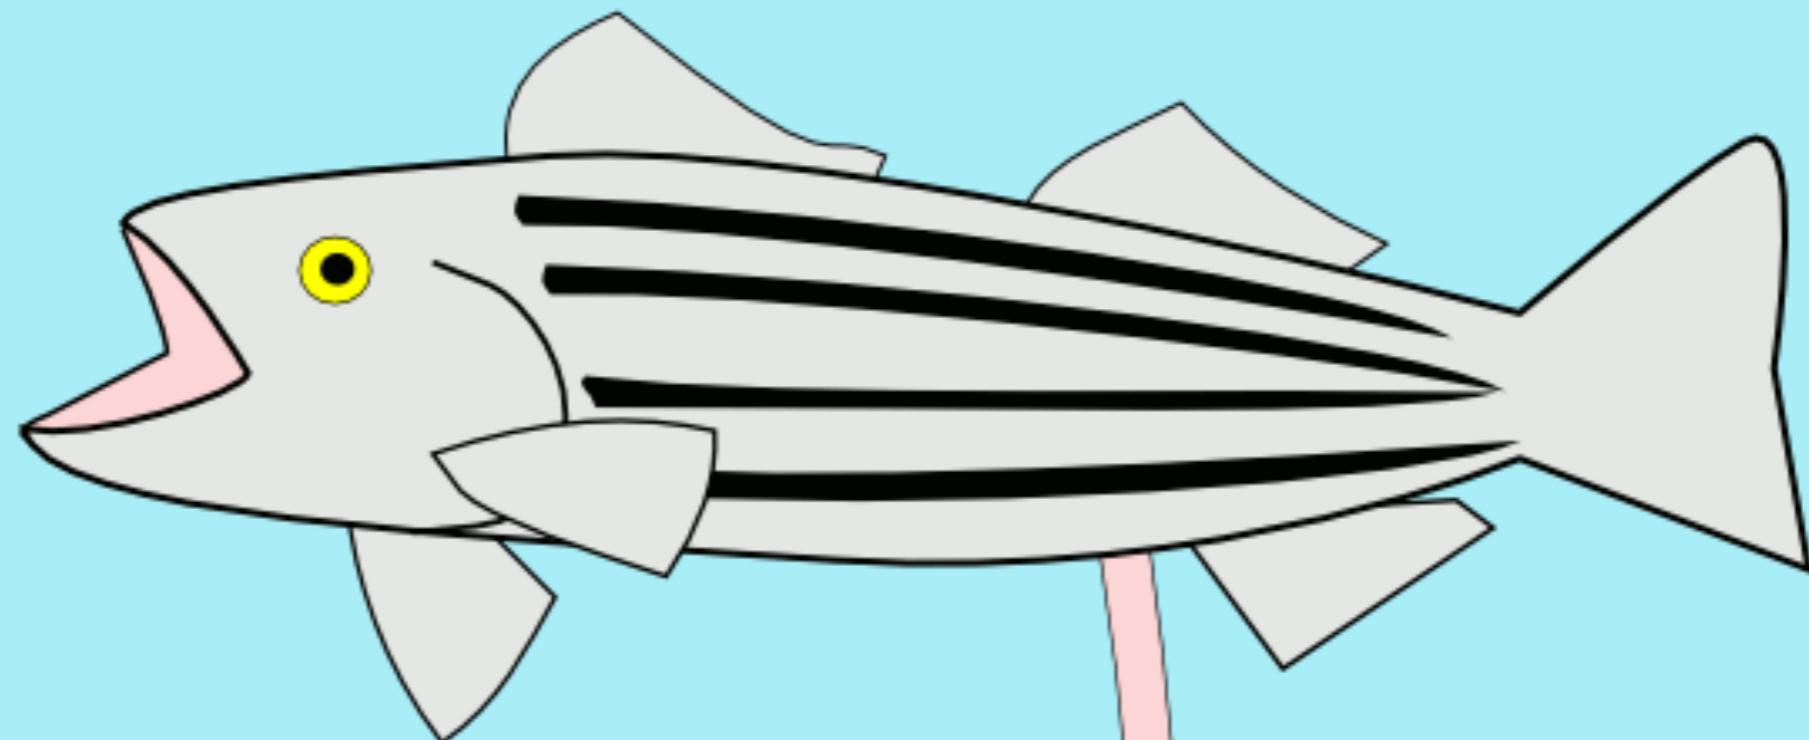


ZUCK's FOLLOWERS

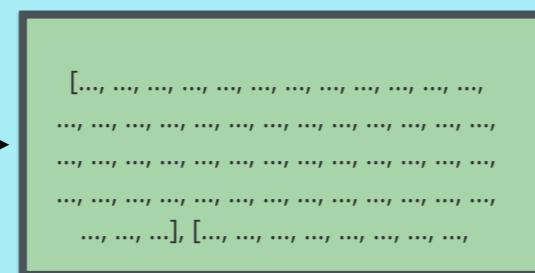


# REPLICATE

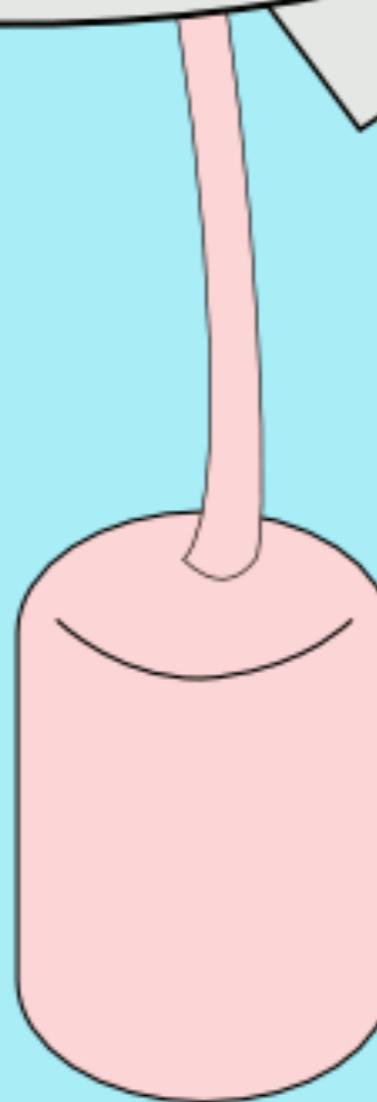




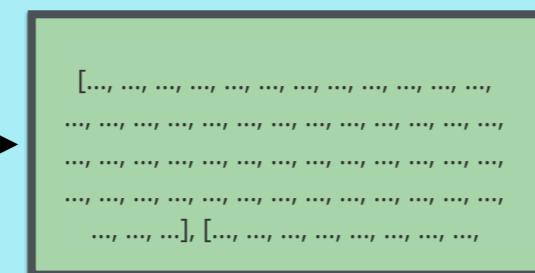
ZUCK's FOLLOWERS

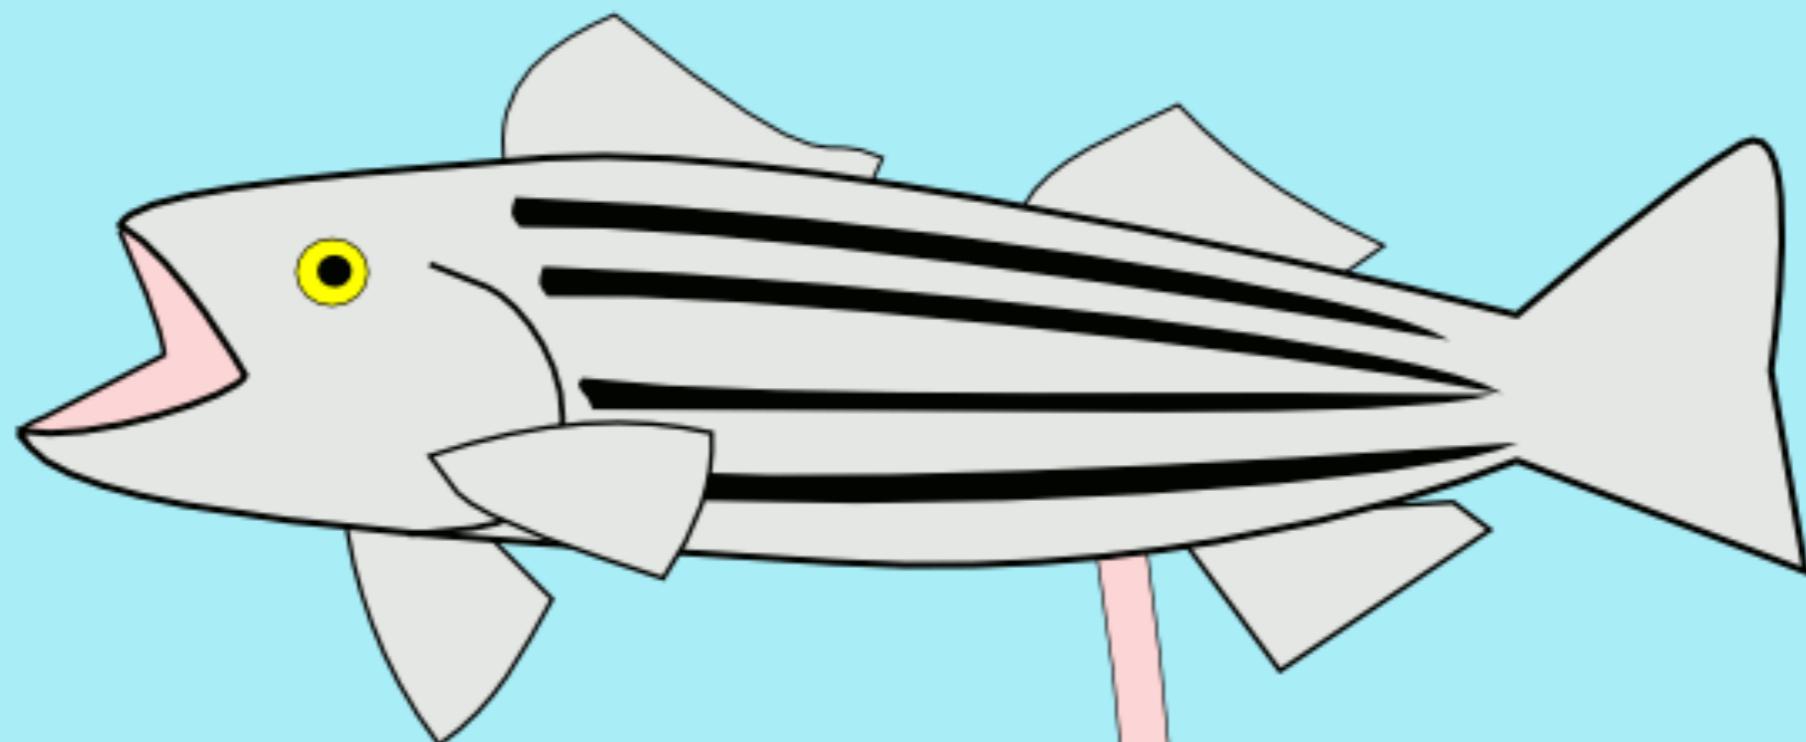


U



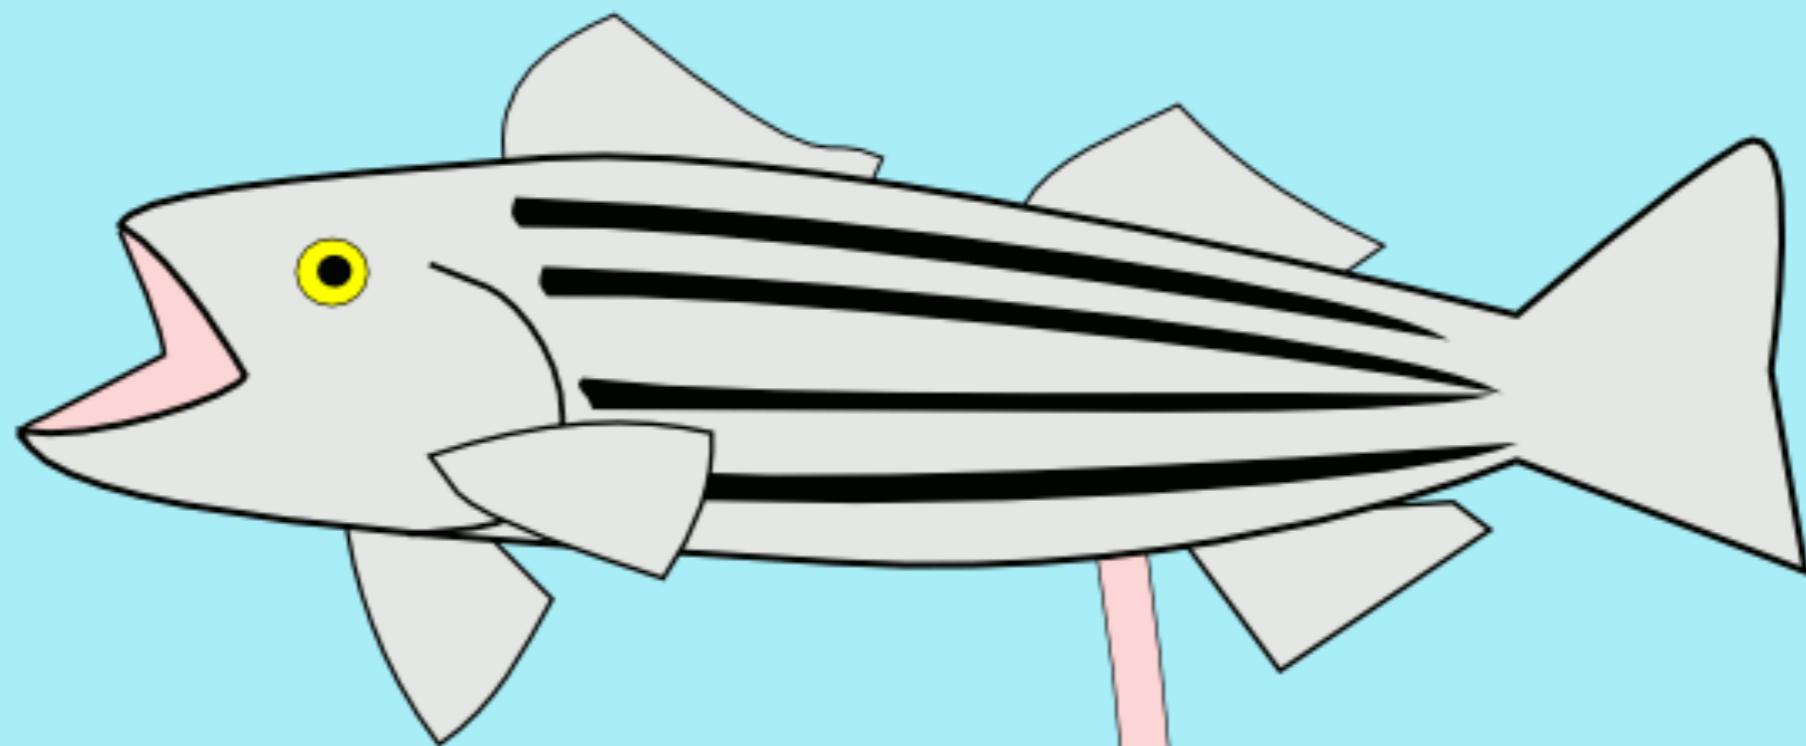
ZUCK's FOLLOWERS



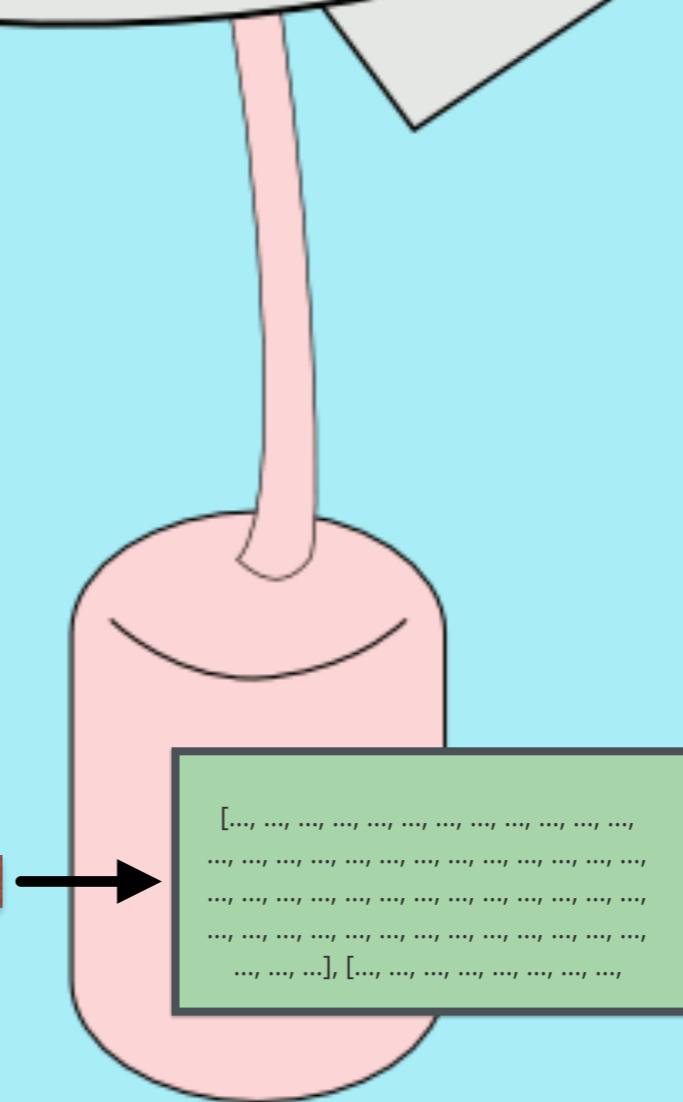


ZUCK's FOLLOWERS

[..., ..., ..., ..., ..., ..., ..., ..., ..., ...  
..., ..., ..., ..., ..., ..., ..., ..., ..., ...  
..., ..., ..., ..., ..., ..., ..., ..., ..., ...  
..., ..., ..., ..., ..., ..., ..., ..., ..., ...  
..., ..., ...], [..., ..., ..., ..., ..., ...]



ZUCK's FOLLOWERS

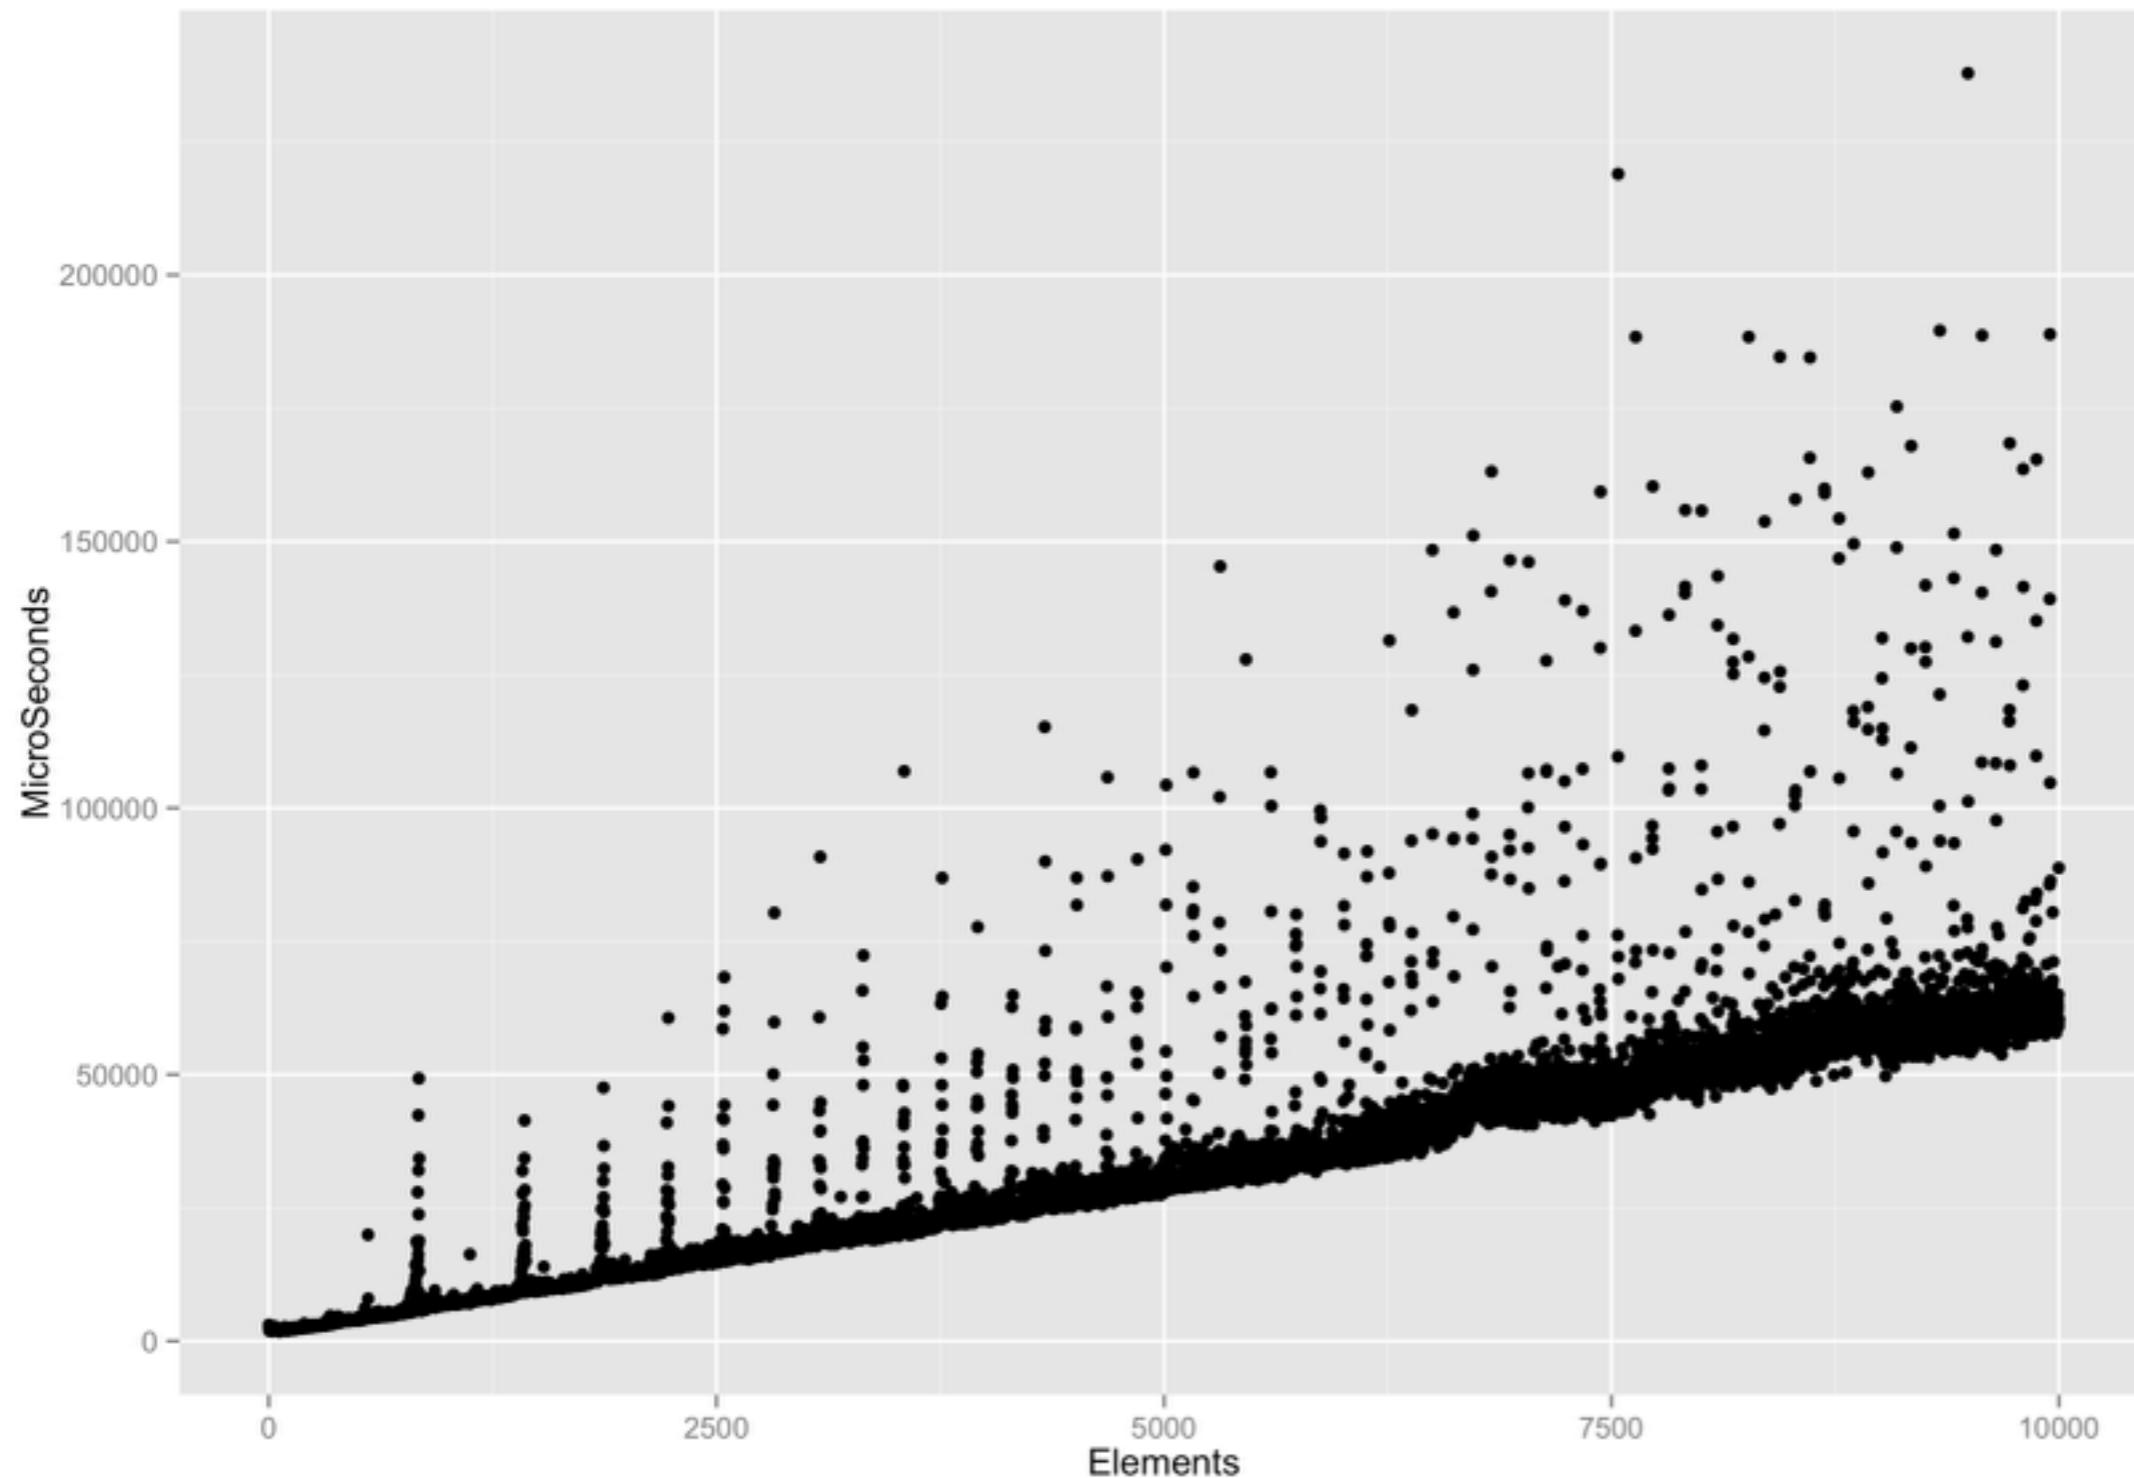


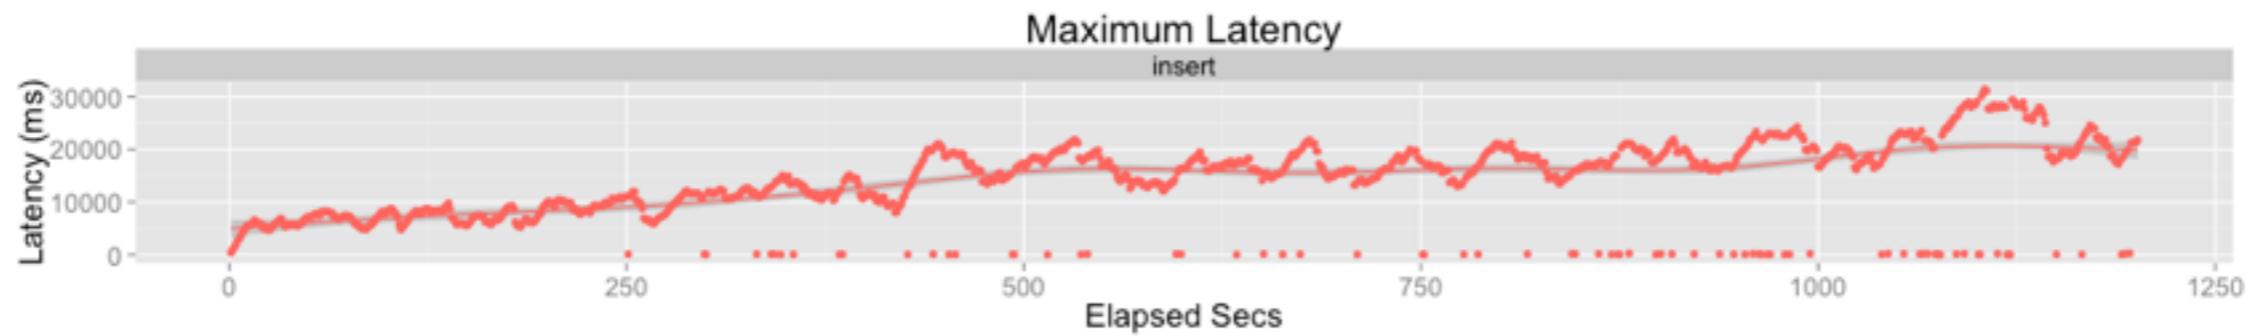
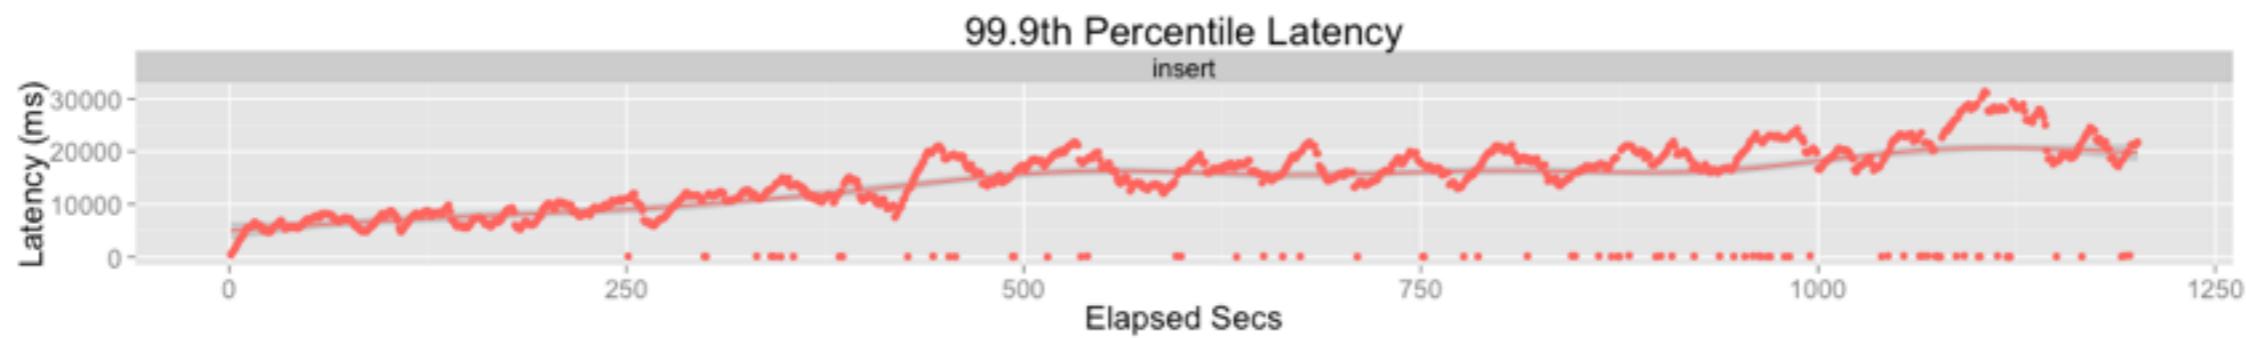
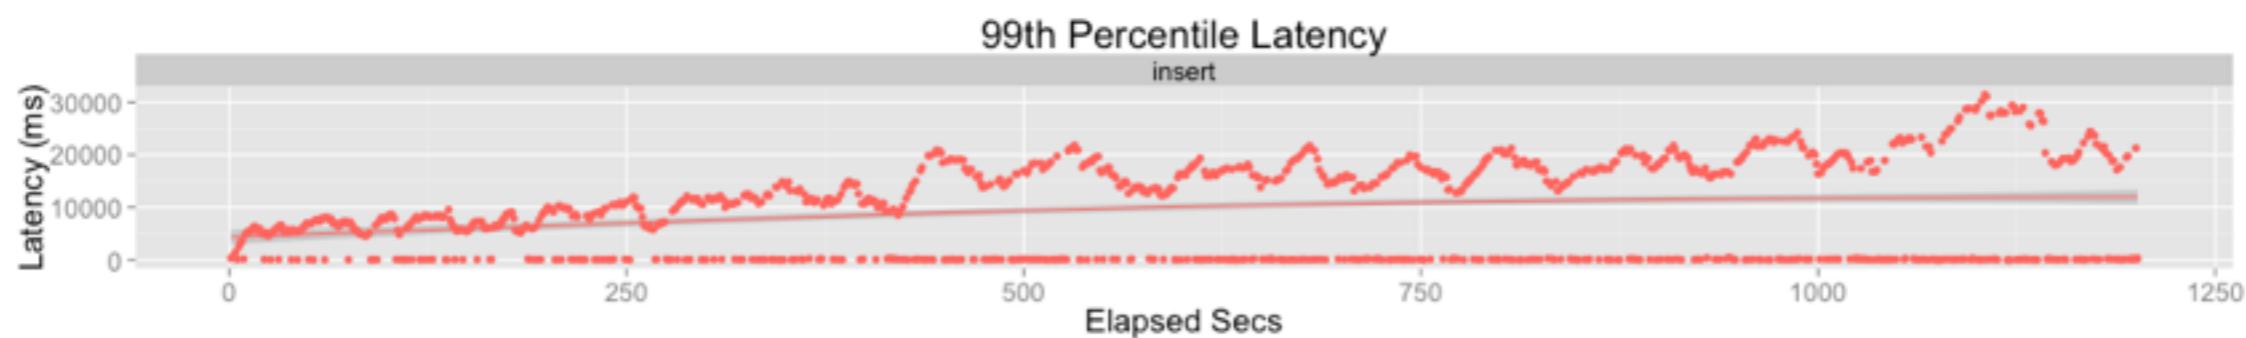
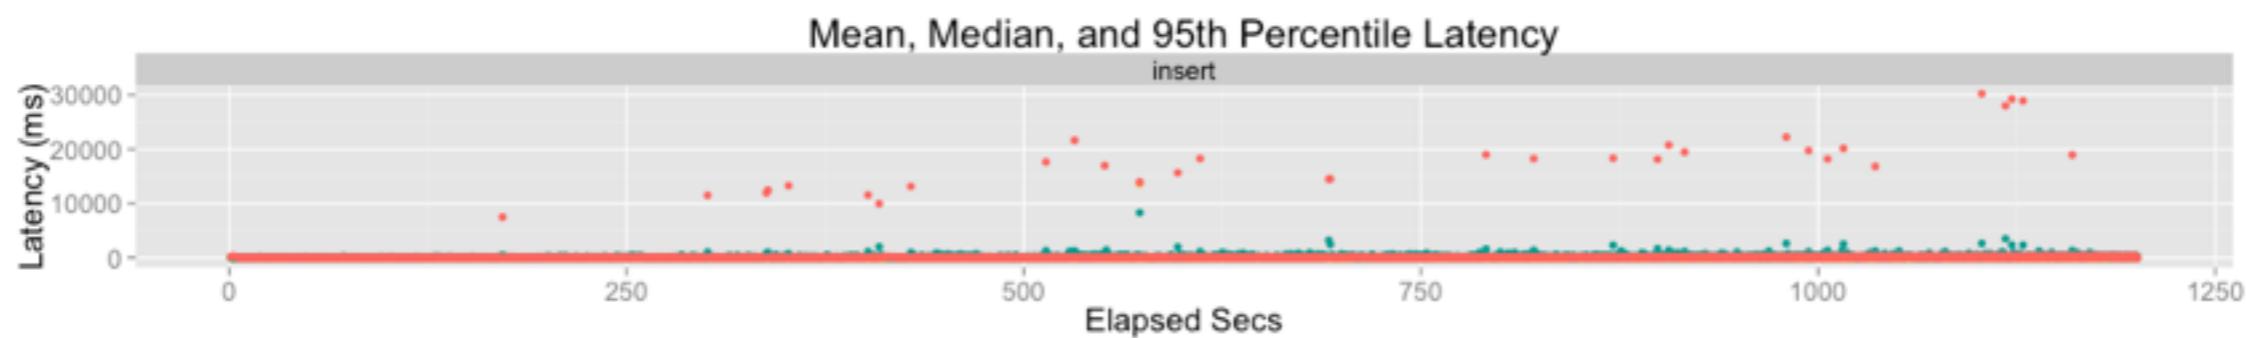
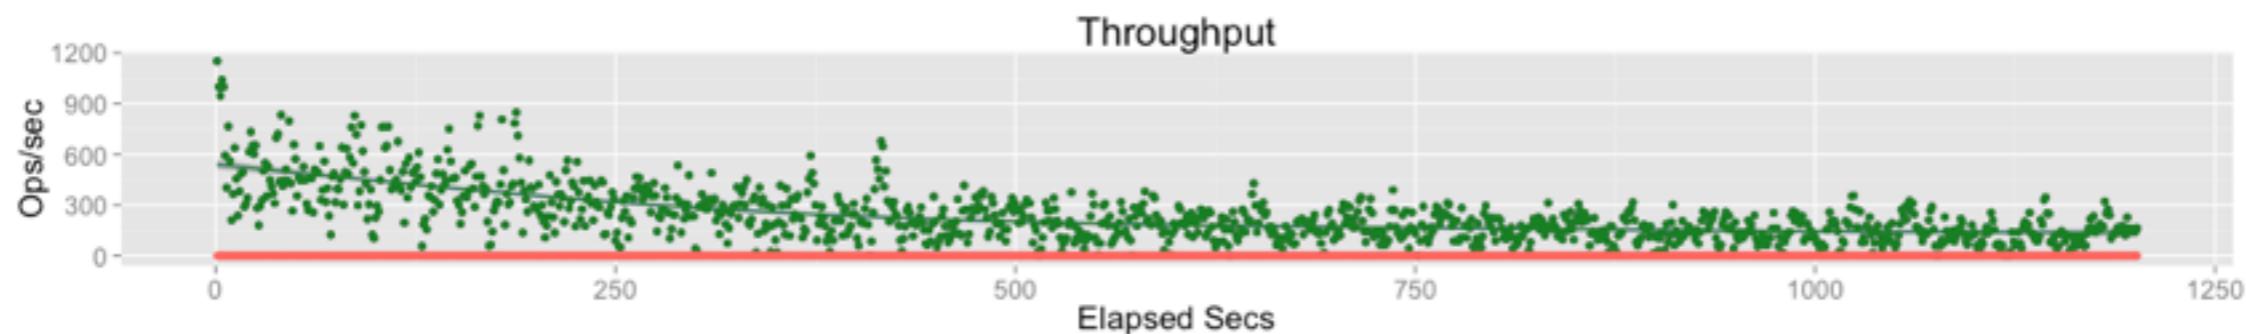
# Problem?

- 1key -> 1 Set
- Poor Write speed
- Can't have “big” sets

Every time we change the  
set we read and write the  
whole set!

# Sets in Riak





10k sets, 100k elements, 50 workers - write

# Sets in Riak

Small : riak object

1MB limit

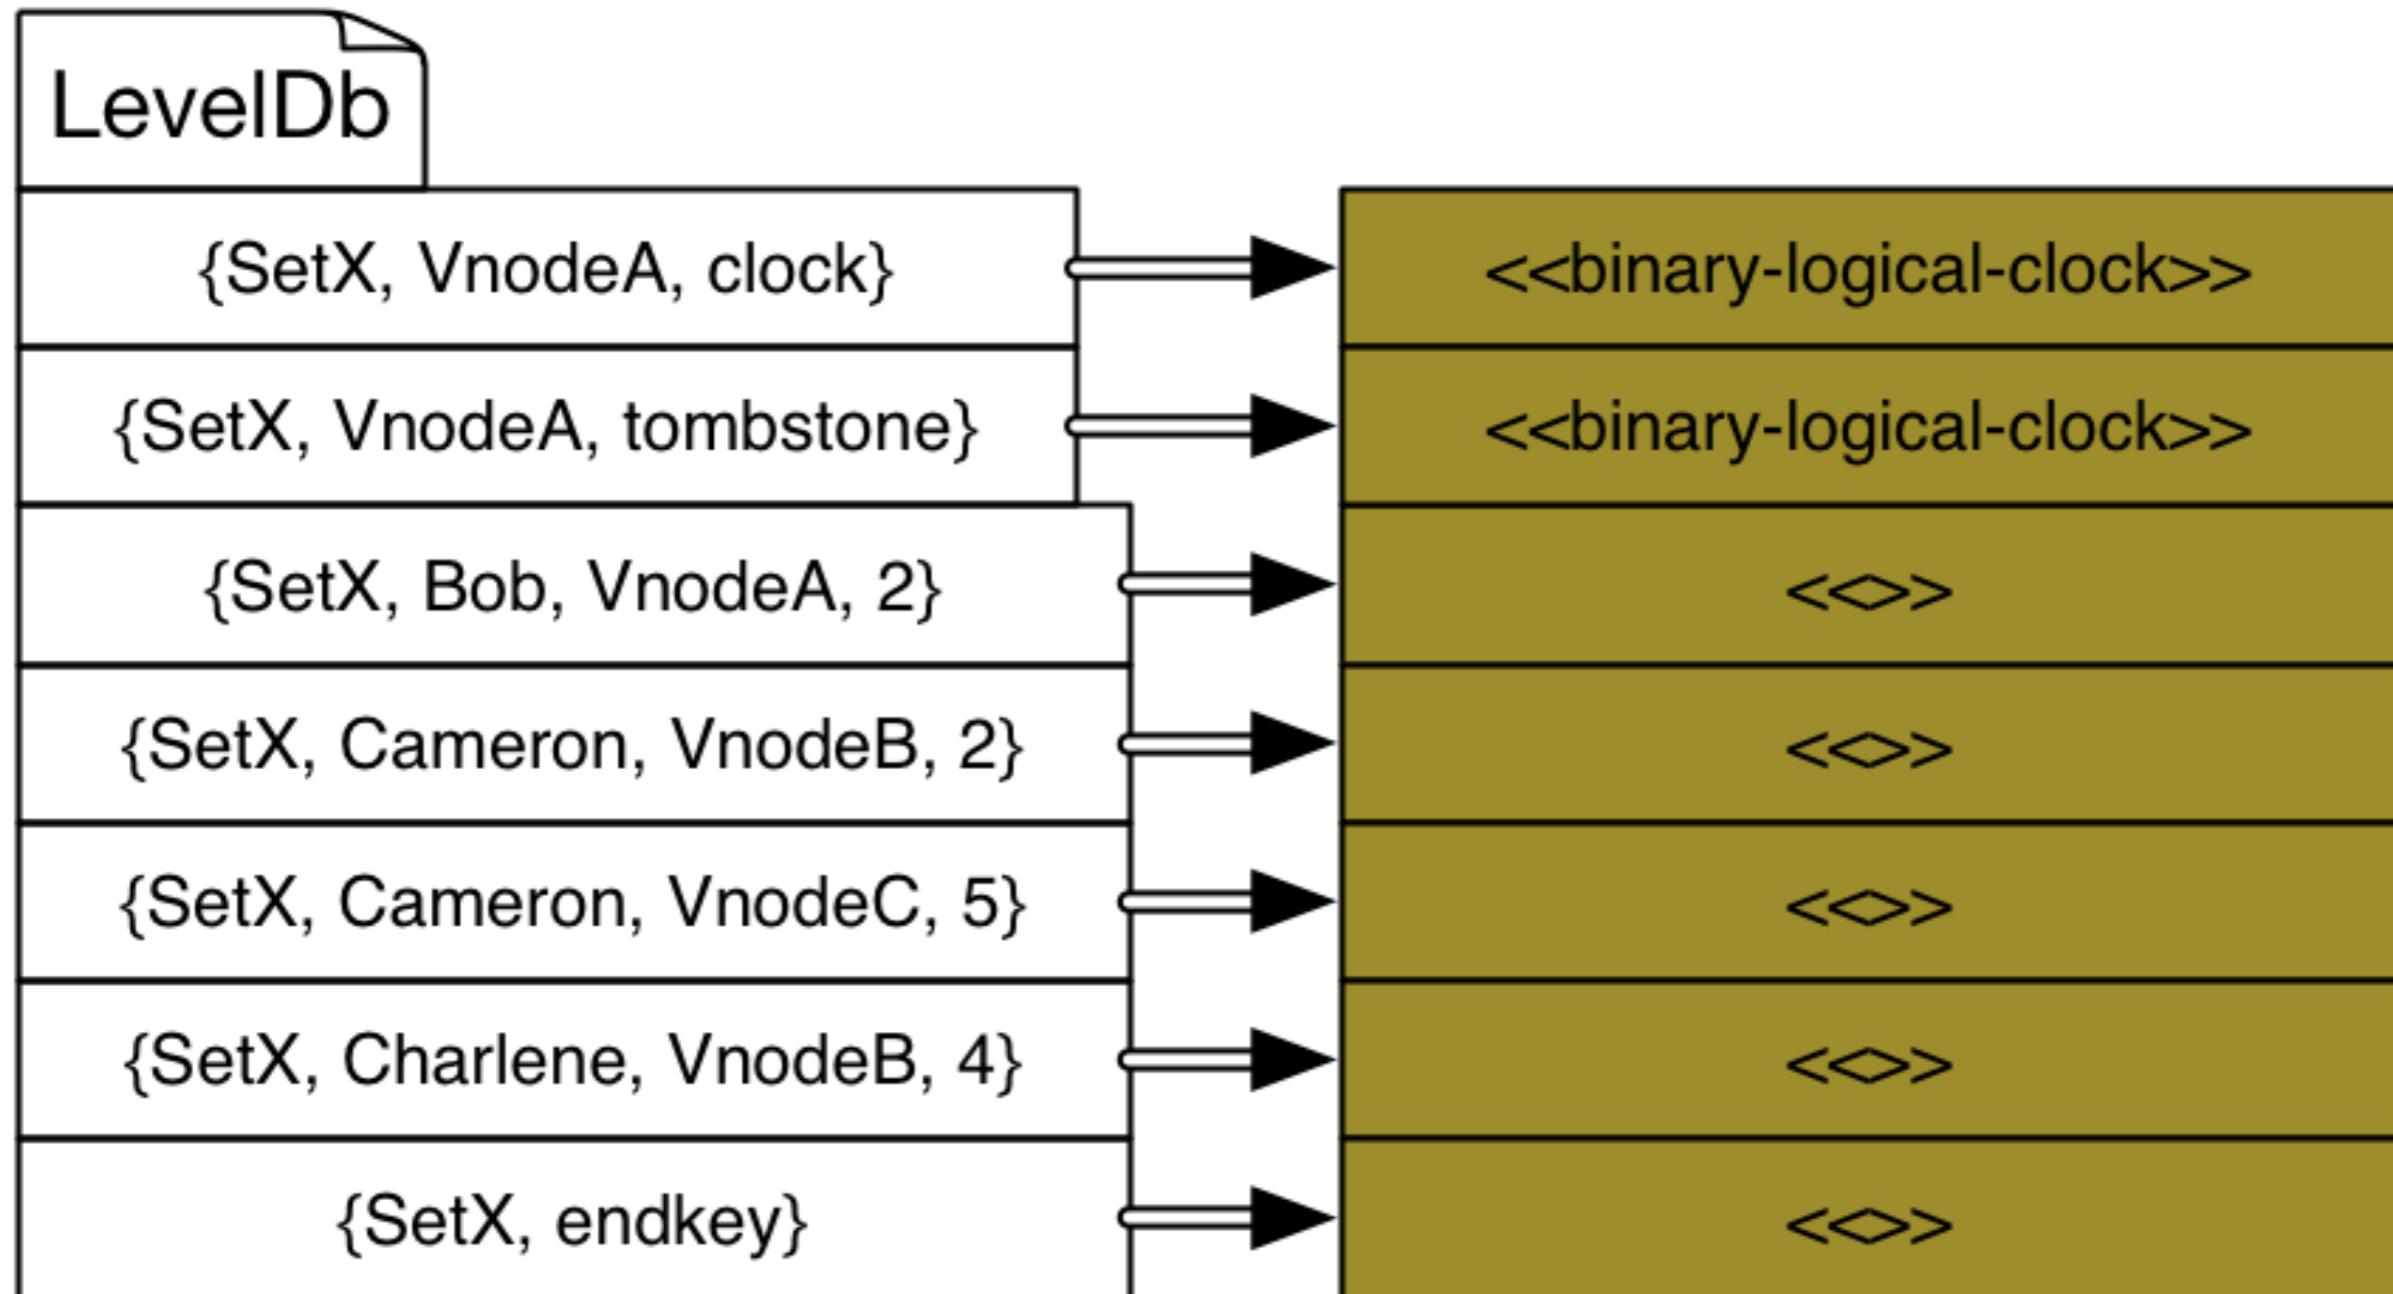
Bigsets:

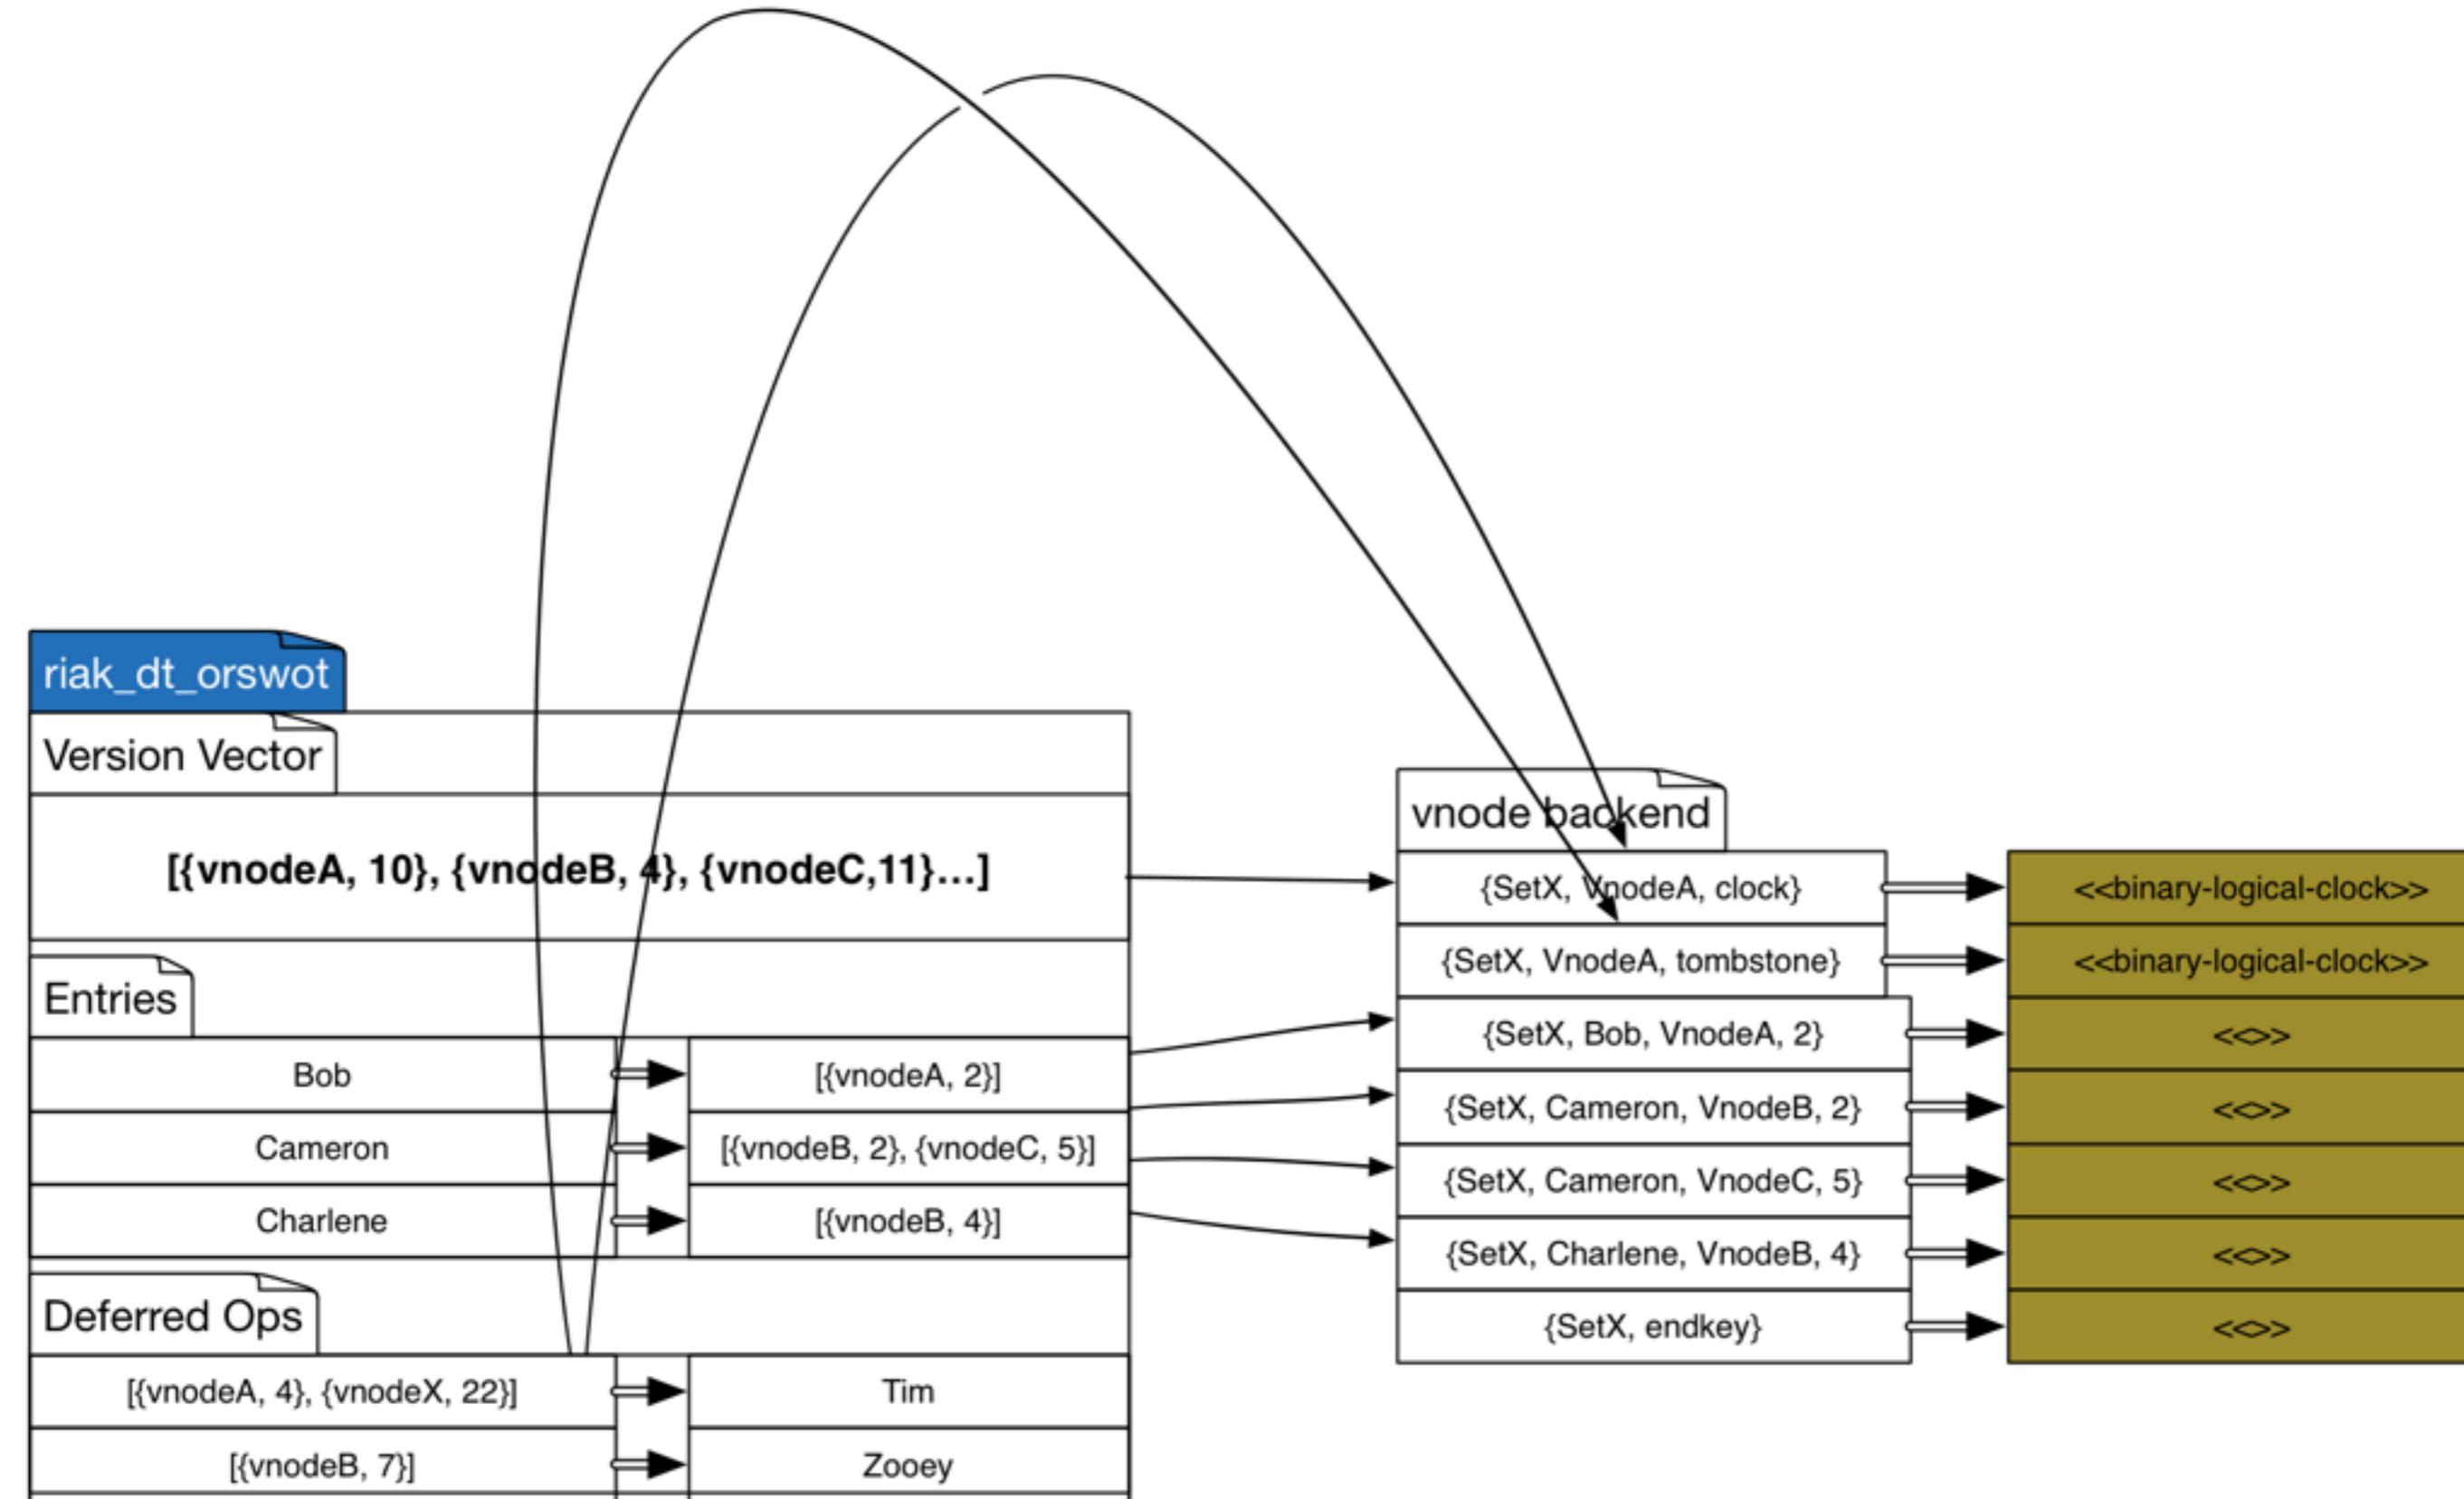
Make writes faster

and

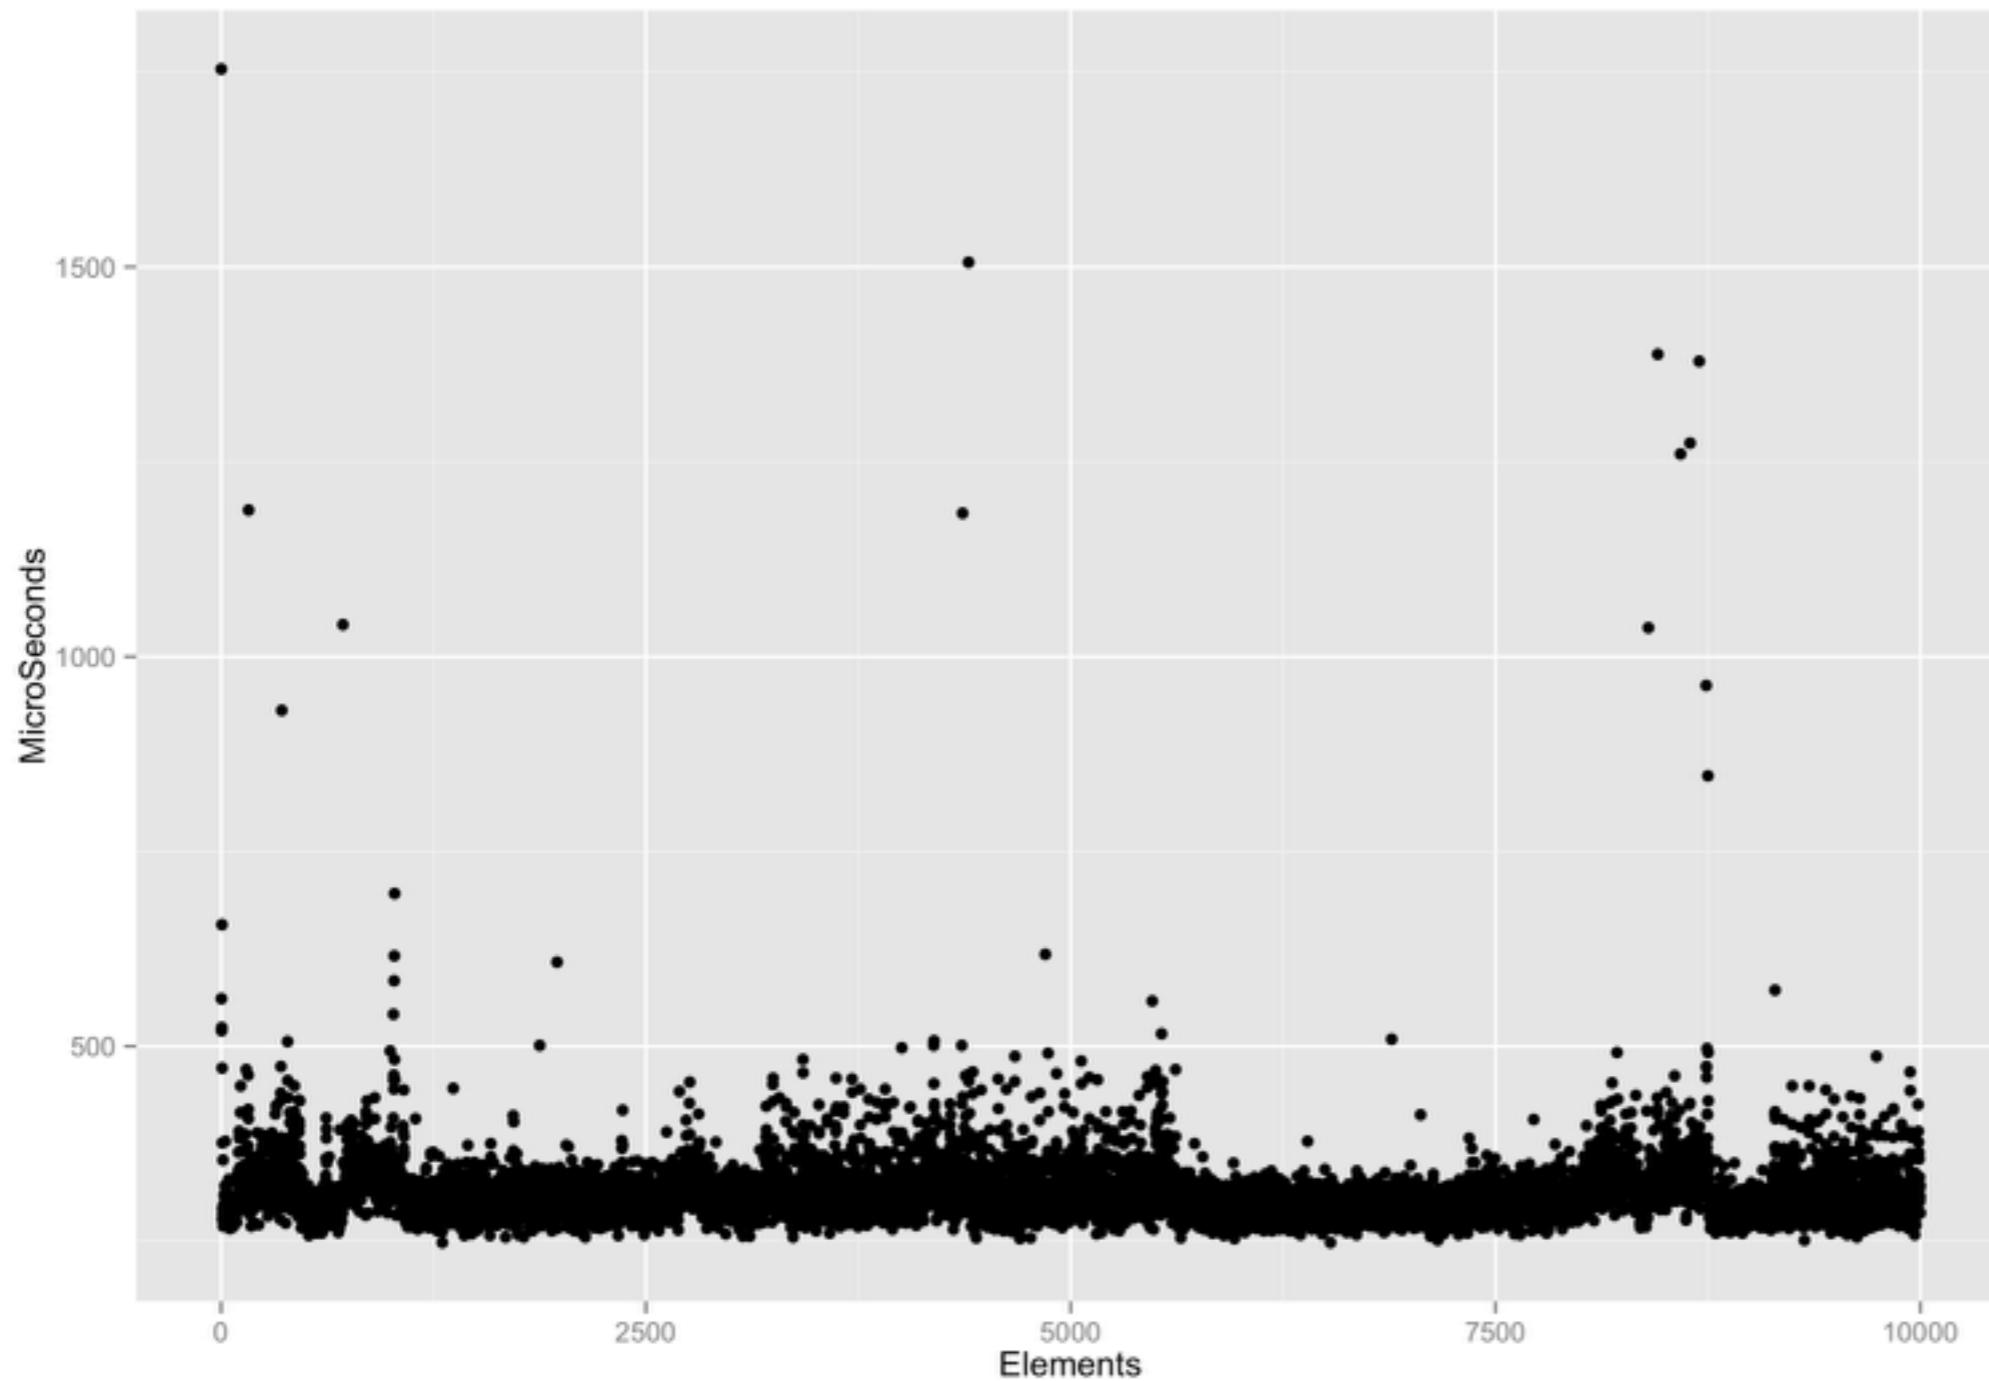
sets bigger

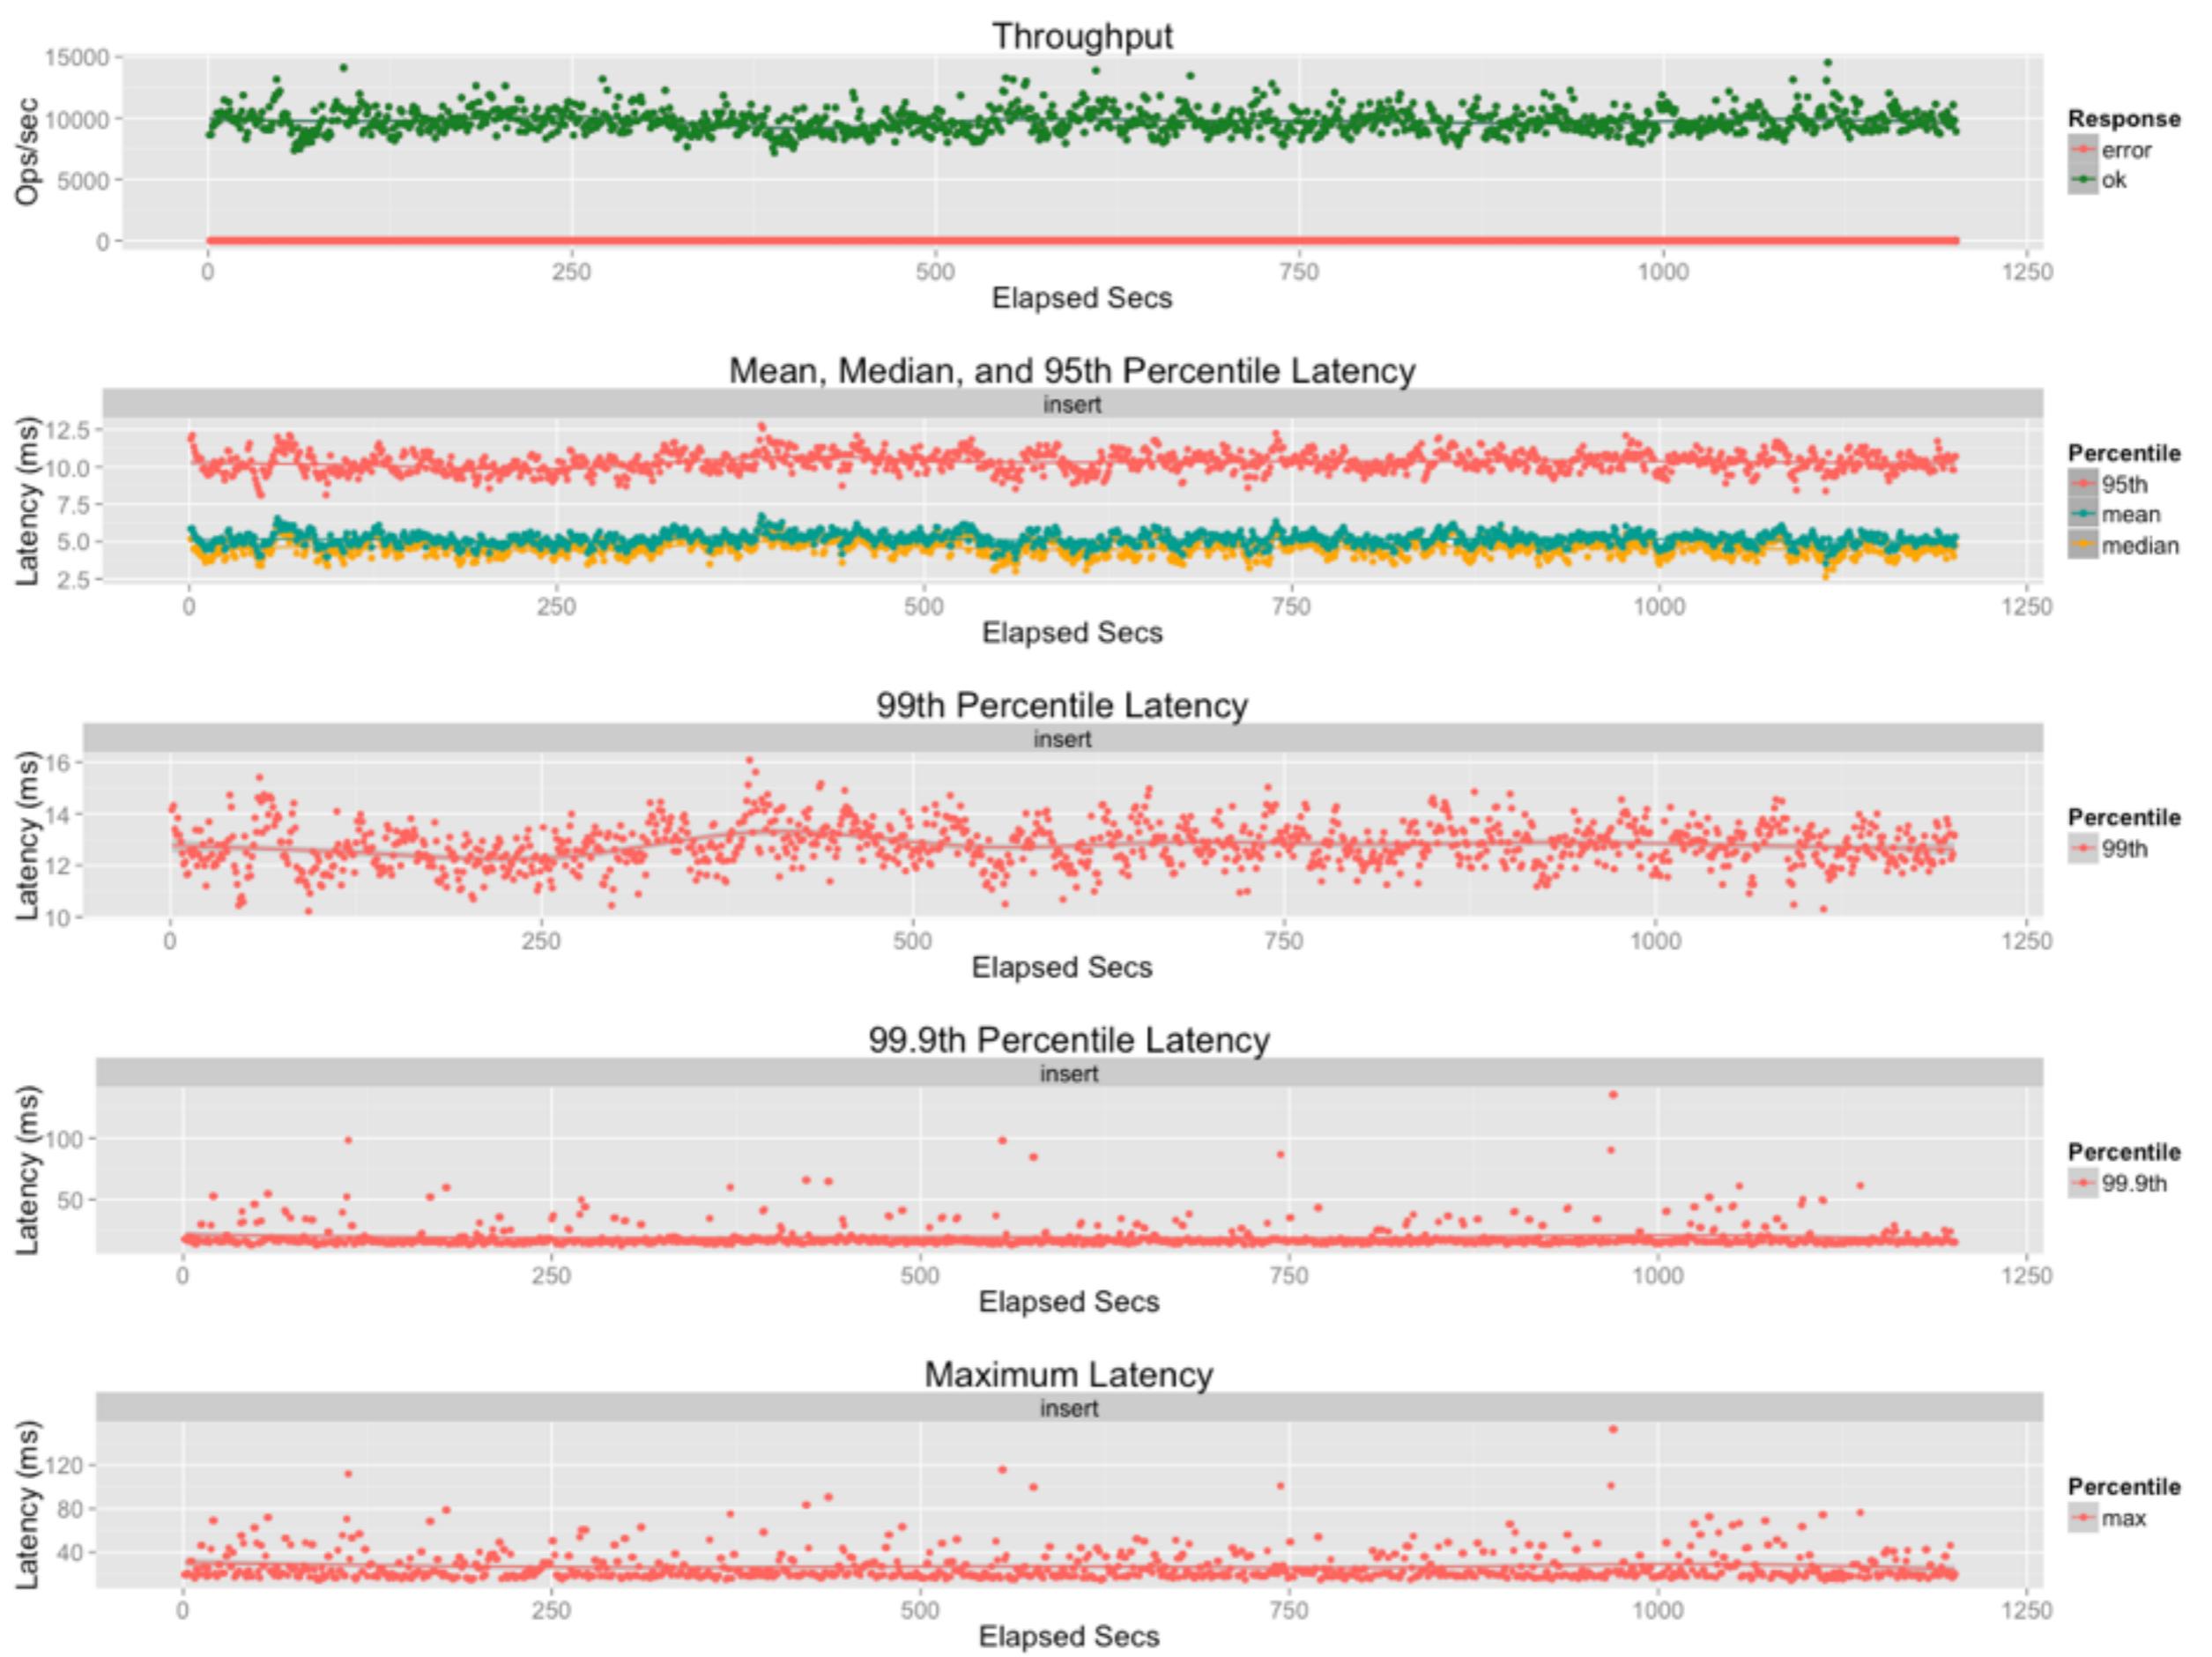
# Bigset Design: Overview



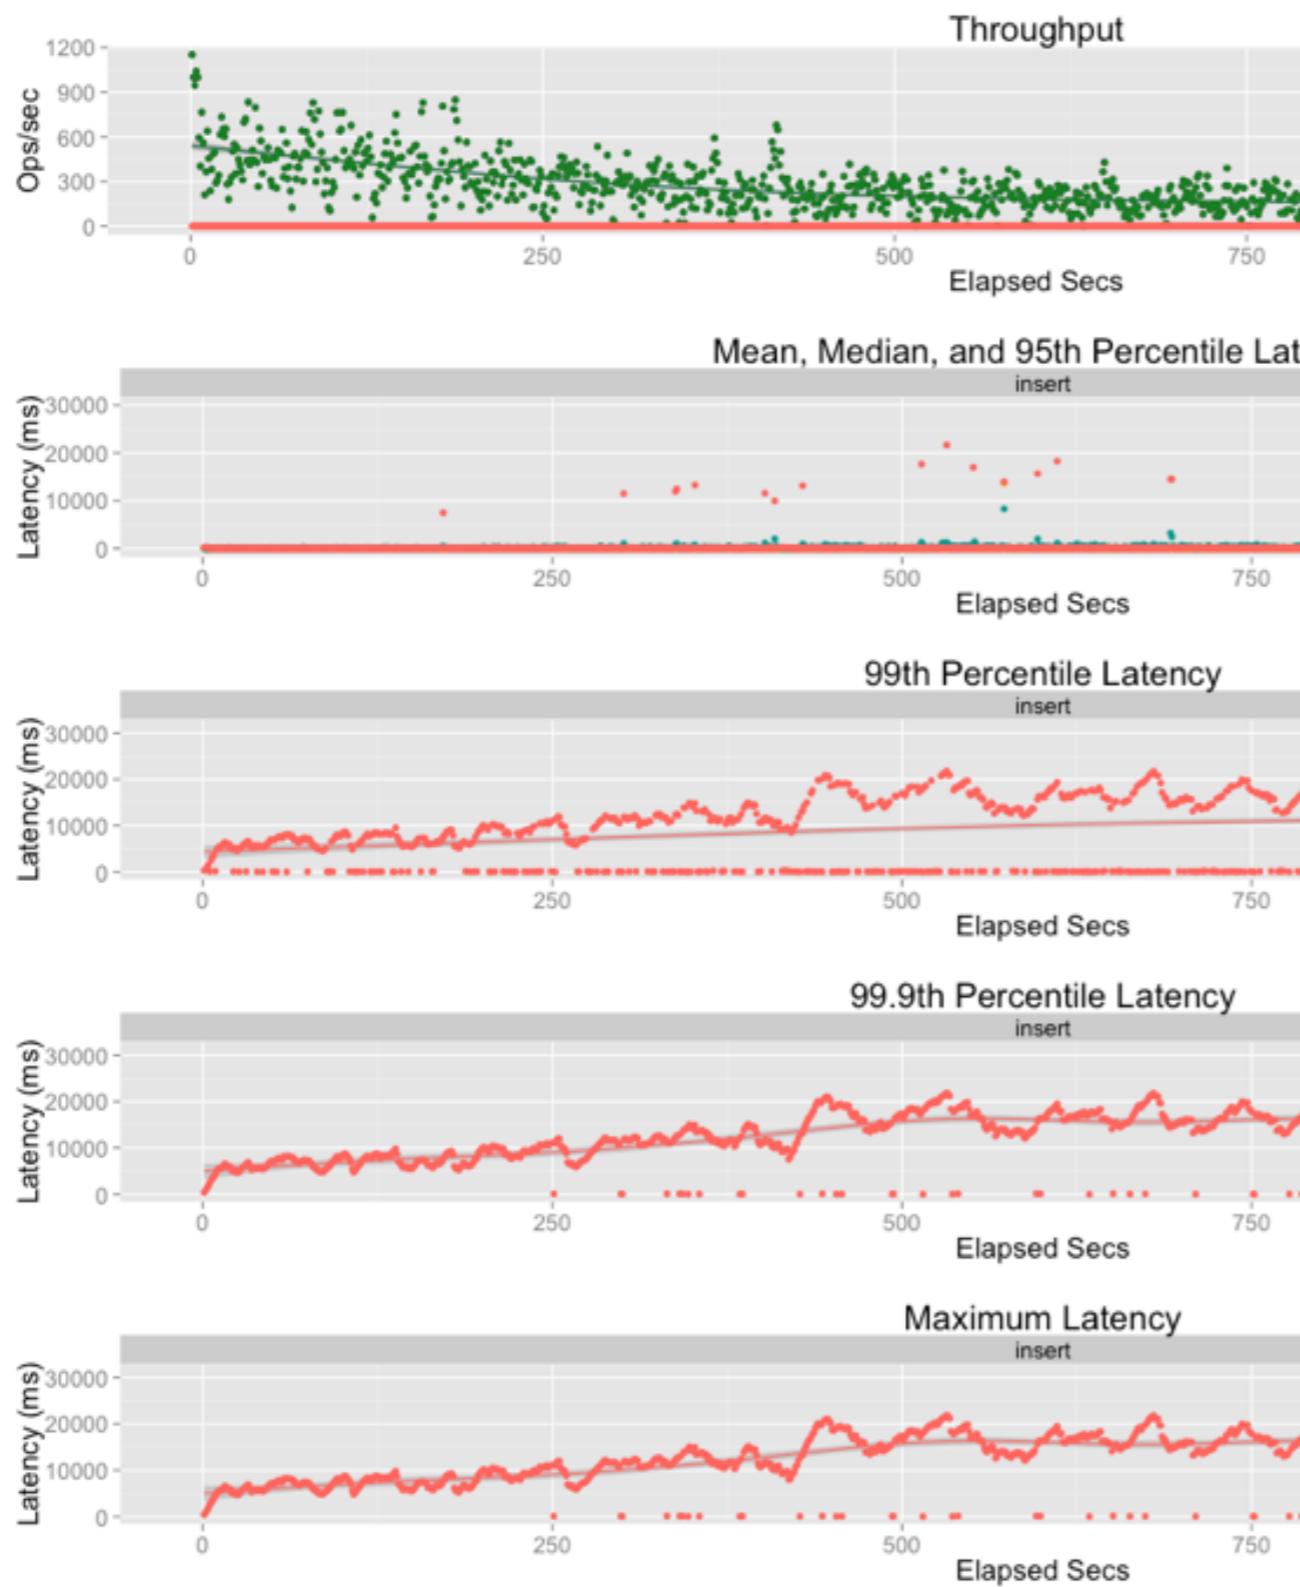


# Initial Results

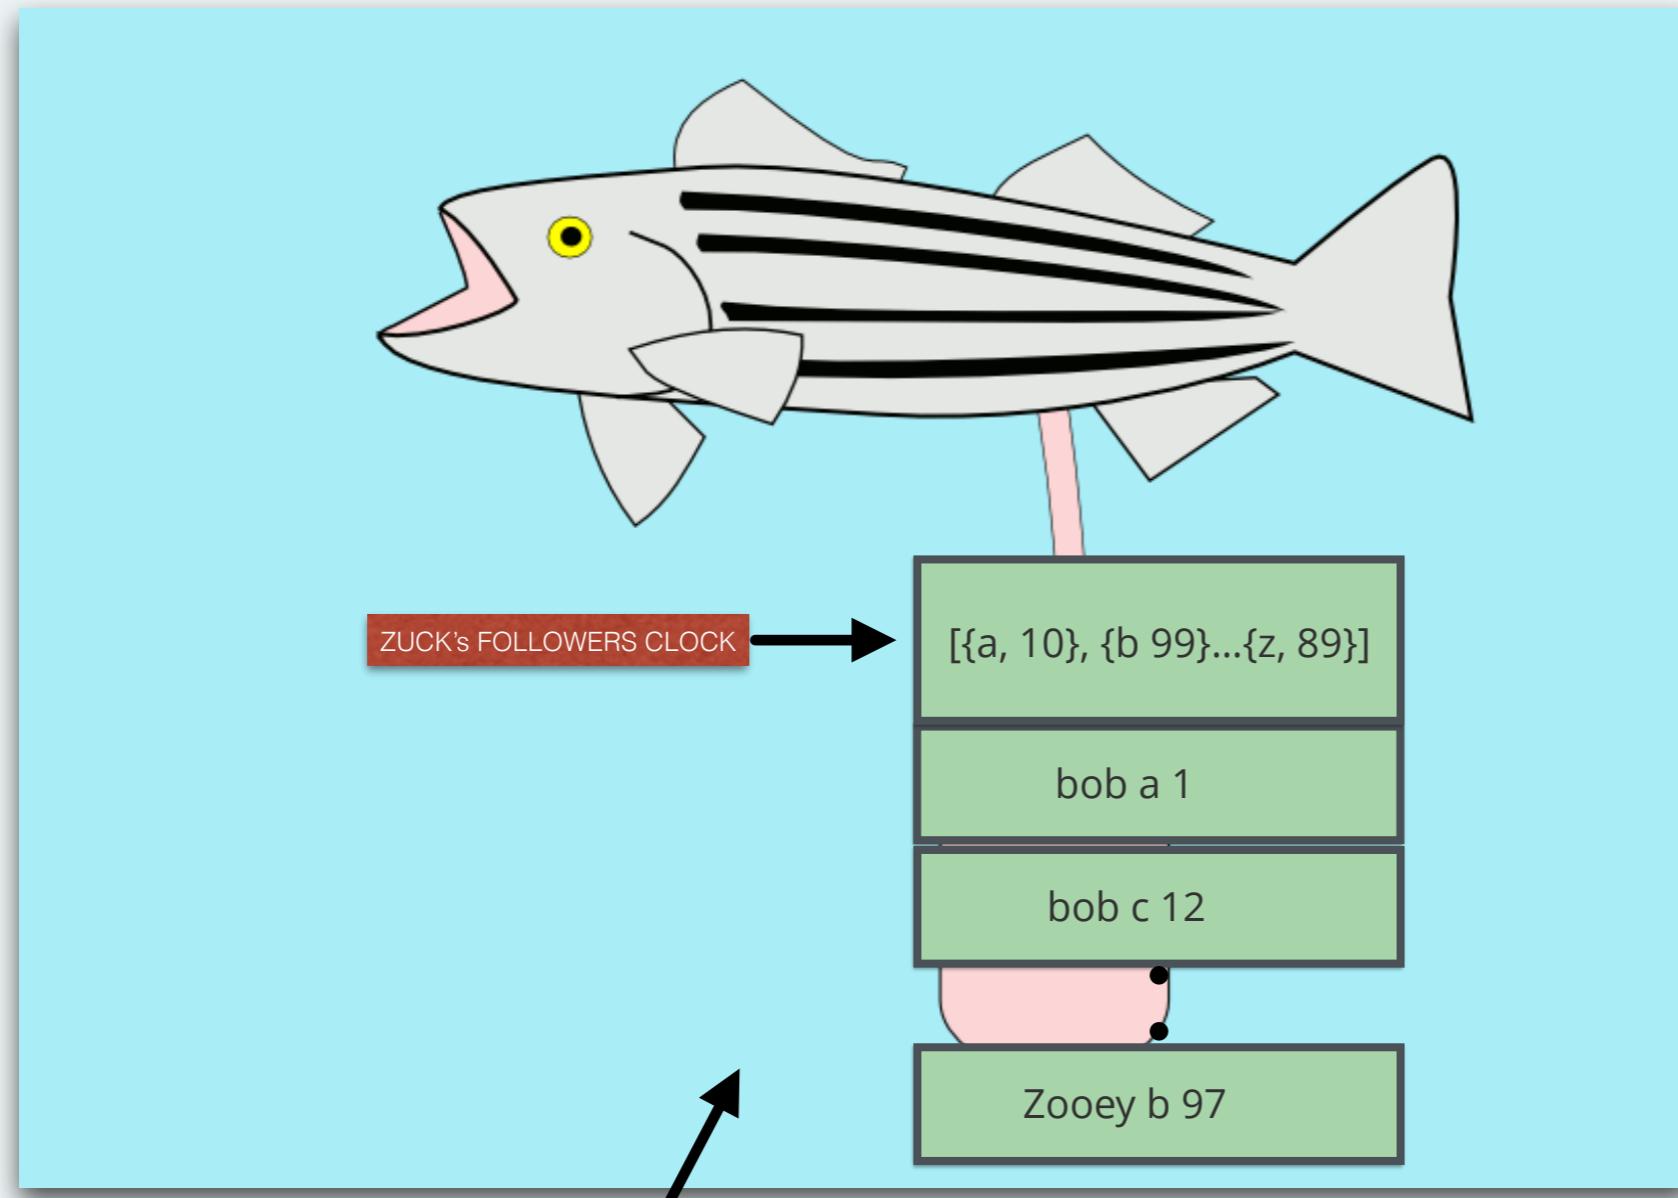




10k sets, 100k elements, 50 workers - write

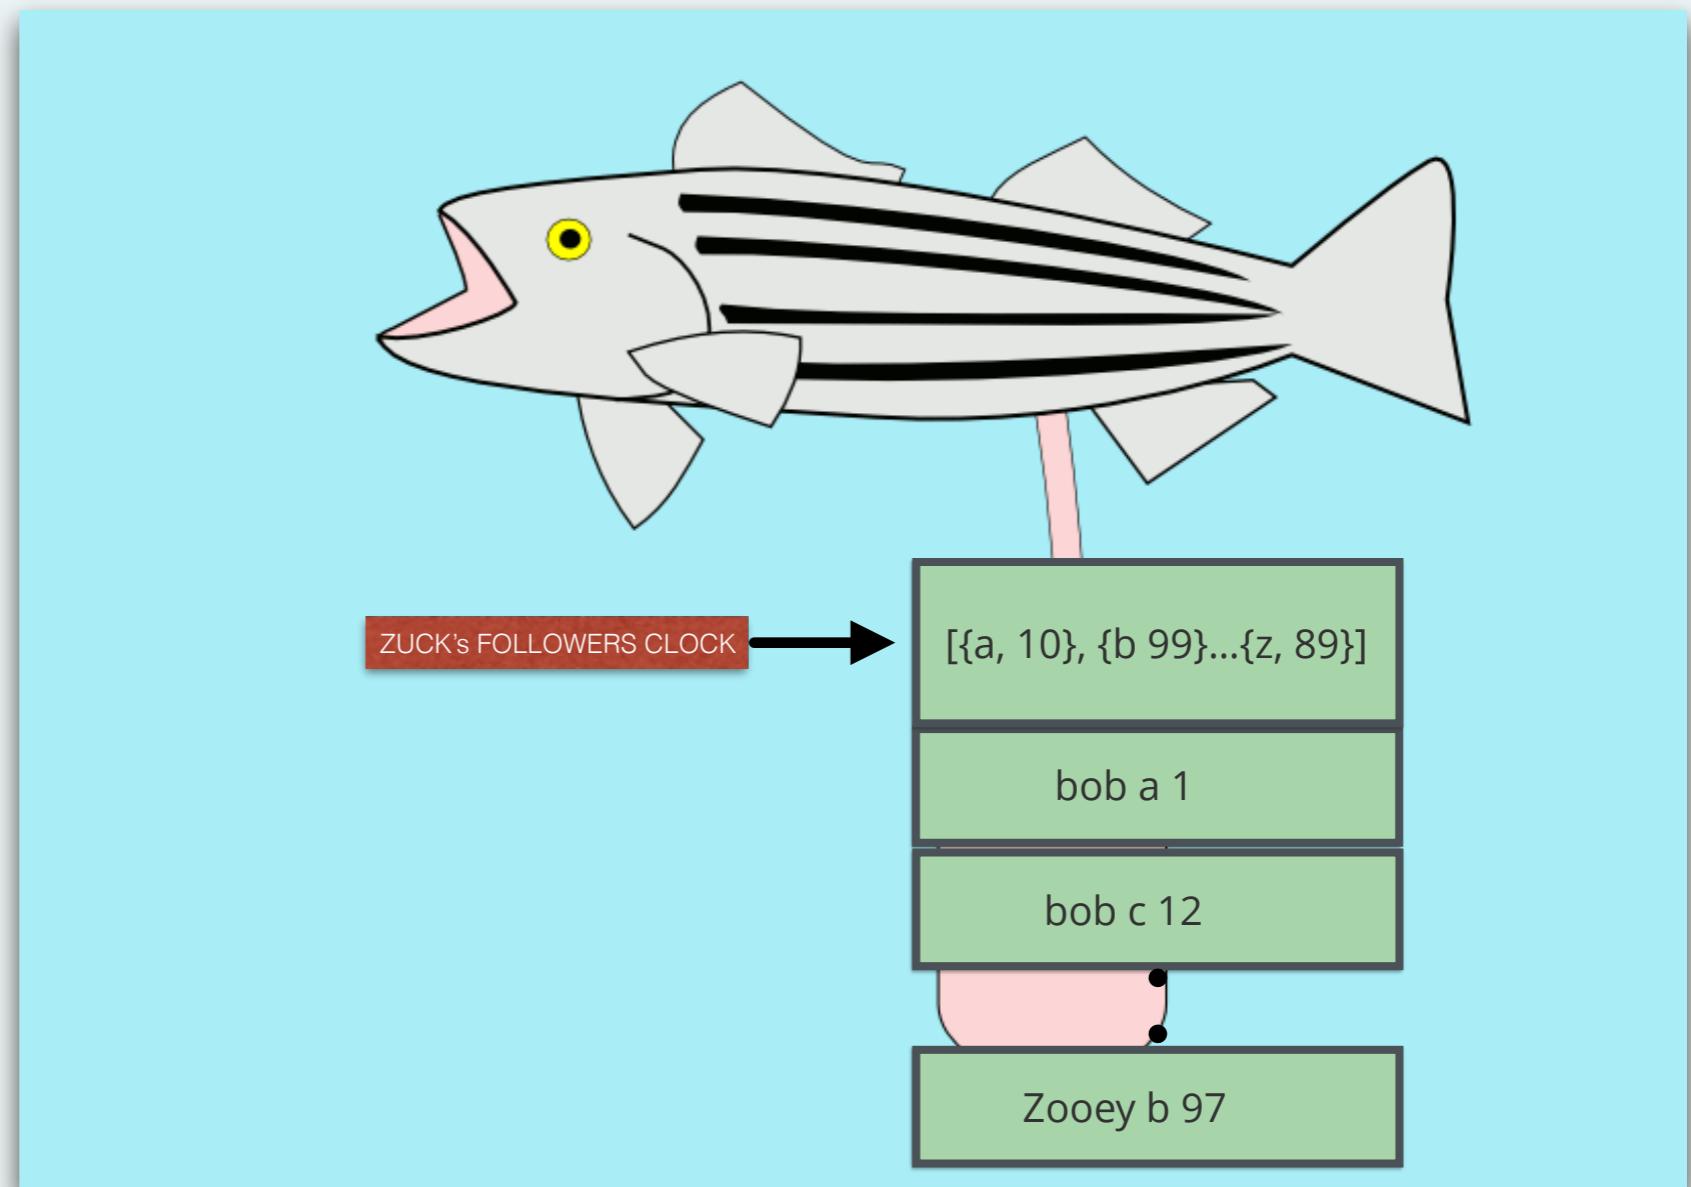


10k sets, 100k elements, 50 workers - write

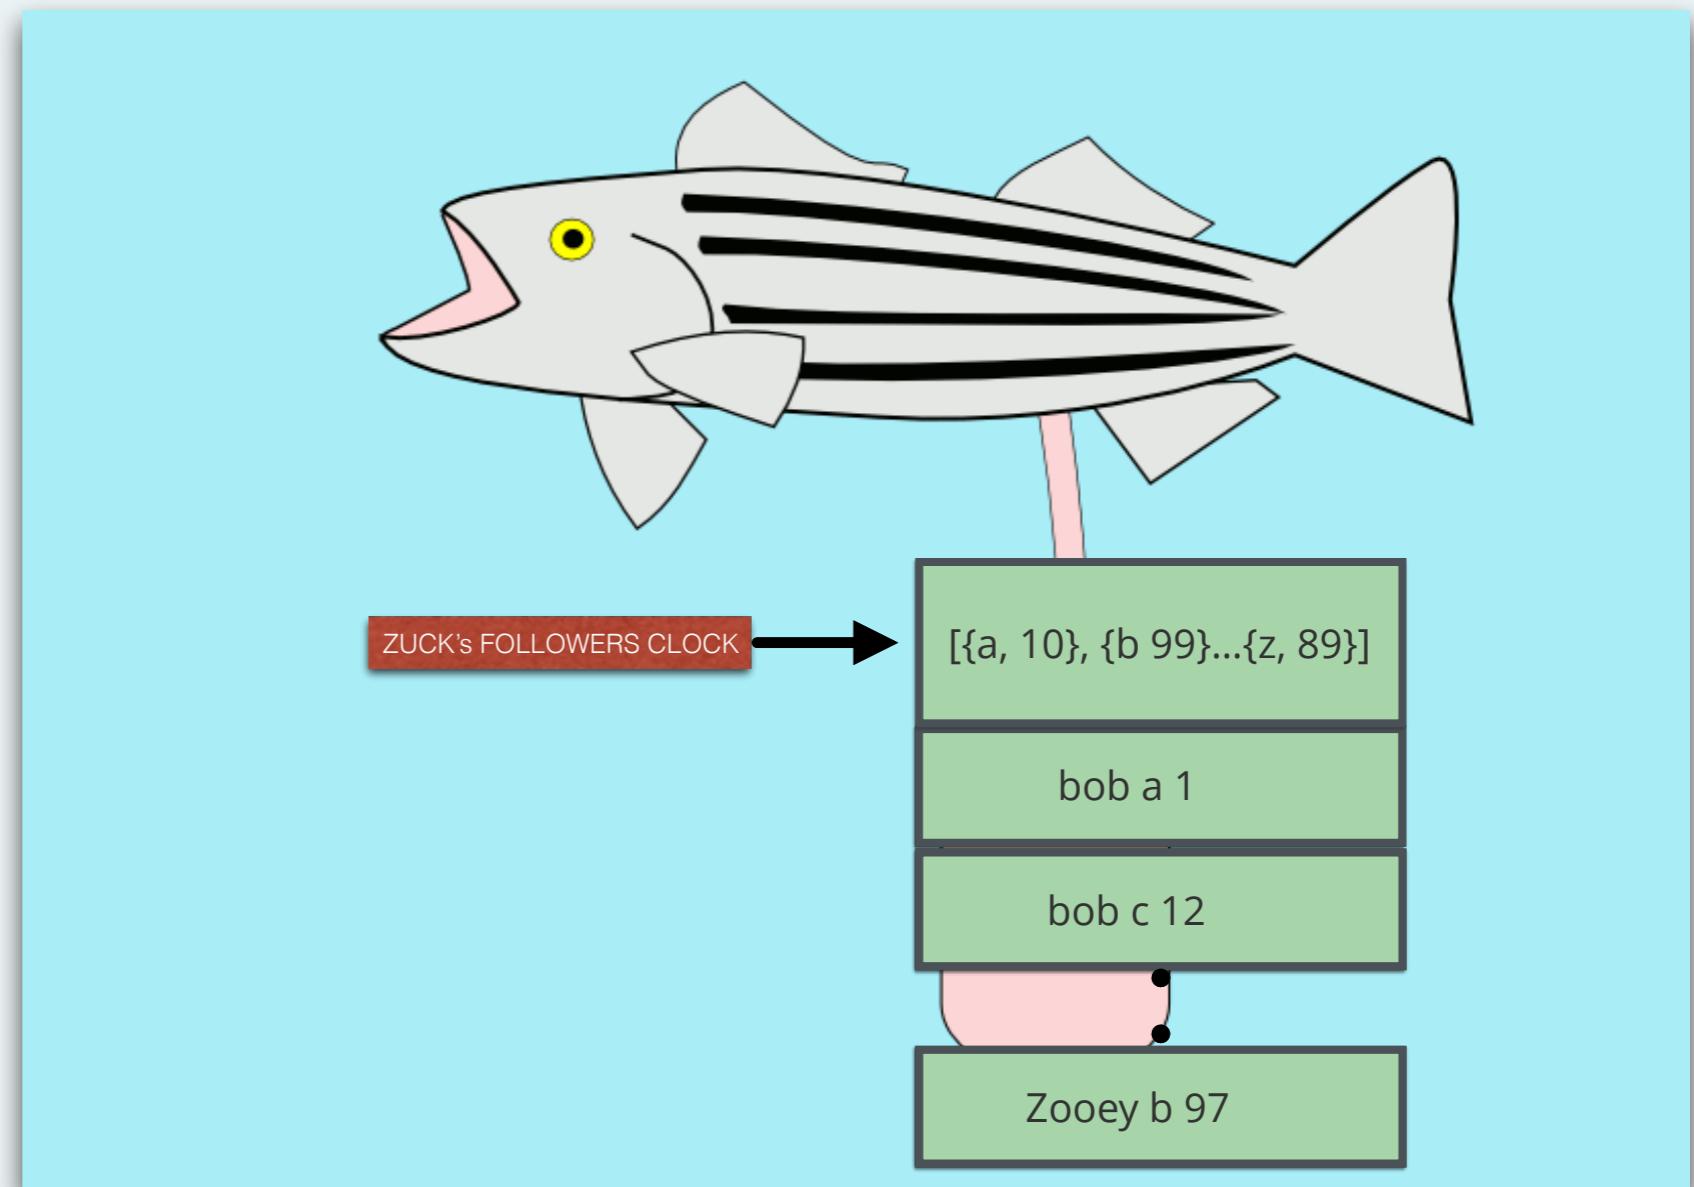


Add “Shelly”

Client X

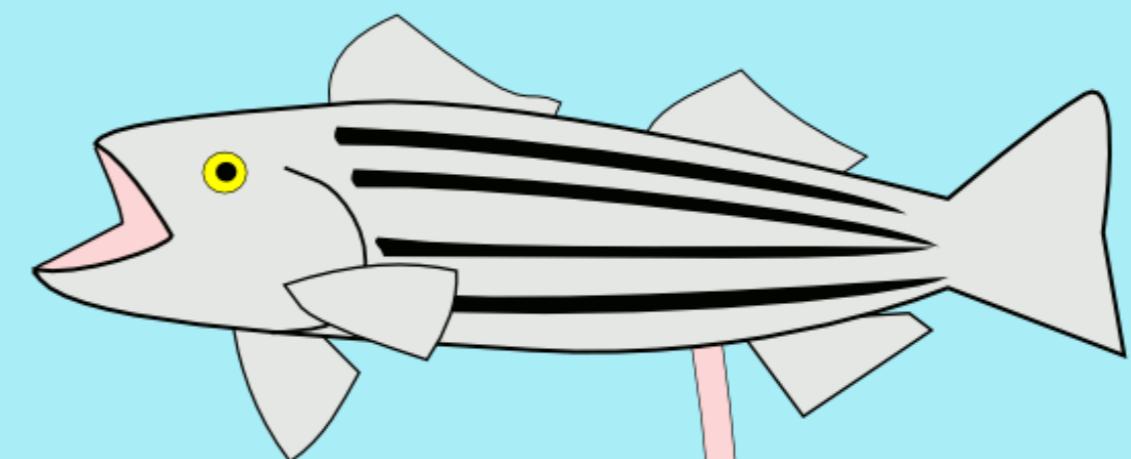


[{a, 10}, {b 99}...{z, 89}]



$[\{a, 11\}, \{b 99\} \dots \{z, 89\}]$

Shelly a 11



ZUCK's FOLLOWERS CLOCK

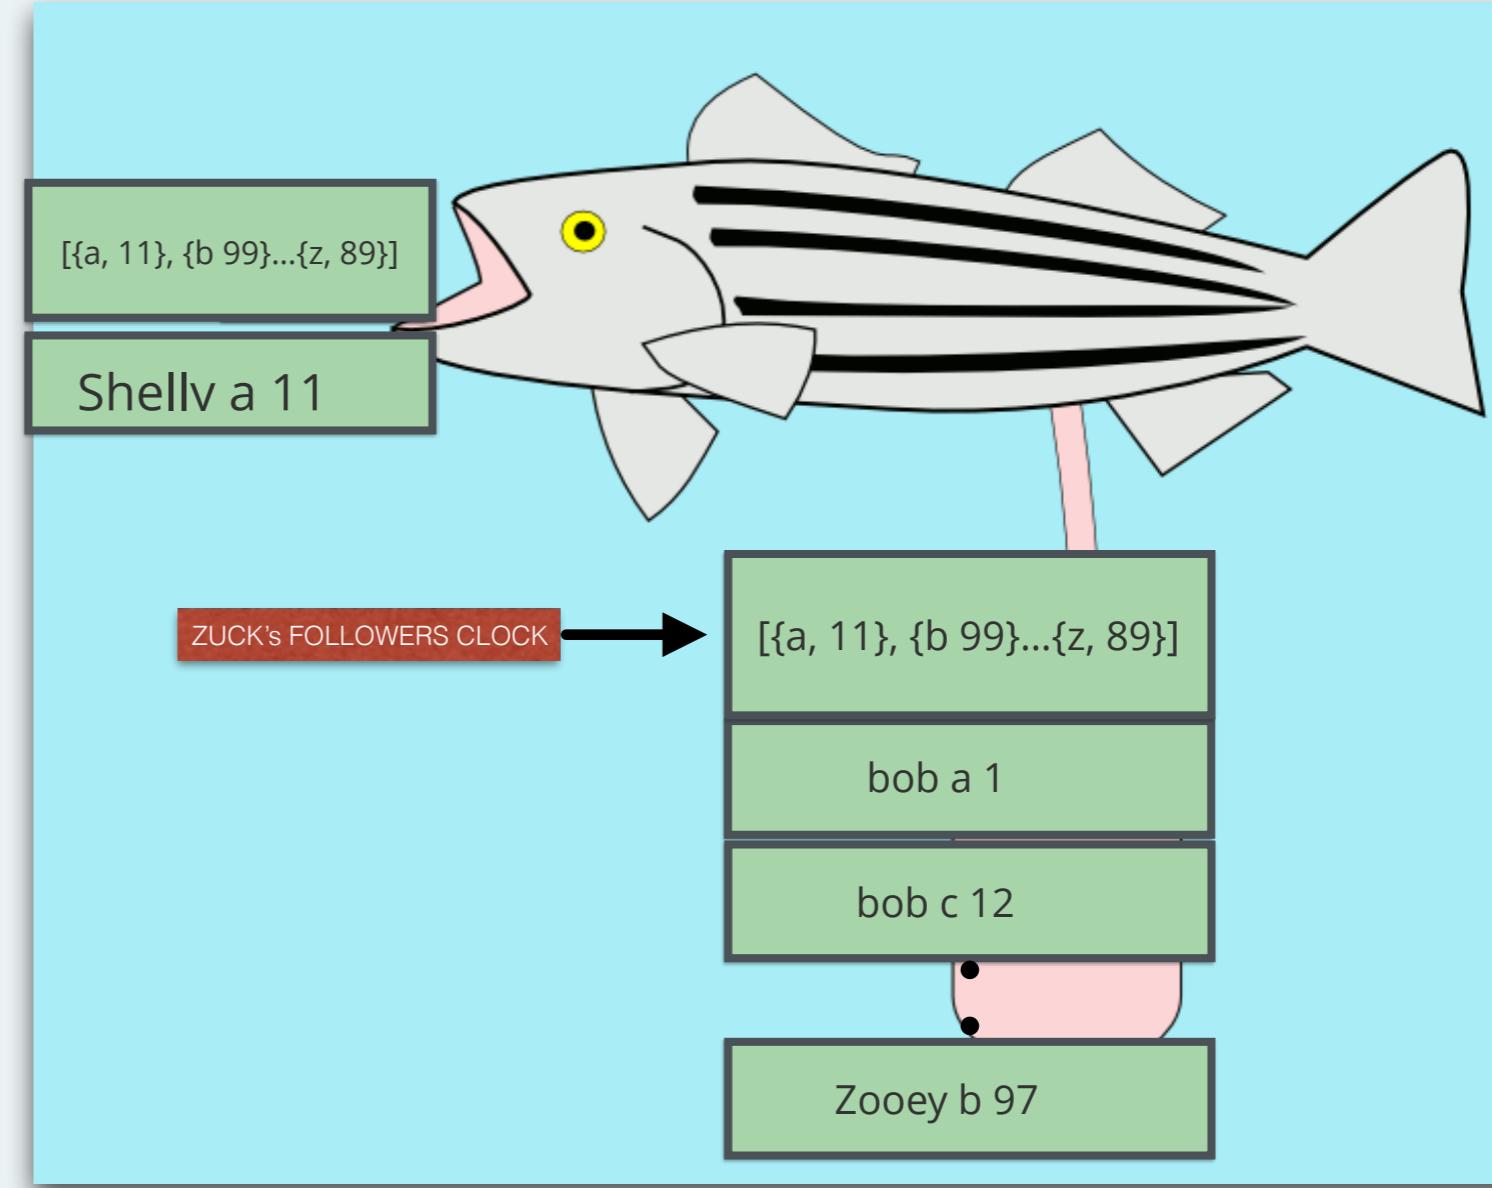


$[\{a, 10\}, \{b 99\} \dots \{z, 89\}]$

bob a 1

bob c 12

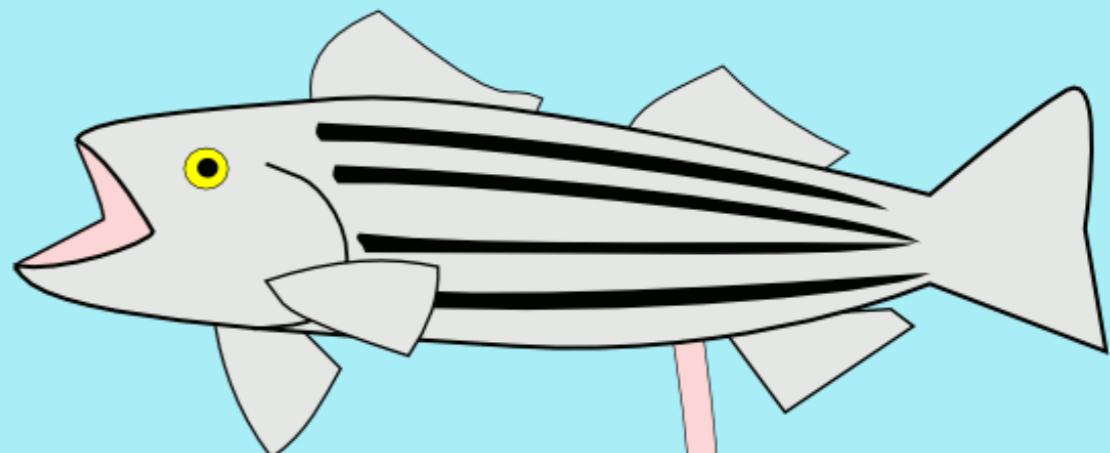
Zooey b 97



Shelly a 11

REPLICATE

ZUCK's FOLLOWERS CLOCK



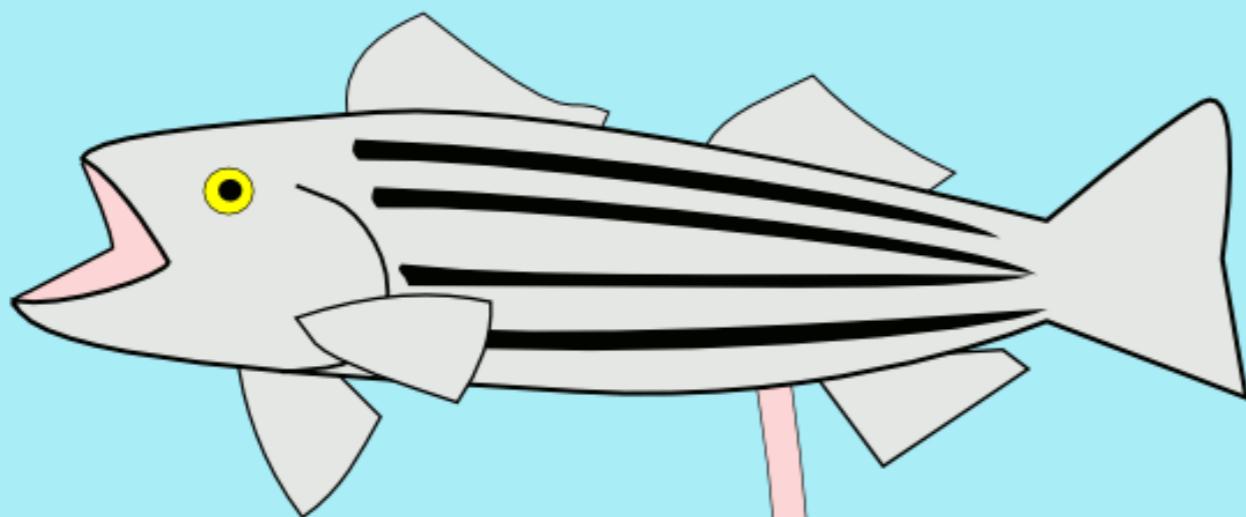
[{a, 11}, {b 99}...{z, 89}]

bob a 1

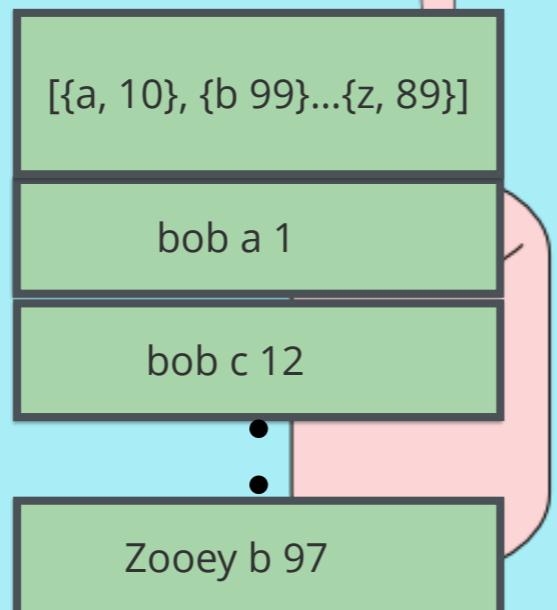
bob c 12

Shelly a 11

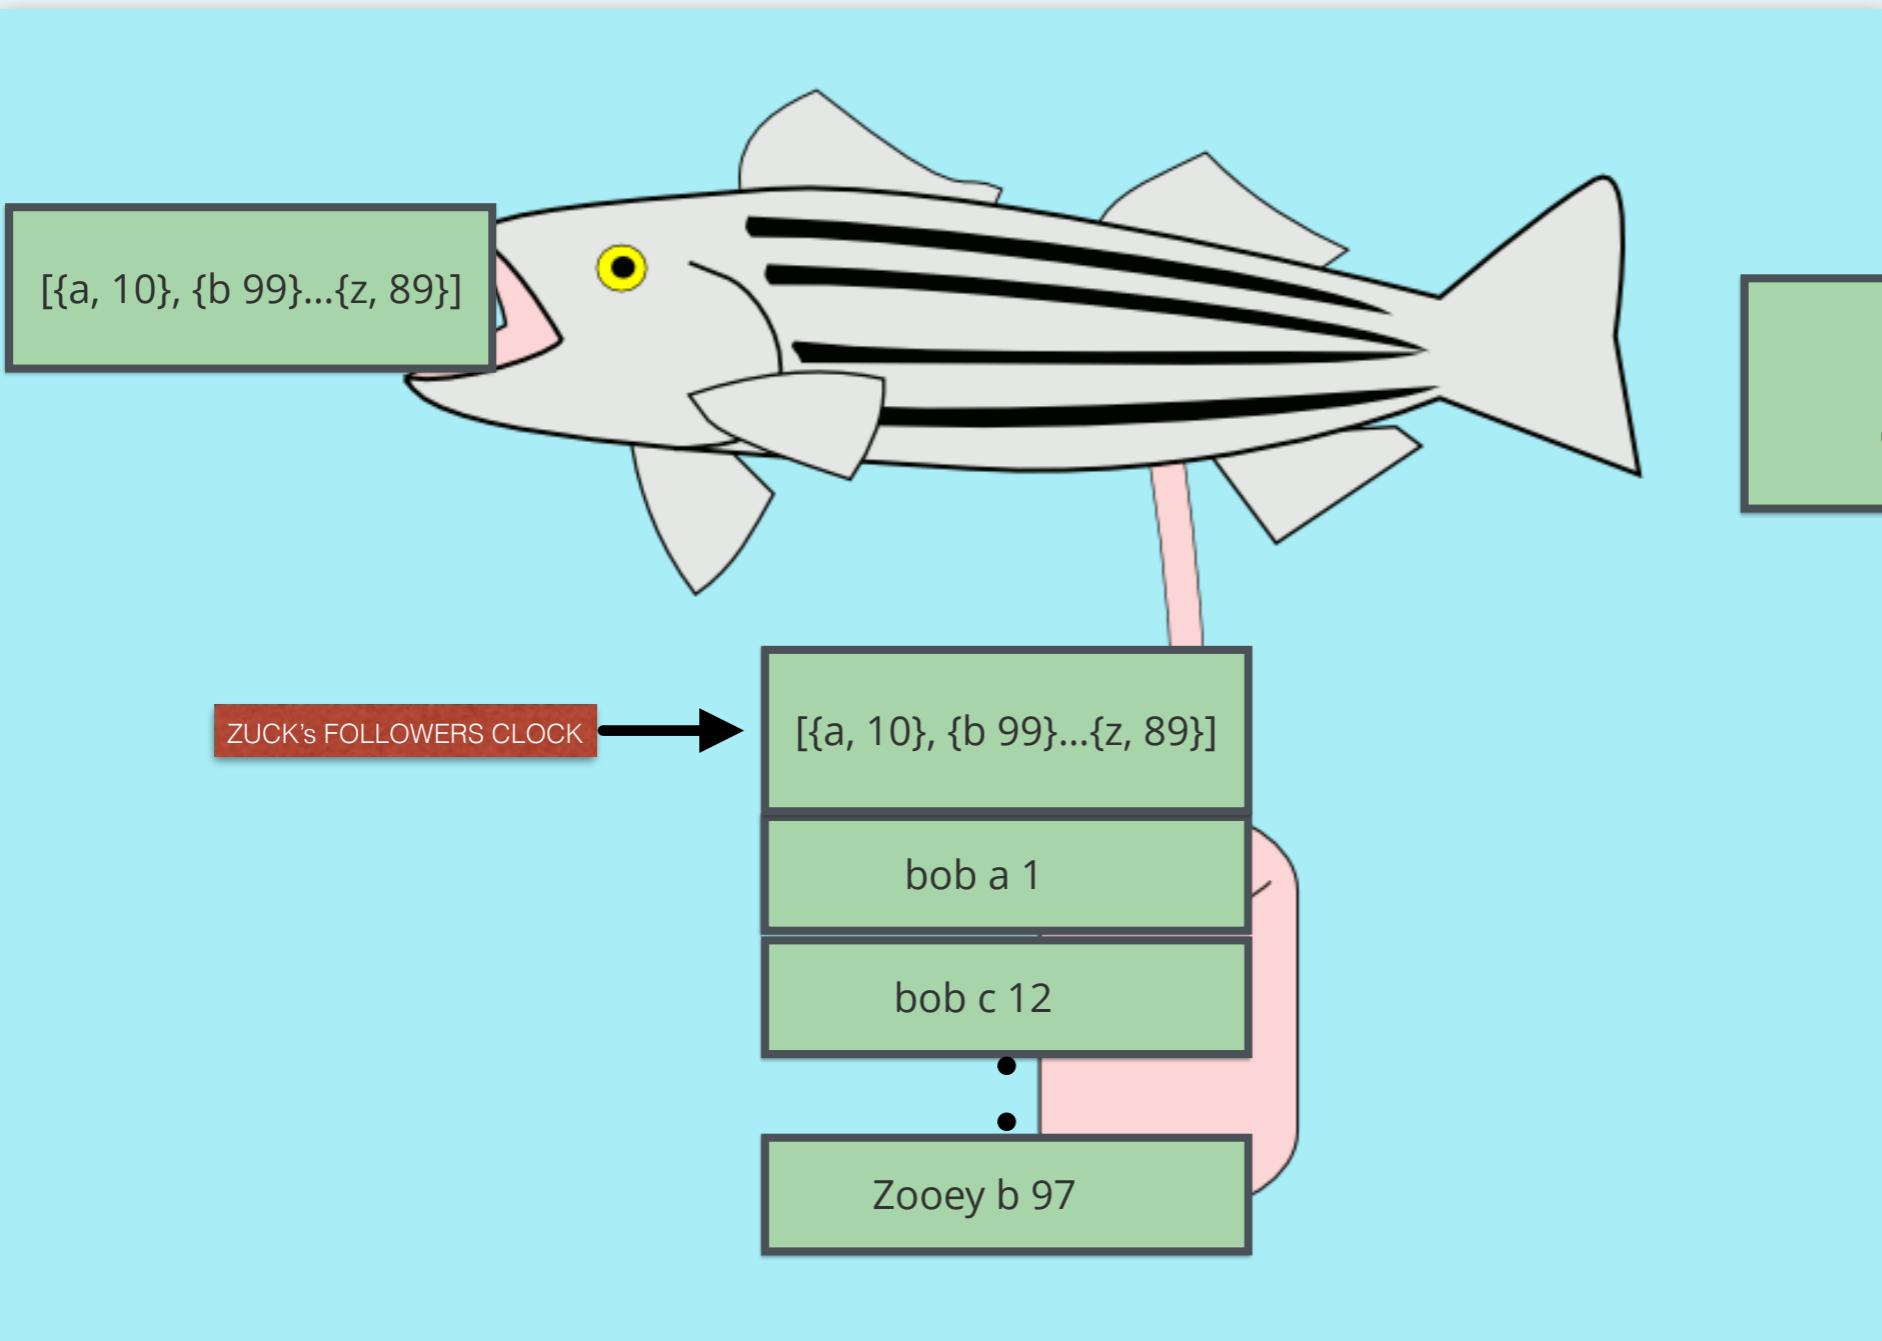
Zooey b 97



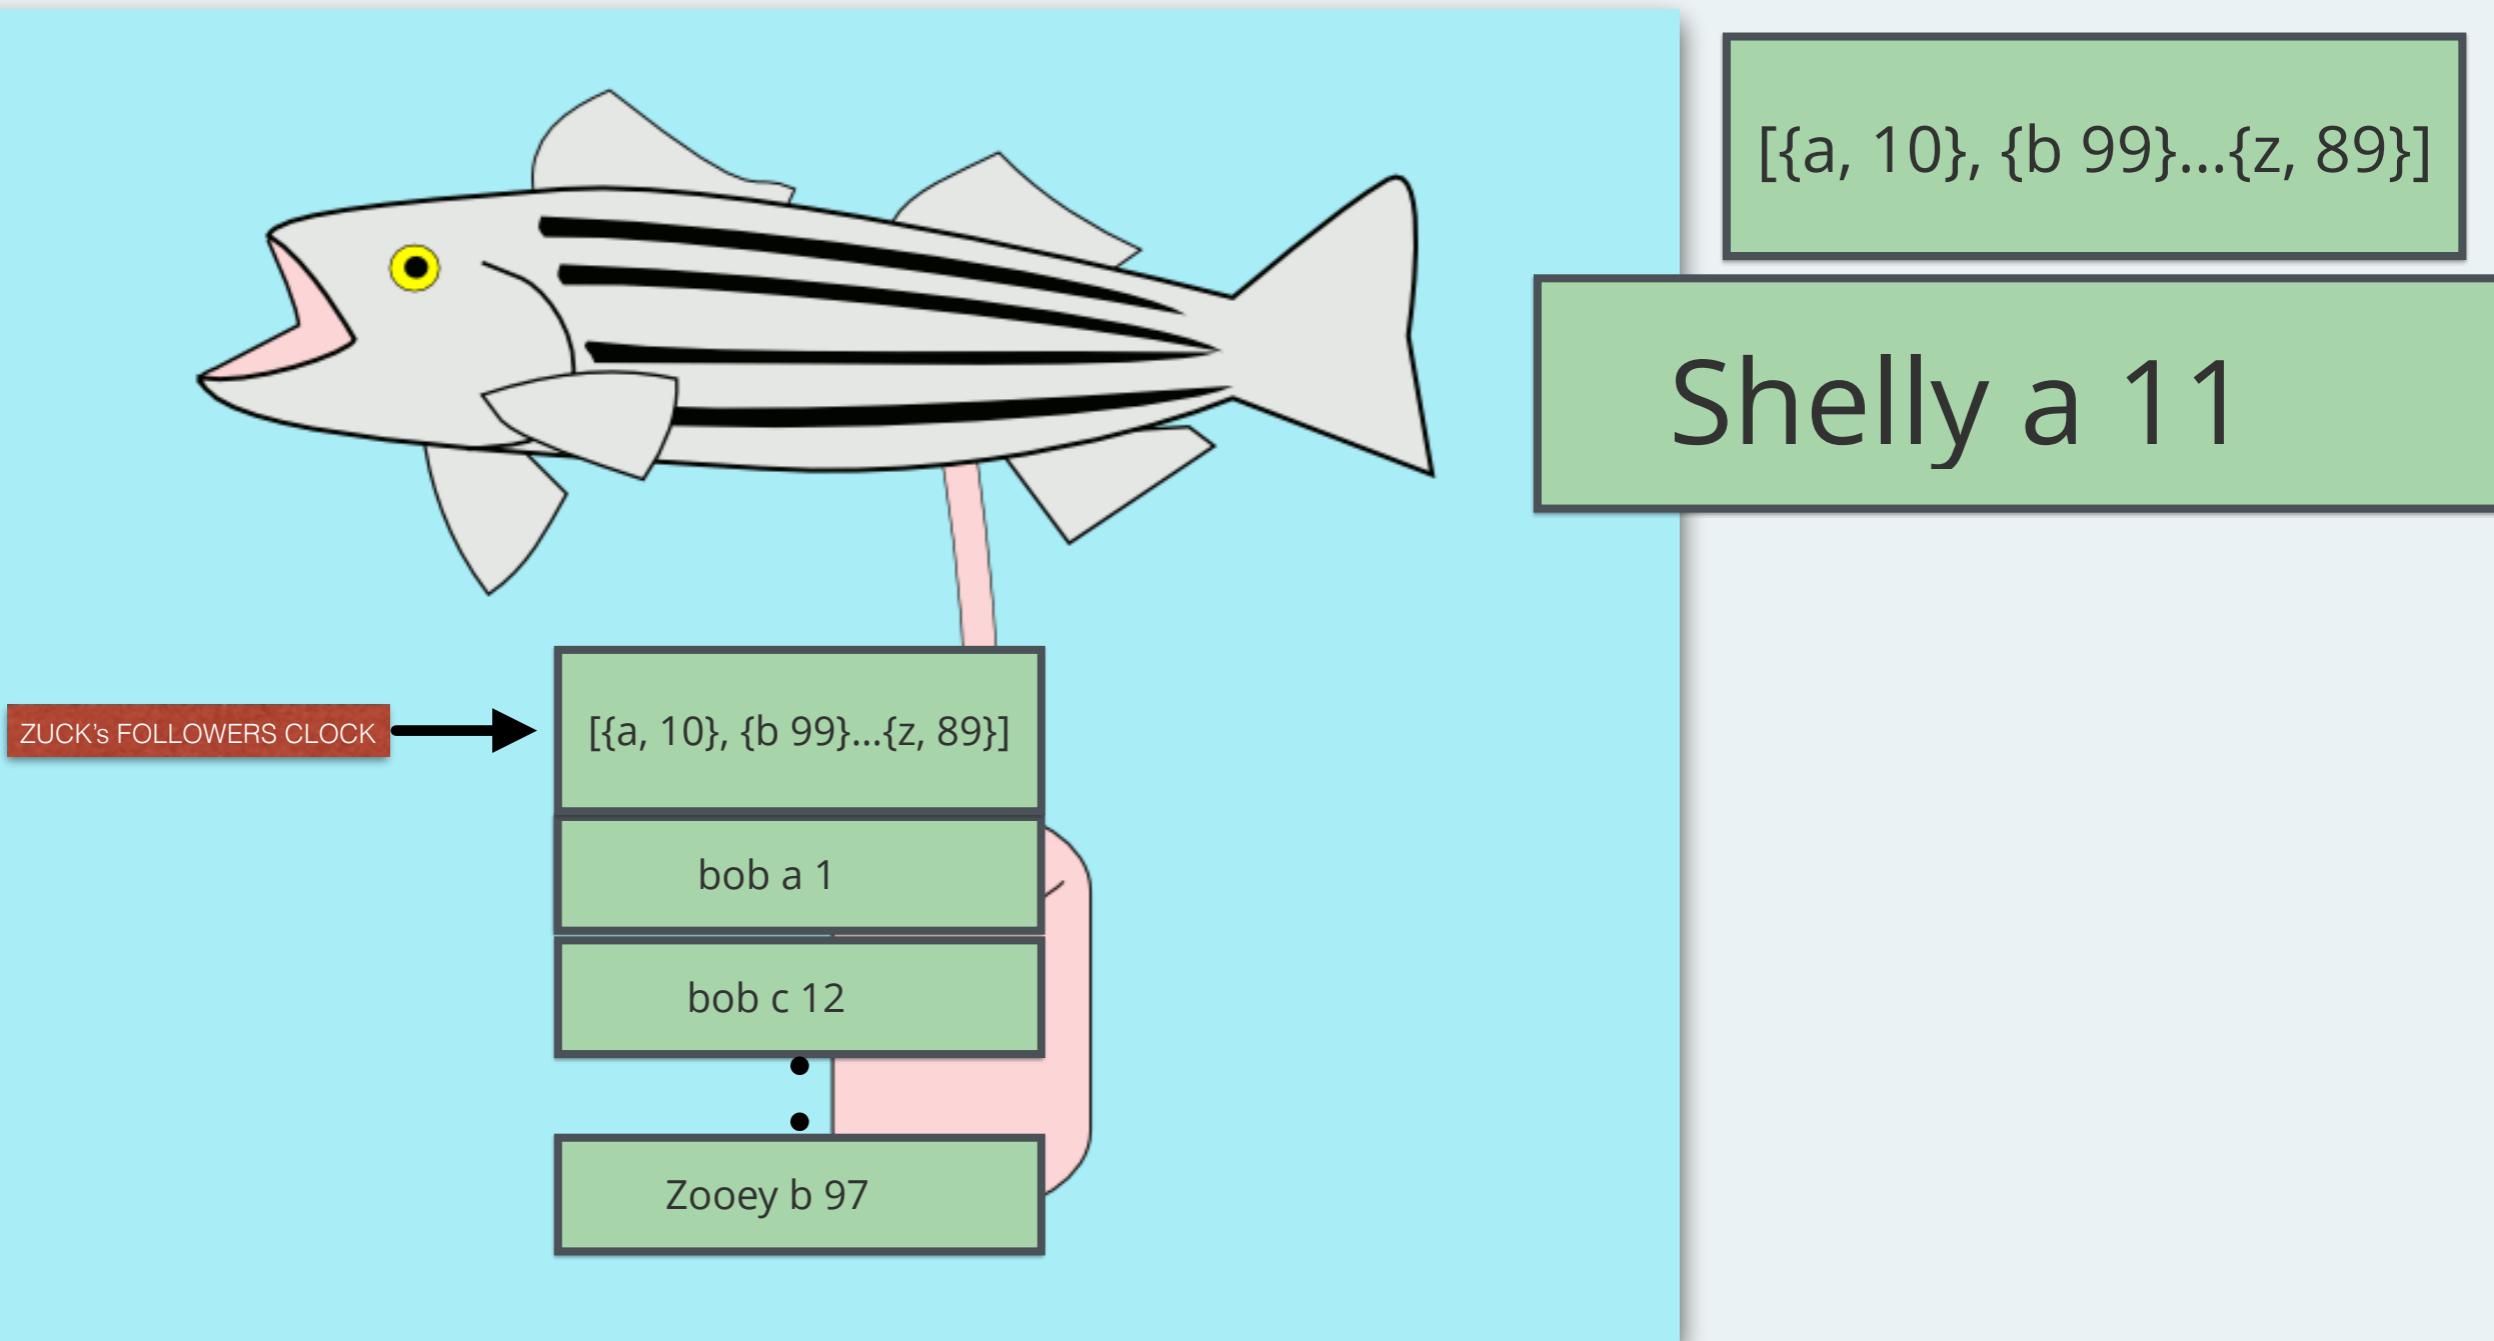
ZUCK's FOLLOWERS CLOCK →

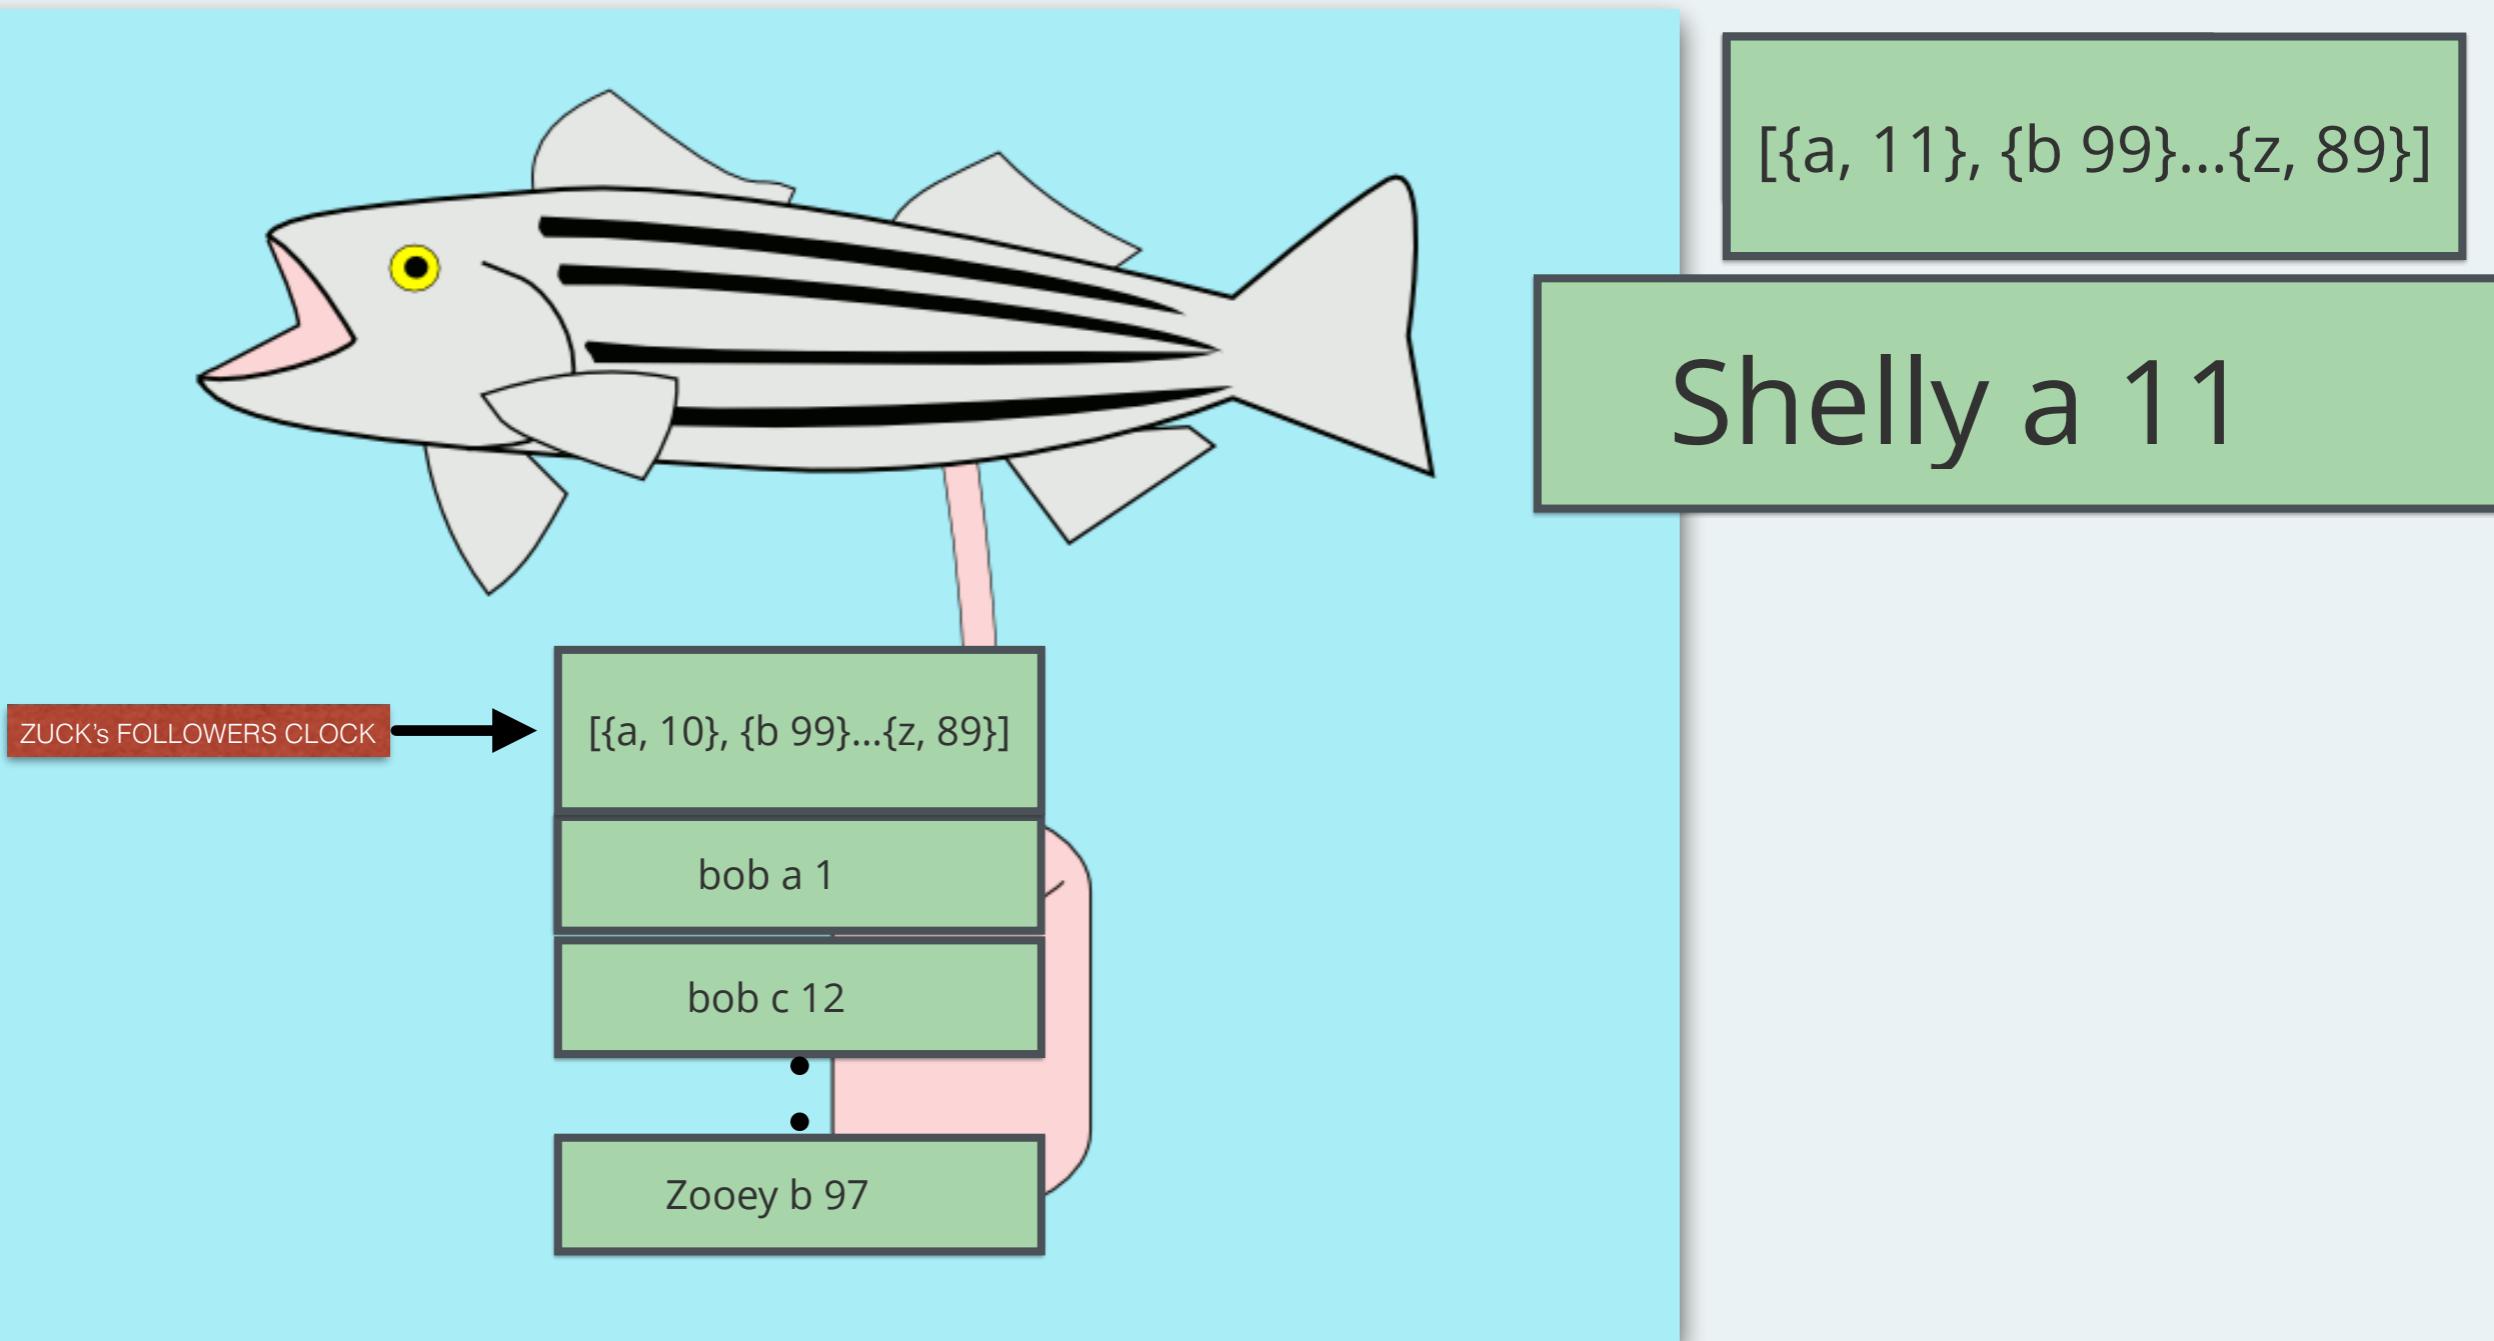


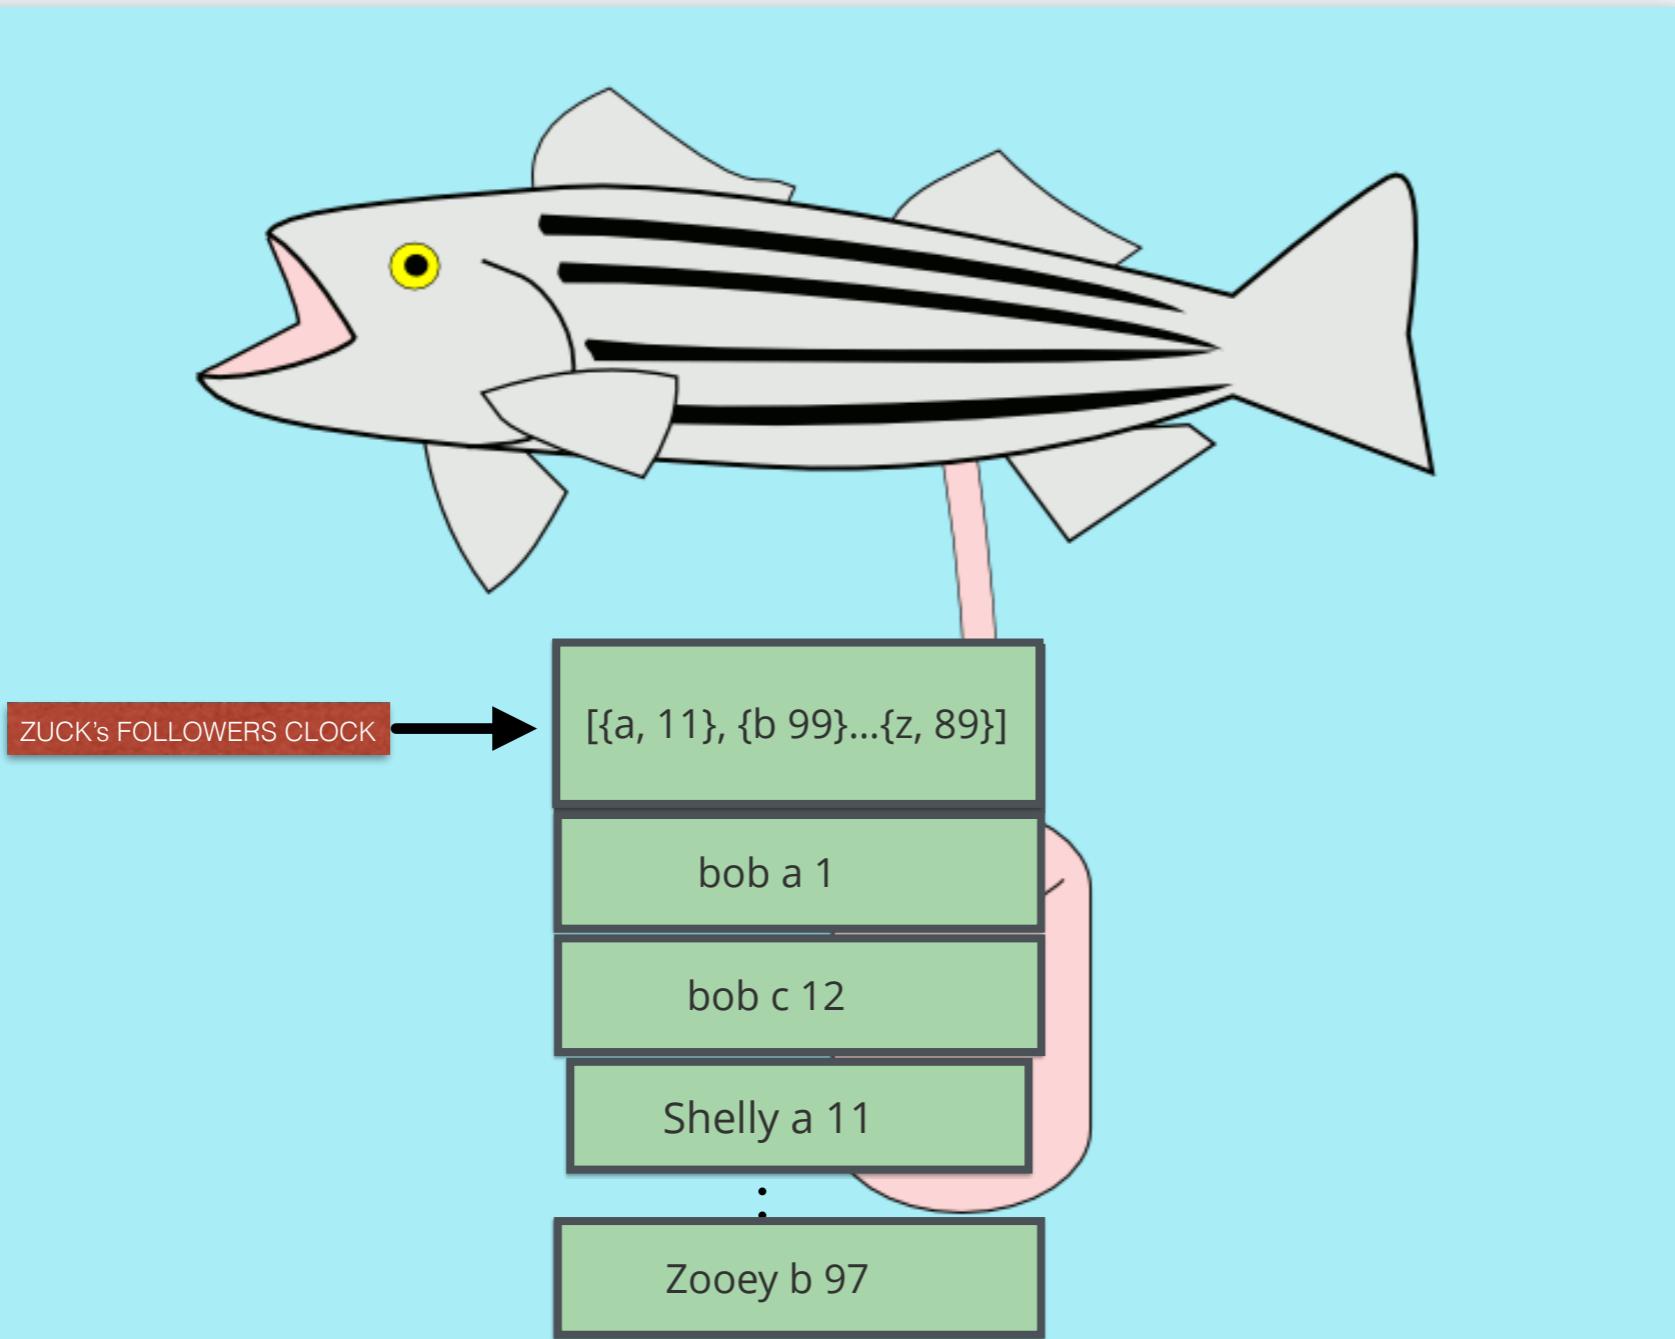
Shelly a 11



Shelly a 11







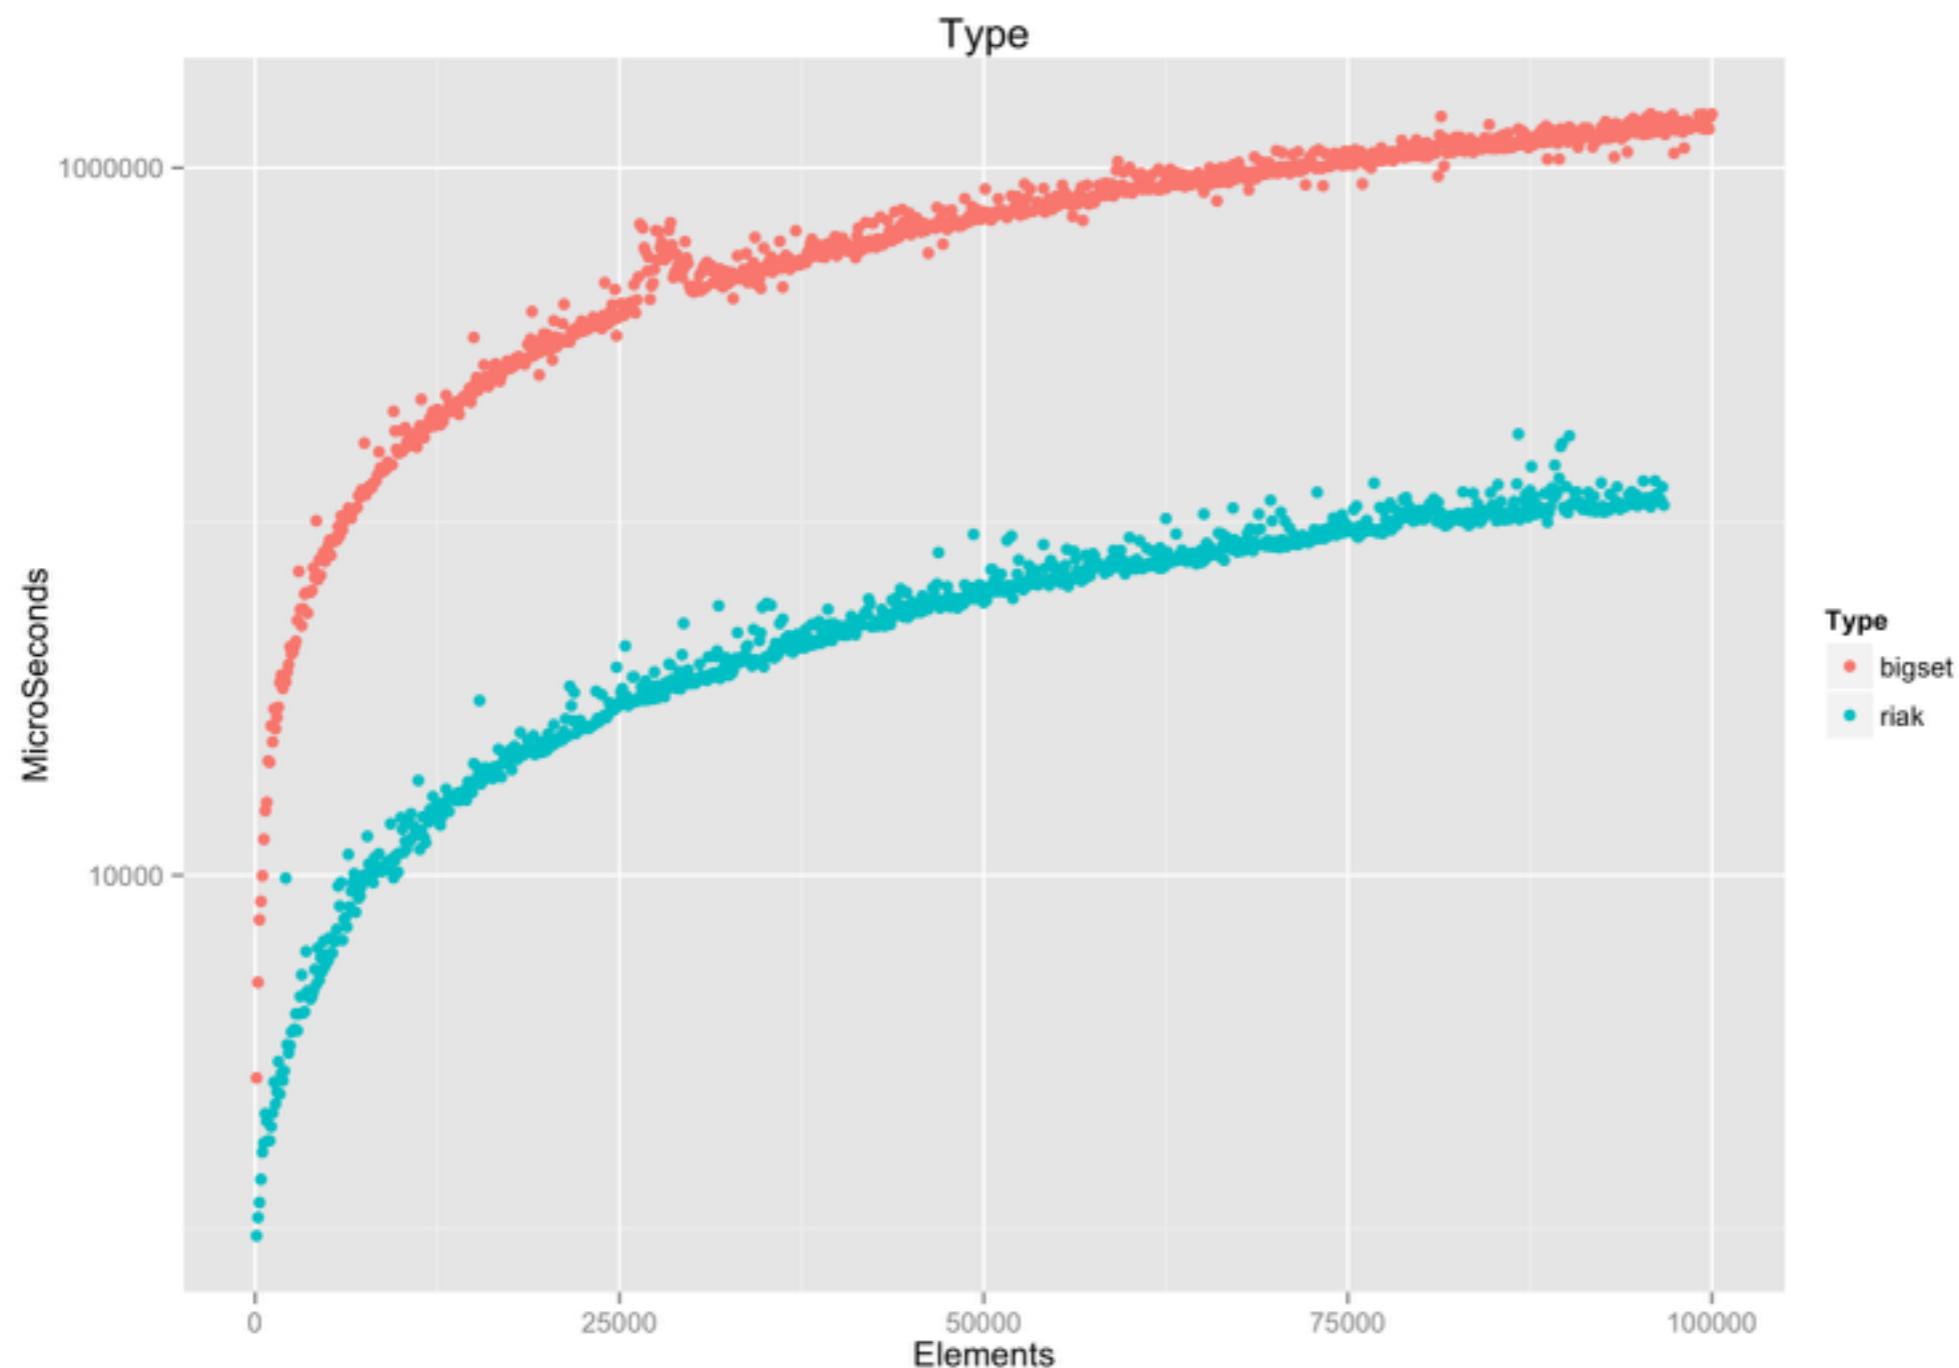
**THAT'S IT!**

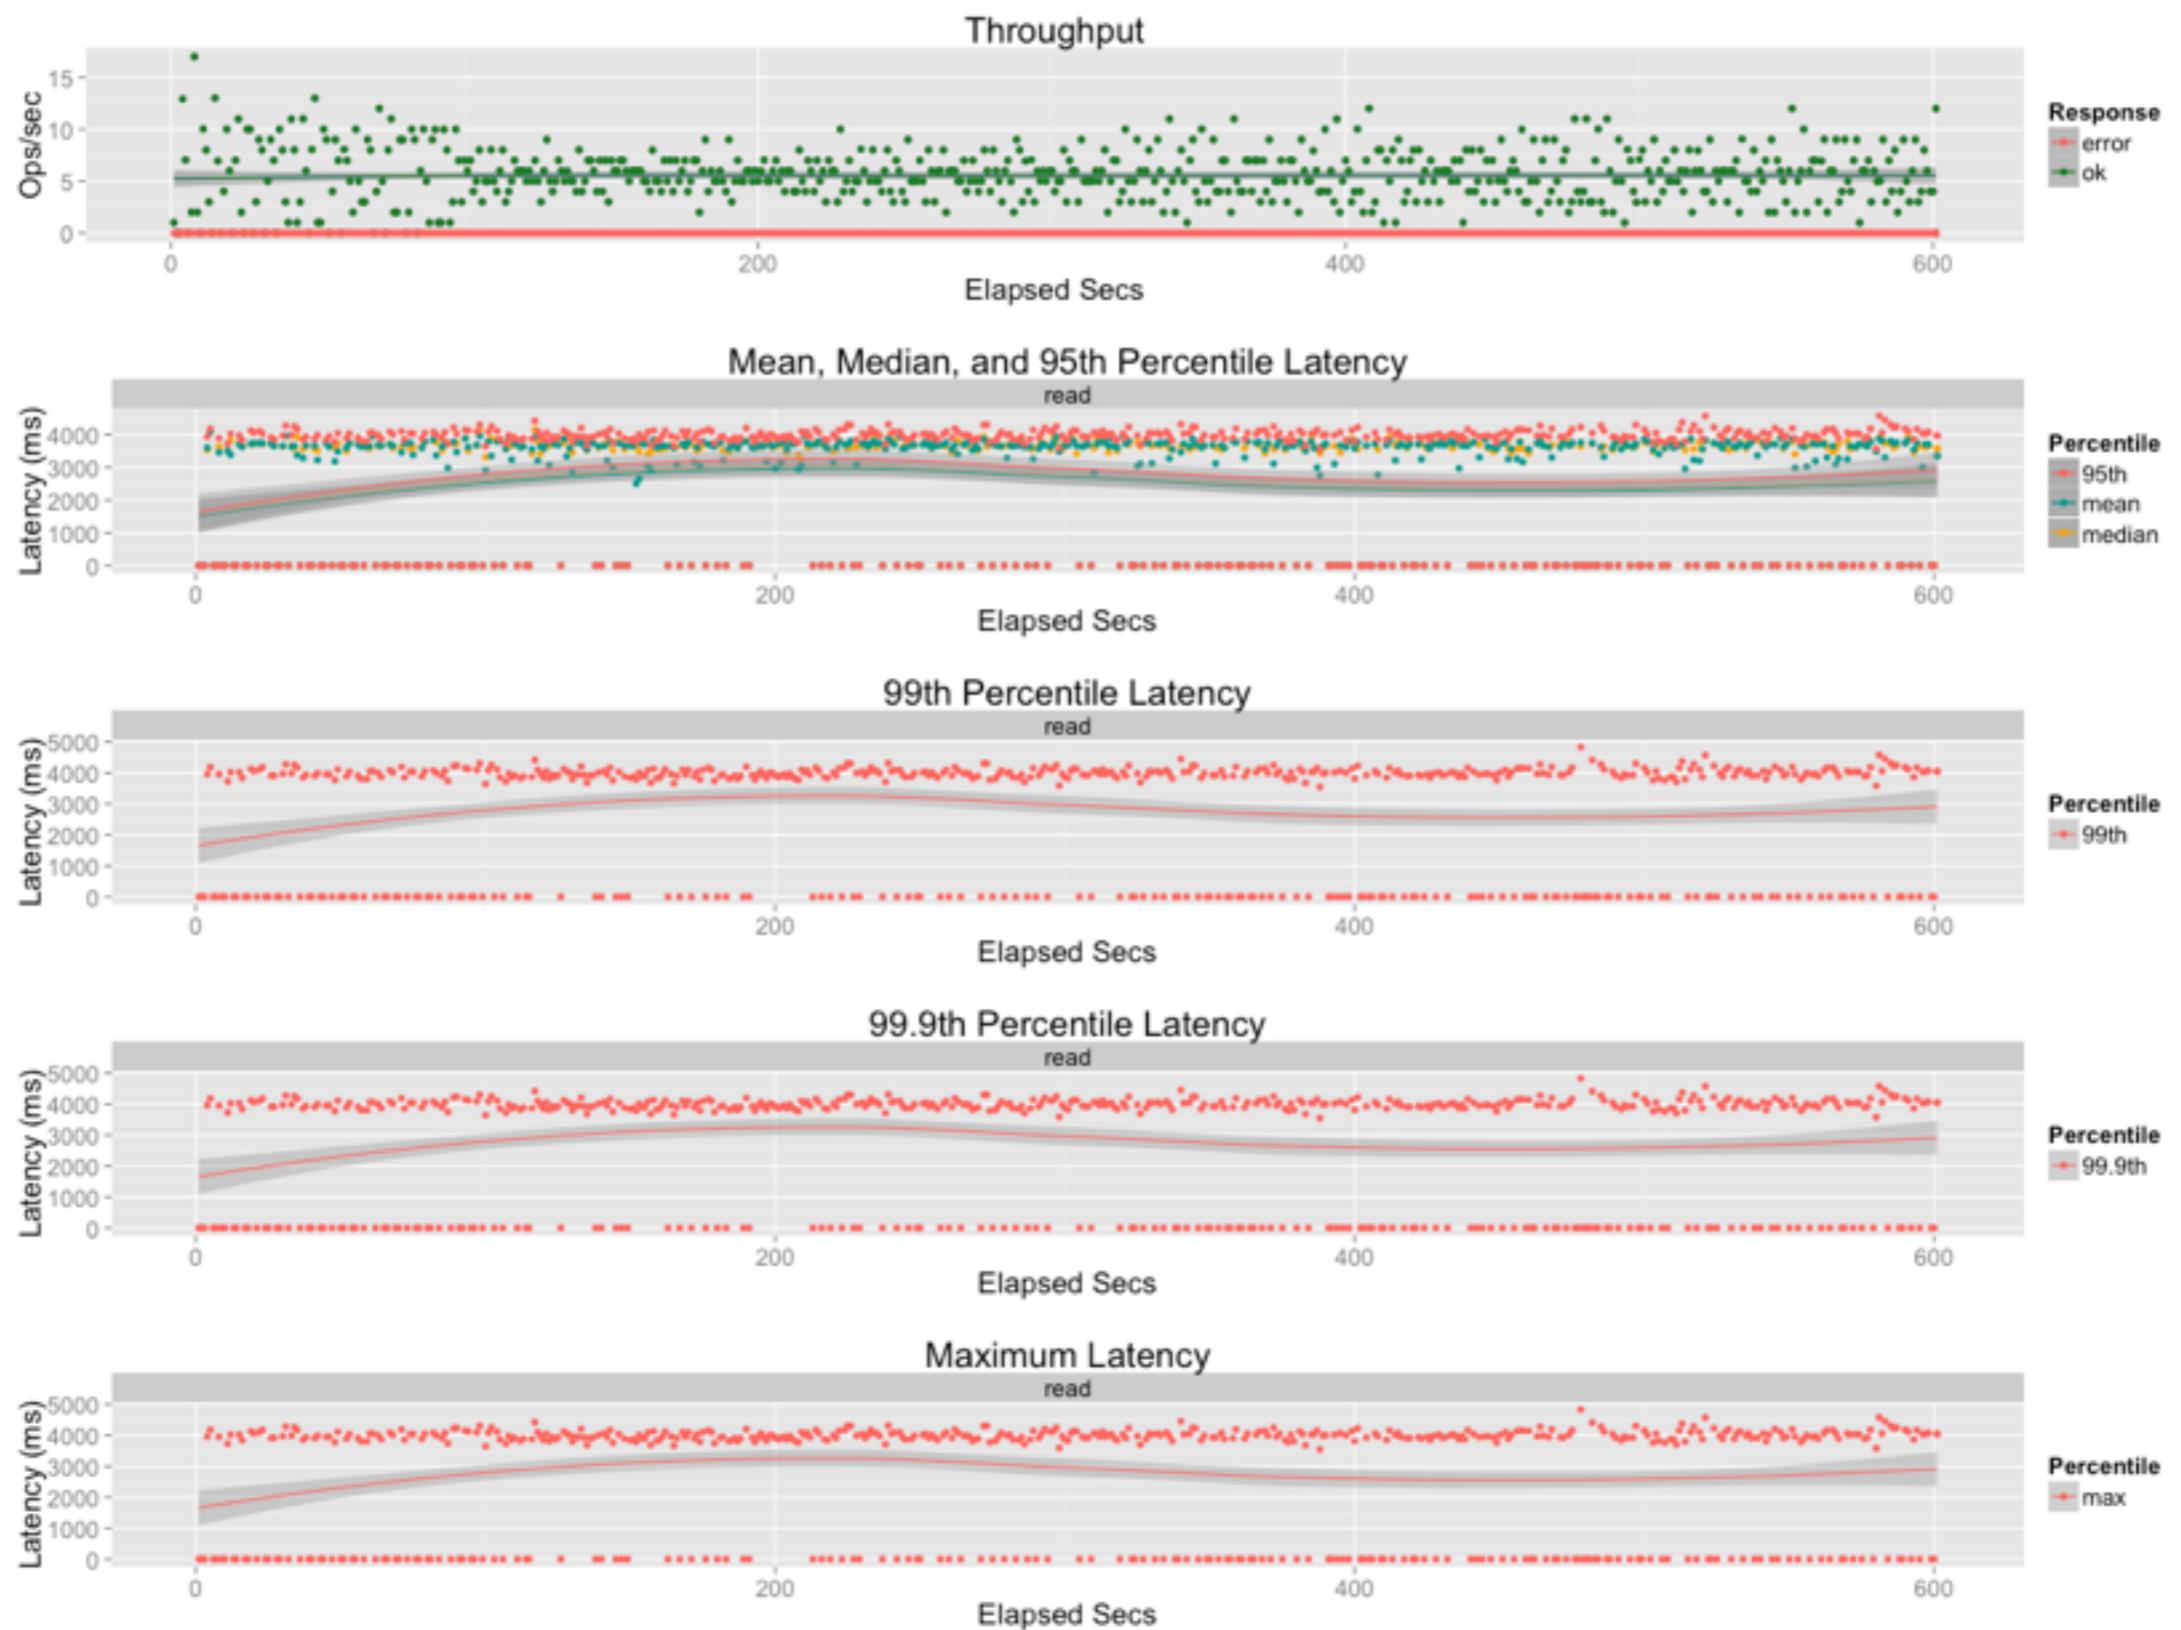
# THAT'S IT?

- Reads!
- Version Vector!
- Hand-Off!
- AAE!

Reads?

# Initial Read Results





10k sets, 100k elements, 20 workers - read

clock

set-tombstone

element-1

•

•

element-N

end\_key

clock

set-tombstone

element-1

•

•

element-N

end\_key

# Read Clock

Iterate keys

riak\_dt\_orswot

Version Vector

[{vnodeA, 10}, {vnodeB, 4}, {vnodeC, 11}...]

Entries

Bob

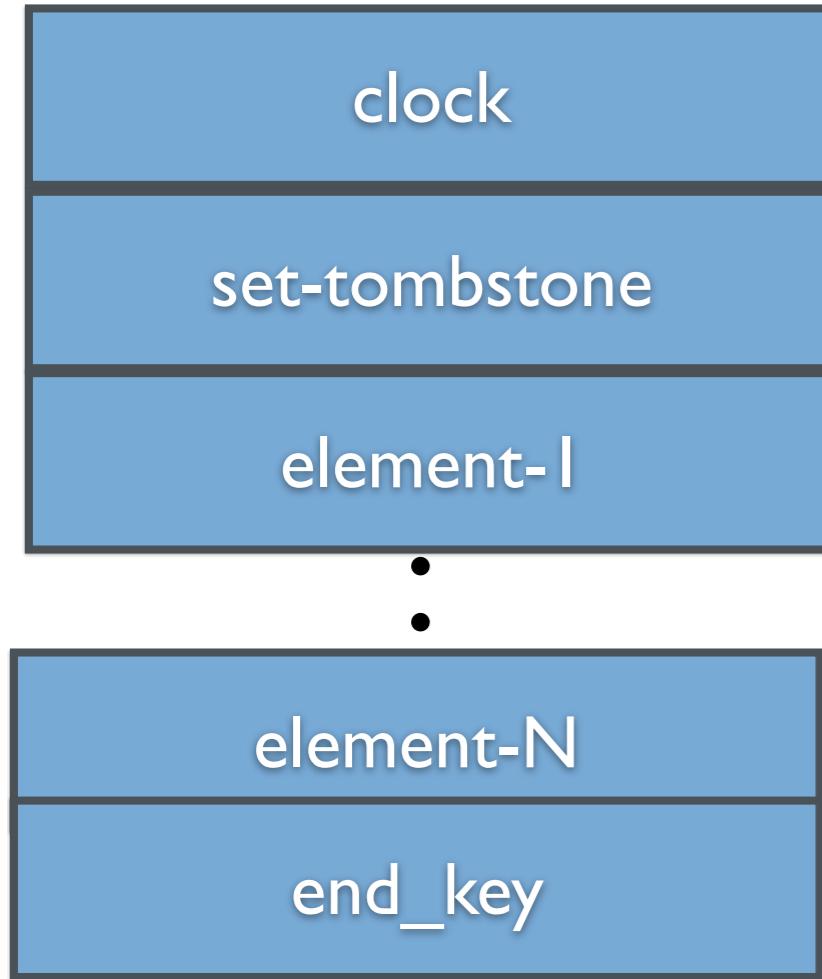
[{vnodeA, 2}]

Cameron

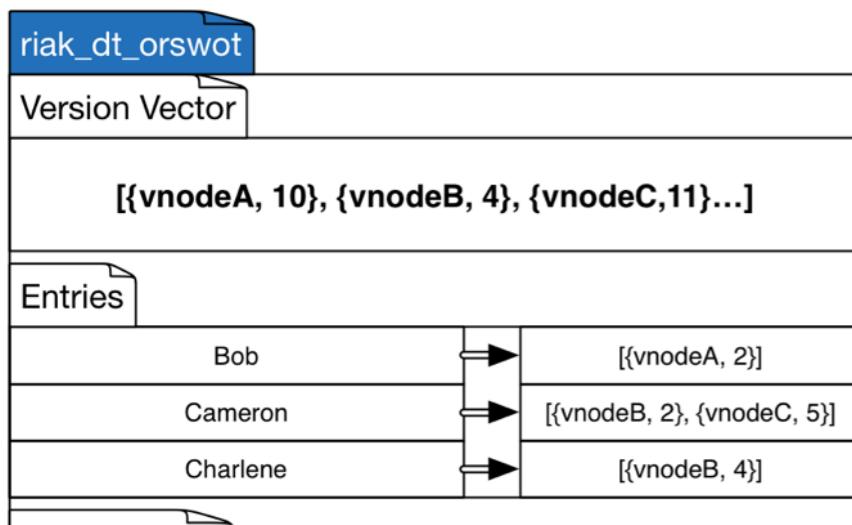
[{vnodeB, 2}, {vnodeC, 5}]

Charlene

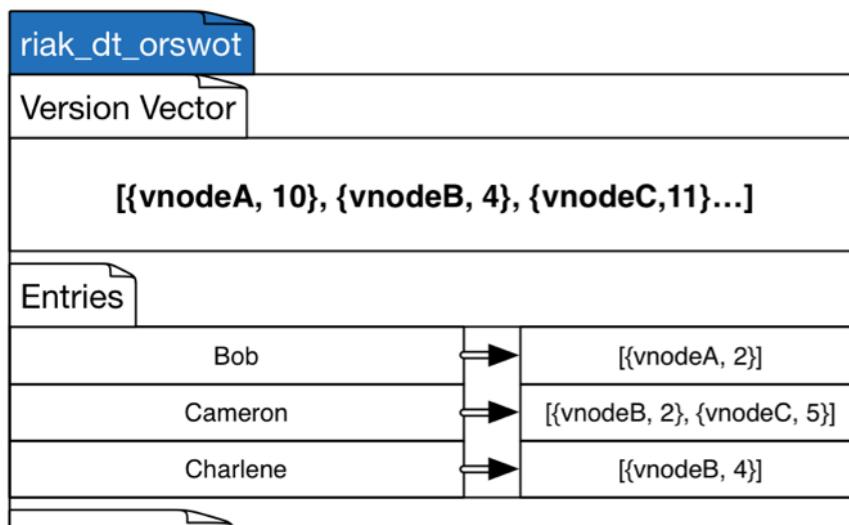
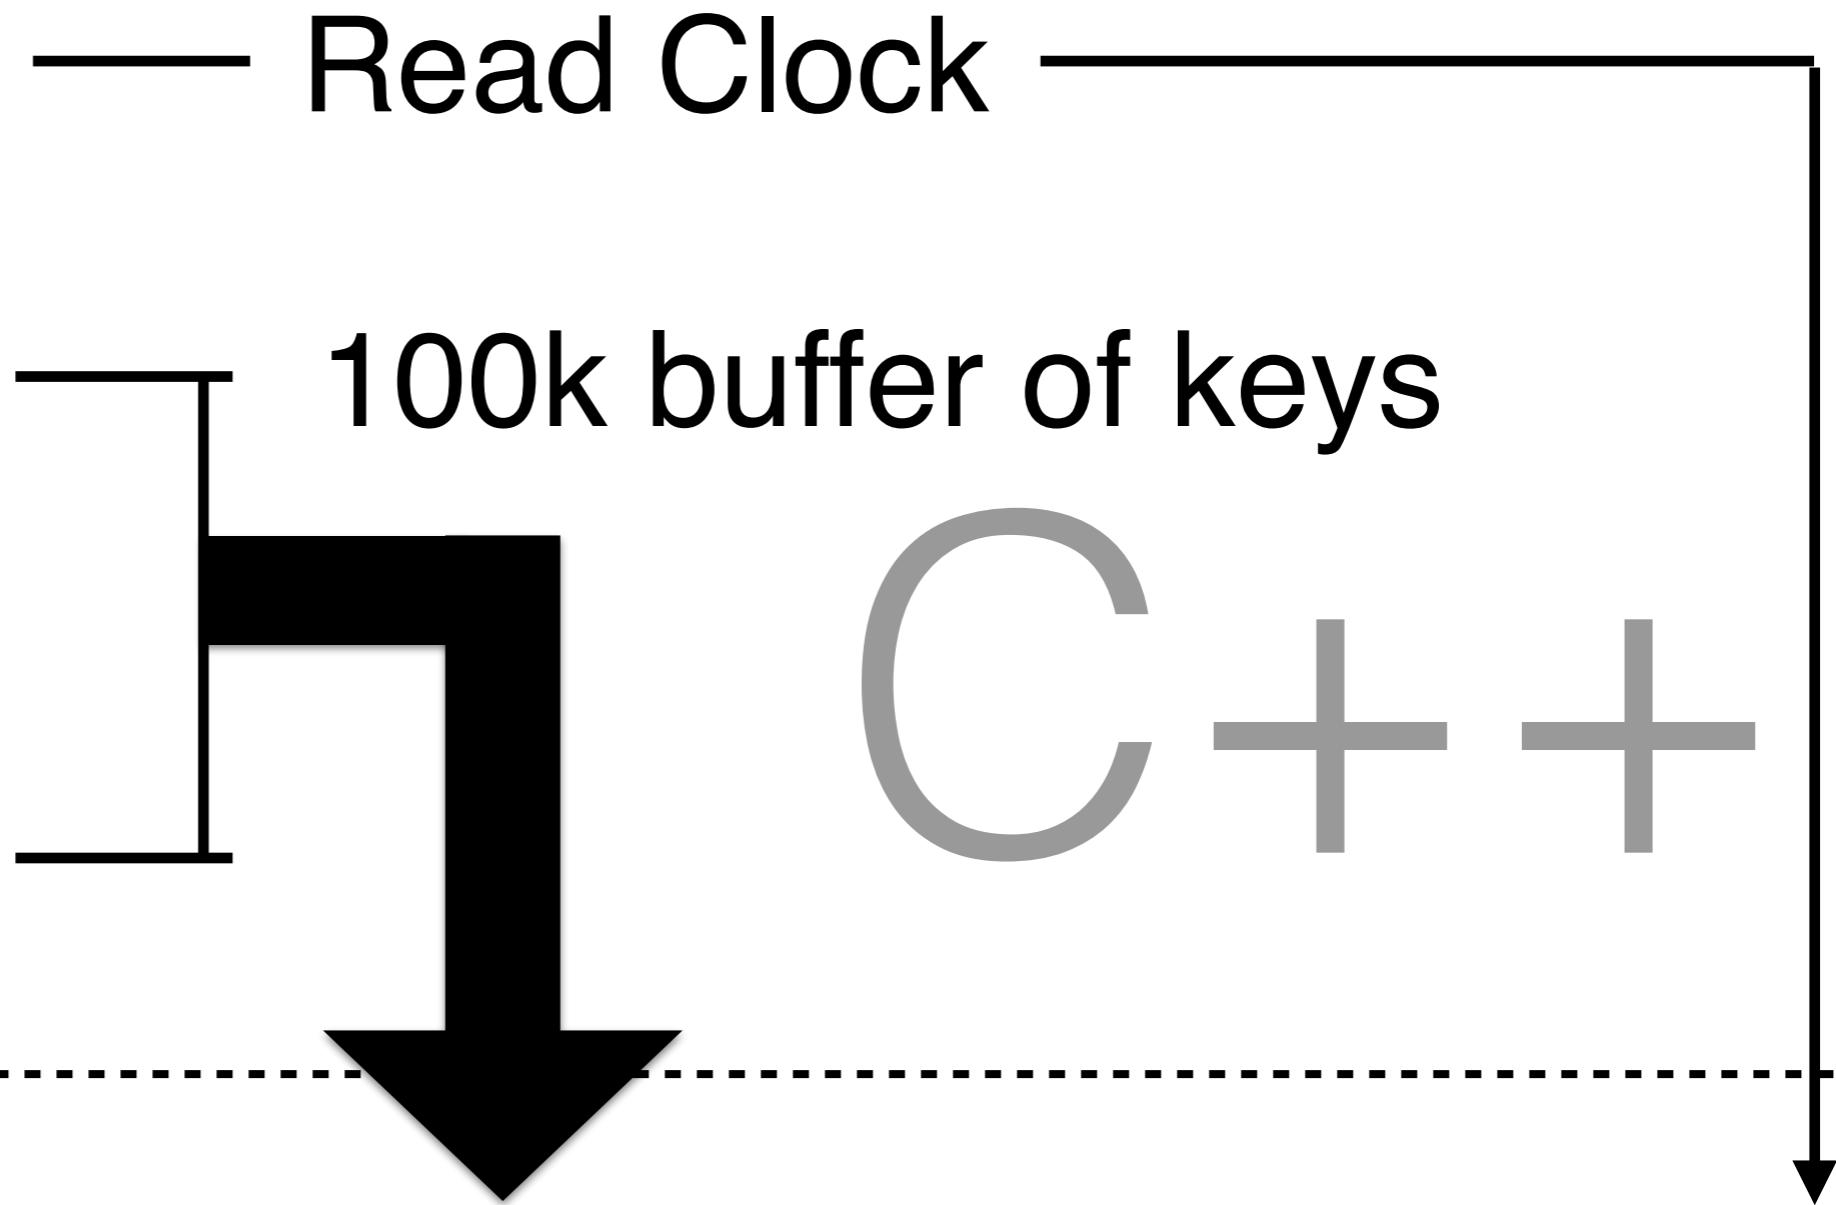
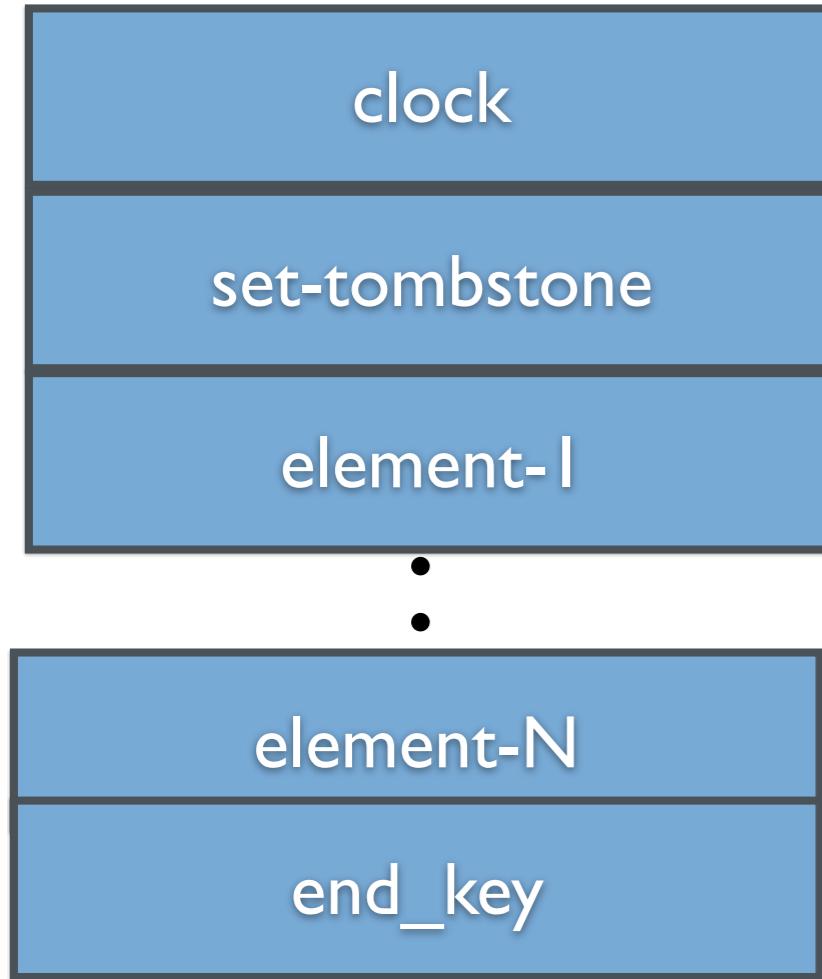
[{vnodeB, 4}]



C + +



# ERLANG



# ERLANG

<<Set, \0, \$c, \0, Actor, \0>>

<<Set, \0, \$e, \0, Element-l, \0,  
Actor, Cnt:/64/big-unsigned-  
integer>>

•  
•  
•  
•  
•

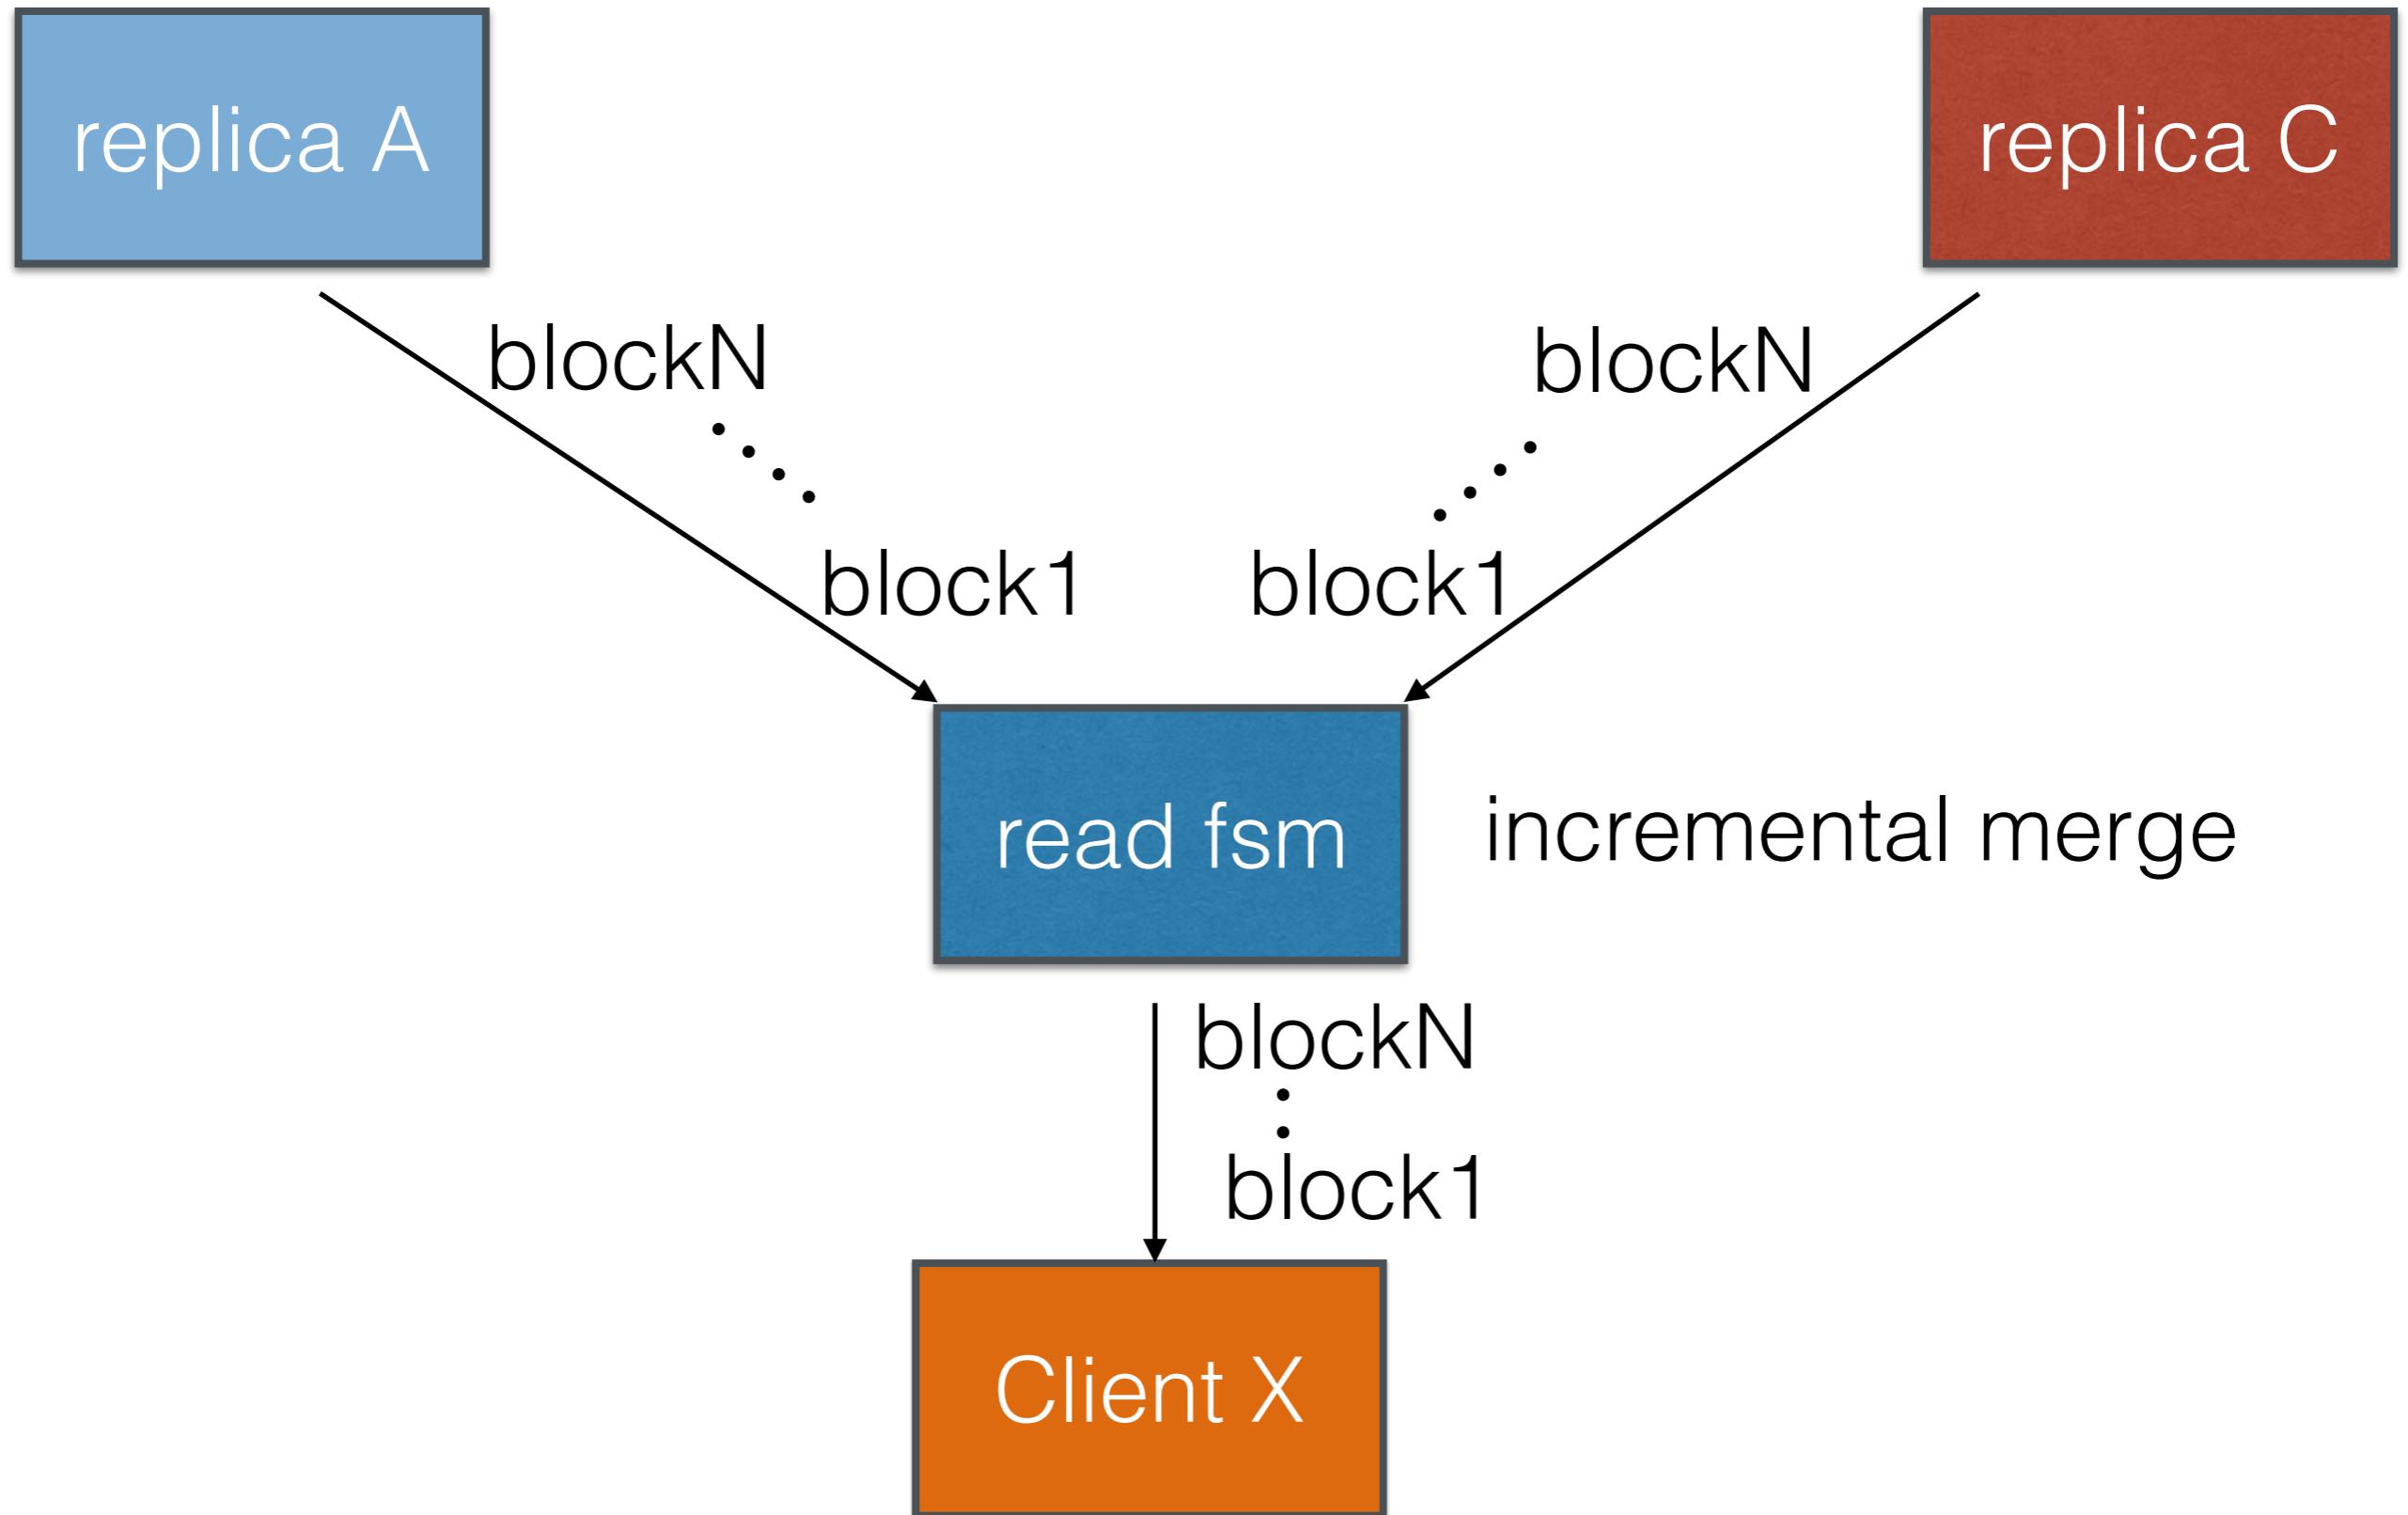
<<Set, \0, \$e, \0, Element-l, \0,  
Actor, Cnt:/64/big-unsigned-  
integer>>

<<Set, \0, %z, \0, \0>>

No Sext

No T2B

# Bigset Design: read

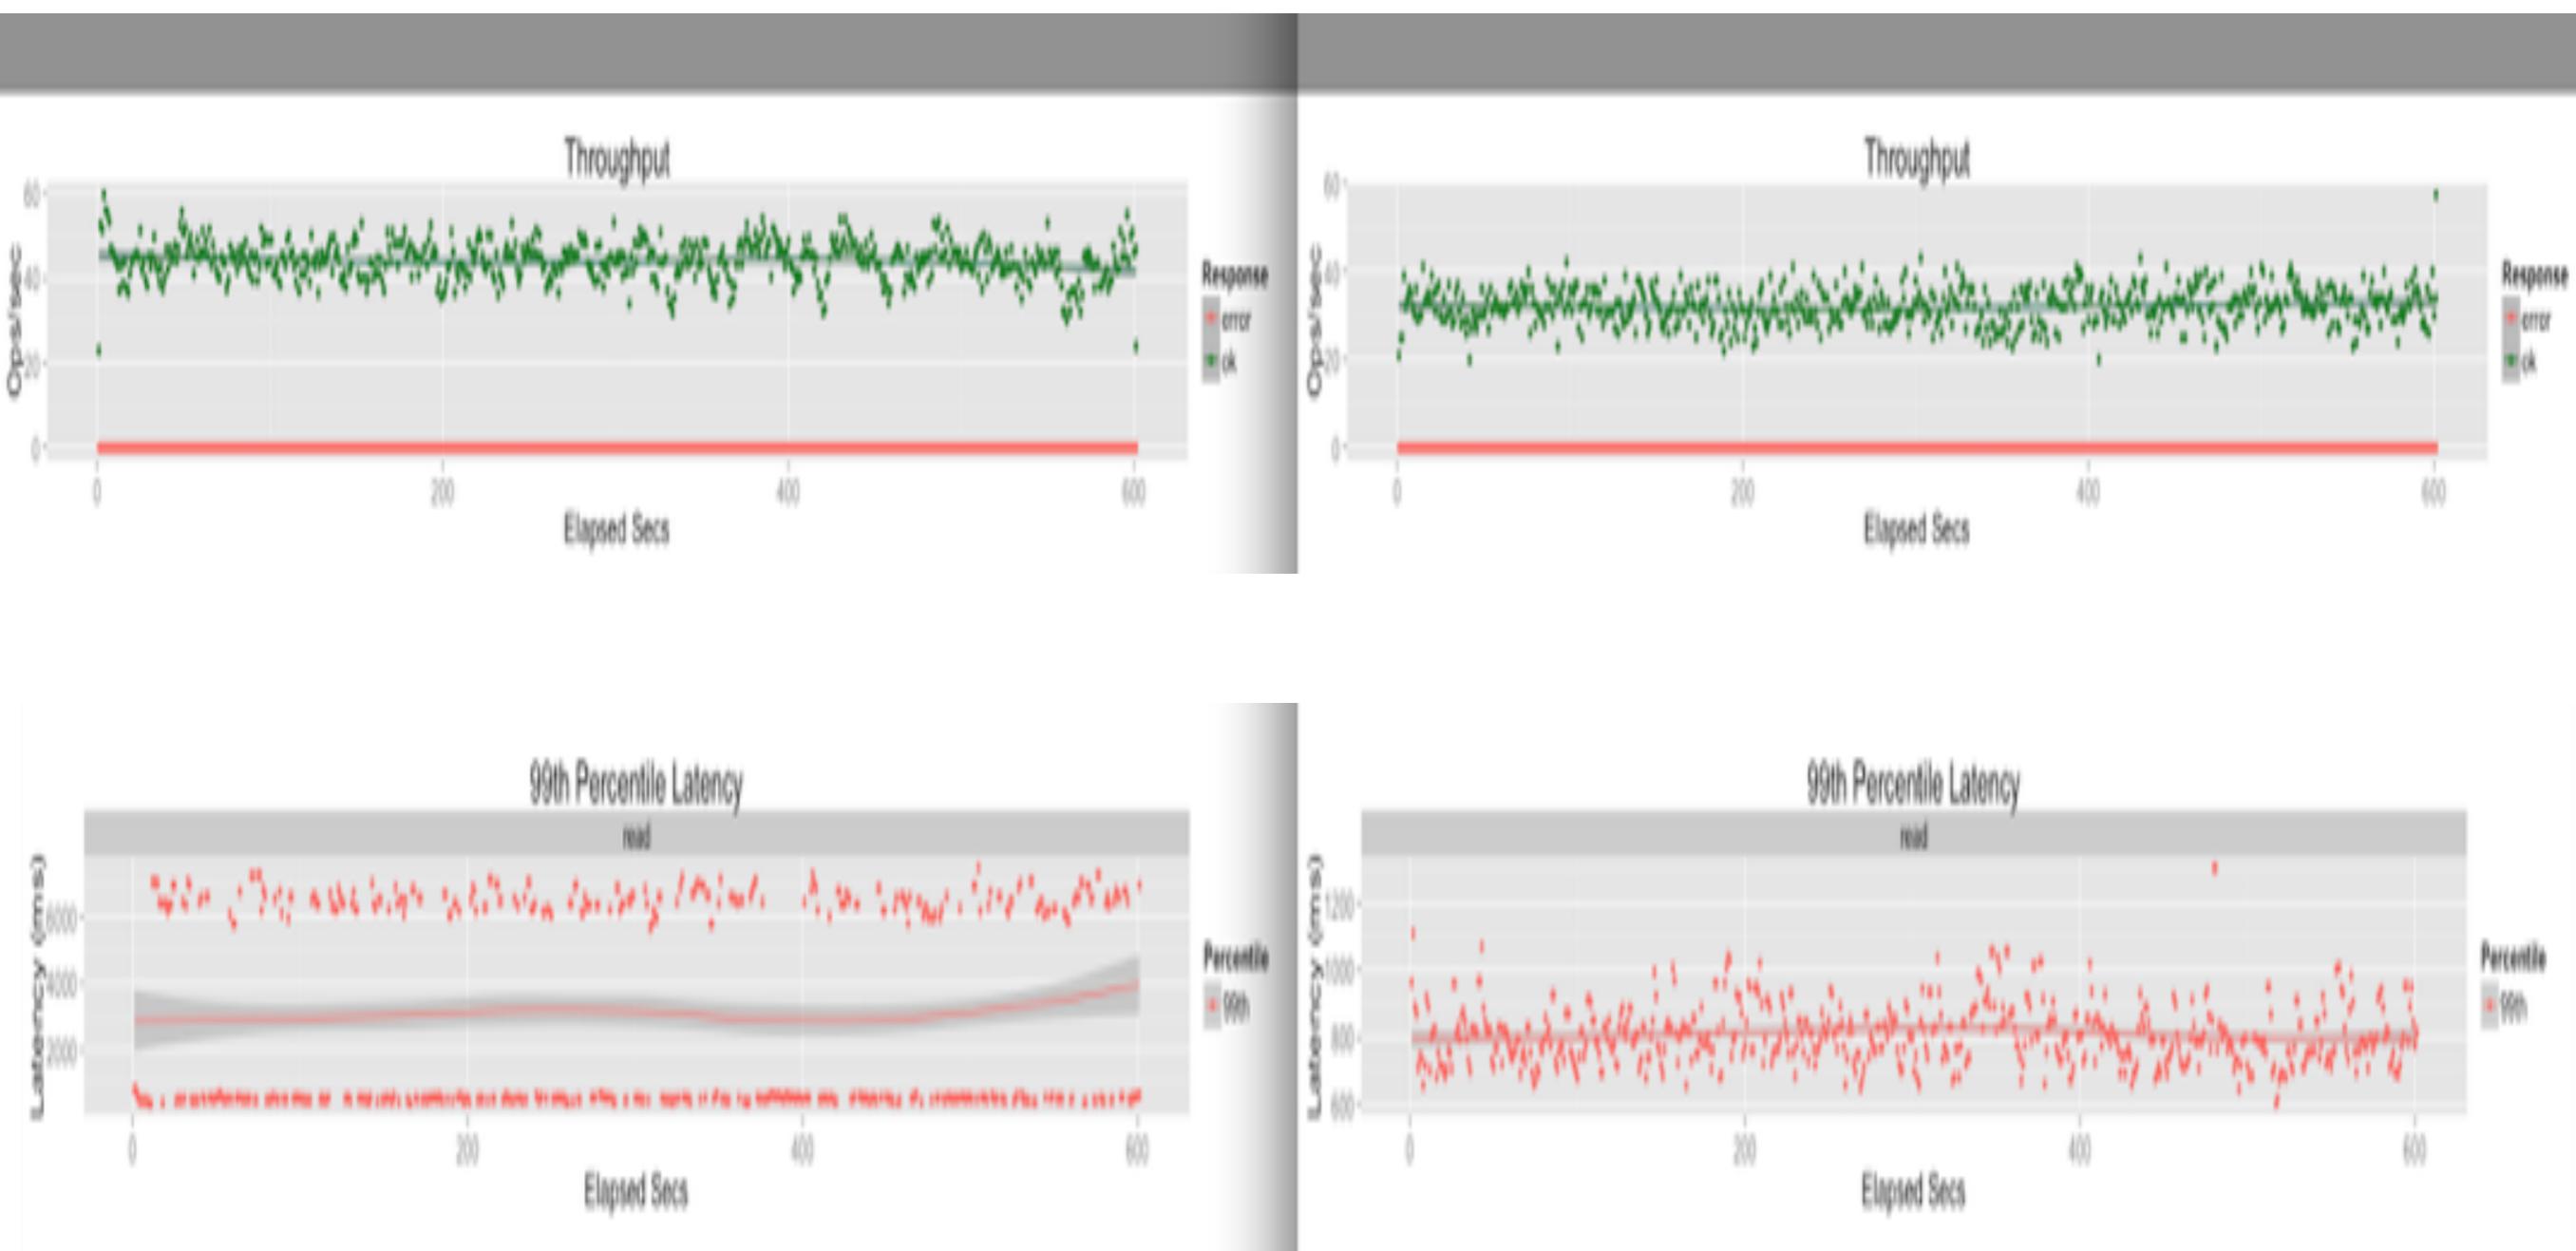


# Reads Today



10k sets, 100k elements, 20 workers - read

# Reads Today



10k sets, 100k elements, 20 workers - read

# Full Set Read or Queries?

Why read the whole set?

‘Cos you HAVE TO!

# Full Set Read or Queries?

## Why read the whole set?

# Full Set Read or Queries?

Why read the whole set?

‘Cos you HAVE TO!

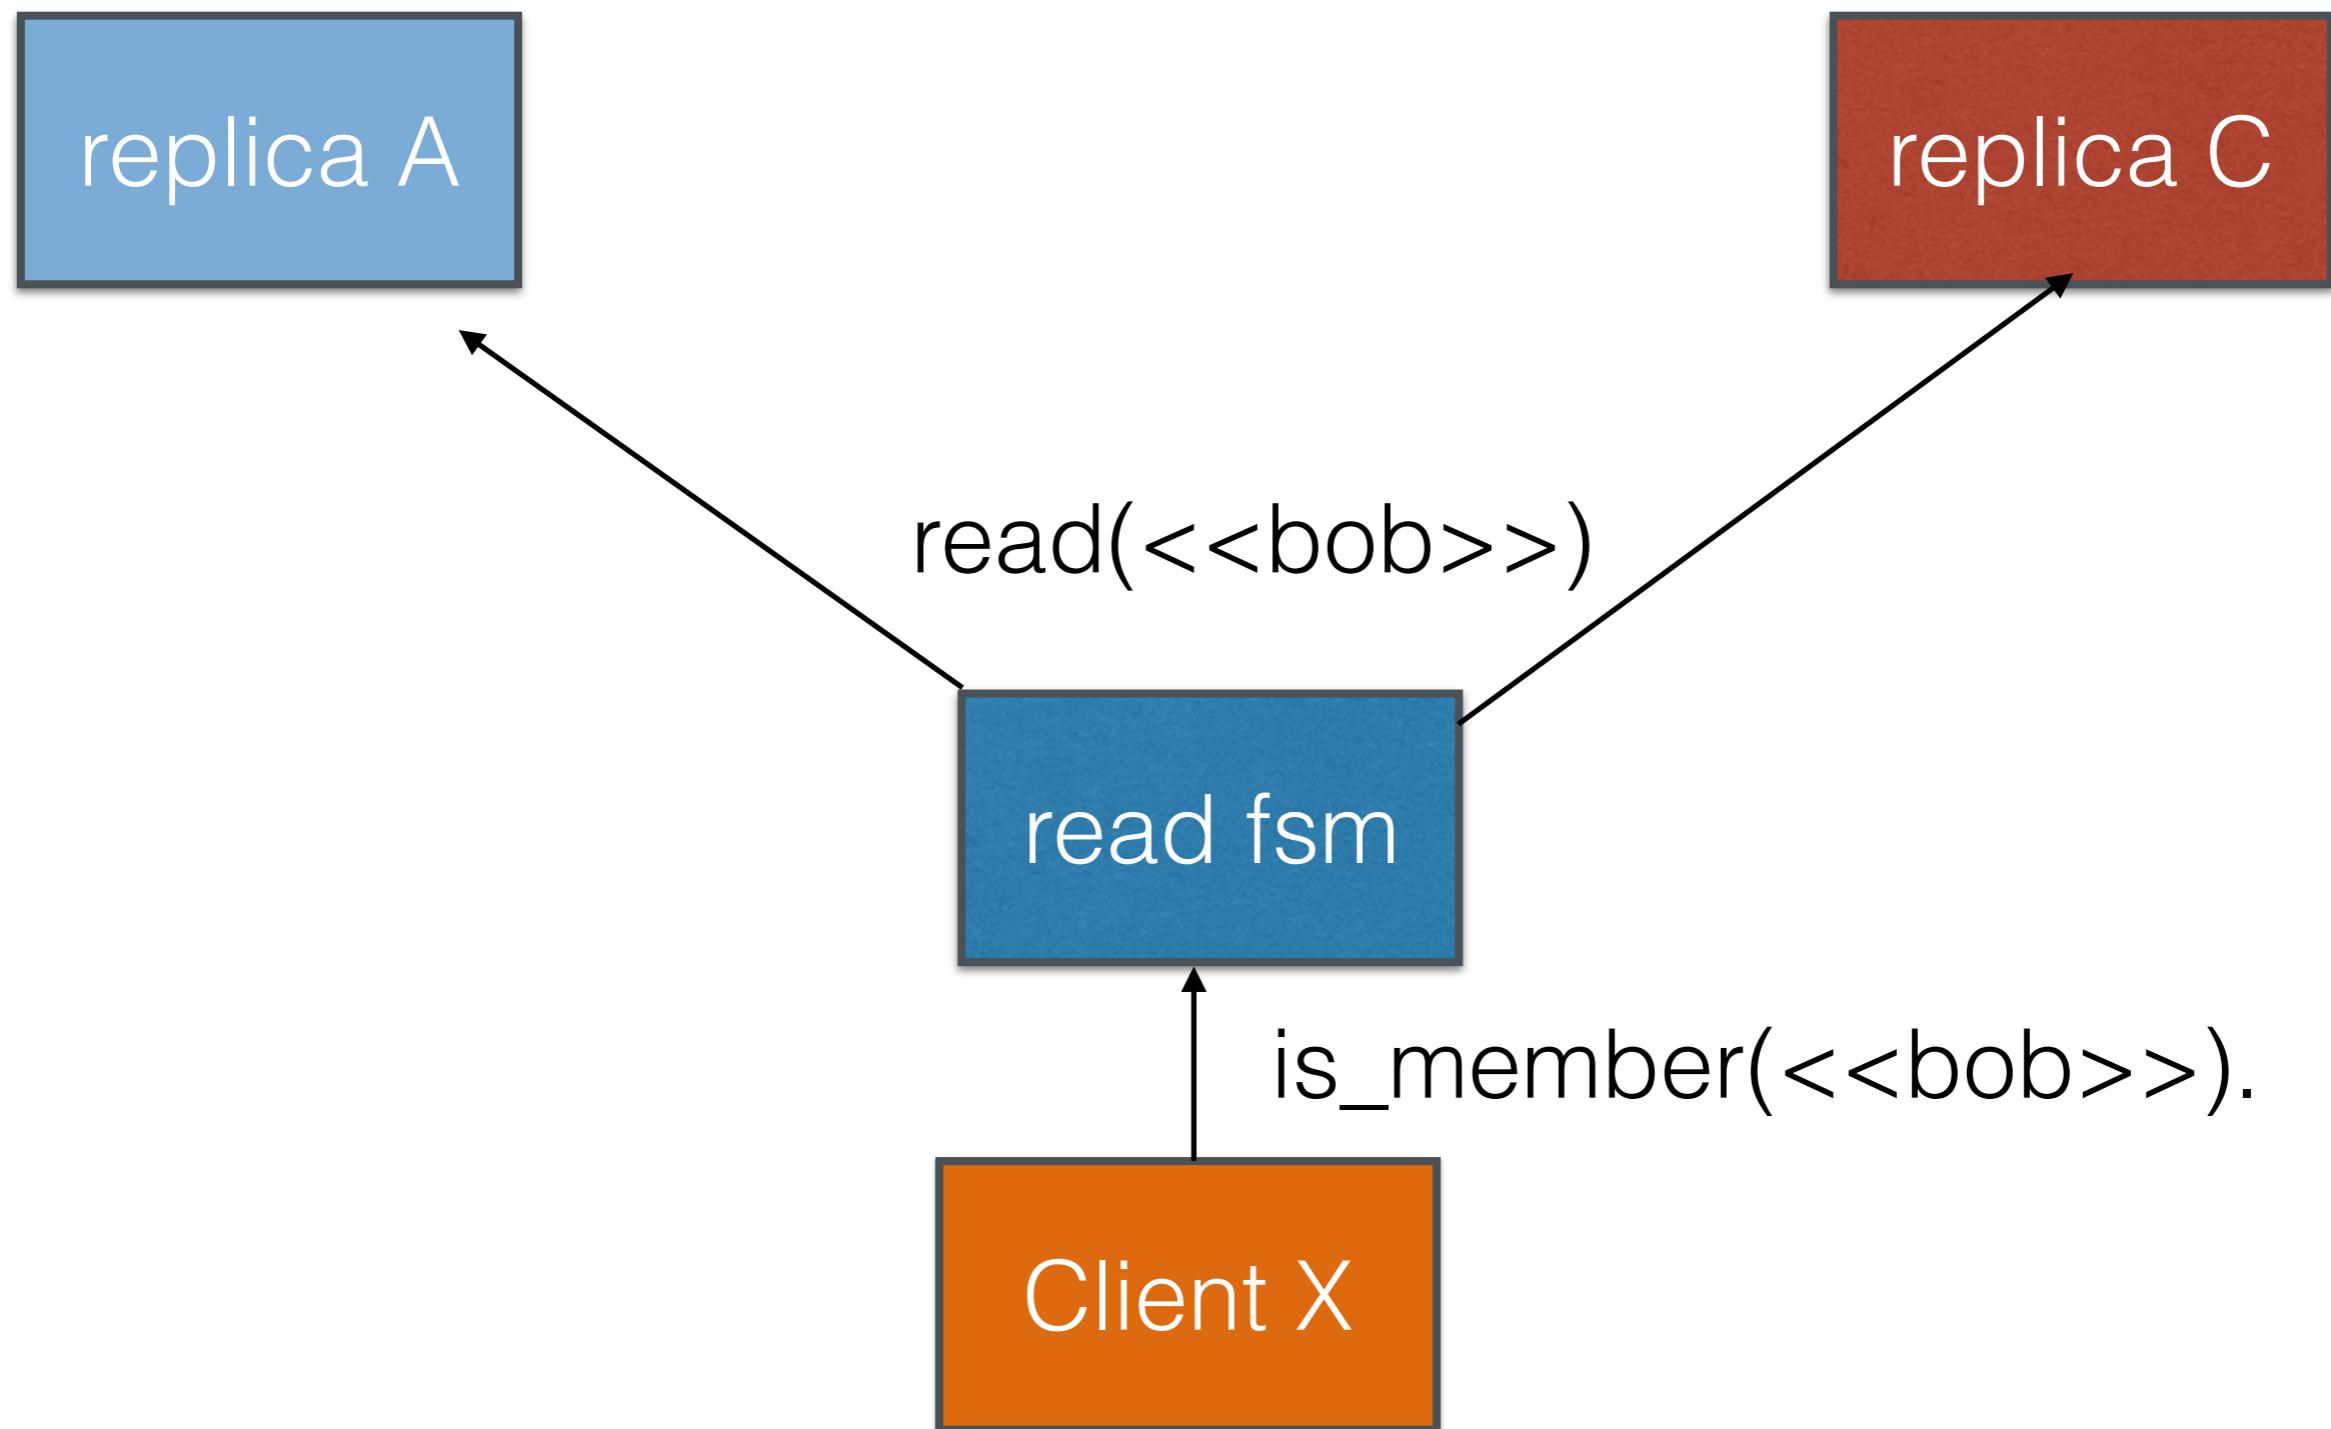
# Bigset Queries

- Subset
- Is Member?
- Range queries SORTED!
- Pagination

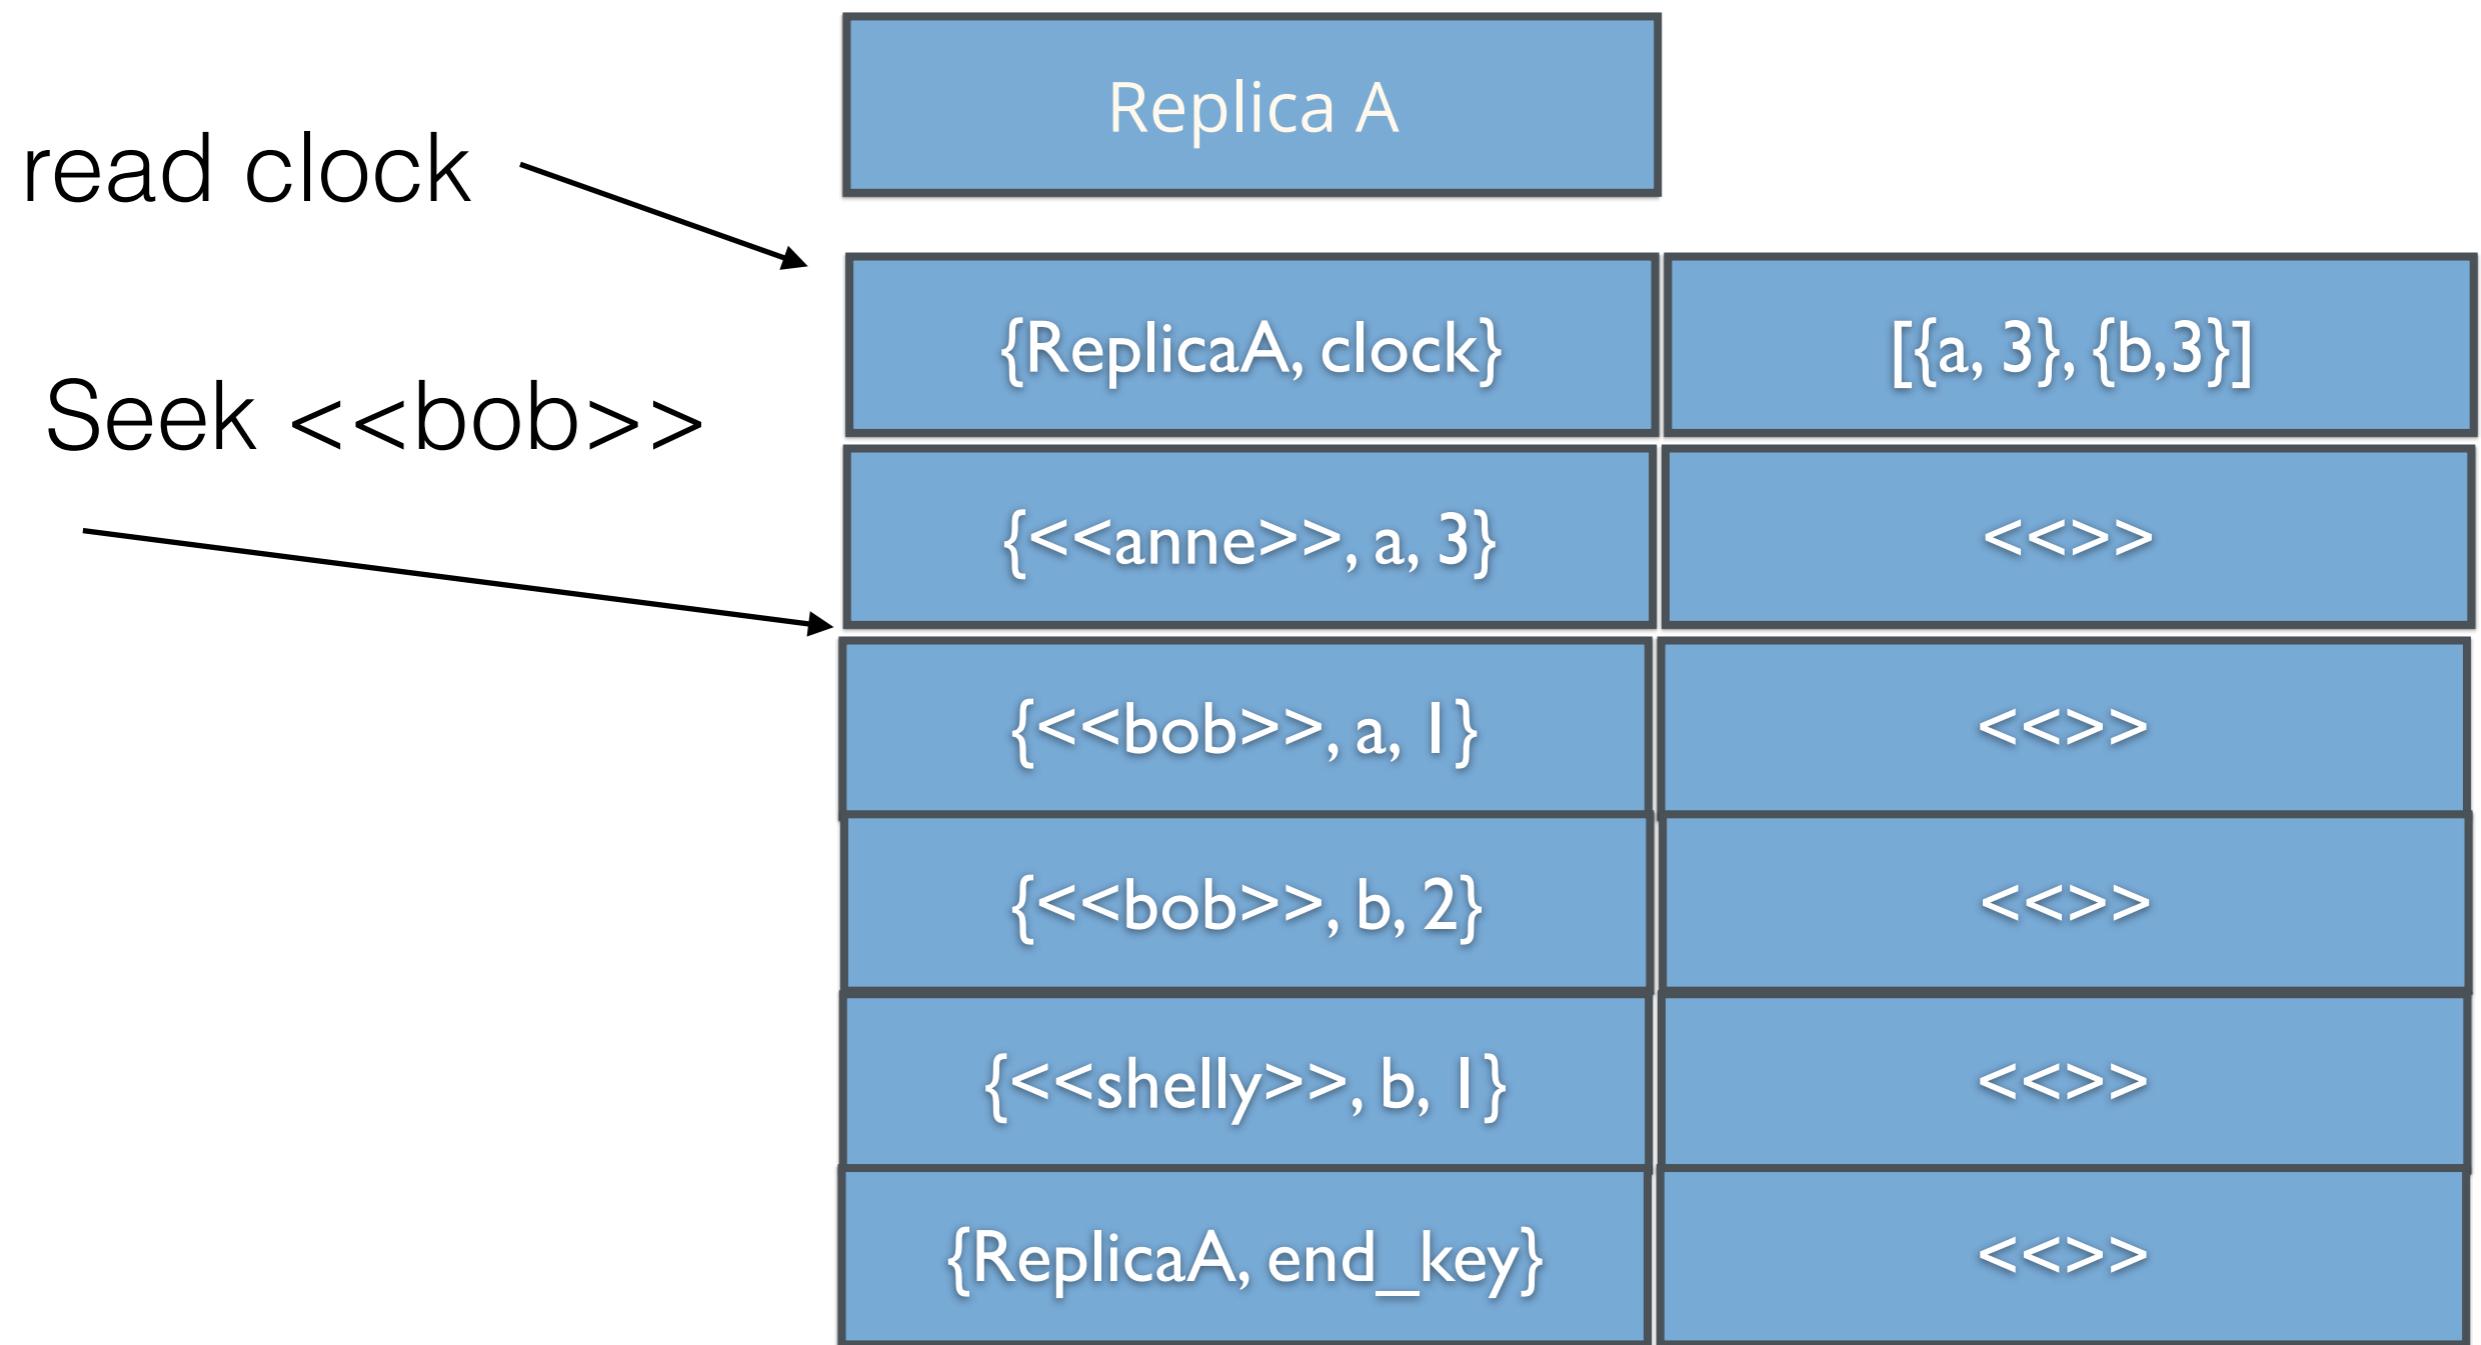
# Removes

- Observed-Remove - context
- Requires some kind of read
  - cheap membership check

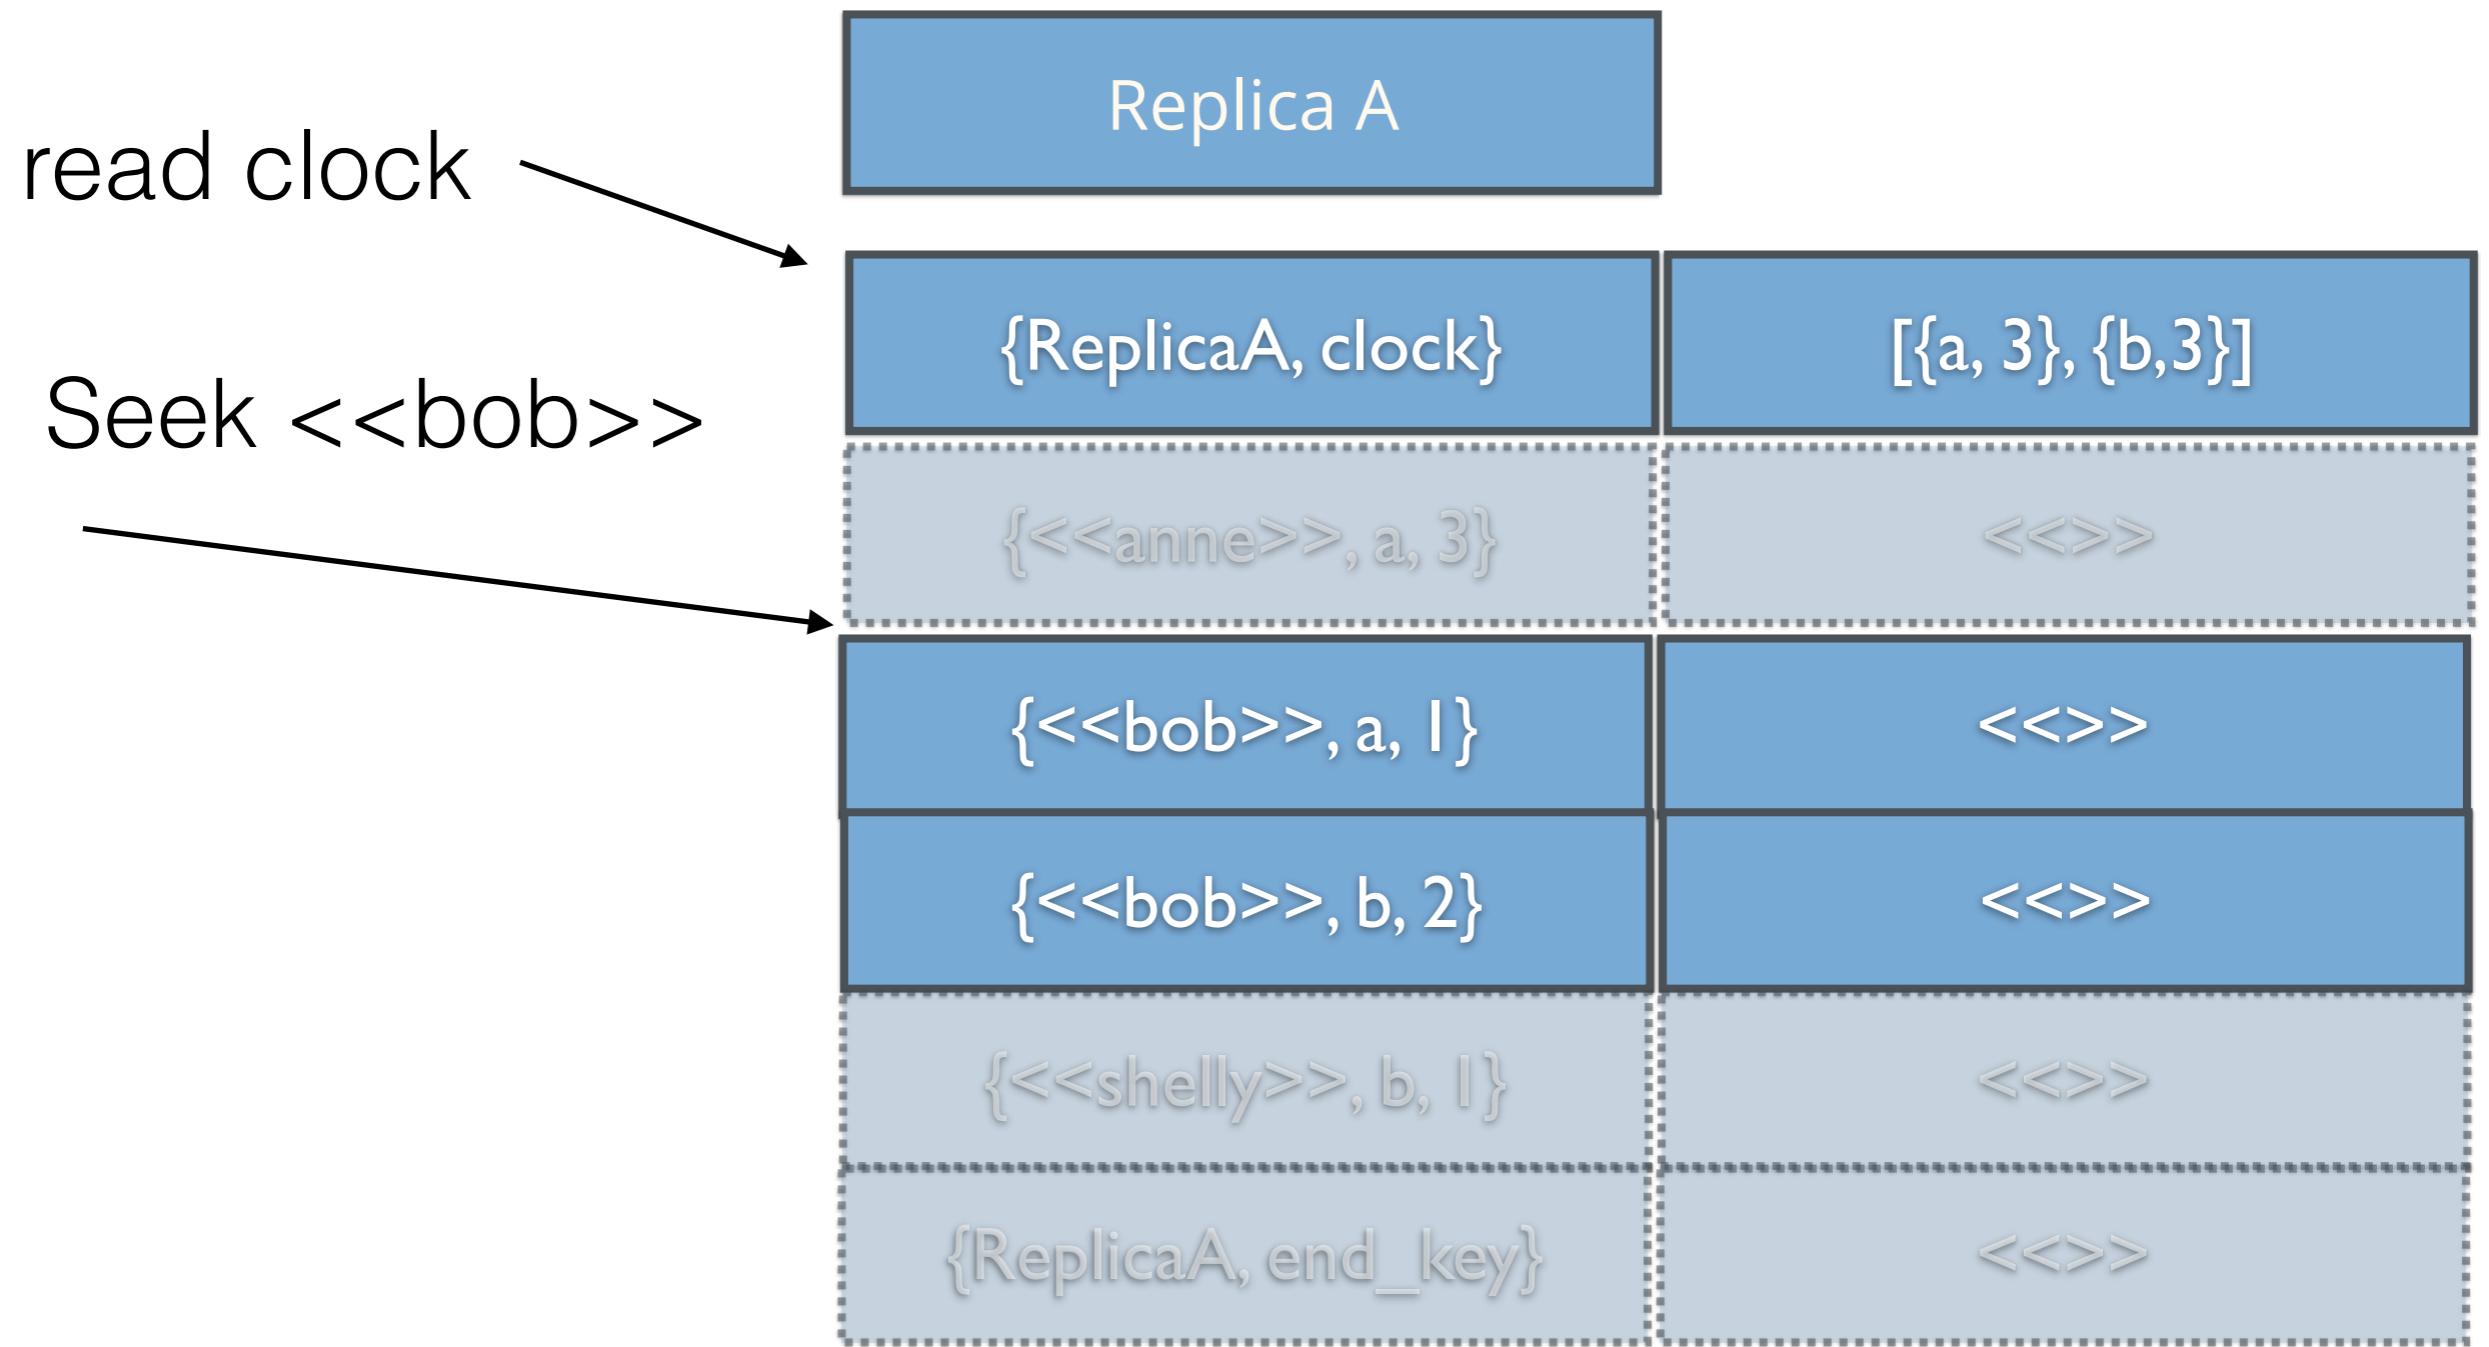
# Is Member(X)



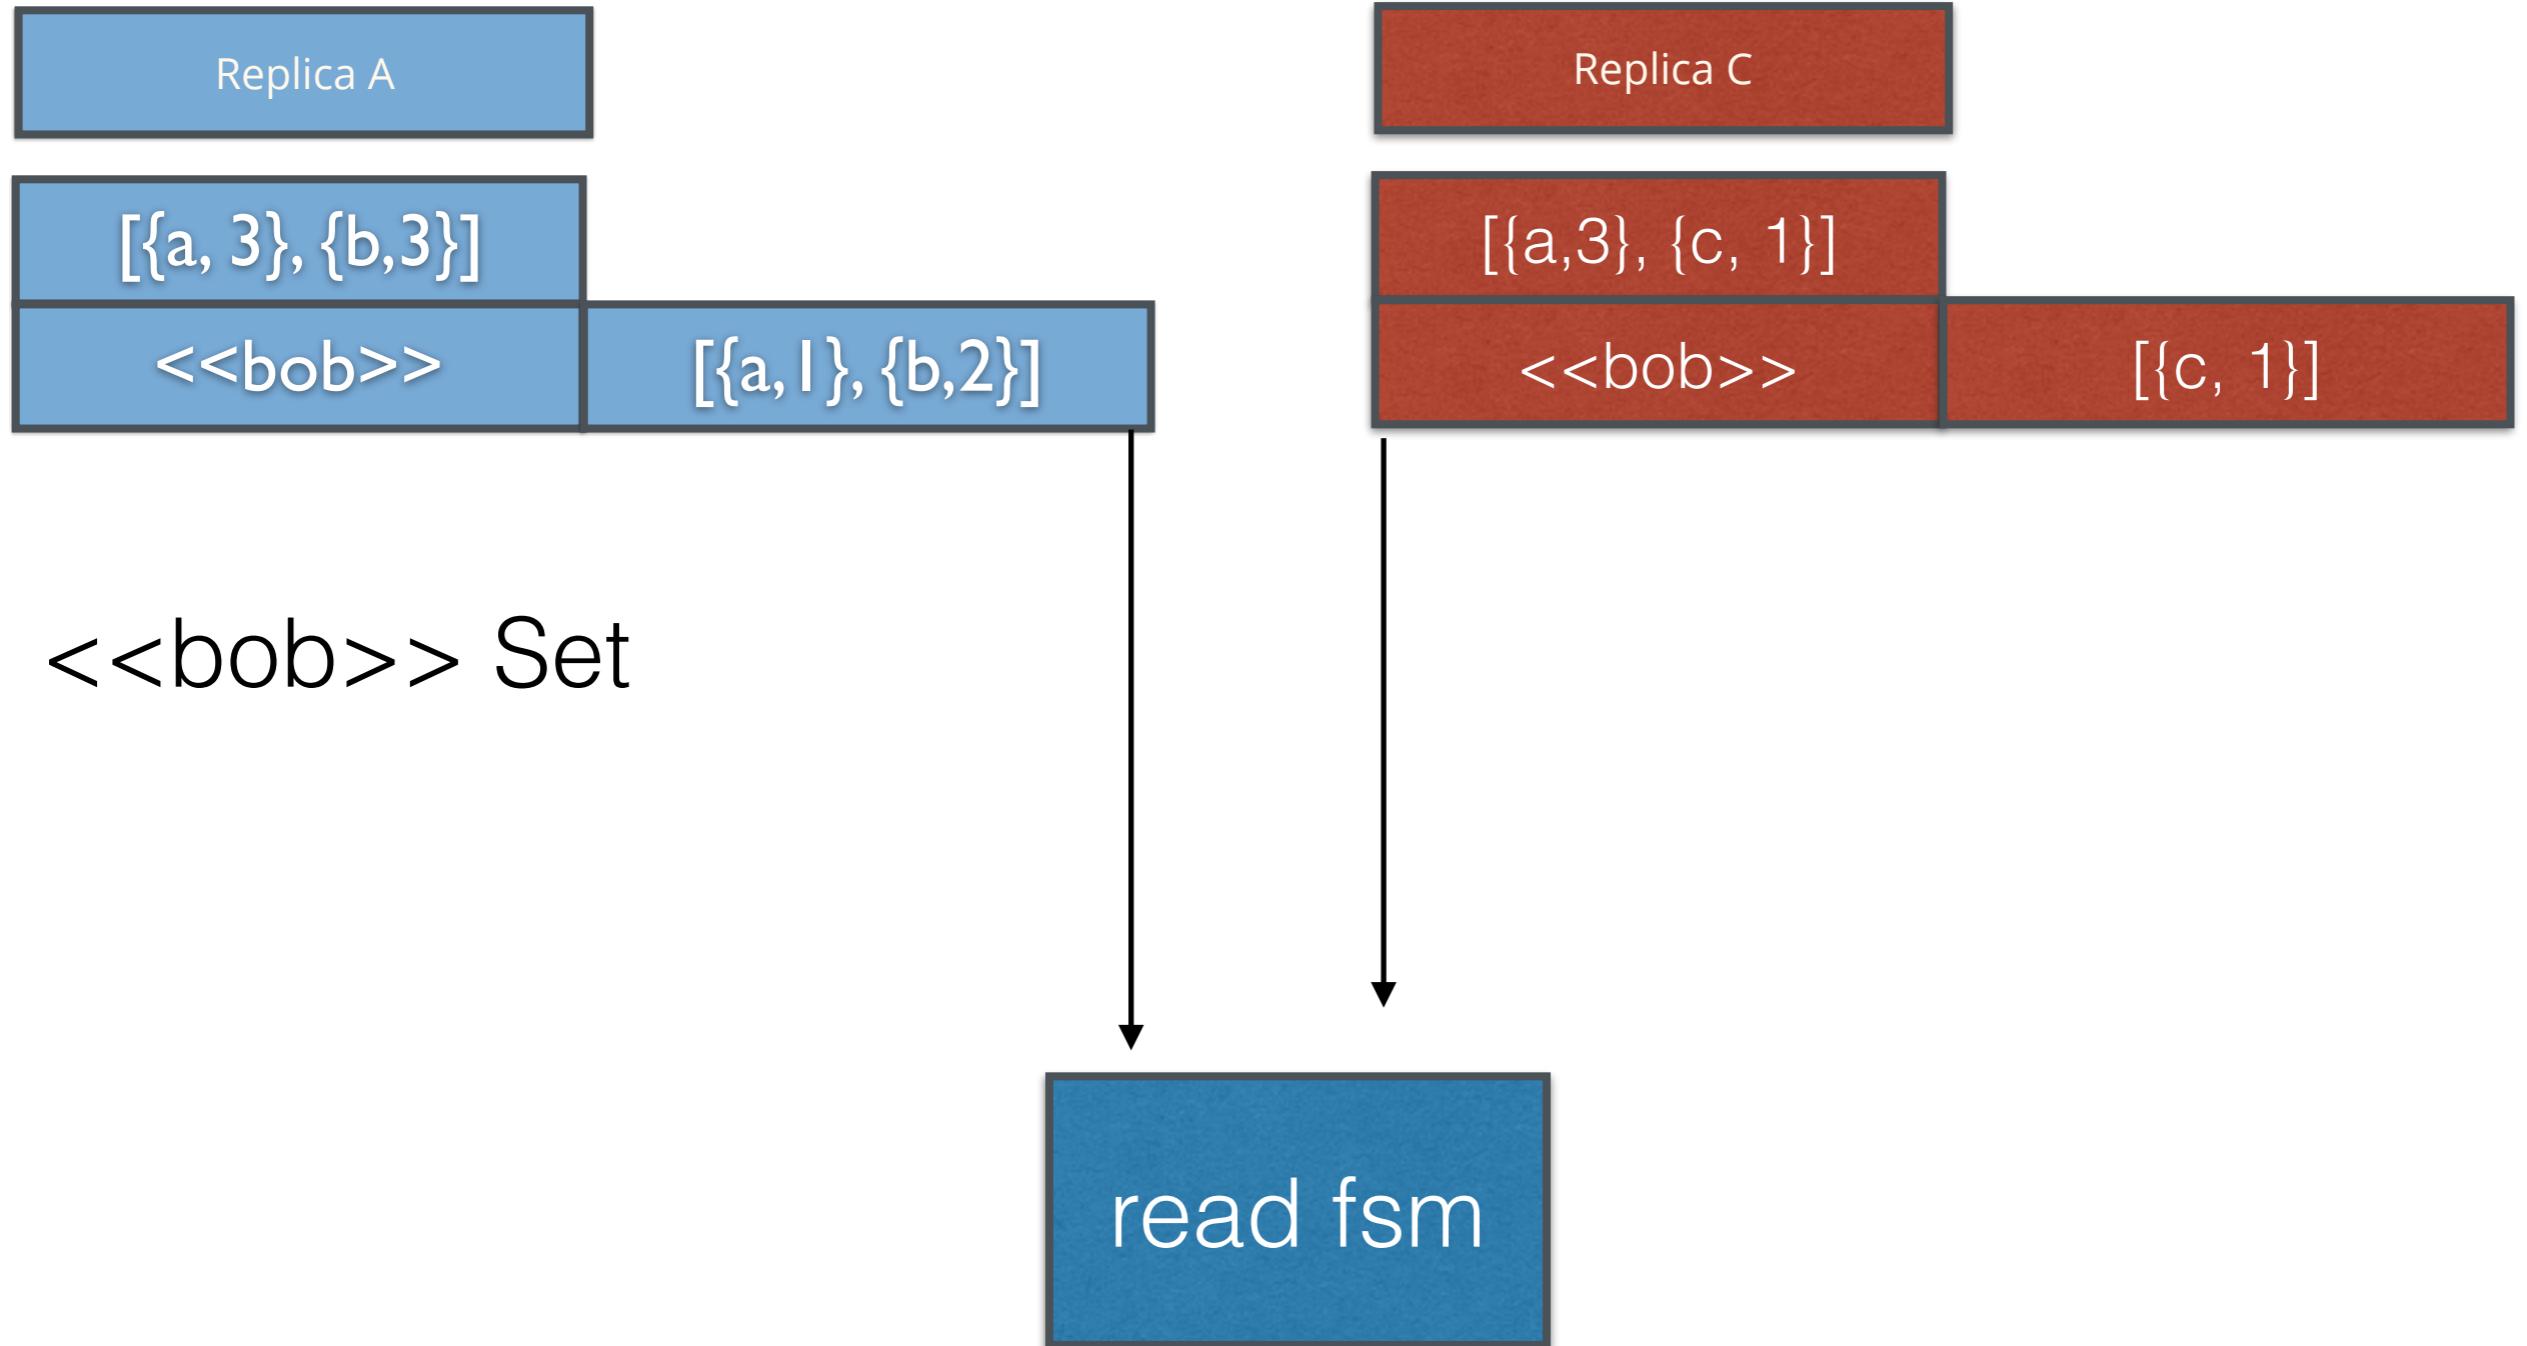
# Is Member(X)



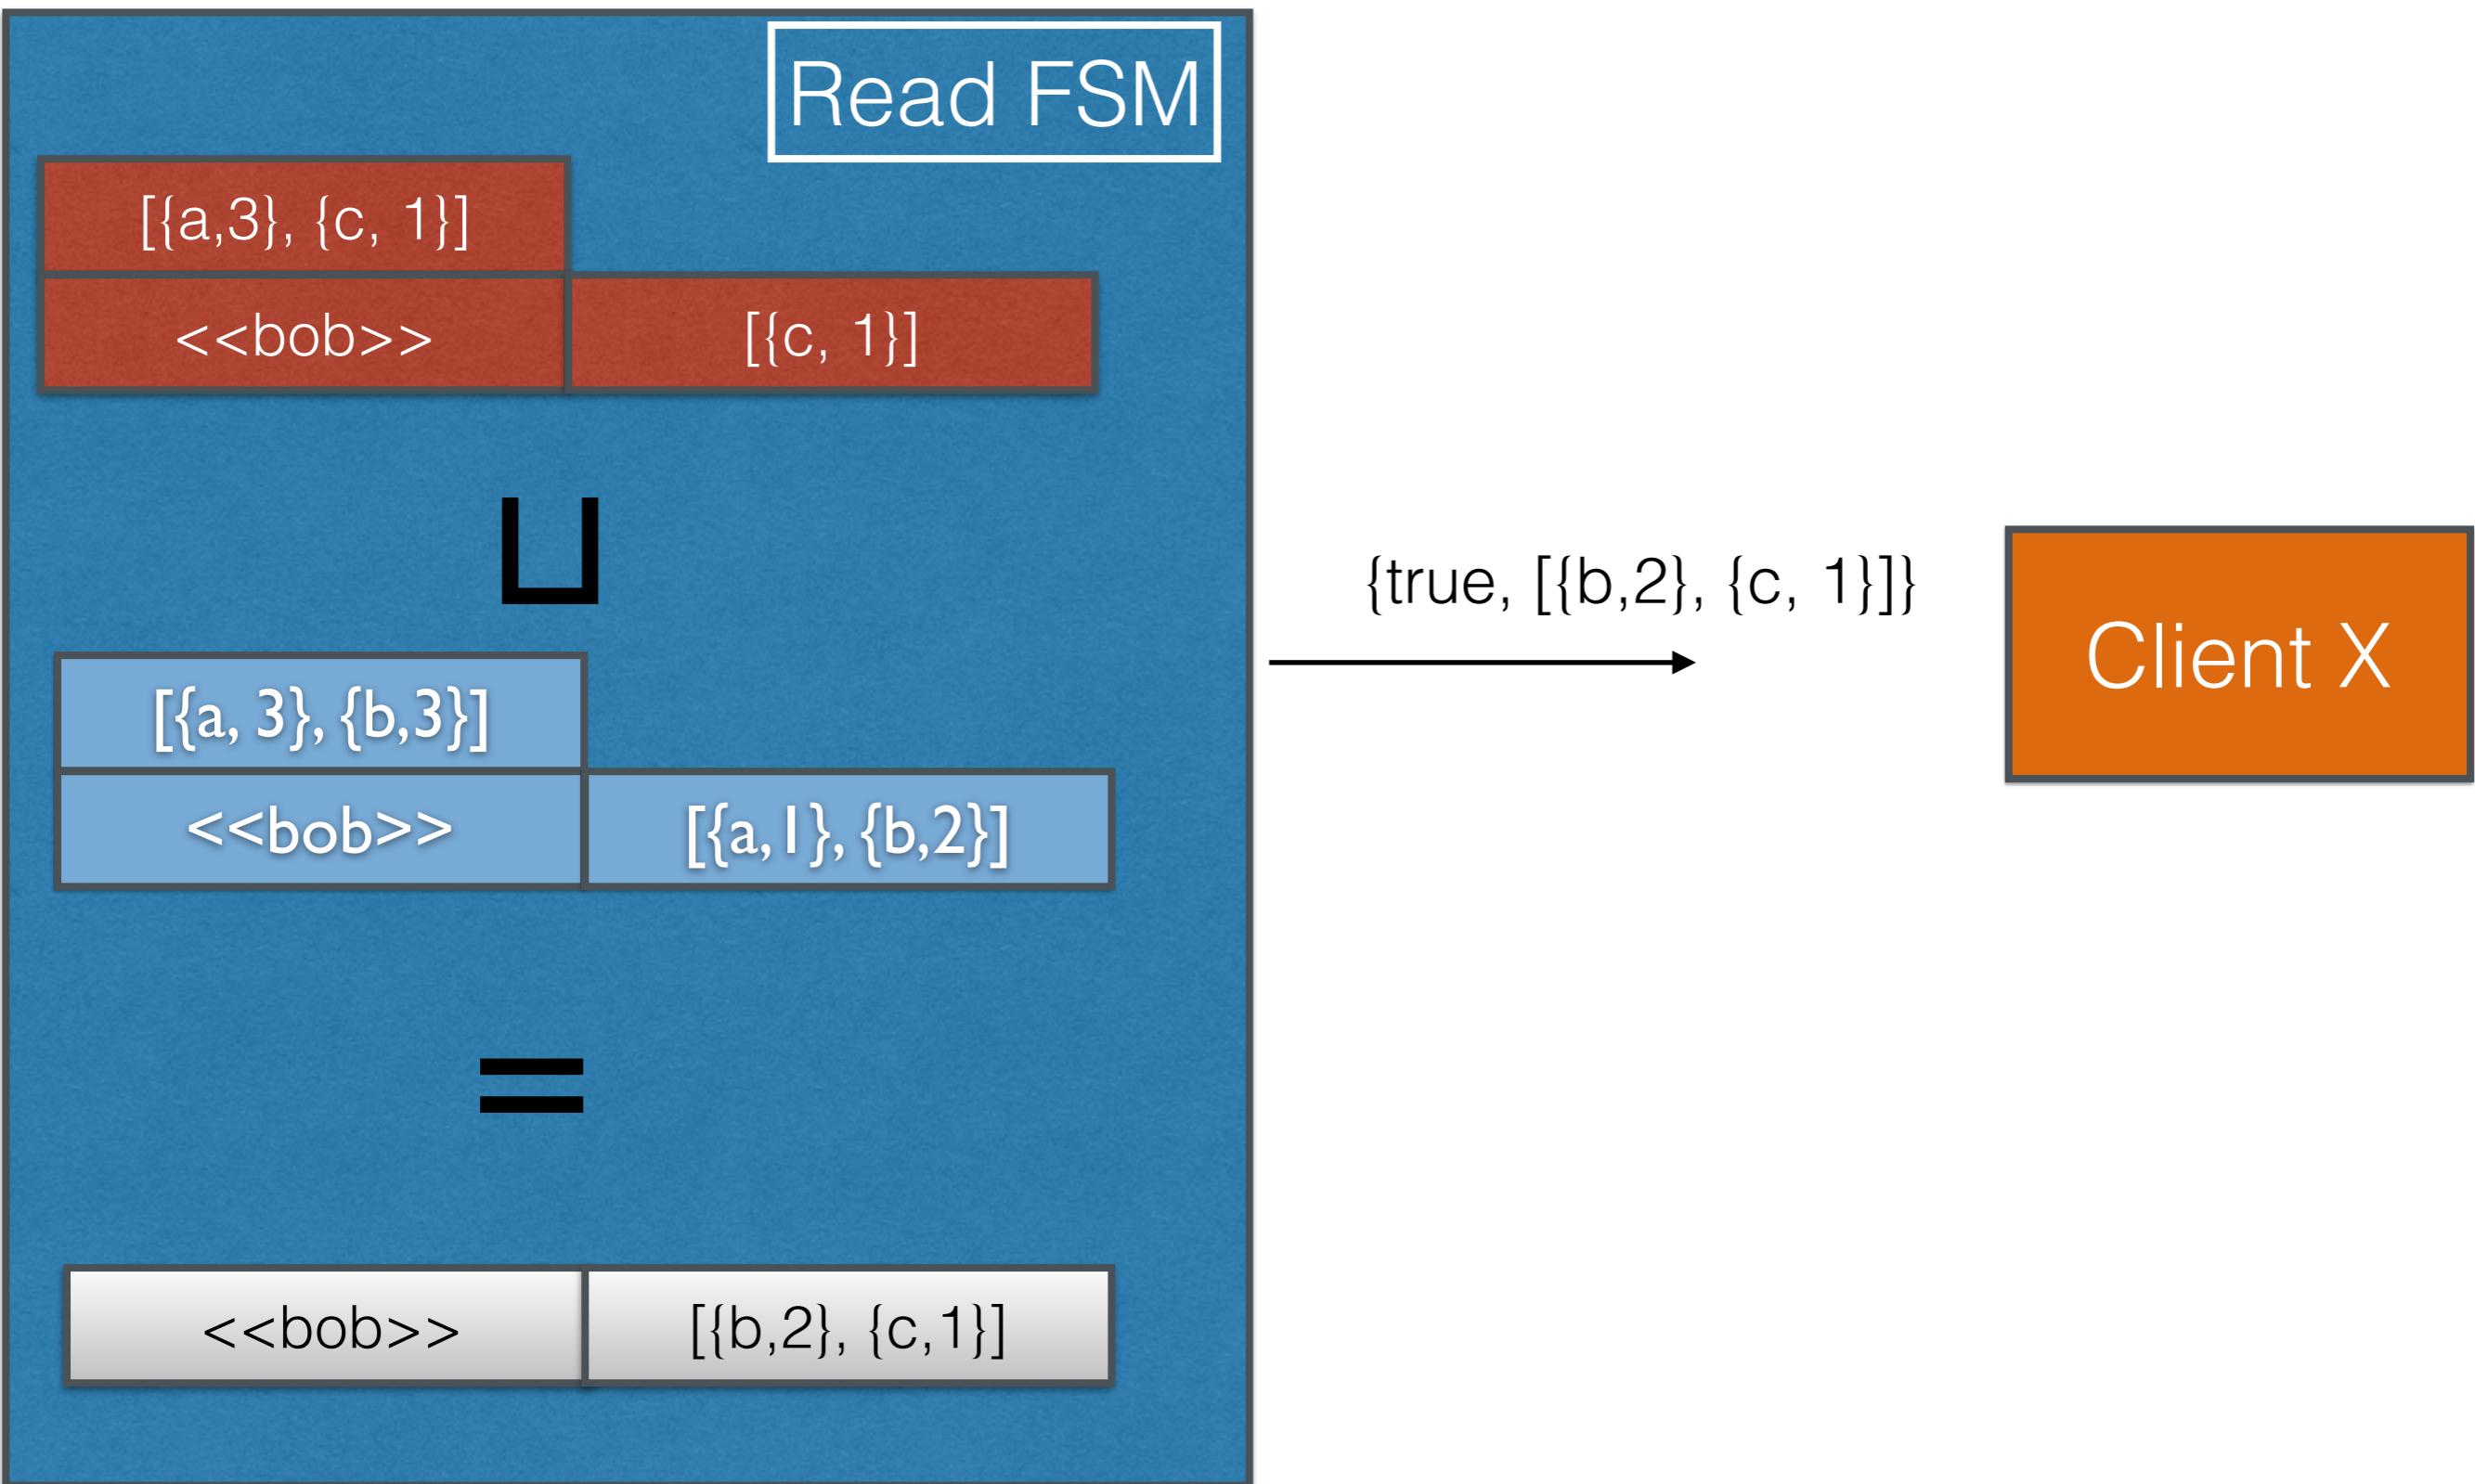
# Is Member(X)



# Is Member(X)



# Is Member(X)



# Next?

- Other “Big” Types - Maps
- Quorum Read Secondary Indexes
- Big Sets of Maps - Tables
- Joins? SQL?

# Summary

- CRDTs make eventual consistency easier on developers
- There exists an Optimised Add-Wins Set...
- It takes more than a lib
- A little engineering goes a long way

# Bigset Paper

<https://arxiv.org/abs/1605.06424>

**THANK YOU!**

---

**WE'RE HIRING!**

UK Client Service Engineer

Developer Advocate EMEA

[bashojobs.theresumator.com](http://bashojobs.theresumator.com)

**VISIT OUR STAND**

Experience our Riak TS demo and be  
entered to win a Scalextric set!

Get your invitation to our IoT Riak TS  
Roadshow

