

# **Bigsets are awesome**

## **My Helsing connection**

**Russell Brown 09/10/2025**

# What Even is a Bigset?

And what has it to do with anything?

CRDTs

CRDT Sets

Real Life vs Papers

Bigsets (2016)

Bigsets->DSon

Bigsets-rs (2025)

# Definition

A conflict-free replicated data type (CRDT) is an abstract data type, with a well defined interface, designed to be replicated at multiple processes and exhibiting the following properties: (i) any replica can be modified without coordinating with another replicas; (ii) when any two replicas have received the same set of updates, they reach the same state, deterministically, by adopting mathematically sound rules to guarantee state convergence.

<https://arxiv.org/abs/1805.06358>

# SOSP 2007 Dynamo

## Amazon Shopping Cart



Created by Amy Schwartz  
from the Noun Project



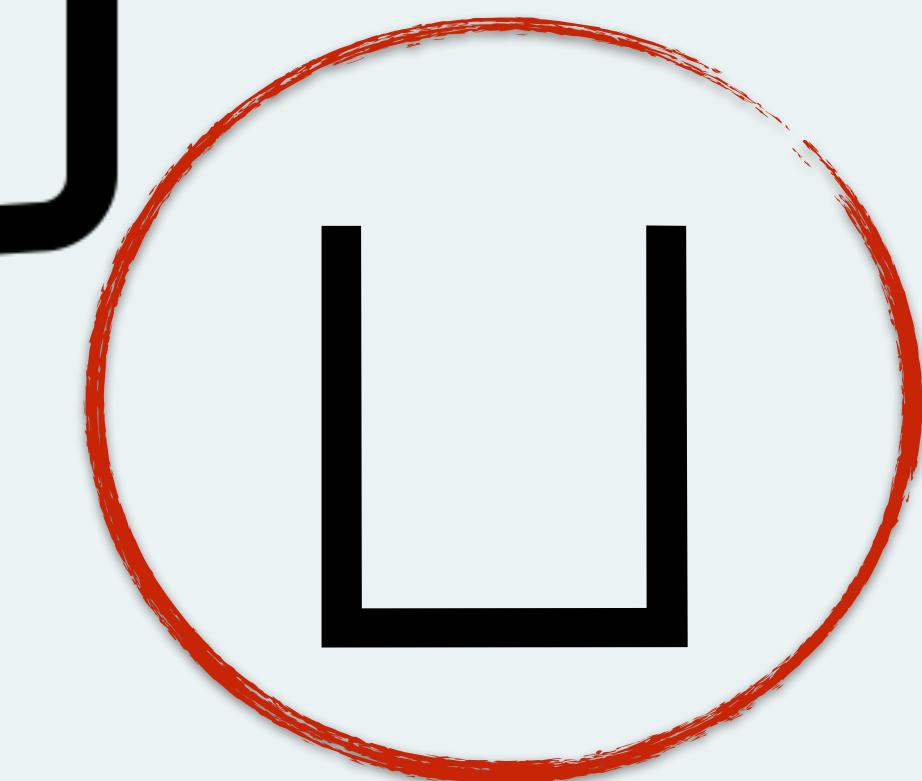
Created by Amy Schwartz  
from the Noun Project



**Created by Amy Schwartz  
from the Noun Project**



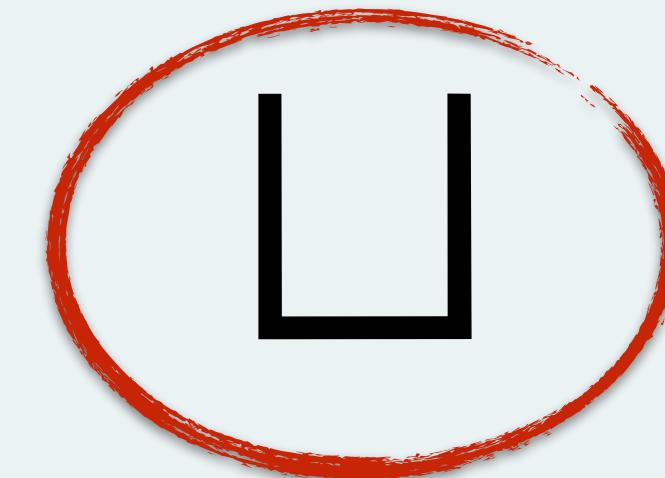
Created by Amy Schwartz  
from the Noun Project



Created by Amy Schwartz  
from the Noun Project



Created by Amy Schwartz  
from the Noun Project



Created by Amy Schwartz  
from the Noun Project



What about removes?

Created by Amy Schwartz  
from the Noun Project

```
/// basic add wins set
#[derive(Clone, Debug)]
struct AddWinsSet {
    adds: BTreeMap<Bytes, BTreeSet<u64>>,
    removes: BTreeMap<Bytes, BTreeSet<u64>>,
}
```

```
fn members(&self) -> BTreeSet<Bytes> {
    let mut members = BTreeSet::new();
    for (elem, tags) in self.adds.iter() {
        if self.removes.get(elem).map_or(true, |r| !tags.is_subset(r)) {
            members.insert(elem.clone());
        }
    }
    members
}
```

```
fn add(&mut self, elem: Bytes, tag: u64) {
    if let Some(tags) = self.adds.get(&elem) {
        self.removes.entry(elem.clone()).or_default().extend(tags);
    }

    self.adds.entry(elem).or_default().insert(tag);
}
```

```
fn remove(&mut self, elem: Bytes) {
    if let Some(tags) = self.adds.get(&elem) {
        self.removes.entry(elem).or_default().extend(tags);
    }
}
```

```
fn merge(&mut self, other: Self) {
    for (elem, tags) in other.removes {
        self.removes.entry(elem.clone()).or_default().extend(tags);
    }

    for (elem, tags) in other.adds {
        self.adds.entry(elem.clone()).or_default().extend(tags);
    }
}
```



Created by Amy Schwartz  
from the Noun Project

Merge



Created by Amy Schwartz  
from the Noun Project



Created by Amy Schwartz

# Optimisations

**Add Wins Sets Without Tombstones**

**Delta Replication (various flavours)**

Papers: One CRDT in memory

# Paper vs Production

One CRDT in Memory vs Many CRDTs on Disk

Deltas are great on the network BUT quadratic on disk

(Read-Deserialise-Mutate-Serialise-Write)

(Send-Delta)

(Read-Deserialise-Apply Delta-Serialise-Write)

# Bigsets

Decompose CRDT elements on Disk

Read minimum  
Generate Delta  
Write minimum

Send Delta  
Read minimum  
Apply Delta

Write minimum

(Trade-Off? Full Set Read is Slow (but enables querying over elements))

# Bigset -> DSON

- Bigsets “paper” <https://dl.acm.org/doi/10.1145/2911151.2911156>
- $\Delta$ -CRDTs: making  $\delta$ -CRDTs delta-based <https://dl.acm.org/doi/10.1145/2911151.2911163>
- Efficient Synchronisation of State-based CRDTs <https://arxiv.org/abs/1803.02750>
- DSON: JSON CRDT using delta-mutations for document stores <https://dl.acm.org/doi/10.14778/3510397.3510403>

# **Bigsets-rs**

<https://github.com/russelldb/bigsets-rs>